Output Form (OF)

Score: 1/5 — Serious concern | Confidence: high

***Benchmark:** CNN/DAILYMAIL* | ***Context:** Ecuador Humanitarian Crisis Tweet Analysis — NGO/IO Analyst Deployment*

Output Form definition: Whether the expected output modality matches regional deployment needs and accessibility.

> Each section below is followed by an evidence table. Three types of evidence are cited: **[QN]** — verbatim quotes extracted from the benchmark paper; **[WEB-N]** — web sources consulted during assessment; **[DN]** / **[DATASET-DN]** — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

Output form is single-token classification accuracy on a closed entity vocabulary [Q34, Q52], with hyperparameter tuning on validation [Q58, Q60, Q61] and attention heat-map interpretability [Q79, Q85, Q86, Q87, Q88]. The deployment requires temporally ordered, multi-sentence narrative timeline summaries with geographic-anchoring precision, institutional-attribution accuracy, temporal coherence, operational-relevance calibration, and bilingual-workflow support per YAML output_task_requirements. The HF dataset's summarization variant produces unordered abstractive bullet lists [DATASET-D1, D5] — also not a temporal timeline. Even at the closest structural overlap (abstractive summarization), the form is misaligned: no temporal sequencing, no phase transitions, no needs-typed multi-sentence narrative output. Memory constraints prevented architectural depth experiments [Q69], and the authors flag the fixed-width hidden vector as an information bottleneck [Q45]. HIGH-priority dimension.

ID	Evidence
Q34	""We propose three neural models for estimating the probability of word type a from document d answering query q: p(a\" (p.4)
Q45	"The fixed width hidden vector forms a bottleneck for this information flow that we propose to circumvent using an attention mechanism inspired by recent results in translation and image recognition [6, 7]." (p.5)
Q52	"Having described a number of models in the previous section, we next evaluate these models on our reading comprehension corpora." (p.6)
Q58	"All model hyperparameters were tuned on the respective validation sets of the two corpora." (p.6)
Q60	"For the Deep LSTM Reader, we consider hidden layer sizes [64, 128, 256], depths [1, 2, 4], initial learning rates [1E−3, 5E−4, 1E−4, 5E−5], batch sizes [16, 32] and dropout [0.0, 0.1, 0.2]. We evaluate two types of feeds. In the cqa setup we feed first the context document and subsequently the question into the encoder, while the qca model starts by feeding in the question followed by the context document. We report results on the best model (underlined hyperparameters, qca setup)." (p.6)
Q61	"For the attention models we consider hidden layer sizes [64, 128, 256], single layer, initial learning rates [1E−4, 5E−5, 2.5E−5, 1E−5], batch sizes [8, 16, 32] and dropout [0, 0.1, 0.2, 0.5]." (p.6)
Q69	"Memory constraints prevented us from experimenting with deeper Attentive Readers." (p.7)
Q79	"We expand on the analysis of the attention mechanism presented in the paper by including visualisations for additional queries from the CNN validation dataset below. We consider examples from the Attentive Reader as well as the Impatient Reader in this appendix." (p.10)
Q85	"To give a better intuition for the behaviour of the Impatient Reader, we use a similar visualisation technique as before. However, this time around we highlight the attention at every time step as the model updates its focus while moving through a given query. Figures 10–13 shows how the attention of the Impatient Reader changes and becomes increasingly more accurate as the model" (p.11)
Q86	"Figure 7: Two more correctly answered validation set queries. The left example (entity ent315) requires correctly attributing the quote, which does not appear trivial with a number of other candidate entities in the vicinity. The right hand side shows our model is not afraid of Chuck Norris (ent164)." (p.12)
Q87	"Figure 8: Attention heat maps from the Attentive Reader for two wrongly answered validation set queries. In the left case the model returns ent85 (correct: ent67), in the right example it gives ent24 (correct: ent64). In both cases the query is unanswerable due to its ambiguous nature and the model selects a plausible answer." (p.12)
Q88	"Note how the attention is distributed fairly arbitraty at first, slowly focussing on the correct entity ent5 only once the question has sufficiently been parsed." (p.12)
D1	A man named as a suspect in the fatal shooting of a Philadelphia, Pennsylvania, police officer last week was captured at a shelter in Miami, Florida — US domestic crime reporting anchored to US cities and police institutions

D5	State-run news agency: China orders an investigation by quality control agencies. Children who swallow the beads can become comatose or have seizures. — Consumer safety story; no humanitarian crisis taxonomy

Strengths

- Multiple evaluation strategies are defined (majority baselines [Q22], word-distance alignment [Q32, Q33], frame-semantic heuristics [Q29], attention models [Q49, Q50, Q51]), and length-sensitivity analysis is reported [Q78] — providing methodological scaffolding that could in principle be adapted, though not the output form itself.

ID	Evidence
Q22	"We define two simple baselines, the majority baseline (maximum frequency) picks the entity most frequently observed in the context document, whereas the exclusive majority (exclusive frequency) chooses the entity most frequently observed in the context but not observed in the query." (p.3)
Q29	"Strategies are ordered by precedence and answers determined accordingly. This heuristic algorithm was iteratively tuned on the validation data set." (p.4)
Q32	"We consider another baseline that relies on word distance measurements. Here, we align the placeholder of the Cloze form question with each possible entity in the context document and calculate a distance measure between the question and the context around the aligned entity." (p.4)
Q33	"We tune the maximum penalty per word (m = 8) on the validation data." (p.4)
Q49	"The representation r of the document d is formed by a weighted sum of these output vectors." (p.5)
Q50	"These weights are interpreted as the degree to which the network attends to a particular token in the document when answering the query:" (p.5)
Q51	"The variable s(t) is the normalised attention at token t." (p.5)
Q78	"To understand how the model performance depends on the size of the context, we plot performance versus document lengths in Figures 4 and 5. The first figure (Fig. 4) plots a sliding window of performance across document length, showing that performance of the attentive models degrades slightly as documents increase in length. The second figure (Fig. 5) shows the cumulative performance with documents up to length N, showing that while the length does impact the models' performance, that effect becomes negligible after reaching a length of ~500 tokens." (p.10)

Checklist

Checklist item definitions:

ID	Assessment Question
OF-1	Does output modality match regional deployment needs?
OF-2	Assess text-to-speech availability for speech output
OF-3	Consider literacy rates and accessibility requirements
OF-4	Document form mismatches harming external validity

Checklist responses:

OF-1: No — expected output modality is single-token classification [Q34, Q52]; deployment requires multi-sentence temporally ordered narrative timeline (YAML output_task_requirements).

OF-2: Not applicable — deployment is text-based; analysts work in office/coordination-hub environments with no speech-output requirement (YAML device_and_connectivity_context_for_analysts).

OF-3: Analyst literacy is high (professional NGO/IO data analysts) per YAML analyst_digital_literacy; not a constraint on output form. However, affected populations (tweet authors) include lower-literacy rural and indigenous communities (~5.4% national illiteracy, higher in rural areas [WEB-11]) — but this is an input concern, not output.

OF-4: The categorical mismatch between single-token classification (or unordered bullet summarization) and temporally ordered narrative timelines violates external validity for the deployment use case.

ID	Evidence

Q34	“"We propose three neural models for estimating the probability of word type a from document d answering query q: p(a\” (p.4)
Q52	“Having described a number of models in the previous section, we next evaluate these models on our reading comprehension corpora.” (p.6)
WEB-11	National illiteracy 5.4% (analyst-side literacy not a constraint) (https://www.statista.com/statistics/1322285/digital-illiteracy-ecuador/)
WEB-19	CrisisFACTS represents the closest existing timeline benchmark — but English-only with no Latin American coverage (https://aclanthology.org/2025.acl-long.783.pdf)

Information Gaps

- Whether any extension of the benchmark task to multi-sentence generation has been published — the original paper does not include such an extension

Requires Expert Verification

- Specific timeline-output format expectations of Red Cross Ecuador and OCHA situation-report consumers

---

Remediation

Gap 1: Single-token classification or unordered bullet summarization mismatched with temporally ordered narrative timelines

Recommendation: Adopt a multi-sentence temporally ordered narrative output form with phase-transition labels and bilingual (Spanish/English) presentation to support UN OCHA inter-agency reporting workflows; evaluate using timeline-coherence metrics adapted from CrisisFACTS [WEB-19] augmented with operational-relevance judgments from the inter-operational team.

ID	Evidence
WEB-19	CrisisFACTS represents the closest existing timeline benchmark — but English-only with no Latin American coverage (https://aclanthology.org/2025.acl-long.783.pdf)