Input Ontology (IO)

Score: 3/5 — Moderate gaps | Confidence: high

Benchmark: *CRISISBENCH* | ***Context:*** *Argentina — Buenos Aires Policy Maker Cohort (Disaster Reporting Reliability)*

Input Ontology definition: Whether the benchmark's test case categories cover the query types expected in deployment.

Each section below is followed by an evidence table. Three types of evidence are cited: **[QN]** — verbatim quotes extracted from the benchmark paper; **[WEB-N]** — web sources consulted during assessment; **[DN]** / **[DATASET-DN]** — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

The top-level taxonomy (informativeness binary; humanitarian multi-class with categories like affected_individual, infrastructure_and_utilities_damage, requests_or_needs, response_efforts, caution_and_advice) is broadly compatible with the international ISCRAM/OCHA-derived frameworks Argentine emergency professionals use as reference points, and elicitation Q1 confirms users find these categories 'broadly sufficient'. However, the benchmark explicitly defers event-type classification [Q22, Q23], does not include sub-labels for Argentine hazard types (pampero, sudestada, ENSO drought, AMBA urban flash flood) or informal-settlement (villa) impact framing, and the absence of a Latin American taxonomy adaptation is confirmed at the field level [WEB-15]. Some humanitarian classes have very low support (e.g., 17 'terrorism related' tweets [Q52]), and categorical granularity varies significantly across constituent datasets [Q76]. Dataset analysis confirms the schema covers the core triage categories Argentine policy makers need [DATASET-D19, D20, D35, D42] but exposes the 'other_relevant_information' class as a heterogeneous catch-all that limits its utility for borderline credibility cases.

ID	Evidence
Q1	“We consolidate eight human-annotated datasets and provide 166.1k and 141.5k tweets for informativeness and humanitarian classification tasks, respectively.” (p.1)
Q22	“We only focused on two tasks for this study and we aim to address event types task in a future study.” (p.2)
Q23	“We consolidate eight datasets that were labeled for different disaster response classification tasks and whose labels can be mapped consistently for two tasks: informativeness and humanitarian information type classification.” (p.3)
Q52	“Only 17 tweets with the label "terrrorism related" are present in CrisisNLP.” (p.5)
Q76	“The motivation of these experiments is to investigate whether model trained with consolidated dataset generalizes well across different sets. For the individual dataset classification experiments, we selected CrisisLex and CrisisNLP as they are relatively larger in size and have a reasonable number of class labels, i.e., six and eleven class labels, respectively.” (p.7)
D19	"Floods kill 3, 75,000 forced from Calgary homes http://t.co/c6OPFjFLkA — AP #news"
D20	As many as 100,000 people could be forced from their homes by heavy flooding in western Canada. Water levels peak today at noon. — Factual caution about Canadian floods
D35	I'M ASKING HELP FOR MY FAMILY WHICH IS VICTIM AND WHICH TAKE REFUGE IN PROVINCE. — Direct help request from disaster victim
D42	#Tacloban has 49 evacuation centers now sheltering over 30,000 people + @care is helping respond to #typhoonhagupit http://t.co/ePE9MsdpTh — Clear response effort tweet — CARE NGO involvement
WEB-15	TREC-IS introduced priority/criticality labels (Low/Medium/High/Critical) as the closest existing benchmark precedent for credibility scoring; covers no Argentine events (https://trecis.github.io/)

Strengths

- Top-level humanitarian categories (affected_individual, infrastructure_and_utilities_damage, requests_or_needs, response_efforts, caution_and_advice, donation_and_volunteering) align with ISCRAM-derived frameworks recognized by Argentine emergency management professionals, supporting first-pass triage utility
- Binary informativeness layer is structurally compatible with the deployment's first-pass signal-triage step and operationalizes a foundation for downstream credibility scoring

Checklist

Checklist item definitions:

ID	Assessment Question
IO-1	Identify test case categories required for regional deployment
IO-2	Check if taxonomy omits regionally relevant categories
IO-3	Check if taxonomy includes irrelevant categories
IO-4	Document category gaps harming content validity

Checklist responses:

IO-1: Required Argentine categories include pampero windstorm, sudestada/Río de la Plata flooding, ENSO-driven drought, urban AMBA flash flood, and villa (informal-settlement) impact framing. Top-level humanitarian categories (infrastructure damage, affected individuals, response efforts) are required and present.

IO-2: Argentine hazard sub-labels and informal-settlement impact framing are omitted. No Latin American adaptation of ISCRAM/OCHA taxonomies addressing pampero, sudestada, or ENSO-drought sub-types exists in the literature [WEB-15]. Event-type classification is deferred entirely [Q22, Q23].

IO-3: Some categories have negligible support (e.g., 'terrorism related' with only 17 tweets [Q52]; 'disease related' only in CrisisNLP [Q53]). The 'other_relevant_information' class is semantically incoherent per dataset analysis (Concern 7), reducing its utility.

IO-4: Documented gaps: (a) no event-type classification implemented; (b) no Argentine hazard sub-types; (c) no informal-settlement impact framing; (d) heterogeneous 'other_relevant_information' catch-all. These constitute construct underrepresentation for the deployment but the user has flagged them as desirable rather than blocking.

ID	Evidence
Q3	“Typical classification tasks in the community include (i) informativeness (i.e., informative vs. not-informative messages), (ii) humanitarian information types (e.g., affected individual reports, infrastructure damage reports), and (iii) event types (e.g., flood, earthquake, fire).” (p.1)
Q22	“We only focused on two tasks for this study and we aim to address event types task in a future study.” (p.2)
Q23	“We consolidate eight datasets that were labeled for different disaster response classification tasks and whose labels can be mapped consistently for two tasks: informativeness and humanitarian information type classification.” (p.3)
Q52	“Only 17 tweets with the label "terrrorism related" are present in CrisisNLP.” (p.5)
Q53	“Similarly, the class "disease related" only appears in CrisisNLP.” (p.5)
Q76	“The motivation of these experiments is to investigate whether model trained with consolidated dataset generalizes well across different sets. For the individual dataset classification experiments, we selected CrisisLex and CrisisNLP as they are relatively larger in size and have a reasonable number of class labels, i.e., six and eleven class labels, respectively.” (p.7)
Q89	“Note that humanitarian task is a multi-class classification problem, which makes it a much more difficult task than the binary informativeness classification.” (p.8)
WEB-15	TREC-IS introduced priority/criticality labels (Low/Medium/High/Critical) as the closest existing benchmark precedent for credibility scoring; covers no Argentine events (https://trecis.github.io/)

Information Gaps

- Whether downstream sub-label extension to Argentine hazard types is feasible given absent training data

Requires Expert Verification

- Whether Argentine emergency management professionals would prioritize villa-impact sub-labels over generic affected_individual labels in operational triage

---

Remediation

Gap 1: No sub-labels for Argentine hazard types (pampero, sudestada, ENSO drought, AMBA urban flash flood) or villa-impact framing

Recommendation: Extend the humanitarian taxonomy with Argentine hazard sub-labels and informal-settlement (villa) impact framing, validated with AMBA emergency management practitioners. Treat as a regional ontology overlay on top of CrisisBench's top-level categories rather than replacing them.