

Validity Analysis: CRISISBENCH

Target context: Argentina — Buenos Aires Policy Maker Cohort (Disaster Reporting Reliability)

Overall risk: **HIGH**

Deployment Context

Use case: Use case: Reliability of disaster reporting in news at national/local level, finding alignment with effective and verified information (e.g. noisy vs real signals). The evaluation data could report or not disaster damage.

Target population: Policy makers and government agencies in Argentina studying trustiness and reliability of disaster communication across different channels are (e.g news, social media post, press). They would be based in Buenos Aires, speakers of Rioplatense Spanish, with an intermediate to advanced level of English, at least a master's degree, and be part of an upper-middle-class household.

Dimension Scores

Dimension	Score	Rating	Confidence
Input Ontology	3	Moderate gaps	high
Input Content 	1	Serious concern	high
Input Form 	3	Moderate gaps	high
Output Ontology 	2	Significant gaps	high
Output Content 	2	Significant gaps	high
Output Form 	3	Moderate gaps	high
Average	2.3		

 = highest concern  = strongest dimension

Dimension Key

Abbr.	Dimension	Definition
IO	Input Ontology	Whether the benchmark's test case categories cover the query types expected in deployment.
IC	Input Content	Whether individual datapoint content is culturally and contextually appropriate for the target region.
IF	Input Form	Whether the input signal encoding (text, audio, image parameters) matches deployment conditions.
OO	Output Ontology	Whether the benchmark's output categories and scoring criteria reflect regionally valid decision boundaries.
OC	Output Content	Whether ground-truth labels align with the judgments of regional stakeholders.
OF	Output Form	Whether the expected output modality matches regional deployment needs and accessibility.

Overall Summary

CrisisBench provides a methodologically rigorous, reproducible foundation for crisis-tweet classification with a clean two-task ontology (binary informativeness + multi-class humanitarian type) that aligns at the top level with international ISCRAM/OCHA-derived frameworks Argentine emergency professionals reference. However, three structural misalignments substantially limit its validity for the Buenos Aires policy-maker deployment: (1) Input Content (HIGH priority) is severely misaligned — 94.46% English [\[Q46\]](#), zero Argentine or Rioplatense Spanish content, and a 2010–2017 temporal window that

predates the dissolution of Télam (July 2024) and creation of AFE (2025); the only Spanish content in the benchmark is Chilean and Venezuelan, not Rioplatense [DATASET-D2, D3, D6, D7]. (2) Output Ontology (HIGH priority) lacks the source credibility / signal-reliability dimension that is the deployment's core requirement; the user acknowledges this as a downstream construction step but the benchmark provides no learnable credibility signal. (3) Output Content shows annotator-population mismatch and empirically observable annotation noise, with evidence of systematic Spanish-language disaster-content under-recognition [DATASET-D9, D8]. Input Form, Input Ontology, and Output Form present moderate but tractable gaps. The benchmark is best treated as a usable signal-triage foundation for the user's planned downstream credibility-aggregation pipeline, not as a complete evaluation framework for the Argentine deployment.

ID	Evidence
Q46	"Among different languages of informativeness tweets, English tweets appear to be highest in the distribution compared to any other language, which is 94.46% of 156,899." (p.5)
D2	RT @FlorrGo: #FuerzaChile ,un 2 de abril ... Que vengan los ingleses a ayudarlos ahora ... A la larga o a la corta ,todo vuelve — Chilean Spanish earthquake solidarity tweet — not Argentine/Rioplatense
D9	Ojo con el exceso de ayuda por los efectos del #TerremotoenChile porque puede producir mucho más obstáculos de lo que se cree. — Chilean Spanish warning about over-donation — labelled not_informative despite substantive crisis-relevant content
D3	Que lindo saber que en mi ciudad bella #Iquique hay gente con un corazón enorme dispuestos a ayudar... Emocionada — Chilean Spanish (Iquique city) solidarity tweet — Chilean, not Argentine
D6	IMPRESIONANTE terremoto Chile camara de seguridad farmacia - NEW VIDEO E...: http://t.co/bkUEIKfUj via @YouTube — Chilean Spanish earthquake tweet — Chilean Spanish, not Rioplatense
D7	RT @marieli_d: Nuevamente se incendia #Amuay http://t.co/AiWlhgy3 — Venezuelan Spanish tweet about refinery explosion
D8	RT @meteomostoles: El tifón #BOPHA,de cat.4 y 947hPa en su centro,va camino de Filipinas.Lleva vientos sostenidos entre 2010 y 249km/h h ... — Spanish tweet about Philippines typhoon — meteorological content labelled not_humanitarian

Practical Guidance

What This Benchmark Measures

CrisisBench measures English-language Twitter content's informativeness (binary) and humanitarian type (multi-class) for a fixed set of 2010–2017 disaster events concentrated in North America and globally reported crises. For the Argentine deployment, it provides a usable first-pass signal-triage capability — separating clear noise from clear crisis content at high accuracy on English content (0.866/0.829 F1) — that can serve as Step 1 of the user's planned downstream credibility-scoring pipeline. The strongest dimensions for the deployment are Input Form (text-only modality match) and the foundational utility of the binary informativeness task; the top-level humanitarian categories are recognizable to Argentine policy professionals via international taxonomy alignment.

Construct Depth

The benchmark probes informativeness and humanitarian type at moderate depth for English North American Twitter content but provides no depth on (a) source credibility, the deployment's core construct, or (b) Spanish-language or Rioplatense register. Cross-dataset F1 degradation of 14.3% [Q87] and ~8% multilingual humanitarian F1 drop [Q99] indicate that even the depth it does provide does not generalize cleanly across event contexts or languages. The dataset analysis reveals annotation noise that further limits the trustworthiness of the F1 numbers as indicators of real-world classification capability, particularly for borderline content where credibility decisions are most consequential.

What Else You Need

The deployment requires three supplementation streams: (1) Argentine/Rioplatense Spanish disaster content for fine-tuning or evaluation — addressing the Input Content gap; possibilities include collecting and annotating Argentine Twitter/X content using the CrisisBench taxonomy, leveraging the cross-lingual approach in arXiv:2209.02139 [WEB-17], or extending TREC-IS-style protocols [WEB-15] to Argentine events. (2) A credibility-scoring extension — addressing the Output Ontology gap; TREC-IS priority labels (Low/Medium/High/Critical) [WEB-15] are the closest precedent and could anchor downstream construction combining classifier outputs with Argentine source metadata (institutional affiliation, post-Télam press source registry, Buenos Aires Province Resolución 4/2025-compliant metadata). (3) Argentine-annotator label validation on a stratified

sample — addressing the Output Content gap; required to quantify expected label disagreement, particularly for Spanish-language content where dataset analysis suggests systematic under-recognition risk.

ID	Evidence
Q87	“The cross dataset (i.e., train on CrisisLex and evaluate on CrisisNLP) results shows that there is a drop in performance. For example, there is 14.3% difference in F1 on CrisisNLP data using the CrisisLex model for the informativeness task.” (p.8)
Q99	“We observe that performance dropped significantly for the humanitarian task compared to English-only dataset. For example, ~8% drop using BERT model.” (p.9)
WEB-17	Sánchez et al. 2022 cross-lingual crisis classification paper does not address Rioplatense register (https://arxiv.org/abs/2209.02139)
WEB-15	TREC-IS introduced priority/criticality labels (Low/Medium/High/Critical) as the closest existing benchmark precedent for credibility scoring; covers no Argentine events (https://trecis.github.io/)