

Validity Analysis: WXIMPACTBENCH

Target context: Argentine Government Disaster Reporting Reliability — Policy Staff

Overall risk: **HIGH**

Deployment Context

Use case: Use case: Reliability of disaster reporting in news at national/local level, finding alignment with effective and verified information (e.g. noisy vs real signals). The evaluation data could report or not disaster damage.

Target population: Policy makers and government agencies in Argentina studying trustiness and reliability of disaster communication across different channels are (e.g news, social media post, press). They would be based in Buenos Aires, speakers of Rioplatense Spanish, with an intermediate to advanced level of English, at least a master's degree, and be part of an upper-middle-class household.

Dimension Scores

Dimension	Score	Rating	Confidence
Input Ontology ✓	3	Moderate gaps	high
Input Content △	1	Serious concern	high
Input Form △	2	Significant gaps	high
Output Ontology △	2	Significant gaps	high
Output Content △	2	Significant gaps	high
Output Form ✓	3	Moderate gaps	medium
Average	2.2		

△ = highest concern ✓ = strongest dimension

Dimension Key

Abbr.	Dimension	Definition
IO	Input Ontology	Whether the benchmark's test case categories cover the query types expected in deployment.
IC	Input Content	Whether individual datapoint content is culturally and contextually appropriate for the target region.
IF	Input Form	Whether the input signal encoding (text, audio, image parameters) matches deployment conditions.
OO	Output Ontology	Whether the benchmark's output categories and scoring criteria reflect regionally valid decision boundaries.
OC	Output Content	Whether ground-truth labels align with the judgments of regional stakeholders.
OF	Output Form	Whether the expected output modality matches regional deployment needs and accessibility.

Overall Summary

WXImpactBench is an English-language, formally-styled, North America-centric benchmark whose six top-level impact categories transfer reasonably to Argentine disaster types but whose content, register, language, annotator pool, and output schema all diverge substantially from the deployment's needs. The most severe validity violations are in input_content (zero Spanish-language, zero Argentine, zero social media content; HIGH priority) and output_ontology (no native credibility scoring — though deliberately handled downstream by user design — and no sub-labels for Argentine-specific granularity; HIGH

priority). Input_form is mismatched on register and length, and the paper itself acknowledges formal historical text is easier for LLMs than informal modern text, meaning benchmark scores likely overstate deployment-relevant capability. Output_content carries zero Argentine annotator representation and exhibits empirical labeling inconsistencies on positively-framed chunks. The benchmark may serve as one input to model triage but cannot stand alone as evidence of deployment-readiness for Argentine government disaster reporting reliability assessment.

Practical Guidance

What This Benchmark Measures

WXImpactBench measures LLM ability to extract six pre-defined societal impact categories from formal long-form English-language historical journalism, and to retrieve relevant articles given a generated question. For the Argentine deployment, this provides indirect evidence of LLM reasoning over formal press content, particularly across two linguistic registers, and exercises some noise-rejection capability through the chunked dataset's heterogeneous off-topic content [DATASET-D2, D7, D23]. Strongest dimensions for the target context are output_form (functional compatibility for upstream feature use) and input_ontology (six top-level categories applicable per elicitation [A1]).

Construct Depth

Construct depth is shallow for the deployment's primary need. The benchmark probes formal English long-form impact classification but does not probe (a) Spanish-language disaster content, (b) social media or short-form informal text, (c) credibility/reliability scoring, (d) Argentine institutional actor recognition, (e) slow-onset disaster framing (ENSO drought), or (f) villa miseria social vulnerability framing. The paper's own observation that formal historical style is easier for LLMs [Q83, Q85] caps the upper bound of construct generalization. Empirical labeling inconsistencies [DATASET-D17, D22, D27] further limit confidence in the convergent validity of the impact-classification construct itself.

What Else You Need

The deployment requires substantial supplementation: (1) a Rioplatense Spanish disaster impact gold-standard test set built with Argentine annotators (no equivalent exists per WEB-14, gap_id 7); (2) a social media short-form evaluation, with HumAID/CrisisBench [WEB-14, WEB-15] as English-language starting points pending cross-lingual transfer validation; (3) a credibility/reliability scoring layer with Argentine source metadata (outlet reach, account verification, AFE/SINAGIR institutional recognition); (4) sub-label refinement for villa miseria infrastructure, agricultural commodity-specific loss, and Argentine partisan-framing political impact; (5) re-annotation of a sample by Argentine emergency management experts to verify whether documented labeling inconsistencies persist under Argentine interpretation.

ID	Evidence
Q83	"The reason might be the structured and formal narrative style used to report disruptive weather events in historical periods, which more explicitly highlights cause-and-effect relationships." (p.8)
Q85	"Though the modern text might dominate within the pre-trained corpus, the language patterns used in historical narrative styles are easier for language models to identify, and thus perform better on classifying disruptive weather impacts." (p.8)
D2	But the real beneficiaries of the work-at-home trend are the employees themselves... She takes a turn at the table in a lay way (5) 29, A good one is of assistance (9) Down 1, Full of ill-will and disease (9)... — Non-weather content: work-from-home editorial + crossword puzzle clues
D7	National Lampoon's Van Wilder 2: The Rise of Taj Starring: Kal Penn, Lauren Cohan and Daniel Percival Playing at: AMC, Brossard, Cavendish, Colossus, LaSalle cinemas. — Film review; no weather content at all
D17	PETER MARTIN, GAZETTE Workers from the Connecticut Light and Power Co were working to reconnect power lines on Oxford Ave in Notre Dame de Grace... Temperatures plummet; today's high minus-15C... more than one million people struggle on without power or heat — 1998 Quebec ice storm coverage describing power outages and extreme cold; labeled all-zero — potential annotation gap
D22	Almost 4 million people across China had been cut off by flood waters, 810,000 homes have collapsed and 2.8 million homes have been damaged in eight provinces... Japan struck at young and old yesterday, killing a schoolgirl and an 85-year-old woman — 1996 China Yangtze floods + Japan E. coli outbreak in same chunk; infrastructural=1 but human_health=0 despite explicit casualties
D23	-2 City res y 5000 5 5 5 Aiexavn 8000 50 50 50 Ckearcdnf 835 825 25..." / "Escenvt 115780 298 J80 289 19 Caiunoil 8000 J9 J9 29... — OCR-corrupted stock market tables; repeated for same article id across multiple rows

D27	A poor lady passenger was dashed to leeward and had her skull fractured. She was landed this morning in a dying state with a coffined child that succumbed last night to its sufferings. — 1881 Atlantic storm with explicit casualties (skull fracture, child death) labeled all-zero
WEB-14	No complementary Argentine Spanish-language benchmark exists for cross-validation (https://crisisnlp.qcri.org/humaid_dataset)
WEB-15	https://ojs.aaai.org/index.php/ICWSM/article/download/18116/17919/21611