Comparative Assessment: HOPE_DAC vs. SWITCHBOARD_1992

Deployment context: Global NLP/AI Research Community — Computational Mental Health (HOPE Benchmark Assessment)

Deployment Context

Use case: An NLP/AI researcher uses the HOPE benchmark dataset to evaluate and improve dialogue-understanding for AI-assistive chatbots. The benchmark provides a evaluation protocol across 12 dialogue act categories in counseling conversations.

Target population: NLP/AI researchers in English speaking countries or willing to work with English data working on computational approaches to mental health, dialogue systems, and psycholinguistics. This includes PhD students, postdoctoral researchers, and industry practitioners building LLM-based systems for clinical NLP and digital mental health applications serving underrepresented regional communities.

Overall Risk Ratings

Benchmark	Risk	Avg. Score
HOPE_DAC	HIGH	2.8 / 5
SWITCHBOARD_1992	HIGH	1.7 / 5

Dimension Scores Comparison

Dimension	HOPE_DAC	Rating	SWITCHBOARD_1992	Rating	Δ
Input Ontology (IO)	3/5	Moderate gaps	1/5	Serious concern	+2
Input Content (IC)	2/5	Significant gaps	1/5	Serious concern	+1
Input Form (IF)	4/5	Minor gaps	3/5	Moderate gaps	+1
Output Ontology (OO)	3/5	Moderate gaps	1/5	Serious concern	+2
Output Content (OC)	2/5	Significant gaps	2/5	Significant gaps	0
Output Form (OF)	3/5	Moderate gaps	2/5	Significant gaps	+1
Average	**2.8**		**1.7**		**+1.2**

Δ = HOPE_DAC score − SWITCHBOARD_1992 score. Positive = regional benchmark scores higher.

Overall Summaries

HOPE_DAC (risk: HIGH)

HOPE provides a methodologically careful, transparent benchmark for dialogue-act classification in dyadic individual counseling — its task framing [Q7–Q9, Q22], 12-label scheme developed with expert consultation [Q27], substantial inter-annotator agreement [Q49], and statistically significant SPARTA results [Q88] make it a reasonable starting point for the elicited use case of evaluating chatbots in individual counseling-style conversations. The strongest validity dimensions are Input Form (text-only English matches deployment exactly) and Input Ontology (well-scoped task framing for individual dyadic counseling). The most significant validity concerns are concentrated in the two HIGH-priority dimensions: Input Content (sessions are explicitly US-clinical [Q98] while downstream chatbots target South Asian, African American, and Latino communities whose communicative norms diverge in well-documented ways [WEB-5, WEB-9, WEB-11]) and Output Content (annotator cultural backgrounds undocumented [Q50], no label-stratified IAA, RoBERTa backbone inherits documented AAVE bias [WEB-8, WEB-9]). The Output Ontology and Output Form dimensions show acknowledged but unresolved limitations — single-label assignment [Q30] mismatches multi-intent reality (per elicitation A3), and no subgroup-disaggregated metrics are reported, which conflicts with emerging FDA equity expectations [WEB-20]. Many of these gaps are field-wide rather than HOPE-specific [WEB-17], but they materially limit the benchmark's external validity for the target downstream populations.

SWITCHBOARD_1992 (risk: HIGH)

Switchboard (1992) is fundamentally misaligned with the deployment context of counseling chatbot dialogue-act evaluation for underrepresented U.S. communities. Three of the four HIGH-priority dimensions (IO, IC, OO) score 1, reflecting that the corpus was designed for speaker authentication and large-vocabulary speech recognition [Q1, Q20] rather than therapeutic discourse, that its content is casual stranger-stranger telephone conversation with no documented racial/ethnic representation [Q3, WEB-1, WEB-5], and that any attached SWBD-DAMSL taxonomy was developed for 'task-free' general speech [WEB-2] without counseling-specific categories or multi-label structure. Output content (annotator quality and cultural validity) and output form (single-label scoring) score 2, with documented annotator-expertise gaps [WEB-3] and metric-form mismatches acknowledged by the user [elicitation A3]. Input form scores 3, partially mitigated by text-based deployment. Counseling-specific published alternatives — HOPE [WEB-7, WEB-12] and AnnoMI [WEB-9, WEB-10] — exist and represent more valid starting points, though they share the cultural coverage gap for target communities.

Per-Dimension Justification Comparison

Each benchmark's full assessment justification and cited evidence, shown side by side.

Input Ontology (IO)

Whether the benchmark's test case categories cover the query types expected in deployment.

HOPE_DAC — 3/5 (Moderate gaps) [medium confidence]	SWITCHBOARD_1992 — 1/5 (Serious concern) [high confidence]
HOPE's task taxonomy targets dialogue-act classification in dyadic counseling, explicitly distinguished from goal-oriented and chit-chat dialogue [Q7, Q8, Q9], with the well-defined objective of assigning one DAC label per utterance [Q22, Q58]. For the elicited use case — individual counseling-style chatbot conversations — this coverage is appropriate, and the elicitation explicitly affirms that 'individual counseling-style conversations' are well covered. However, the taxonomy is restricted to individual dyadic sessions [Q26], with no sub-task disaggregation for group therapy, crisis intervention, or culturally adapted modalities (e.g., spiritually grounded counseling, collectivist distress framing). The elicitation marks IO as MODERATE priority and treats cultural adaptation as deferred future work, so this is a known but accepted gap rather than a fundamental misalignment. The baseline architecture comparison and ablations [Q79, Q80, Q115] cover utterance, local context, global context, and speaker-awareness signals, which is structurally appropriate for the construct.	The Switchboard input ontology is explicitly oriented toward speaker authentication and large-vocabulary speech recognition [Q1, Q20], with depth and breadth of coverage calibrated for speaker verification [Q23]. There is no dialogue-act or therapeutic-move taxonomy in the original paper, and the SWBD-DAMSL scheme later attached to it was developed for 'task-free' casual telephone speech to improve speech recognition, not for therapeutic discourse [WEB-2, WEB-4]. For a counseling chatbot deployment requiring categories such as empathic reflection, psychoeducation, and motivational interviewing techniques, the input task taxonomy is fundamentally misaligned. Counseling-specific alternatives (HOPE, AnnoMI) exist but are structurally incompatible with SWBD-DAMSL [WEB-7, WEB-9]. Given the HIGH priority assigned to IO in elicitation and the absence of any therapeutic category coverage, the misalignment is severe.

HOPE_DAC evidence	SWITCHBOARD_1992 evidence
[Q7] 'Unlike general goal-oriented dialogues, a conversation between a patient and a therapist is considerably implicit, though the objective of the conversation is quite apparent.' (p.1)	[Q1] 'SWITCHBOARD is a large multispeaker corpus of conversational speech and text which should be of interest to researchers in speaker authentication and large vocabulary speech recognition.' (p.1)

[Q22] 'we annotated ~ 12.9*K* utterances with 12 dialogue-act labels carefully designed to cater to the requirements of a counseling session.' (p.3)	[Q20] 'The design of the SWITCHBOARD corpus emphasizes the importance of both depth and breadth of coverage, especially for speaker verification research.' (p.3)
[Q26] 'all of them are dyadic conversations only.' (p.3)	[Q23] 'The design assumes that 25 target speakers should suffice to get statistically reliable estimates of the performance of a speaker verification algorithm under development...' (p.3)
[Q79] 'we categorize the baselines into three groups – utterance-driven, utterance+global context driven, and utterance + global context + speaker-aware driven.' (p.7)	[WEB-2] SWBD-DAMSL designed for task-free casual telephone speech to improve speech recognition, not therapeutic discourse
	[WEB-4] Stolcke et al. 2000 confirms SWBD-DAMSL purpose was computational DA modeling for conversational speech recognition
	[WEB-7] HOPE counseling-specific 12-label DA hierarchy developed with therapist input as alternative
	[WEB-9] AnnoMI MI behavioral coding scheme aligned with MISC as alternative counseling DA taxonomy

Input Content (IC)

Whether individual datapoint content is culturally and contextually appropriate for the target region.

HOPE_DAC — 2/5 (Significant gaps) [high confidence]	SWITCHBOARD_1992 — 1/5 (Serious concern) [high confidence]
Input content is sourced from 212 publicly available YouTube counseling videos [Q3, Q13, Q23], and the paper explicitly states that 'the majority of the sessions in HOPE belong to the mental health professionals and patients based in United States' [Q98], with explicit acknowledgment that 'the effectiveness of SPARTA on data from other geographical or demographical regions may vary' [Q99]. The elicitation marks IC as HIGH priority precisely because the downstream chatbot deployment targets South Asian, African American, and Latin American communities whose communicative norms diverge from US clinical baseline (indirect disclosure, somatic idioms, AAVE pragmatics, familismo, personalismo). No participant demographic metadata is documented — gap_id 6 confirmed [WEB-17]. Latino adults are 28% less likely than US adults overall to receive mental health treatment [WEB-11], and Asian Americans utilize services at roughly half the general-population rate [WEB-5] — meaning the populations the chatbot targets are precisely those least represented in HOPE's source content. AAVE-related transformer bias is documented in the literature backbone (RoBERTa) [WEB-8, WEB-9]. Utterance length distributions (~103/124 words) [Q56, Q57] further diverge from short-turn chatbot deployment patterns. Construct-irrelevant variance risk is high.	The corpus consists of casual prompted telephone conversations between paid stranger volunteers from American English dialect regions in 1990–1992 [Q2, Q3, Q4]. Demographic metadata covers age, sex, education, and dialect region [Q25] but does not include race/ethnicity, and no demographic audit of speaker composition for race/ethnicity/SES has been published [WEB-1, WEB-5]. The content is contextually and normatively remote from clinical counseling sessions, no community validation exists, and target downstream communities (Black, Latino, South Asian, Indigenous) are not documented as proportionally represented. With IC marked HIGH priority and substantial construct-irrelevant variance from both the casual stranger-conversation context and the demographic gap, the score is 1.

HOPE_DAC evidence	SWITCHBOARD_1992 evidence
[Q98] 'the majority of the sessions in HOPE belong to the mental health professionals and patients based in United States.' (p.8)	[Q2] 'About 2500 conversations by 500 speakers from around the U.S. were collected automatically over T1 lines at Texas Instruments.' (p.1)

[Q99] 'Hence, the effectiveness of SPARTA on data from other geographical or demographical regions may vary.' (p.8)	[Q3] '...2500 conversations of three to ten minutes' duration, carried on by about 500 paid volunteers both hosts from every major dialect of American English.' (p.1)
[Q23] 'One of major hurdles we faced in the process of data collection was the unavailability of public counseling sessions...we carefully explored the web and collected publicly-available pre-recorded counselling videos on YouTube.' (p.3)	[Q4] '...over 250 hours of speech and nearly 3 million words of text.' (p.1)
[Q56] 'the dialogue sessions in HOPE are usually lengthy (~ 59 utterances per session).' (p.5)	[Q25] 'Demographic information about the speakers is entered in an Oracle data base when they register to participate. This includes their age, sex, level of education, and the geographically-defined dialect area where they grew up.' (p.4)
[Q57] 'the average length of utterance of a patient being 103 words, whereas for therapist, it is ~ 124 words.' (p.5)	[WEB-1] LDC catalog confirms 543 speakers (302M/241F), dialect region logged but no racial/ethnic breakdown
[WEB-17] Translational Psychiatry (2023) systematic review: only 40 of 102 NLP-MHI studies report patient demographic information	[WEB-5] OLAC record confirms no demographic audit of race/ethnicity/SES exists
[WEB-11] HHS Office of Minority Health: Latino adults 28% less likely than US adults overall to receive mental health treatment in 2024	
[WEB-5] PMC review: only 9% of Asian Americans utilized mental health services vs. 18% of general US population	
[WEB-8] ResearchGate study quantifying AAVE bias in transformer-based language models (including RoBERTa)	
[WEB-9] npj Digital Medicine (2025): LLMs provide differential treatment recommendations when AAVE dialect cues indicate patient race	

Input Form (IF)

Whether the input signal encoding (text, audio, image parameters) matches deployment conditions.

HOPE_DAC — 4/5 (Minor gaps) [high confidence]	SWITCHBOARD_1992 — 3/5 (Moderate gaps) [medium confidence]
Input form is text-only English derived from speech via OTTER ASR with manual correction [Q25], with names systematically masked [Q97]. The benchmark language matches the deployment modality exactly (text-only English chatbot evaluation), with no script or primary-modality mismatch — the regional context marks IF as LOWER priority. Standard ML compute infrastructure (Tesla V100, 32GB GPU, PyTorch, HuggingFace) [Q108, Q64] is accessible for the global researcher cohort. The only meaningful concern is that the speech-to-text transfer 'causes some information loss' [Q94], removing prosodic and paralinguistic cues that may carry communicative-intent signal in counseling — but this is acknowledged. Utterance length distribution (~103/124 words) [Q56, Q57] is a mild form-distribution mismatch with short-turn chatbot deployment, though this is more of a content-distribution issue than a signal-encoding issue.	Switchboard inputs are 1992 telephone-channel digital audio with time-aligned word-level transcriptions produced via supervised phone-based recognition [Q6, Q7, Q8, Q17]. For a text-based chatbot deployment, the audio mismatch is partially mitigated [elicitation infrastructure note], but transcription conventions (disfluency marking, turn segmentation tailored to ASR benchmarking) may diverge from contemporary chatbot text distributions. No empirical study comparing Switchboard transcription artifacts to modern LLM preprocessing has been found [WEB-13, WEB-17]. Given MODERATE priority and partial mitigation, a midrange score is warranted.

HOPE_DAC evidence	SWITCHBOARD_1992 evidence

[Q25] 'we extract the transcriptions of each video using OTTER...an automatic speech recognition tool. Subsequently, we correct transcription errors to remove any noise.' (p.3)	[Q6] 'A time-aligned word for word transcription accompanies each recording.' (p.1)
[Q94] 'The automatic transfer of utterance from the speech modality to the text causes some information loss, though we tried our best to recover them through manual intervention.' (p.8)	[Q7] 'all-digital 4-wire format' (p.1)
[Q97] 'The names of the patients and therapists involved in these sessions have been systematically masked.' (p.8)	[Q8] 'detailed transcription and time alignment of all conversations' (p.1)
[Q108] 'GPU: TeslaV100, Memory: 32GB, Linux 64 Bit: Ubuntu 18.04 LTS' (p.10)	[Q17] 'The SWITCHBOARD conversations are also time aligned at the word level. The time alignment is accomplished using supervised phone-based recognition...' (p.3)
[Q56] 'the dialogue sessions in HOPE are usually lengthy (~ 59 utterances per session).' (p.5)	[WEB-13] arXiv NDAP benchmark notes Switchboard has highly skewed transition distributions vs counseling corpora
	[WEB-17] Switchboard Dysfluency Annotation Stylebook documents transcription conventions

Output Ontology (OO)

Whether the benchmark's output categories and scoring criteria reflect regionally valid decision boundaries.

HOPE_DAC — 3/5 (Moderate gaps) [high confidence]	SWITCHBOARD_1992 — 1/5 (Serious concern) [high confidence]
The 12-label DA taxonomy [Q2, Q14, Q17] was developed in consultation with therapist and counseling experts [Q27], explicitly distinguished from general schemes (Switchboard) on domain specificity grounds [Q27], and organized hierarchically [Q29, Q32]. For the elicited use case, the elicitation (A2) affirms that 'the dialogue-act taxonomy is designed to capture intent and conversational flow in a way that is generally applicable across settings' and treats the taxonomy as useful for guiding chatbot dialogue trajectory. However, three structural concerns remain: (1) Single-label assignment is acknowledged by the authors to be a simplification — 'sometimes, an utterance can have multiple dialog-acts' [Q30] — and the elicitation (A3) confirms this is a real concern for downstream chatbot use; (2) The authors themselves note that speaker-initiative/responsive label pairs 'seem complementary in nature' [Q51] with confirmed systematic confusion in error analysis (OD:ID 43%, ID:ACK 28%, YNQ:IRQ 26%) [Q91, Q92]; (3) The taxonomy is grounded in US clinical practice [Q98] and may systematically misclassify communicative acts prevalent in target downstream communities (call-and-response → potentially miscoded as ACK or GC; personalismo opening → possibly GT/GC instead of therapeutically functional rapport-building; spiritual coping → possibly OD or GC). No multi-label counseling DA benchmark currently exists [WEB-15, WEB-18], and no culturally adapted DA taxonomy has been identified [WEB-23], so HOPE is not exceptional but the gap is real.	The original paper documents no output ontology beyond word-level transcription [Q6, Q8]; any dialogue-act taxonomy is external (SWBD-DAMSL) [WEB-2, WEB-3]. SWBD-DAMSL provides ~42 collapsed categories from ~220 unique combinations [WEB-3, WEB-4], designed for 'task-free' casual telephone conversation [WEB-2], with no therapeutic move categories. The deployment requires multi-label/ranked classification per utterance for multi-intent counseling utterances, but SWBD-DAMSL enforces single-label categorical structure [elicitation A3]. Counseling-specific alternatives (HOPE 12-label, AnnoMI MI behavioral codes) are structurally incompatible with SWBD-DAMSL [WEB-7, WEB-9]. With OO marked HIGH priority and both category coverage and structural mismatch problems, the score is 1.

HOPE_DAC evidence	SWITCHBOARD_1992 evidence

[Q27] 'we, in consultation with therapist and counselling experts, design a set of 12 dialog-act labels that are arranged in a hierarchy.' (p.3)	[Q6] 'A time-aligned word for word transcription accompanies each recording.' (p.1)
[Q30] 'Sometimes, an utterance can have multiple dialog-acts; however, they are rare in the annotated dataset. Hence, for simplicity, we consider only one (primary) label for each utterance.' (p.3)	[Q8] 'detailed transcription and time alignment of all conversations' (p.1)
[Q51] 'The broad definitions of speaker-initiative and speaker-responsive dialogue-act label pairs, IRQ & ID; ORQ & OD, and CRQ & CD, seem complementary in nature.' (p.4)	[WEB-2] Jurafsky et al. 2000 explicitly notes Switchboard domain is task-free with few external constraints on DA category definitions
[Q91] 'We observe three pairs with significant error-rates (≥ 25%) – YNQ:IRQ (26%), OD:ID (43%), ID:ACK (28%).' (p.7)	[WEB-3] SWBD-DAMSL Coder's Manual: ~50 basic tags, ~220 combinations, collapsed to 42 categories
[WEB-23] JMIR Mental Health (2025) confirms 'limited attention to linguistic or cultural adaptation' in LLM mental health studies	[WEB-4] Stolcke et al. 2000 confirms 42-tag scheme for computational DA modeling
[WEB-15] AnnoMI uses MI-specific behavior codes; no multi-label counseling DA benchmark identified	[WEB-7] HOPE 12-label counseling DA hierarchy with therapist input
[WEB-18] Scientific Reports (2025): MultiWD covers multi-label wellness dimensions but for social media, not counseling dialogue acts	[WEB-9] AnnoMI MI behavioral codes (reflection/question/input) aligned with MISC

Output Content (OC)

Whether ground-truth labels align with the judgments of regional stakeholders.

HOPE_DAC — 2/5 (Significant gaps) [high confidence]	**SWITCHBOARD_1992 — 2/5 (Significant gaps) [high confidence]**
Three expert linguists annotated the corpus [Q45], operating from a hierarchical scheme developed with mental health professional and linguist consultation [Q16, Q95], with calibration phase [Q46] and expert-therapist-moderated discussion [Q47]. Aggregate Cohen's Kappa is 0.7234 ('substantial') [Q48, Q49]. The authors transparently acknowledge that 'the annotator's bias cannot be ruled out completely' [Q96]. However, multiple critical gaps remain for the elicited HIGH-priority use case: (1) Annotator cultural backgrounds, native languages, and exposure to target downstream communities are entirely undocumented — the only reported demographics are age 25–35 and 2–10 years professional experience [Q50]; (2) No label-stratified IAA scores are reported, masking whether agreement varies on culturally loaded labels (GC, ACK, GT, ID); (3) The institutional affiliations (BITS Pilani, MAIT Delhi, IIIT-Delhi) [Q12] suggest annotators are likely South Asian by background — itself a confound when annotating US-clinical content for downstream targeting of US underrepresented communities; (4) Empirical evidence shows LLMs and transformer models (including RoBERTa, HOPE's backbone) exhibit racial and dialectal bias in clinical contexts [WEB-8, WEB-9], grounding the OC concern empirically. AnnoMI [WEB-15] demonstrates that richer annotator documentation is achievable in this field, making HOPE's gap a methodological choice rather than a constraint.	The original paper describes automated, computer-controlled collection [Q9, Q12] with supervised phone-based time alignment [Q17] but no human dialogue-act annotation phase, no annotator demographics, and no inter-annotator agreement statistics. The external SWBD-DAMSL annotators were computational linguists and speech researchers from Stanford, Colorado, and JHU [WEB-3] with no documented clinical or cross-cultural counseling expertise; annotators labeled from transcripts alone (~30 min/conversation) without listening to audio. Reported human accuracy is 84% on the 42-tag scheme [WEB-4, WEB-6] with no published kappa. For underrepresented community utterances, no published study validates standard English DA schemes against AAE, Chicano/Latino English, South Asian English, or Indigenous community norms [WEB-14, WEB-15, gap_id 2]. With OC marked MODERATE and the user partially trusting taxonomic generality, a low-mid score of 2 reflects substantial concerns short of complete invalidation.

HOPE_DAC evidence	SWITCHBOARD_1992 evidence
[Q45] 'We employed three annotators who are experts in linguistics.' (p.4)	[Q9] 'automated collection' (p.1)
[Q49] 'We obtain the inter-rater agreement score of 0.7234 – which falls under the "substantial" category.' (p.4)	[Q12] 'The conversations in SWITCHBOARD were collected under computer control, without human intervention.' (p.1)
[Q50] 'Annotators are in the age group of 25-35, with 2-10 years of professional experience.' (p.4)	[Q17] 'The time alignment is accomplished using supervised phone-based recognition...' (p.3)
[Q96] 'However, the annotator's bias cannot be ruled out completely.' (p.8)	[WEB-3] SWBD-DAMSL annotators were computational linguists/speech researchers labeling from transcripts alone, no clinical/cross-cultural expertise documented
[Q47] 'every annotation was discussed in the presence of the annotators and an expert therapist as moderator to ensure consistency.' (p.4)	[WEB-4] Stolcke et al. 2000 reports 84% human accuracy on 42-tag scheme
[Q12] 'BITS Pilani, Goa, India, Maharaja Agrasen Institute of Technology, New Delhi, India, IIIT-Delhi, India.' (p.1)	[WEB-6] Cornerstone MNSU repository confirms 84% human accuracy figure
[WEB-9] npj Digital Medicine (2025): LLMs more racially biased in mental health than other health domains; AAVE dialect cues trigger inferior treatment recommendations	[WEB-14] TACL Culturally Aware NLP survey confirms cultural validation of DA schemes is an open research gap
[WEB-8] ResearchGate: RoBERTa shows complex bias toward AAVE in masked language modeling — bias inherited by SPARTA	[WEB-15] PMC AAVE clinical NLP study recommends cultural expert annotation
[WEB-15] AnnoMI uses 10 annotators with post-annotation survey, demonstrating richer annotator documentation is achievable	

Output Form (OF)

Whether the expected output modality matches regional deployment needs and accessibility.

HOPE_DAC — 3/5 (Moderate gaps) [high confidence]	SWITCHBOARD_1992 — 2/5 (Significant gaps) [high confidence]
Performance is reported via macro-F1, weighted-F1, and accuracy [Q63], with label-wise F1 [Q76] and 3-fold cross-validation [Q87]. SPARTA-TAA achieves 60.29 macro-F1 / 64.53 weighted-F1 / 64.75% accuracy [Q73, Q85] with statistical significance vs. CASA [Q88]. Strengths include label-wise reporting that surfaces ACK weakness (F1=47.86%) [Q77] and proper statistical testing. However, the OF dimension (MODERATE priority) faces a confirmed mismatch: evaluation is single-label classification only [Q63], while the elicited downstream chatbot use case benefits from multi-label or ranked outputs (A3 explicitly confirms 'multi-label or ranked outputs would yield more contextually appropriate downstream responses'). The benchmark itself acknowledges multi-intent utterances [Q30] but defers the form change as future work. The emotion annotation experiment was abandoned due to imbalance [Q126–Q129], leaving no affective dimension in the output form. No subgroup-disaggregated performance metrics are reported, which becomes a regulatory issue under emerging FDA equity expectations [WEB-20, WEB-27]. Confusion analysis [Q90, Q91, Q92] is reported but not disaggregated by speaker role or cultural subgroup.	The original paper specifies no concrete evaluation metrics or scoring rubrics — it describes infrastructure for 'multiple testing runs and a variety of technical approaches' [Q10] supported by a relational database [Q11] but no scoring schema. SWBD-DAMSL implementations enforce single-label categorical output per utterance, with single-label accuracy as the primary metric [region YAML evaluation_protocol_notes; WEB-4]. The user explicitly acknowledges that counseling utterances often carry multiple simultaneous communicative intents [elicitation A3] and that single-label scoring is an oversimplification. ICSI-MRDA reintroduced multi-label assignment that SWBD-DAMSL eliminated, and AnnoMI-full preserves multi-attribute labels [WEB-10, WEB-16]. With OF marked MODERATE and the form mismatch acknowledged but treated as deferred, score 2 reflects significant but not critical concerns.

HOPE_DAC evidence	SWITCHBOARD_1992 evidence
[Q63] 'we compute macro-F1, weighted-F1, and accuracy scores.' (p.6)	[Q10] 'design for multiple testing runs and a variety of technical approaches' (p.1)
[Q73] 'The SPARTA-TAA system obtains the best scores of 60.29 macro-F1, 64.53 weighted-F1' (p.6)	[Q11] 'an underlying relational database system' (p.1)
[Q77] 'SPARTA consistently yields good scores for the majority of the dialogue-acts, except for the Acknowledgement (ACK) where it records F1-score of merely 47.86%.' (p.7)	[WEB-4] Stolcke et al. 2000 establishes single-label accuracy convention on SWBD-DAMSL
[Q88] 'our results are significant with > 95% confidence across macro-F1 (p-value= 0.009), weighted-F1 (p-value= 0.014), and accuracy (p-value= 0.048).' (p.7)	[WEB-10] AnnoMI GitHub preserves 10-annotator multi-attribute labels per utterance
[Q129] 'Due to such imbalance, this data was not used in the final version of our proposed architecture.' (p.11)	[WEB-16] ICSI-MRDA reintroduced multi-label assignment that SWBD-DAMSL eliminated
[WEB-20] Orrick (Nov 2025): FDA DHAC signals sponsors will need 'evidence of equitable performance across populations and languages' for AI mental health devices	
[WEB-27] FDA document on generative AI in digital mental health medical devices	