

Output Form (OF)

Score: 4/5 — Minor gaps | Confidence: high

Benchmark: MILU | Context: North Indian Hindi-Medium Postgraduate Teacher — MCQ Evaluation Deployment

Definition: Whether the expected output modality matches regional deployment needs and accessibility.

Justification

Output form aligns well with the deployment. Accuracy is the primary metric [Q3, Q15], with log-likelihood scoring for non-API models via LM-EVALUATION-HARNESS [Q52-Q55] and structured-JSON generative scoring for API models [Q56, Q57]. Both approaches produce a single answer-label output that maps directly to the deployment's binary correct/incorrect requirement. The deployment requires binary correct/incorrect feedback, which this output form supports natively. Reporting at 0/1/5-shot levels [Q50, Q100] gives operational flexibility. The authors explicitly acknowledge that log-likelihood evaluation may yield different results from generation-based or chain-of-thought approaches [Q85] — a caveat that matters if the deployment uses generative inference. API models are evaluated zero-shot only due to cost [Q58], potentially limiting comparability with few-shot deployment configurations.

Strengths

- Accuracy-based binary scoring directly maps to deployment's positive/negative marking requirement [Q3, Q15, Q16].
- Reproducible evaluation via LM-EVALUATION-HARNESS [Q52].
- Structured JSON generative output for API models matches programmatic-integration needs [Q57].
- 0/1/5-shot reporting enables deployment-config sensitivity analysis [Q50, Q100].

Checklist

Checklist item definitions:

ID	Assessment Question
OF-1	Does output modality match regional deployment needs?
OF-2	Assess text-to-speech availability for speech output
OF-3	Consider literacy rates and accessibility requirements
OF-4	Document form mismatches harming external validity

Checklist responses:

OF-1: Output modality (single answer-label) matches deployment's binary correct/incorrect feedback requirement [Q15, Q55, Q57].

OF-2: TTS not required — deployment is text-only Hindi mobile/desktop application per deployment_modality field.

OF-3: Target population is postgraduate-qualified teachers (high literacy assumed); accessibility considerations beyond Devanagari rendering not flagged in elicitation. INSUFFICIENT DOCUMENTATION on benchmark-side accessibility analysis but not relevant for this deployment.

OF-4: Minor mismatch: log-likelihood scoring [Q53-Q55] differs from generation-based deployment inference; the paper acknowledges this as a limitation [Q85]. API-evaluation zero-shot constraint [Q58] may limit comparability with deployments running few-shot.

Information Gaps

- Whether the deployment will use log-likelihood or generation-based inference; benchmark caveat suggests results may differ across these modes.

Remediation

Gap 1: Headline accuracy figures are computed via log-likelihood scoring, which the authors note may diverge from generation-based or chain-of-thought evaluation [Q85]; deployment may use generative inference.

Recommendation: Re-run MILU Hindi evaluation in the deployment's actual inference mode (generative, with structured JSON output) on at least the candidate model to ensure accuracy estimates reflect deployed behaviour.