Output Content (OC)

Score: 2/5 — Significant gaps | Confidence: high

Benchmark: *MRBENCH* | ***Context:*** *Metropolitan India — Grades 1–8 Mathematics Tutoring (Student Population)*

Output Content definition: Whether ground-truth labels align with the judgments of regional stakeholders.

> Each section below is followed by an evidence table. Three types of evidence are cited: **[QN]** — verbatim quotes extracted from the benchmark paper; **[WEB-N]** — web sources consulted during assessment; **[DN]** / **[DATASET-DN]** — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

Output content is MODERATE-priority per elicitation, with A3 indicating partial alignment expected between Indian student satisfaction and benchmark annotators but anticipating divergence on tone, warmth, and cultural fit. MRBench's annotators are four post-graduate Computer Science students proficient in English [Q31] with no disclosed nationality, cultural background, or teaching experience [Q32]; teaching experience was explicitly NOT required [Q32]. Inter-annotator agreement is substantial (Fleiss' kappa 0.65 pilot, Cohen's kappa 0.71 full-scale) [Q39, Q61], indicating internal label reliability — but reliability does not address whether labels reflect Indian stakeholder perspectives. No Indian educators, parents, or students contributed to ground-truth labels, and no cross-cultural validation has been performed; the BEA 2025 Shared Task extended MRBench without addressing this [WEB-17, WEB-18]. The deployment's actual judges (human tutor co-authors and student engagement signals per A3) are structurally different from the benchmark's annotator pool. LLM-as-judge alternatives are explicitly unreliable [Q107, Q109, Q114], a finding corroborated by 2024–2025 literature [WEB-19, WEB-20], so they cannot substitute for human re-annotation.

ID	Evidence
Q31	“The annotation team consisted of two male and two female annotators, with all four annotators holding at least a post-graduate degree in Computer Science and being proficient in English.” (p.5)
Q32	“We note that for this study, we do not require annotators to have direct teaching experience, as understanding of the mathematical tasks at the middle school level and being able to judge the responses from the perspective of a potential user of such AI tutors (or a student), rather than specifically a teacher, is sufficient.” (p.5)
Q39	“To measure inter-annotator agreement, we computed Fleiss' kappa value, which for this annotation experiment equals 0.65, indicating substantial agreement.” (p.5)
Q61	“The annotators reached an average Cohen's kappa score of 0.71, which indicates substantial inter-annotator agreement (McHugh, 2012).” (p.6)
Q107	“Across both LLMs, it can be observed that most of the correlation scores are negative (except for the human-likeness dimension), indicating that the annotations from the LLMs are unreliable for the challenging pedagogical dimensions.” (p.8)
Q109	“We believe both LLMs have a limited understanding of rich pedagogical concepts, as they were not specifically trained on pedagogically rich datasets.” (p.8)
Q114	“We also assess the feasibility of LLMs as evaluators in this context by correlating their judgements with human annotations, indicating that they are often unreliable.” (p.9)
WEB-17	BEA 2025 Shared Task extended MRBench without Indian rater validation (https://arxiv.org/html/2507.10579v1)
WEB-18	NAACL 2025 BEA proceedings: no participating system documented Indian adaptation (https://aclanthology.org/2025.naacl-long.57.pdf)
WEB-19	2025 MDPI: LLM evaluator inconsistency confirmed for educational evaluation (https://www.mdpi.com/2076-3417/15/2/671)
WEB-20	2025 study: LLM and human evaluator scores share little common signal (https://arxiv.org/pdf/2508.02442)

Strengths

- Strong inter-annotator agreement (Fleiss' kappa 0.65, Cohen's kappa 0.71) demonstrates internal label reliability [Q39, Q61]
- In-house annotation with comprehensive training and validation pilot avoids crowdsourcing quality issues [Q33, Q34, Q35]

- Expert and Novice human responses included as comparison baselines, with Expert treated as gold standard [Q68, Q78]
- Authors explicitly assess and report the unreliability of LLM critics [Q107, Q109, Q114] — providing honest evidence that automated re-annotation is not a substitute

ID	Evidence
Q33	"To control the annotation workflow and ensure quality, we opted not to use public annotation outsourcing platforms such as Prolific or MTurk, which allowed us to implement rigorous training protocols and a robust validation mechanism for the annotations." (p.5)
Q34	"First, we provided all annotators with comprehensive training, including an interactive training document (see Section C for more details) and oral instructions." (p.5)
Q35	"Following this, we conducted validation pilot study to evaluate the annotation quality and the annotators' understanding of the instructions before rolling out the large-scale human evaluation detailed in Section 5.2." (p.5)
Q39	"To measure inter-annotator agreement, we computed Fleiss' kappa value, which for this annotation experiment equals 0.65, indicating substantial agreement." (p.5)
Q61	"The annotators reached an average Cohen's kappa score of 0.71, which indicates substantial inter-annotator agreement (McHugh, 2012)." (p.6)
Q68	"We consider human-based evaluations as gold standard." (p.7)
Q78	"We also investigated the pedagogical value of human responses for both Novice and Expert." (p.7)
Q107	"Across both LLMs, it can be observed that most of the correlation scores are negative (except for the human-likeness dimension), indicating that the annotations from the LLMs are unreliable for the challenging pedagogical dimensions." (p.8)
Q109	"We believe both LLMs have a limited understanding of rich pedagogical concepts, as they were not specifically trained on pedagogically rich datasets." (p.8)
Q114	"We also assess the feasibility of LLMs as evaluators in this context by correlating their judgements with human annotations, indicating that they are often unreliable." (p.9)

Checklist

Checklist item definitions:

ID	Assessment Question
OC-1	Do ground-truth labels reflect regional stakeholder perspectives?
OC-2	Assess potential annotator-regional population disagreement
OC-3	Review annotator demographics from documentation
OC-4	Consider label re-annotation by regional annotators
OC-5	Review aggregation methods for minority perspective erasure
OC-6	Document label issues harming convergent/external validity

Checklist responses:

OC-1: No — labels reflect the perspectives of four Computer Science post-graduates [Q31] with undisclosed nationality and no teaching experience [Q32]. There is no documented Indian stakeholder representation in the annotator pool.

OC-2: Disagreement is anticipated, particularly on Tutor Tone and human-likeness dimensions where cultural register expectations differ. A3 indicates partial alignment expected; A4 confirms divergence on cultural appropriateness. The paper itself notes Tutor Tone's neutral/encouraging distinction collapses trivially [Q90], suggesting the dimension lacks discriminative power on the very axis where Indian register differences would matter.

OC-3: Annotator demographics partially disclosed: gender (2 male, 2 female), education (post-graduate Computer Science), and English proficiency [Q31]. Nationality, cultural background, and teaching experience are NOT disclosed; teaching experience explicitly not required [Q32]. This is below the level of detail recommended in Datasheets/Data Statements for cross-cultural validity assessment.

OC-4: Re-annotation by representative regional annotators (Indian educators, Indian metropolitan students or parents) is needed for the deployment, particularly on tone, human-likeness, and any added cultural-appropriateness dimension. No such re-annotation exists [WEB-17].

OC-5: Aggregation method (majority across four annotators with Cohen's kappa reported [Q57, Q61]) is reasonable for internal reliability but provides no mechanism to detect minority cultural perspectives; with all four annotators from a single demographic profile, minority perspectives are structurally absent.

OC-6: Documented harms to convergent validity (labels not validated against Indian raters) and external validity (judgments may not generalize to Indian metropolitan student satisfaction). The deployment's reliance on human tutor co-authorship partially mitigates this for the deployed system, but the benchmark itself cannot be relied on as a satisfaction proxy.

ID	Evidence
Q31	“The annotation team consisted of two male and two female annotators, with all four annotators holding at least a post-graduate degree in Computer Science and being proficient in English.” (p.5)
Q32	“We note that for this study, we do not require annotators to have direct teaching experience, as understanding of the mathematical tasks at the middle school level and being able to judge the responses from the perspective of a potential user of such AI tutors (or a student), rather than specifically a teacher, is sufficient.” (p.5)
Q39	“To measure inter-annotator agreement, we computed Fleiss' kappa value, which for this annotation experiment equals 0.65, indicating substantial agreement.” (p.5)
Q57	“A total of 192 instances were annotated, with 40 of those annotated independently by two annotators (10 instances from Bridge and 30 from MathDial) allowing us to calculate pairwise inter-annotator agreement.” (p.6)
Q61	“The annotators reached an average Cohen's kappa score of 0.71, which indicates substantial inter-annotator agreement (McHugh, 2012).” (p.6)
Q90	“When we combine these two labels into "Non-offensive", the DAMR score reaches 100% as we observe no offensive responses from any LLMs or humans.” (p.8)
Q107	“Across both LLMs, it can be observed that most of the correlation scores are negative (except for the human-likeness dimension), indicating that the annotations from the LLMs are unreliable for the challenging pedagogical dimensions.” (p.8)
WEB-17	BEA 2025 Shared Task extended MRBench without Indian rater validation (https://arxiv.org/html/2507.10579v1)
WEB-18	NAACL 2025 BEA proceedings: no participating system documented Indian adaptation (https://aclanthology.org/2025.naacl-long.57.pdf)
WEB-19	2025 MDPI: LLM evaluator inconsistency confirmed for educational evaluation (https://www.mdpi.com/2076-3417/15/2/671)
WEB-20	2025 study: LLM and human evaluator scores share little common signal (https://arxiv.org/pdf/2508.02442)

Information Gaps

- Annotator nationality and cultural background not disclosed in the paper
- No cross-cultural validation study with Indian raters has been published
- No published comparison of MRBench labels to Indian human-tutor co-author judgments

Requires Expert Verification

- Re-annotation of a sample of MRBench instances by Indian educators and metropolitan Grade 1–8 students to quantify divergence on tone and human-likeness
- Calibration of the deployment's human tutor co-author judgments against MRBench gold labels

---

Remediation

Gap 1: Ground-truth labels reflect four Computer Science post-graduates with undisclosed nationality and no teaching experience; no Indian rater validation in MRBench or BEA 2025 Shared Task.

Recommendation: Recruit a panel of Indian educators (NCERT-familiar primary and middle-school math teachers) and a smaller panel of metropolitan Grade 1–8 students or parents to re-annotate at least the 40 double-annotated instances [Q57] across all dimensions, then quantify divergence from MRBench gold labels and report dimension-specific recalibration.

ID	Evidence

Q57	“A total of 192 instances were annotated, with 40 of those annotated independently by two annotators (10 instances from Bridge and 30 from MathDial) allowing us to calculate pairwise inter-annotator agreement.” (p.6)