

Validity Analysis: BLURB

Target context: EU/French Pharmaceutical Regulatory NLP — BLURB Assessment

Overall risk: **HIGH**

Deployment Context

Use case: In a pharmaceutical setting, a regulatory specialist uses document management software to ensure compliance in French drug labeling and scientific abstracts by applying a model evaluated on semantic textual similarity and fine-grained named-entity recognition. The system identifies over a dozen types of precise entities, such as chemical compounds and dosage modes, while verifying that safety warnings across different leaflets or patents are semantically consistent. These results determine whether pharmacological documentation meets strict professional standards for official submission or requires manual revision to prevent regulatory non-compliance.

Target population: Regulatory affairs specialists, pharmacologists, and legal experts at pharmaceutical companies or government health agencies.

Dimension Scores

Dimension	Score	Rating	Confidence
Input Ontology	2	Significant gaps	high
Input Content $\triangle$	1	Serious concern	high
Input Form $\triangle$	1	Serious concern	high
Output Ontology $\triangle$	2	Significant gaps	high
Output Content $\triangle$	2	Significant gaps	high
Output Form $\checkmark$	3	Moderate gaps	medium
Average	1.8		

$\triangle$ = highest concern $\checkmark$ = strongest dimension

Dimension Key

Abbr.	Dimension	Definition
IO	Input Ontology	Whether the benchmark's test case categories cover the query types expected in deployment.
IC	Input Content	Whether individual datapoint content is culturally and contextually appropriate for the target region.
IF	Input Form	Whether the input signal encoding (text, audio, image parameters) matches deployment conditions.
OO	Output Ontology	Whether the benchmark's output categories and scoring criteria reflect regionally valid decision boundaries.
OC	Output Content	Whether ground-truth labels align with the judgments of regional stakeholders.
OF	Output Form	Whether the expected output modality matches regional deployment needs and accessibility.

Overall Summary

BLURB is a well-constructed benchmark for English-language PubMed-based biomedical NLP, but it is fundamentally misaligned with the EU/French pharmaceutical regulatory NLP deployment across five of six validity dimensions. The language gap (English-only [Q1, Q126] vs French-language deployment), document-genre gap (PubMed abstracts vs SmPCs/PILs/CTD

modules), entity-ontology gap (no INN/ATC/posology/excipient label spaces), STS-construct gap (general biomedical proximity [Q74] vs legally-determinative regulatory equivalence), and annotator-norm gap (biomedical researchers vs EMA/ANSM-aligned regulatory experts) are each rated HIGH severity by the elicitation and confirmed by web research [WEB-9, WEB-10, WEB-12, WEB-23]. Output form is the only dimension with reasonable representational alignment, and even there, the absence of calibration/threshold-sensitivity reporting limits direct utility for a human-review-triggering compliance pipeline. No remediation can be done within BLURB itself; bespoke French regulatory benchmark development against the EMA AI Reflection Paper [WEB-6] normative framework is required.

ID	Evidence
Q1	"we compile a comprehensive biomedical NLP benchmark from publicly-available datasets." (p.1)
Q74	"The Sentence Similarity Estimation System for the Biomedical Domain [54] contains 100 pairs of PubMed sentences each of which is annotated by five expert-level annotators with an estimated similarity score in the range from 0 (no relation) to 4 (equivalent meanings)." (p.9)
Q126	"For biomedical domain-specific pretraining, we generate the vocabulary and conduct pretraining using the latest collection of PubMed abstracts: 14 million abstracts, 3.2 billion words, 21 GB." (p.12)
WEB-6	EMA AI Reflection Paper requires human-oversight mechanisms and quality review for AI in product information drafting — implies threshold/calibration characterization the deployment must satisfy (https://www.ema.europa.eu/en/documents/scientific-guideline/reflection-paper-use-artificial-intelligence-ai-medicinal-product-lifecycle_en.pdf)
WEB-9	2025 Frontiers SmPC IDMP study — ATC codes and excipient dosages 'frequently missing or misclassified' by LLMs (https://www.frontiersin.org/journals/medicine/articles/10.3389/fmed.2025.1598979/full)
WEB-10	EMA 2024 AI Observatory Report — internal regulatory NLP tools developed under EMA staff oversight, not represented in any public benchmark (https://www.ema.europa.eu/en/documents/report/2024-ai-observatory-report_en.pdf)
WEB-12	No regulatory-expert-annotated pharmaceutical NLP dataset (French or English) identified; DART Italian corpus does not use regulatory-expert annotators (https://arxiv.org/html/2510.18475)
WEB-23	Practitioner review confirms NLP for EU SmPC/prescribing information remains hybrid with no validated regulatory-specific benchmark (https://intuitionlabs.ai/articles/nlp-regulatory-labeling)

Practical Guidance

What This Benchmark Measures

BLURB measures the relative quality of English biomedical pretrained encoders on PubMed-derived NER, PICO extraction, relation extraction, sentence similarity (BIOSSES), document classification (HoC), and yes/no QA tasks [Q48, Q150]. For the EU/French regulatory deployment it can serve, at most, as an upper-bound sanity check on a French-translated or French-equivalent regulatory pipeline's underlying biomedical NLP competence — and even this only if a French biomedical PLM (DrBERT, CamemBERT-bio, AliBERT [WEB-14, WEB-15, WEB-16]) is the model under evaluation. None of the deployment's HIGH-priority constructs (regulatory NER ontology, regulatory equivalence STS, regulatory-expert annotation norms) is directly measurable through BLURB.

Construct Depth

BLURB probes generic biomedical NER, RE, STS, and QA constructs at moderate depth on English PubMed text. It does not probe regulatory-document understanding, French-language regulatory phrasing, regulatory-equivalence scoring, cross-document consistency, or calibration at decision boundaries. The five HIGH-priority dimensions in this deployment (IO, IC, IF, OO, OC) are essentially unaddressed by BLURB, and the MODERATE-priority OF dimension is only partially addressed (representation matches; calibration/threshold reporting does not).

What Else You Need

A complete assessment requires: (1) a French-language regulatory document corpus seeded by QUAERO EMEA [WEB-19, WEB-20] and informed by the DART Italian SmPC precedent [WEB-12]; (2) a regulatory-specific entity ontology (INN, ATC, excipient, posology, EMA-template qualifiers) with annotation guidelines anchored in EMA QRD v10.4 [WEB-1, WEB-2] and ANSM circulars; (3) a regulatory-equivalence STS dataset annotated by EMA/ANSM-trained regulatory affairs specialists, with critical-mismatch sub-categories (dose, population, indication scope); (4) calibration and threshold-sensitivity evaluation

framed against the EMA AI Reflection Paper [WEB-6] and EU AI Act high-risk requirements; (5) an IAA study comparing biomedical NLP annotators with regulatory-expert annotators on shared regulatory text to quantify the OC gap.

ID	Evidence
Q48	“BLURB is comprised of a comprehensive set of biomedical NLP tasks from publicly available datasets, including named entity recognition (NER), evidence-based medical information extraction (PICO), relation extraction, sentence similarity, document classification, and question answering.” (p.7)
Q150	“The BLURB score is the macro average of average test results for each of the six tasks (NER, PICO, relation extraction, sentence similarity, document classification, question answering).” (p.14)
WEB-1	EMA QRD SmPC Template v10.4 (Feb 2024) — defines the formulaic template language the deployment must process (https://www.ema.europa.eu/en/human-regulatory-overview/marketing-authorisation/product-information-requirements/product-information-qrd-templates-human)
WEB-2	https://www.ema.europa.eu/en/documents/template-form/qrd-product-information-template-version-104-highlighted_en.pdf
WEB-6	EMA AI Reflection Paper requires human-oversight mechanisms and quality review for AI in product information drafting — implies threshold/calibration characterization the deployment must satisfy (https://www.ema.europa.eu/en/documents/scientific-guideline/reflection-paper-use-artificial-intelligence-ai-medicinal-product-lifecycle_en.pdf)
WEB-12	No regulatory-expert-annotated pharmaceutical NLP dataset (French or English) identified; DART Italian corpus does not use regulatory-expert annotators (https://arxiv.org/html/2510.18475)
WEB-14	DrBERT — French biomedical PLM pretrained from scratch (ACL 2023), not part of BLURB (https://arxiv.org/abs/2304.00958)
WEB-15	CamemBERT-bio — French biomedical continual-pretraining model (LREC-COLING 2024) (https://arxiv.org/abs/2306.15550)
WEB-16	AliBERT — French biomedical PLM with Unigram tokenizer (BioNLP 2023) (https://aclanthology.org/2023.bionlp-1.19/)
WEB-19	QUAERO French Medical Corpus — only public French regulatory-document NLP resource, very small EMEA subcorpus (https://quaerofrenchmed.limsi.fr/)
WEB-20	QUAERO HuggingFace dataset documentation — 4 EMEA documents segmented into 15 sub-documents (https://huggingface.co/datasets/DrBenchmark/QUAERO)