

Output Form (OF)

Score: 4/5 — Minor gaps | Confidence: high

Benchmark: DRBENCHMARK | Context: EU Pharmaceutical Regulatory NLP — French

Definition: Whether the expected output modality matches regional deployment needs and accessibility.

Justification

Output representations (IOB2 NER labels, multi-label classification, multiple-choice indices, continuous 0–5 STS scores [Q52, Q53, Q54]) match the deployment's consumption modalities. The deployment ingests entity spans and STS scores; both are produced by the benchmark in compatible formats. Statistical reporting (averaging over four runs, Student's t-test [Q51], majority-class baseline [Q56]) supports interpretable results. SeqEval with first-token prediction [Q52, Q53] provides tokenizer-agnostic evaluation. The minor concerns are (a) the EDRM metric's specificity to DEFT-2020 [Q54] without documented calibration for regulatory micro-differences, and (b) the ~2,500 GPU-hour reproducibility cost [Q92, Q93] for users without HPC access. OF is LOWER priority per the elicitation.

Strengths

- Standard NER output via SeqEval IOB2 with tokenizer-agnostic first-token prediction [Q52, Q53]
- Continuous 0–5 STS scores in a format directly consumable by the deployment's compliance flagging logic [Q41, Q54]
- Robust statistical reporting: four runs averaged, Student's t-test [Q51], majority-class baseline [Q56]
- Best-model selection via validation metric mitigates overfitting [Q100]
- Hyperparameters documented for reproducibility [Q49, Q99]

Checklist

Checklist item definitions:

ID	Assessment Question
OF-1	Does output modality match regional deployment needs?
OF-2	Assess text-to-speech availability for speech output
OF-3	Consider literacy rates and accessibility requirements
OF-4	Document form mismatches harming external validity

Checklist responses:

OF-1: Output modalities (IOB2 entity spans, classification labels, continuous STS scores) match deployment needs [Q52, Q54]. The deployment ingests entity spans for NER and continuous scores for STS — both are produced by the benchmark.

OF-2: Not applicable — the deployment is text-only enterprise document processing; no text-to-speech requirements (deployment infrastructure_notes).

OF-3: Not applicable — the target population is highly literate professional regulatory affairs specialists; no accessibility constraints relevant to output form.

OF-4: Minor form concerns: (a) EDRM metric specific to DEFT-2020 task design [Q54] has no documented calibration for regulatory equivalence, but this is a calibration concern (OO) rather than a representation mismatch; (b) reproducibility burden of 2,500 GPU hours [Q92, Q93] limits independent re-running but does not affect output-form compatibility.