

Validity Analysis: DRBENCHMARK

Target context: EU Pharmaceutical Regulatory NLP — French

Overall risk: **HIGH**

Deployment Context

Use case: In a pharmaceutical setting, a regulatory specialist uses document management software to ensure compliance in French drug labeling and scientific abstracts by applying a model evaluated on semantic textual similarity and fine-grained named-entity recognition. The system identifies over a dozen types of precise entities, such as chemical compounds and dosage modes, while verifying that safety warnings across different leaflets or patents are semantically consistent. These results determine whether pharmacological documentation meets strict professional standards for official submission or requires manual revision to prevent regulatory non-compliance.

Target population: Regulatory affairs specialists, pharmacologists, and legal experts at pharmaceutical companies or government health agencies.

Dimension Scores

Dimension	Score	Rating	Confidence
Input Ontology 	2	Significant gaps	high
Input Content 	2	Significant gaps	high
Input Form 	4	Minor gaps	high
Output Ontology 	1	Serious concern	high
Output Content 	2	Significant gaps	high
Output Form 	4	Minor gaps	high
Average	2.5		

 = highest concern  = strongest dimension

Dimension Key

Abbr.	Dimension	Definition
IO	Input Ontology	Whether the benchmark's test case categories cover the query types expected in deployment.
IC	Input Content	Whether individual datapoint content is culturally and contextually appropriate for the target region.
IF	Input Form	Whether the input signal encoding (text, audio, image parameters) matches deployment conditions.
OO	Output Ontology	Whether the benchmark's output categories and scoring criteria reflect regionally valid decision boundaries.
OC	Output Content	Whether ground-truth labels align with the judgments of regional stakeholders.
OF	Output Form	Whether the expected output modality matches regional deployment needs and accessibility.

Overall Summary

DrBenchmark is the most comprehensive French biomedical NLU benchmark currently available [Q16, Q83], and parts of its content (QUAERO EMEA, Mantra-GSC fr_emea, DEFT2020 drug-leaflet pairs, PxCorpus posology schema) provide partial signal for the EU pharmaceutical regulatory deployment. However, the deployment's HIGH-priority dimensions — input

ontology, input content, output ontology, and output content — show severe misalignment. The benchmark is calibrated for clinical and research French biomedical NLP, not for EMA/ANSM regulatory affairs workflows: (i) regulatory document genres (SmPCs, CTD modules, EU-RMPs) are absent; (ii) regulatory-specific entity types (INN as distinct from excipient, ATC codes, EMA-templated posology, contraindication qualifiers) are not annotated as distinct classes; (iii) STS scoring is calibrated for general semantic proximity and demonstrably tolerates legally significant micro-differences (2×–60× numeric divergences score 3.0–4.0); (iv) annotators are biomedical NLP / clinical experts, not regulatory affairs specialists, and elicitation explicitly anticipates systematic disagreement on borderline cases. Lower-priority dimensions (input form, output form) are well aligned. Web search confirms several of these gaps are field-level, not just benchmark-level [WEB-11, WEB-12, WEB-16].

ID	Evidence
Q16	“To our knowledge, aside the multilingual benchmark BigBio (Fries et al., 2022) which includes only 4 corpora for French and is initially intended for generative text completion under zero-shot scenario, no large benchmark specialized in the French biomedical field exists.” (p.2)
Q83	“In this paper, we introduced DrBenchmark, the first large language understanding benchmark tailored for the French biomedical domain.” (p.9)
WEB-11	Kadi et al. 2025 — confirms field-level gap; no regulatory-grade annotation rubric for SmPC extraction is established (https://pmc.ncbi.nlm.nih.gov/articles/PMC12380897/)
WEB-12	Bayer et al. 2021 — used pharmacovigilance evaluators for ADE extraction but not regulatory equivalence scoring; closest comparator confirms gap (https://link.springer.com/article/10.1007/s40264-020-00996-3)
WEB-16	IQVIA fact sheet — IDMP requires structured extraction from unstructured SmPC text; ~70% of required attributes reside in unstructured text; this task category is unaddressed by DrBenchmark (https://www.iqvia.com/-/media/iqvia/pdfs/library/fact-sheets/nlp-text-analytics-for-regulatory-affairs.pdf)

Practical Guidance

What This Benchmark Measures

DrBenchmark provides credible measurement of general French biomedical NLU capabilities in clinical and research register: clinical NER (UMLS Semantic Groups), POS, MeSH-axis classification, ICD-10 chapter classification, MCQA over pharmacy exam knowledge, and clinical STS. For the deployment, it offers partial signal on (a) coarse pharmaceutical entity recognition over EMEA leaflet text via QUAERO EMEA and Mantra-GSC fr_emea, (b) French-language STS capability over biomedical sentence pairs including some drug-leaflet content via CLISTER and DEFT2020, and (c) compositional posology recognition via PxCorpus and DEFT2021. Output form (IO-2) is fully aligned, and input form (IF) is mostly aligned.

Construct Depth

Construct probing for the deployment's core needs is shallow. Regulatory-equivalence STS — the deployment's central decision rule — is not measured at all; CLISTER and DEFT2020 measure general semantic proximity, and dataset analysis empirically shows the calibration is incompatible (CLISTER-D12, DEFT2020-D16). Regulatory entity NER is measured only as conflated CHEM/SUBSTANCE classes; INN/ATC/excipient sub-distinctions are unevaluated. Annotation provenance for the most deployment-relevant subsets (QUAERO, Mantra-GSC, CLISTER) is undocumented, so even successful benchmark scores carry unquantified construct-validity risk for borderline regulatory cases.

What Else You Need

To complete the assessment, the deployment team would need: (1) a regulatory-affairs-annotated SmPC/PIL/CTD evaluation set covering INN, ATC, excipient, posology, and contraindication-qualifier entity types — addressing Input Ontology, Input Content, and Output Ontology gaps; (2) a regulatory-equivalence STS rubric and dataset with annotator guidelines that penalize dose-threshold and population-qualifier divergences — addressing Output Ontology calibration and Output Content (annotator alignment); (3) a re-annotation of QUAERO EMEA and Mantra-GSC fr_emea spans by regulatory affairs specialists to quantify systematic disagreement — addressing Output Content; (4) integration of EMA QRD v10.4 [WEB-7, WEB-18] and IDMP [WEB-16, WEB-17] structural extraction targets as evaluation tasks. Given that web search confirms these are field-level gaps [WEB-11, WEB-12], supplementation will require new resource creation, not just resource discovery.

ID	Evidence
----	----------

CLISTER-D12	L'AHG urinaire est à 10 mmol/l." / "L'AHG urinaire est à 20 mmol/l. — 2× quantitative difference scored 4.0 (very similar)
DEFT2020-D16	Ce produit peut provoquer un syndrome de sevrage opiacé s'il est administré à un toxicomane moins de 4 heures après la dernière prise de stupéfiant. — High similarity score despite addition of specific drug names in target — legally significant detail unmarked
WEB-7	EMA QRD-templated documents have formulaic typographic conventions (section numbering, tabular posology) that may stress sentence-splitting at boundaries (https://www.ema.europa.eu/en/human-regulatory-overview/marketing-authorisation/product-information-requirements/product-information-qrd-templates-human)
WEB-11	Kadi et al. 2025 — confirms field-level gap; no regulatory-grade annotation rubric for SmPC extraction is established (https://pmc.ncbi.nlm.nih.gov/articles/PMC12380897/)
WEB-12	Bayer et al. 2021 — used pharmacovigilance evaluators for ADE extraction but not regulatory equivalence scoring; closest comparator confirms gap (https://link.springer.com/article/10.1007/s40264-020-00996-3)
WEB-16	IQVIA fact sheet — IDMP requires structured extraction from unstructured SmPC text; ~70% of required attributes reside in unstructured text; this task category is unaddressed by DrBenchmark (https://www.iqvia.com/-/media/iqvia/pdfs/library/fact-sheets/nlp-text-analytics-for-regulatory-affairs.pdf)
WEB-17	https://www.linguamatics.com/blog/countdown-idmp-compliance-are-you-ready
WEB-18	EMA QRD v10.4 template PDF — current operational template structure (https://www.ema.europa.eu/en/documents/template-form/qrd-product-information-template-version-104-highlighted_en.pdf)