

Comparative Assessment: LTZGLUE vs. GLUE

Deployment context: Luxembourg Public Administration — Civil Servant NLU Deployment

Deployment Context

Use case: In a public administration setting, a civil servant uses LLM-powered software to manage Luxembourgish citizen feedback by applying a model evaluated on topic classification, named entity recognition, intent detection, and sentiment analysis. The system automatically categorizes messages by domain and identifies key entities while classifying user intent and gauging sentiment to help detect frustrated users. These outputs enable the automated routing of requests to specific departments and the immediate prioritization of disgruntled or urgent correspondence, ensuring administrative actions are triggered accurately based on the content and tone of the message.

Target population: Civil servants in Luxembourgish government agencies who process digital correspondence and citizen inquiries written in the diverse orthographic styles of the Luxembourgish language.

Overall Risk Ratings

Benchmark	Risk	Avg. Score
LTZGLUE	HIGH	2.2 / 5
GLUE	HIGH	1.2 / 5

Dimension Scores Comparison

Dimension	LTZGLUE	Rating	GLUE	Rating	Δ
Input Ontology (IO)	2/5	Significant gaps	1/5	Serious concern	+1
Input Content (IC)	2/5	Significant gaps	1/5	Serious concern	+1
Input Form (IF)	3/5	Moderate gaps	1/5	Serious concern	+2
Output Ontology (OO)	1/5	Serious concern	1/5	Serious concern	0
Output Content (OC)	2/5	Significant gaps	1/5	Serious concern	+1
Output Form (OF)	3/5	Moderate gaps	2/5	Significant gaps	+1
Average	2.2		1.2		+1

Δ = LTZGLUE score – GLUE score. Positive = regional benchmark scores higher.

Overall Summaries

LTZGLUE (risk: HIGH)

LTZGLUE is a credible first GLUE-style NLU benchmark for Luxembourgish and a valuable resource for general LTZ NLP capability assessment. However, for the specific deployment of routing citizen correspondence in Luxembourg public administration, validity is materially limited across the four HIGH-priority dimensions (IO, IC, OO, OC). The most severe gap is output ontology: LTZGLUE produces single hard labels across taxonomies (RTL editorial topics, voice-assistant intents, ternary sentiment) that do not encode the administrative routing categories, multi-label structure, confidence scores, urgency scaling, or language-of-input axis the deployment requires. Input content reflects predominantly formal/institutional registers, with the authors themselves warning this does not represent informal and multilingual contexts — precisely the distribution of real citizen correspondence given Luxembourg's ~47% frontier workforce and entrenched code-switching. Annotation quality is moderate ($\kappa=0.45$ on sentiment) and lacks documented administrative-domain expertise. Input form and output form mismatches are

moderate: text-to-label alignment is preserved, but length/language filtering and absence of calibration/multi-label metrics weaken external validity. The benchmark's authors explicitly caution against using performance as evidence of cultural or demographic coverage [Q121], and this caution applies directly here.

GLUE (risk: HIGH)

GLUE is fundamentally invalid as an evaluation benchmark for a Luxembourgish public-administration citizen-correspondence routing system. Five of six dimensions score 1 (major validity violations), with OF scoring 2 only because the surface text-in/label-out modality matches. The mismatches are total along every HIGH-priority dimension: (IO) GLUE's nine English NLU tasks contain zero administrative-routing categories and no national-vs-communal or frontier-vs-resident classifications [Q3, Q18, WEB-16, WEB-19]; (IC) all 409 sampled examples are English with US-political/cultural dominance, while the deployment input is trilingual lb/fr/de code-switched correspondence with non-standardized Luxembourgish orthography [WEB-11, WEB-12, DATASET-Concern 1]; (IF) GLUE assumes standardized English text input, while Luxembourgish input requires trilingual tokenization and ~21% residual orthographic noise even after best-available normalization [WEB-12]; (OO) hard single-label classification cannot represent multi-label routing with confidence scoring [Q12, DATASET-Concern 4]; (OC) annotator pools are undocumented English-speaking populations with no connection to Luxembourgish administrative register [Q26, Q36, Q73]; (OF) macro-average over nine irrelevant tasks [Q60] is uninformative as a deployment proxy. The strongest reusable value is methodological — GLUE's multi-task structure, diagnostic minimal-pair design [Q65, Q66], and human-baseline reporting [Q73, Q74] are transferable design patterns, not transferable evaluation signal.

Per-Dimension Justification Comparison

Each benchmark's full assessment justification and cited evidence, shown side by side.

Input Ontology (IO)

Whether the benchmark's test case categories cover the query types expected in deployment.

LTZGLUE — 2/5 (Significant gaps) [high confidence]	GLUE — 1/5 (Serious concern) [high confidence]
LTZGLUE provides a GLUE-style suite of eight tasks covering binary and multi-class sentence- and token-level classification [Q11, Q67], which is a credible general-purpose NLU foundation for Luxembourgish. However, for the specific deployment — routing citizen correspondence to government departments — the taxonomies are misaligned. Topic classification covers only five RTL editorial categories (SPORTS, CULTURE, TECHNOLOGY, BUSINESS, ANIMALS) [Q48, Q153] with no administrative categories such as cross-border worker issues, housing, transport, tax, social security, or communal vs. national competency routing — all flagged as required in the elicitation and regional context. Intent detection labels are exclusively voice-assistant commands (alarm, reminder, weather) [Q154], with no administrative intents (complaint, document request, escalation). The benchmark contains no multi-label task formulation [Q67], whereas the deployment requires multi-label outputs [elicitation Q4]. The authors themselves acknowledge no unified benchmark currently exists for LTZ NLU [Q14]. Given IO is HIGH priority and the gaps span both omission (administrative categories) and irrelevance (voice-assistant intents) for this deployment, this constitutes significant taxonomic misalignment.	GLUE's task ontology consists of nine English sentence-understanding tasks — sentiment, paraphrase, entailment, QA, acceptability, similarity [Q3, Q18] — with no administrative-domain categories. The deployment requires routing categories that GLUE does not contain: cross-border worker (frontalier) classification, national vs. communal (État vs. commune) competency routing, housing urgency, social security, immigration, taxation, and public-transport queries (deployment_specific_categories in YAML). The diagnostic set's 'World Knowledge' [Q195] and 'Common Sense' [Q196] categories are general English-language reasoning probes with no jurisdictional specificity. IO is marked HIGH priority and the gap is total — the dataset analysis confirms zero administrative-domain content across 409 sampled examples.
LTZGLUE evidence	GLUE evidence

[Q11] 'In this section, we introduce the eight tasks for LTZGLUE. The set spans binary and multi-class sentence and token-level classification tasks.' (p.2)	[Q3] 'GLUE is a collection of NLU tasks including question answering, sentiment analysis, and textual entailment, and an associated online platform for model evaluation, comparison, and analysis.' (p.1)
[Q14] 'No unified benchmark currently exists to evaluate LTZ language understanding consistently, a gap we aim to fill.' (p.2)	[Q18] 'GLUE is centered on nine English sentence understanding tasks, which cover a broad range of domains, data quantities, and difficulties.' (p.3)
[Q48] 'From the available categories, we focused on five principal domains: SPORTS, CULTURE, TECHNOLOGY, BUSINESS, and ANIMALS.' (p.4)	[Q195] 'World Knowledge In this category we focus on knowledge that can clearly be expressed as facts, as well as broader and less common geographical, legal, political, technical, or cultural knowledge.' (p.20)
[Q67] 'Together, the eight tasks in LTZGLUE form a broad and balanced evaluation suite, covering four binary and four multi-class settings, sentence- and document-level inputs, as well as a token-level sequence-labelling task.' (p.5)	
[Q153] 'topic_classification: Classify topic of the document by title and text. Allowed category_names: sports, animals, business, culture, technology.' (p.15)	
[Q154] 'slot_intent_detection: Detect the intent for the text given. Allowed intents: reminder/show_reminders, weather/find, reminder/set_reminder...' (p.15)	

Input Content (IC)

Whether individual datapoint content is culturally and contextually appropriate for the target region.

LTZGLUE — 2/5 (Significant gaps) [high confidence]	GLUE — 1/5 (Serious concern) [high confidence]
Benchmark inputs are drawn predominantly from RTL news articles, the Luxembourgish Online Dictionary, Wikipedia/Wikidata, and a translated English voice-assistant corpus [Q12, Q16, Q24, Q30, Q34, Q46, Q51]. The authors explicitly warn that 'most data sources reflect formal writing or institutional usage and therefore do not fully represent informal and multilingual contexts' [Q112], and that LTZ data 'often originate from a limited set of public domains' [Q119]. The deployment, however, expects orthographically variable, code-switched (Luxembourgish/French/German) citizen correspondence [elicitation Q2; regional context languages.code_switching_patterns]. Only the human-annotated NER sub-corpus from RTL user comments provides partial coverage of informal and code-mixed writing [Q41]. The intent detection corpus uses a voice-assistant register that has no LTZ reference, so terminology could not be systematically verified [Q55, Q56]. Given that ~47% of Luxembourg's workforce are frontaliers (mostly French-speaking) [WEB-5, WEB-2] and the language law requires the administration to respond in the citizen's chosen language [WEB-14], the benchmark's content distribution is materially narrower than the deployment input distribution. IC is HIGH priority and the documented gaps are substantial.	GLUE's input content is entirely English, drawn from American/British movie reviews [Q26], US/UK news [Q29, Q48], Wikipedia [Q41], Quora [Q31], linguistics journals [Q21], and English-language fiction and transcribed speech [Q38]. The dataset analysis confirms across 409 sampled examples that 'every single example reviewed across all 12 configs is in English' with US-political (Clinton, Bush, Gingrich, Ryan, Tillerson — DATASET-D10, D13, D14, D15, D22, D34), India-specific (DATASET-D16, D18), and US-pop-culture (DATASET-D3) cultural references dominating. The deployment input is trilingual Luxembourgish/French/German code-switched citizen correspondence with non-standardized Luxembourgish orthography [WEB-12, WEB-11], from a population that is 47.3% non-Luxembourgish nationals [WEB-4] with 47% of the workforce being cross-border workers from France, Belgium, and Germany [WEB-6, WEB-7]. The cultural and linguistic distance is total. IC is marked HIGH priority.
LTZGLUE evidence	GLUE evidence

[Q12] 'Unless stated otherwise, the textual data used across most tasks stems from two main sources' (p.2)	[Q21] 'The Corpus of Linguistic Acceptability (Warstadt et al., 2018) consists of English acceptability judgments drawn from books and journal articles on linguistic theory.' (p.3)
[Q41] 'It covers a wider range of text types and registers, including informal and code-mixed writing' (p.4)	[Q26] 'The Stanford Sentiment Treebank (Socher et al., 2013) consists of sentences from movie reviews and human annotations of their sentiment.' (p.3)
[Q55] 'The source segments consist of user commands for a voice-controlled AI assistant, representing a specialised spoken register for which there is no equivalent reference corpus in LTZ' (p.5)	[Q29] 'The Microsoft Research Paraphrase Corpus (Dolan & Brockett, 2005) is a corpus of sentence pairs automatically extracted from online news sources...' (p.3)
[Q56] 'Due to the lack of LTZ references in this register, it was not possible to systematically verify the translated terminology.' (p.5)	[Q38] 'The premise sentences are gathered from ten different sources, including transcribed speech, fiction, and government reports.' (p.4)
[Q112] 'most data sources reflect formal writing or institutional usage and therefore do not fully represent informal and multilingual contexts' (p.9)	[WEB-4] 47.3% foreign nationals as of 1 January 2024; over 170 nationalities — confirms input population's linguistic diversity
[Q119] 'LTZ is a small language community, and linguistic data often originate from a limited set of public domains.' (p.10)	[WEB-6] ~47% of workforce are cross-border workers (~231,290 frontaliers in 2024) — confirms French/German linguistic input share
[WEB-5] ~231,290 frontaliers (~47% of employment) — citizen-input distribution will be heavily multilingual	[WEB-11] LuxBorrow corpus documents systematic LU/DE/FR/EN code-switching in Luxembourgish text — empirical evidence of trilingual input characteristic
[WEB-14] Law of 24 February 1984 Article 4 requires reply in citizen's chosen language — input language detection is legally salient	[WEB-12] Best neural normalizer for Luxembourgish achieves only ~78.8% accuracy — empirical evidence of orthographic non-standardization absent from GLUE
	DATASET-D10 : 'I did not mention Monica in my lecture, but the first question I was asked was how President Clinton could do his job...' — US political scandal dependency
	DATASET-D16 : 'What is the effect of black money on India's macro economy?' — confirms even non-US content is non-Luxembourgish

Input Form (IF)

Whether the input signal encoding (text, audio, image parameters) matches deployment conditions.

LTZGLUE — 3/5 (Moderate gaps) [medium confidence]	GLUE — 1/5 (Serious concern) [high confidence]
Form mismatch is moderate. Both benchmark and deployment are text-only in Latin script, so there is no modality or script-direction violation [regional context infrastructure_notes.deployment_modality; writing_systems]. Tokenization (BPE, vocab 50,368, max sequence 1,024) [Q128, Q129] is suitable for typical citizen-correspondence lengths (40–400 words covers many but not all messages [regional context infrastructure_notes.expected_input_length, Q47]). However, signal-distribution mismatch persists at the form level: task inputs are filtered to remove non-Luxembourgish text via OpenLID and constrained to 40–400 words [Q47], which removes precisely the multilingual and length-extreme inputs the deployment will encounter. The benchmark's text register (formal news, dictionary examples, transcribed political speeches) [Q74, Q16, Q30] differs from informal, orthographically variable citizen messages. IF is MODERATE priority; the form-level alignment is partial because text-text matching is achieved, but distributional filtering and register narrowness reduce external validity.	GLUE inputs are standardized written English text in a single script with consistent orthography, formatted as single sentences or sentence pairs [Q4, Q15]. The deployment input is text in three languages with frequent code-switching at lexical, phrasal, and sentence level [WEB-11], non-standardized Luxembourgish spelling where 'large amount of orthographic variation' and homograph density are documented challenges [WEB-12, WEB-15], and a register requiring trilingual tokenization and normalization that English NLP pipelines do not provide [WEB-2, WEB-13, WEB-14]. GLUE's preprocessing transformations (recasting SQuAD into NLI [Q42], converting Winograd [Q52]) are English-specific and provide no transferable methodology for non-standardized multilingual signals. IF is HIGH priority and the gap is fundamental.
LTZGLUE evidence	GLUE evidence
[Q47] 'we removed articles identified as non-Luxembourgish by OpenLID (Burchell et al., 2023), as well as those containing fewer than 40 words or more than 400 words.' (p.4)	[Q4] 'GLUE does not place any constraints on model architecture beyond the ability to process single-sentence and sentence-pair inputs and to make corresponding predictions.' (p.1)
[Q74] 'A large portion of the data stems from RTL... transcribed radio interviews... user comments... transcribed podcasts... transcribed political speeches and debates' (p.6)	[Q42] 'We convert the task into sentence pair classification by forming a pair between each question and each sentence in the corresponding context, and filtering out pairs with low lexical overlap...' (p.4)
[Q128] 'a BPE-based tokenizer, which we train on the entire pre-training set, using a minimum frequency of two and a vocabulary size of 50,368.' (p.12)	[Q52] 'To convert the problem into sentence pair classification, we construct sentence pairs by replacing the ambiguous pronoun with each possible referent.' (p.4)
[Q129] 'a max sequence length of 1024.' (p.12)	[WEB-11] LuxBorrow corpus with token-level LID and code-switching annotation across LU/DE/FR/EN — confirms code-switching as systematic input characteristic
	[WEB-12] Neural Text Normalization for Luxembourgish — best ByT5 base achieves ~78.8% accuracy; documents 'large amount of orthographic variation'
	[WEB-15] spellux Python package documents homograph density and orthographic variation as ongoing Luxembourgish NLP challenges
	[WEB-13] LuxemBERT (2022) outperformed mBERT — confirms standardized English models are inadequate baselines for Luxembourgish
	[WEB-2] CTIE operates MyGuichet.lu under Ministry for Digitalisation

Output Ontology (OO)

Whether the benchmark's output categories and scoring criteria reflect regionally valid decision boundaries.

LTZGLUE — 1/5 (Serious concern) [high confidence]	GLUE — 1/5 (Serious concern) [high confidence]
Severe structural misalignment between benchmark output ontology and deployment requirements. All eight LTZGLUE tasks produce single hard labels [Q67] across narrow label spaces: binary headline acceptability [Q148], three-class sentiment [Q149], binary/four-class linguistic acceptability [Q33, Q151], BIO-tagged NER over PER/ORG/LOC/MISC/DATE [Q44, Q152], five-class topic [Q48, Q153], nine voice-assistant intents [Q154], and binary entailment [Q155]. The deployment explicitly requires multi-label classification with confidence scores for routing and human-fallback escalation [elicitation Q4], and demands administrative topic categories (cross-border worker, housing, transport, tax, communal vs. national) and urgency/frustration scales — none of which appear in any task. Sentiment's three-class scheme [Q149] cannot encode urgency or formal-but-understated frustration cues that are operationally significant per the elicitation [elicitation Q3] and regional cultural notes. EU AI Act considerations [WEB-12] and the language law's requirement to detect citizen language choice [WEB-14] introduce further required output dimensions (language-of-input, possibly risk flags) absent from the benchmark. OO is HIGH priority. This is the most severe dimension violation.	GLUE's output ontology is hard single-label classification (binary or three-class) plus one regression task [Q12], scored via accuracy/MCC/F1/Pearson [Q20, Q23, Q30, Q33, Q35] aggregated as a macro-average [Q59, Q60]. The deployment requires multi-label classification (a single citizen message may concern multiple departments simultaneously) plus confidence scoring to flag uncertain cases for human review (elicitation A4). The dataset analysis confirms 'every classification task in the sampled data uses exactly one label per example' and 'no example has multiple labels, soft labels, or confidence scores.' Additionally, the output categories themselves (grammatical/ungrammatical, positive/negative, entailment/neutral/contradiction) have no correspondence to administrative routing categories: no État-vs-commune label, no frontalier-vs-resident label, no department-routing label, no urgency flag. OO is HIGH priority.
LTZGLUE evidence	GLUE evidence
[Q33] 'The binary dataset distinguishes between correct (1) and incorrect (0)' (p.4)	[Q12] 'All tasks are single sentence or sentence pair classification, except STS-B, which is a regression task. MNLI has three classes; all other classification tasks have two.' (p.2)
[Q44] 'the tag set is harmonised by merging the GPE and LOC categories into a single location label, while retaining PER, ORG, and MISC unchanged.' (p.4)	
[Q67] 'Together, the eight tasks in LTZGLUE form a broad and balanced evaluation suite, covering four binary and four multi-class settings' (p.5)	
[Q149] 'sentiment_analysis: Classify sentiment of the text. Allowed labels: positive, neutral, negative.' (p.15)	
[Q152] 'ner: Perform Named Entity Recognition... Allowed Tags: O, B-LOC, I-LOC, B-PER, I-PER, B-DATE, I-DATE, B-ORG, I-ORG, B-MISC, I-MISC.' (p.15)	
[Q153] 'topic_classification: ... Allowed category_names: sports, animals, business, culture, technology.' (p.15)	
[Q154] 'slot_intent_detection: ... Allowed intents: reminder/show_reminders, weather/find, reminder/set_reminder, ...' (p.15)	
[WEB-12] EU AI Act Annex III Category 5(a) high-risk classification potentially relevant — implies need for documented decision rules and oversight outputs not in benchmark	

[WEB-14] Language law requires response in citizen's chosen language — implies language-of-input output dimension absent from benchmark

Output Content (OC)

Whether ground-truth labels align with the judgments of regional stakeholders.

LTZGLUE — 2/5 (Significant gaps) [high confidence]	GLUE — 1/5 (Serious concern) [high confidence]
Annotation provenance and quality are mixed and weakly tied to civil-administration norms. Sentiment was annotated by two native LTZ speakers with Cohen's Kappa 0.45 — moderate agreement — and disagreements resolved by consensus [Q25, Q26, Q27, Q28]; no documentation of administrative-register expertise. JUDGEWEL NER relies on LLM-as-judge with 'minimal human verification' [Q37]. RTE used ChatGPT-5-mini to verify labels, finding nearly 10% incorrect labels [Q63, Q64], with manual correction afterwards [Q65]. The intent detection translation lacked LTZ register references and could not be systematically verified [Q56]. Authors explicitly caution against 'using benchmark performance as evidence of cultural or demographic coverage' [Q121] and note 'models may reproduce dominant norms while under-representing regional, sociolectal, or multilingual practices' [Q120]. For a deployment where sentiment must capture formal-but-understated frustration in civil correspondence [elicitation Q3, regional context cultural_norms_notes], $\kappa=0.45$ with consumer-facing news comments as source material is insufficient. OC is HIGH priority.	Annotator demographics for the main inherited GLUE tasks are not documented in the paper [Q6 establishes the inherited-dataset reliance; no demographic disclosure for SST-2 [Q26] or MNLI crowdworkers [Q36]]. The diagnostic set's annotators are explicitly six NLP researchers [Q73] with high inter-annotator agreement (Fleiss's $\kappa = 0.73$ [Q74], $R3 = 0.80$ [Q75]) — but this is a specialized English-speaking population with no documented connection to Luxembourgish administrative culture or continental-European formal correspondence register. Ground-truth sentiment in SST-2 reflects English movie-review judgments (DATASET-D1 through D5, D29, D30); ground-truth entailment in MNLI/RTE reflects English crowdsourced judgments. The dataset analysis shows SST-2 fragments often labelled in ways that depend entirely on film-criticism context (DATASET-D2, D29, D30). Luxembourg's understated formal administrative register, where frustration is expressed indirectly (cultural_norms_notes), is not represented in any annotator pool. The user explicitly flagged OC as a 'live risk' (elicitation A3). Even ItzGLUE [WEB-10] and the only Luxembourgish sentiment corpus [WEB-25] do not cover administrative register — the latter sources from RTL.lu informal news comments.
LTZGLUE evidence	GLUE evidence
[Q25] 'we extract 4,583 sentences, which are then annotated by two native speakers of LTZ.' (p.3)	[Q6] 'None of the datasets in GLUE were created from scratch for the benchmark; we rely on preexisting datasets...' (p.1)
[Q27] 'We calculated Cohen's Kappa at 0.45.' (p.3)	[Q26] 'The Stanford Sentiment Treebank (Socher et al., 2013) consists of sentences from movie reviews and human annotations of their sentiment.' (p.3)
[Q28] 'For the final set, the annotators agree on a label in cases of label disagreement.' (p.3)	[Q36] 'The Multi-Genre Natural Language Inference Corpus (Williams et al., 2018) is a crowdsourced collection of sentence pairs with textual entailment annotations.' (p.4)
[Q37] 'The resulting sentences are then evaluated using LLMs acting as judges, with minimal human verification to calibrate quality thresholds.' (p.4)	[Q73] 'To establish human baseline performance on the diagnostic set, we have six NLP researchers annotate 50 sentence pairs (100 entailment examples) randomly sampled from the diagnostic set.' (p.6)
[Q56] 'Due to the lack of LTZ references in this register, it was not possible to systematically verify the translated terminology.' (p.5)	[Q74] 'Inter-annotator agreement is high, with a Fleiss's κ of 0.73.' (p.6)
[Q64] 'Nearly 10% of the labels were false.' (p.5)	[WEB-25] RTL Corpus Sentiment Annotation Tool (STRIPS) — only known Luxembourgish sentiment resource, sourced from informal RTL news comments rather than administrative register

[Q120] 'models may reproduce dominant norms while under-representing regional, sociolectal, or multilingual practices.' (p.10)	[WEB-10] ItzGLUE provides Luxembourgish sentiment (positive/neutral/negative) but not in administrative register
[Q121] 'We therefore caution against using benchmark performance as evidence of cultural or demographic coverage.' (p.10)	DATASET-D2: 'at its best moments' labelled positive — illustrates that SST-2 labels depend on film-critic register absent from administrative correspondence
	DATASET-D29: 'the horrors' labelled negative — context-dependent label inappropriate for administrative register
	DATASET-D30: 'heroes' labelled positive — single-word fragment with film-review context dependency

Output Form (OF)

Whether the expected output modality matches regional deployment needs and accessibility.

LTZGLUE — 3/5 (Moderate gaps) [medium confidence]	GLUE — 2/5 (Significant gaps) [high confidence]
Output form is text-classification labels evaluated by macro-F1 with results averaged over three runs and standard deviations reported [Q93, Q96]. This is a standard, well-documented evaluation methodology and matches the deployment's text-to-label format at surface level. However, the deployment requires confidence scores for routing thresholds and multi-label outputs for messages spanning multiple departments [elicitation Q4]. LTZGLUE reports only macro-F1 with no calibration metrics (e.g., ECE, Brier score), no per-class precision/recall thresholds, and no multi-label evaluation methodology. Invalid LLM outputs are simply discarded [Q90, Q91], which masks the operational reality that malformed outputs in deployment must be handled (escalated). OF is MODERATE priority. The form is text-to-text-label as required, but the metric set and label cardinality do not match the deployment's operational form.	GLUE evaluates text-based outputs only via hard single-label scoring on private test sets [Q57, Q58], with per-task scores aggregated as a macro-average [Q59, Q60] subject to two-submissions-per-day rate limiting [Q131]. The deployment is also text-in/label-out, so the surface-level format aligns — both are written-text pipelines, and OF is correctly marked MODERATE priority (lower than HIGH for the other dimensions). However, GLUE's output form lacks the multi-label structure and confidence scoring the deployment requires (elicitation A4), and the submission-based privately-held test data model [Q7, Q58] prevents the iterative, task-specific calibration a live civil-service deployment requires. The macro-average [Q60] also rewards balanced performance across nine irrelevant English NLU tasks, making the aggregate GLUE score uninformative as a deployment proxy. Score is 2 (not 1) because the underlying text-based output modality matches and the per-task evaluation infrastructure is methodologically reusable, but the specific output-form requirements (multi-label, confidence) are unsatisfied.
LTZGLUE evidence	GLUE evidence
[Q90] 'Prompted LLMs do not always produce well-formed outputs and may return an incorrect number of predictions for a given task; such outputs are discarded prior to evaluation.' (p.7)	[Q7] 'Four of the datasets feature privately-held test data, which will be used to ensure that the benchmark is used fairly.' (p.1)
[Q93] 'Table 6 shows F1 scores for all models across all tasks' (p.7)	[Q57] 'The GLUE benchmark follows the same evaluation model as SemEval and Kaggle.' (p.5)
[Q96] 'Encoder results are averaged over three runs with standard deviations as subscripts. Prompted LLMs were evaluated once; we report macro-F1 only.' (p.8)	[Q58] 'To evaluate a system on the benchmark, one must run the system on the provided test data for the tasks, then upload the results to the website gluebenchmark.com for scoring.' (p.5)
	[Q59] 'The benchmark site shows per-task scores and a macro-average of those scores to determine a system's position on the leaderboard.' (p.5)

	[Q60] 'For tasks with multiple metrics (e.g., accuracy and F1), we use an unweighted average of the metrics as the score for the task when computing the overall macro-average.' (p.5)
	[Q131] 'The GLUE website limits users to two submissions per day in order to avoid overfitting to the private test data.' (p.14)