

Output Form (OF)

Score: 4/5 — Minor gaps | Confidence: high

Benchmark: LTZGLUE | Context: Luxembourgish Professional Newsroom

Definition: Whether the expected output modality matches regional deployment needs and accessibility.

Justification

Output form aligns well with deployment needs (LOWER priority dimension). All tasks produce discrete label outputs evaluated by macro-F1 [Q93, Q96], directly compatible with the deployment's per-task significance-score-and-threshold pipeline [A4]. Encoder results are averaged over three runs with reported standard deviations [Q89, Q96], enabling stability assessment. The label-output (not MCQA) paradigm [Q84] mirrors the deployment's classification framing. Class-imbalanced tasks use class-balanced loss [Q92]. Caveats: LLM outputs failing label format are discarded [Q90], inflating reported scores; LLMs are evaluated only once with macro-F1 only [Q96] — meaning LLM performance is less stable, and the benchmark itself flags this as 'indicative rather than directly comparable' [Q86]. The deployment's threshold-flagging logic [A4] requires confidence scores per task, which encoder fine-tuning supports natively but prompted LLM evaluation may not.

Strengths

- Discrete label outputs evaluated by macro-F1 [Q93, Q96] directly support the deployment's per-task threshold flagging [A4]
- Encoder results averaged over three runs with standard deviations [Q89, Q96] permit stability assessment for production deployment
- Class-balanced loss applied to imbalanced tasks (LA, SA) [Q92] addresses real distributional skew
- Label-output paradigm rather than MCQA [Q84] matches the deployment's classification framing

Checklist

Checklist item definitions:

ID	Assessment Question
OF-1	Does output modality match regional deployment needs?
OF-2	Assess text-to-speech availability for speech output
OF-3	Consider literacy rates and accessibility requirements
OF-4	Document form mismatches harming external validity

Checklist responses:

OF-1: Discrete label outputs [Q93, Q148, Q149, Q150, Q151, Q152, Q153] match the deployment's per-task scoring needs [A4]; the deployment's threshold flagging is a downstream wrapper compatible with macro-F1 evaluation.

OF-2: Not applicable — deployment is text-only [deployment_modality]; speech output is not required.

OF-3: Target users are professional journalists/editors with high literacy and trilingual proficiency [WEB-14, WEB-15]; no accessibility constraints introduced by the output form.

ID	Evidence
WEB-14	Luxembourg trilingual education system produces newsroom professionals with high literacy in Luxembourgish/French/German plus English
WEB-15	EU Education and Training Monitor confirms strong language education context for target users

OF-4: Minor mismatches: (a) prompted LLM scores are single-run macro-F1 [Q96] less directly comparable to encoder runs [Q86]; (b) discarded malformed LLM outputs [Q90] inflate apparent performance — relevant if the deployment uses prompted

LLMs rather than fine-tuned encoders.

Information Gaps

- Whether confidence scores from fine-tuned encoders are well-calibrated for the deployment's threshold-based flagging logic

Requires Expert Verification

- Calibration of model confidence outputs against editorial flagging thresholds in production