

Comparative Assessment: AFRIDOC_MT vs. NLLB

Deployment context: Amhara Region Agricultural & Microfinance Document Translation — Cooperative and Rural MFI End-Users

Deployment Context

Use case: Document-level translation of agricultural input subsidy notices, land-use policy updates, and credit-scheme terms from federal and regional bureaus into Amharic

Target population: Farmer cooperative leaders and rural microfinance clients (e.g., ACSI, ADCSI borrowers) in Amhara interpreting government and lender document

Overall Risk Ratings

Benchmark	Risk	Avg. Score
AFRIDOC_MT	HIGH	2.3 / 5
NLLB	HIGH	2.7 / 5

Dimension Scores Comparison

Dimension	AFRIDOC_MT	Rating	NLLB	Rating	Δ
Input Ontology (IO)	2/5	Significant gaps	2/5	Significant gaps	0
Input Content (IC)	1/5	Serious concern	2/5	Significant gaps	-1
Input Form (IF)	4/5	Minor gaps	4/5	Minor gaps	0
Output Ontology (OO)	2/5	Significant gaps	2/5	Significant gaps	0
Output Content (OC)	2/5	Significant gaps	2/5	Significant gaps	0
Output Form (OF)	3/5	Moderate gaps	4/5	Minor gaps	-1
Average	2.3		2.7		-0.3

Δ = AFRIDOC_MT score – NLLB score. Positive = regional benchmark scores higher.

Overall Summaries

AFRIDOC_MT (risk: HIGH)

AFRIDOC-MT is a high-quality document-level MT benchmark for five African languages including Amharic in Ethiopic script, but its fitness for evaluating Amharic agricultural and microfinance document translation in Amhara region is severely limited by three convergent gaps: (1) total domain absence — no agricultural, land-tenure, cooperative governance, or credit-scheme content exists in the benchmark, confirmed empirically across all sampled examples and corroborated by web research showing no other Amharic MT benchmark covers these domains either; (2) source-language direction mismatch — the benchmark translates from English to Amharic and exhibits acknowledged translationese, whereas the deployment operates on documents natively authored in Amharic proclamation register; and (3) annotator-authority mismatch — the four named Amharic translators were prepared for health/IT domains with no documented agricultural or legal Amharic expertise, while the user-specified preferred ground-truth authority is federal bureau translators. Output scoring is calibrated for narrative discourse coherence with no metric for terminological consistency — the critical coherence requirement for legal-administrative documents. The benchmark's strongest alignment is on Input Form (text-only Amharic in Ethiopic script with documented tokenization accommodations) and Output Form (text output modality match). Of the three HIGH-priority dimensions (IO, IC, OC), all score 1-2.

NLLB (risk: HIGH)

Flores-200/NLLB provides credible language- and script-level support for Amharic — Ge'ez script is included with a dedicated SentencePiece vocabulary [Q382], and the benchmark's text-only output form aligns with the deployment's accepted format [elicitation A4]. These are real strengths for this text-only Amharic deployment. However, the four HIGH-priority dimensions (Input Ontology, Input Content, Output Ontology, Output Content) show substantial validity gaps. The benchmark's task taxonomy excludes agricultural, land-tenure, and microfinance genres entirely [Q95, Q166-Q169]; its content is exclusively Wikimedia-derived [Q95, Q151] with the authors themselves acknowledging cultural-coverage limitations [Q151]; its output ontology offers only general translation-quality scalars [Q455, Q648] that cannot detect the high-consequence error types (numerical-threshold alteration, conditional-logic distortion, terminology drift) that matter for binding administrative documents; and its annotator workforce is professional generalist translators [Q97-Q103], not the federal bureau translators with Ethiopian government-domain expertise that the deployer identifies as the validation authority [elicitation A3]. Independent 2025 research [WEB-10] empirically confirms that Flores+ scores fail to predict out-of-domain translation quality and that chrF++/BLEU can reward superficial named-entity copying. Web search found no Ethiopian agricultural/administrative MT benchmark exists [WEB-3, WEB-4] and no public ACSI/ADCSI Amharic terminology resource [WEB-15, WEB-16], meaning the gaps are real and unaddressed by the broader research community. Aggregate NLLB scores on Flores-200 should be treated as a coarse general-domain capability signal for Amharic, not as evidence of fitness for binding subsidy, land-policy, or credit documents.

Per-Dimension Justification Comparison

Each benchmark's full assessment justification and cited evidence, shown side by side.

Input Ontology (IO)

Whether the benchmark's test case categories cover the query types expected in deployment.

AFRIDOC_MT — 2/5 (Significant gaps) [high confidence]	NLLB — 2/5 (Significant gaps) [high confidence]
The benchmark's task taxonomy (document-level and sentence-level MT between English and five African languages, including Amharic) is structurally relevant — the deployment also needs document-level Amharic MT, and the user agreed that document-level structural demands are broadly similar across genres. However, the test-case category distribution covers only health (WHO) and IT news (Techpoint Africa) genres [Q1, Q2, Q80], with zero coverage of agricultural input subsidy notices, land-use policy updates, credit-scheme terms, or cooperative governance documents — the four genres specified for the deployment. The only legally-binding-text instance in the entire dataset is a translation of the WHO International Health Regulations [DATASET-D17], which does not approach the register of Ethiopian federal proclamations. Web research confirms this is a sector-wide gap: no Amharic MT benchmark covers Ethiopian government proclamation, agricultural subsidy, cooperative governance, or legal-administrative domains [WEB-14, WEB-22, WEB-15]. Given the user's HIGH priority weight on IO and the total absence of deployment-relevant test categories, the score is low — but not 1, because the abstract task type (document-level Amharic MT) is correctly represented and supports both sentence-level and pseudo-document subtasks at multiple chunk sizes [Q3, Q35, Q49] that broadly match deployment document length variation.	The benchmark's task taxonomy is organized along resource level (high/low/very low) and a narrow set of domains (news, scripted talks, chat, WHO health) tested only in NLLB-MD for six language directions [Q159, Q166-Q169], none of which include Amharic. The deployment's core document genres — agricultural input subsidy notices, land-use policy updates, and credit-scheme terms — fall entirely outside the evaluation taxonomy. The paper itself frames the generalization question as 'How well do models generalize to non-Wikimedia domains?' [Q159] but does not extend domain testing to administrative, agricultural, or financial genres. Independent 2025 evidence confirms Flores+ scores do not reliably predict out-of-domain translation quality [WEB-10]. Given the HIGH priority assigned to this dimension by the elicitation, the absence of any agricultural/bureaucratic/financial category in the benchmark taxonomy is a significant content-validity concern.
AFRIDOC_MT evidence	NLLB evidence

[Q1] 'This paper introduces AFRIDOC-MT, a document-level multi-parallel translation dataset covering English and five African languages: Amharic, Hausa, Swahili, Yorùbá, and Zulu.' (p.1)	[Q159] 'How well do models generalize to non-Wikimedia domains?' (p.24)
[Q2] 'The dataset comprises 334 health and 271 information technology news documents, all human-translated from English to these languages.' (p.1)	[Q166] 'We translate the English side of the WMT21 English-German development set' (p.25)
[Q80] 'We introduce AFRIDOC-MT, a document-level translation dataset in the health and tech domains for five African languages.' (p.9)	[Q169] 'one World Health Organisation report... combined with sentences translated from the English portion of the TAUS Corona Crisis Report' (p.25)
[Q3] 'We conduct document-level translation benchmark experiments by evaluating the ability of neural machine translation (NMT) models and large language models (LLMs) to translate between English and these languages, at both the sentence and pseudo-document levels' (p.1)	[WEB-10] 2025 'Languages Still Left Behind' study: NLLB models fine-tuned on naturalistic domain data outperformed Flores+ on domain test sets but underperformed on Flores+ — direct evidence that Flores+ scores do not predict domain-specific quality
[WEB-14] EthioNLP survey confirms no Amharic MT benchmark covers legal-administrative domains	
[WEB-22] EthioMT parallel corpus is general-domain only	
[WEB-15] EthioLLM/EthioBenchmark covers classification, NER, sentiment, hate speech — not legal-domain MT	
DATASET-D17 : International Health Regulations is the closest legal-register item but is WHO international law, not Ethiopian proclamation	

Input Content (IC)

Whether individual datapoint content is culturally and contextually appropriate for the target region.

AFRIDOC_MT — 1/5 (Serious concern) [high confidence]	NLLB — 2/5 (Significant gaps) [high confidence]
The input content gap is total. All English source documents are scraped from Techpoint Africa and WHO [Q18, Q22], and the Amharic translations were produced from this English source material with explicit acknowledgment of translationese risk [Q103, Q104, Q105]. The deployment requires Amharic content drawn from Ethiopian federal and regional proclamations, agricultural subsidy notices, land-tenure documents, MFI loan agreements, and cooperative governance materials — none of which appear anywhere in the corpus. Critical proclamation-derived term categories (cooperative tier names, land-tenure category labels, subsidy scheme names, MFI credit-scheme terms) are entirely absent, confirmed empirically across all sampled examples [DATASET-D9, DATASET-D4]. The source-language direction is also mismatched: the benchmark translates FROM English [Q103], whereas the deployment operates on documents natively authored in Amharic administrative register [WEB-10]. No public Amharic legal-administrative glossary exists [WEB-14, WEB-15] to serve as an external terminology anchor. Given the user's HIGH priority weight on IC and the empirical confirmation of total content absence, this dimension is at the floor.	Flores-200 and NLLB-Seed inputs are exclusively sampled from English-language Wikimedia projects [Q95, Q146], and the paper explicitly acknowledges 'the content reflects what Wikipedia editors find is relevant for English Wikipedia, and likely does not cover diverse content from different cultures' [Q151]. Low-resource bitext for African languages is 'often opportunistically drawn from known collections of multilingual content, such as the Christian Bible or publications of multinational organizations' [Q284], 'often limited in quantity and domain' [Q285]. The deployment requires Ethiopian-specific land-tenure proclamation vocabulary, ACSI/ADCSI financial terminology, and regional cooperative administrative phrasing — none of which is documented as present in NLLB training or evaluation data. The Hugging Face NLLB model card confirms 'Our model has been tested on the Wikimedia domain with limited investigation on other domains' [WEB-11]. Web search found no public ACSI/ADCSI Amharic terminology resources [WEB-15, WEB-16], confirming the absence of any reference content for the deployment's core terminology. Given HIGH priority and binding-document consequences, the content gap is severe.

AFRIDOC_MT evidence	NLLB evidence
[Q18] 'We scraped English articles from the websites of Techpoint Africa and the World Health Organization (WHO).' (p.3)	[Q95] '3001 sentences sampled from English-language Wikimedia projects... Wikinews, Wikijunior, and Wikivoyage' (p.19)
[Q22] 'we selected 334 and 271 English articles/documents from the health and tech domains respectively, which represents 10k sentences each per domain.' (p.3)	[Q146] 'Wikimedia's List of articles every Wikipedia should have... 11 categories such as People, History, Philosophy and Religion, Geography' (p.23)
[Q103] 'A potential limitation of our dataset is the influence of translationese... Since all source material translated originates in English, translated sentences in African languages may exhibit patterns such as unnatural syntax or overly literal phrasing.' (p.10)	[Q151] 'the content reflects what Wikipedia editors find is relevant for English Wikipedia, and likely does not cover diverse content from different cultures' (p.24)
[Q105] 'AFRIDOC-MT may reflect a bias toward English in terms of structure, semantics, and cultural framing.' (p.10)	[Q284] 'often opportunistically drawn from known collections of multilingual content, such as the Christian Bible or publications of multinational organizations' (p.37)
[WEB-10] Ministry of Justice provides authoritative Amharic terminology for cooperative tiers	[Q285] 'often limited in quantity and domain' (p.37)
[WEB-14] No public Amharic legal-administrative glossary exists	[WEB-11] NLLB Hugging Face model card: 'Our model has been tested on the Wikimedia domain with limited investigation on other domains'
[WEB-15] EthioLLM resources do not include legal-administrative MT data	[WEB-15] ACSI official website provides institutional background only — no Amharic terminology resources publicly available
DATASET-D9 : Crypto/fintech vocabulary (Nestcoin, FTX, pre-seed) — zero overlap with cooperative or subsidy terminology	[WEB-16] FinDev Gateway ACSI profile — no terminology guide
DATASET-D4 : Antimicrobial medical terminology — zero overlap with land-tenure or cooperative tier vocabulary	

Input Form (IF)

Whether the input signal encoding (text, audio, image parameters) matches deployment conditions.

AFRIDOC_MT — 4/5 (Minor gaps) [high confidence]	NLLB — 4/5 (Minor gaps) [high confidence]
Input form is the most aligned dimension and the user assigned it LOWER priority. Both benchmark and deployment are text-only, and Amharic is explicitly handled in Ethiopic (Ge'ez) script with documented accommodations: token limits were specifically increased for Amharic [Q162] because its non-Latin script produced larger token counts than English [Q127]. The benchmark adopts pseudo-document chunking with k=10 to fit Amharic within model context limits [Q47, Q48, Q128], a practical choice that the deployment will likely also need. Empirically, all sampled Amharic content is rendered consistently in Ethiopic script across all configurations [DATASET-D1, DATASET-D5, Strength 1]. The remaining concern is that the deployment may evolve toward format-preserving or multimodal translation (scanned PDFs, structured documents with tables and signature blocks) — a possibility flagged but unconfirmed in the regional context. Some annotation-workflow infrastructure (CSV/Google Sheets with paragraph-marking empty rows) [Q119, Q122] partially preserves document structure at the corpus level. Score of 4 reflects strong technical alignment on script and tokenization, with a minor gap on structural fidelity that becomes relevant only if deployment expands to format-preserving use cases.	The benchmark's input form is well-aligned with the deployment at the modality and script levels. All inputs are text-only encoded via BCP 47 with ISO 639-3 + ISO 15924 script subtags [Q79], the Ethiopic (Ge'ez) script is explicitly supported with a dedicated 8k SentencePiece vocabulary [Q382], and the deployer confirmed text-only prose input is acceptable [elicitation A4]. Average sentence length (~21 words) [Q133] reflects general-domain prose. The deployment is HIGH on text-only operation, and the benchmark matches. Residual concerns are: (a) sentence-level evaluation architecture does not assess document-level structures (numbered clauses, conditional logic), and (b) Flores-200 sentences are general-domain prose, whereas administrative source documents have higher punctuation density and clause structure — though deployer accepted prose output [elicitation A4]. IF was assigned LOWER priority by elicitation, consistent with this score.
AFRIDOC_MT evidence	NLLB evidence
[Q47] 'An initial analysis showed that some models were unable to process entire documents due to input length limits, which were exceeded by token counts in some languages (Amharic and Yorùbá).' (p.5)	[Q79] 'BCP 47 tag sequence using a three-letter ISO 639-3 code as the base subtag, which we complement with ISO 15924 script subtags' (p.17)
[Q127] 'Amharic and Yorùbá—languages with unique characteristics such as non-Latin scripts and diacritics, respectively—had the largest average token counts across the tokenizers.' (p.19)	[Q133] 'On average, sentences are approximately 21 words long' (p.22)
[Q128] 'To accommodate both languages in our experiments, we chose pseudo-documents with k=10.' (p.19)	[Q382] 'a specific 8k SentencePiece vocabulary to better support the Ge'ez script' (p.45)
[Q162] 'Therefore, we increased the token limit specifically for Amharic.' (p.21)	
[Q119] 'Please do not delete double empty rows, as they serve to separate paragraphs.' (p.18)	
DATASET-D1: Amharic Ethiopic-script document rendered consistently across multi-paragraph health content	
DATASET-D5: k=10 pseudo-document chunk in Amharic confirms working configuration	

Output Ontology (OO)

Whether the benchmark's output categories and scoring criteria reflect regionally valid decision boundaries.

AFRIDOC_MT — 2/5 (Significant gaps) [high confidence]	NLLB — 2/5 (Significant gaps) [high confidence]
The output ontology is calibrated for news discourse: GPT-4o-judged fluency (1–5), content errors, lexical cohesion errors, and grammatical cohesion errors [Q54, Q55, Q180, Q191], plus d-chrF and d-BLEU [Q51, Q52] and human direct assessment on a 0–100 scale [Q195]. The benchmark explicitly defines lexical cohesion as 'vocabulary usage, incorrect or missing synonyms' [Q189] — narrative discourse flow, not legal terminological consistency. No rubric component evaluates whether named entities such as cooperative tier designations, land-tenure category labels, or subsidy scheme names are rendered identically across all clauses of a document — the critical coherence requirement for the deployment. Web research confirms terminology-consistency metrics exist (Semenov & Bojar WMT 2022; HHI) but have not been validated for Amharic [WEB-25, WEB-26]. Compounding this, the benchmark itself reports that GPT-4o-as-judge is unreliable for African language translation directions [Q67, Q85, Q99, Q219], leaving human DA as the residual signal — which had Krippendorff's alpha of only 0.46 for Amharic [Q71, Q211]. Given the user's MODERATE priority on OO and the partial match (d-chrF captures morphological richness [Q53], cohesion definitions touch on relevant phenomena), score is 2 rather than 1.	The benchmark's output scoring ontology is a uniform scalar quality score (chrF++, spBLEU, or calibrated XSTS) applied to all directions and sentence types [Q455, Q648, Q661]. There is no domain-specific rubric for correctly rendering numerical eligibility conditions, conditional legal clauses, or financial terminology — the most consequential error types for the deployment. The XSTS scale measures 'meaning preservation' on a 5-point scale [Q654] but treats sentences uniformly. For toxicity, the taxonomy distinguishes source-present from added toxicity [Q705, Q706] with categories covering vulgar/profane/pornographic/hate-speech/bullying [Q712–Q714]; the paper acknowledges this is 'culturally sensitive' [Q715] with no Ethiopian regional calibration. A chrF++ improvement of +0.5 corresponds to detectable human improvement [Q624, Q631] but cannot distinguish silently altered subsidy thresholds from generally good prose. Independent 2025 evidence [WEB-10] shows mock translations reproducing named entities scored non-trivially on BLEU/ChrF++, confirming these metrics reward superficial overlap. Given OO is HIGH priority for binding documents, this is a structural-validity violation.
AFRIDOC_MT evidence	NLLB evidence
[Q53] 'we report d-chrF scores... as chrF better captures the morphological richness of African languages' (p.5)	[Q455] 'We use the chrF++ metric to compare the model performance' (p.53)
[Q55] 'we assess each translated document's fluency, content errors (CE), and cohesion errors—specifically lexical (LE) and grammatical (GE) errors—using GPT-4o' (p.5)	[Q648] 'we primarily evaluate using chrF++ using the settings from sacrebleu' (p.73)
[Q189] 'Lexical Cohesion Mistakes: Issues with vocabulary usage, incorrect or missing synonyms, or overuse of certain words that disrupt the flow.' (p.23)	[Q654] 'XSTS rates each source sentence and its machine translation on a five-point scale' (p.74)
[Q67] 'These inconsistencies raise concerns about GPT-4o's reliability.' (p.7)	[Q661] 'we take the majority XSTS score for each sentence and average these majority scores over all evaluated sentences' (p.74)
[Q99] 'GPT-4o was employed to assess and rate the translation outputs... While its ratings were consistent for translations into English, the same was not observed for translations into African languages, likely due to the model's limited understanding of these languages.' (p.10)	[Q624] 'chrF++ improvements of +0.5 chrF++ usually correlate to statistically significant human evaluation improvements' (p.70)
[Q219] 'These results suggest that using GPT-4o as a translation judge is not yet well-suited for low-resource languages.' (p.26)	[Q705] 'In the context of translation, toxicity may originally be present in the source text or it can be generated de-novo in the target text (added toxicity)' (p.79)
[WEB-25] Semenov & Bojar (WMT 2022) terminology-consistency metric exists but not validated for African languages	[Q715] 'Toxicity is culturally sensitive' (p.80)

[WEB-26] HHI applied to legal corpora terminology consistency — not validated for Amharic

[WEB-10] 2025 study: mock translations merely reproducing named entities achieved non-trivial BLEU and ChrF++ scores, confirming metrics can reward superficial overlap rather than translational adequacy

Output Content (OC)

Whether ground-truth labels align with the judgments of regional stakeholders.

AFRIDOC_MT — 2/5 (Significant gaps) [high confidence]	NLLB — 2/5 (Significant gaps) [high confidence]
Reference Amharic translations were produced by four named native-speaker translators [Q23, Q109] under a controlled workflow with QE filtering at threshold 0.65 [Q28, Q29] and language-coordinator manual inspection [Q27]. However, the user's preferred ground-truth authority is domain-trained federal bureau translators [elicitation Q3], with regional bureau staff or extension workers as acceptable surrogates — and no documentation of institutional affiliation or agricultural/legal domain expertise exists for the four Amharic translators [Q109; gap_3 web search NOT FOUND]. The translators were prepared for health/IT domains via a terminology workshop [Q26], not for proclamation-derived agricultural or financial vocabulary. Inter-annotator agreement for Amharic was Krippendorff's alpha = 0.46 [Q71, Q211], described by the authors as 'relatively low' [Q71], even within the health domain in which translators were specifically prepared. Reference translations also exhibit translationese due to English source provenance [Q103, Q104]. Human evaluation used three native-speaker annotators per language [Q194], primarily drawn from the translation team [Q194], 80 documents each with 30 shared [Q197]. No demographics on annotator regional origin (Amhara region vs. other Amharic-speaking areas) are documented. Given the user's HIGH priority weight on OC and the multiple convergent gaps (annotator domain mismatch, low IAA, translationese, no federal-bureau-translator anchor), score is 2.	Flores-200 reference translations were produced through a multi-step workflow [Q93] with translators requiring 2-5 years of experience [Q97], 18-month re-testing [Q98], and reviewer native-speaker status with typically 5+ years of experience [Q99, Q100]. For low-resource languages, qualification thresholds were relaxed to accept broader linguistics degrees or extensive commercial experience [Q103]. No Amharic-specific annotator provenance, domain expertise, or inter-annotator agreement figures are documented in the registry. The community interview study informing research priorities included only 12 African-language speakers [Q31], a majority being immigrants in US/Europe with ~1/3 tech workers [Q32], whose perspectives 'may not consummately capture the sentiments of their communities back home' [Q39]. The deployer's required validation authority — federal bureau translators with Ethiopian agricultural/financial domain expertise [elicitation A3] — is not represented in the Flores-200 annotator workforce. Web search confirms post-publication corrections exist for Hausa/Sepedi/Xitsonga/isiZulu Flores-200 splits [WEB-14] but not Amharic, leaving Amharic reference quality unaudited. Common errors documented include 'mistranslation and unnatural translation errors' from non-English language influences and 'lower levels of standardization' [Q136-Q138]. Given OC is HIGH priority, this is a substantial concern.
AFRIDOC_MT evidence	NLLB evidence
[Q23] 'We translated the extracted 10k English sentences to the 5 African languages through 4 expert translators per language.' (p.3)	[Q93] 'professional translators and reviewers aligned on language standards... If the quality assessment indicated that the quality is above 90 percent' (p.19)
[Q26] 'we conducted a translation workshop, during which three translation experts shared their experiences in creating terminologies' (p.3)	[Q97] 'Translators are required to have at least two to three years of translation experience' (p.20)
[Q71] 'We obtained Krippendorff's alpha values of ≥ 0.40, which are relatively low due to the fine granularity of the evaluation scale.' (p.8)	[Q99] 'Flores-200 reviewers are also required to be native speakers of the target language' (p.20)
[Q103] 'translated sentences in African languages may exhibit patterns such as unnatural syntax or overly literal phrasing.' (p.10)	[Q103] 'we modified the qualification process to accept applications from reviewers with more general language degrees' (p.20)
[Q109] 'Amharic: Bereket Tilahun, Hana M. Tamiru, Biniam Asmlash, Lidetewold Kebede' (p.11)	[Q31] '12 in Africa' (p.8) — community study's African-speaker count

[Q194] 'three native speakers of the African languages—primarily translators involved in the dataset creation—were recruited.' (p.24)	[Q32] 'majority of our participants are immigrants living in the U.S. and Europe, and about a third of them (n = 17) identify as tech workers' (p.8)
[Q211] 'We obtained a scores of 0.46, 0.57, 0.40, and 0.81, and 0.54 for Amharic, Hausa, Swahili, Yorùbá, and Zulu respectively.' (p.26)	[Q39] 'their perspectives may not consummately capture the sentiments of their communities back home' (p.8)
	[Q136] 'Mistranslation and unnatural translation errors were still the most common errors' (p.22)
	[WEB-14] Post-publication FLORES-200 corrections exist for Hausa, Sepedi, Xitsonga, isiZulu — but not Amharic, leaving Amharic reference unaudited

Output Form (OF)

Whether the expected output modality matches regional deployment needs and accessibility.

AFRIDOC_MT — 3/5 (Moderate gaps) [medium confidence]	NLLB — 4/5 (Minor gaps) [high confidence]
Output form is text-only translation [Q8, Q44, Q201], matching the deployment's text-only requirement. Document-level outputs are produced by realigning sentence-level or pseudo-document translations into full documents [Q51, Q57], partially preserving paragraph structure via maintained empty rows [Q119]. However, no metric assesses structural fidelity (preservation of numbered clauses, tables, signature blocks) [regional context: structure_preservation_gap]. The benchmark also reports specific output-form pathologies for African-language translations: under-generation, over-generation, repetition of words and phrases, off-target translations [Q6], and inconsistent diacritics [Q77] — these failure modes would be especially harmful in a legal-administrative deployment where exact term repetition matters. Embedding-based document-level evaluation for African languages remains unavailable [Q50, Q96, Q175, Q176]. Greedy decoding is used [Q155]. Given the user's MODERATE priority on OF and the partial match (text-only ✓; structural fidelity gap; no plain-language accessibility scoring for variable-literacy end-users), score is 3.	The benchmark evaluates text outputs only [Q5] using sentence-level chrF++ and spBLEU [Q455, Q645, Q648] plus 5-point XSTS [Q654]. The deployer confirmed text-only prose output is acceptable for this NLLB deployment [elicitation A4], and this matches the benchmark's output form. However, no document-level coherence, terminology-consistency, or structural-preservation metrics are reported, all evaluation is sentence-level on general-domain content [Q5, Q336], and aggregate sentence-level chrF++ improvements may not indicate fitness for binding administrative documents requiring consistent rendering across clauses. National Ethiopia adult literacy is ~51.8% [WEB-5] with rural Amhara likely lower, meaning translated documents are often relayed orally through cooperative leaders — a use case that further increases the importance of consistent terminology, which sentence-level scoring does not assess. OF was rated LOWER priority by elicitation; given form alignment at the modality level, a moderately high score is warranted, with the residual concern being absence of document-level scoring.
AFRIDOC_MT evidence	NLLB evidence
[Q6] 'some LLMs exhibit issues such as under-generation, over-generation, repetition of words and phrases, and off-target translations, specifically for translation into African languages.' (p.1)	[Q5] 'evaluated the performance of over 40,000 different translation directions using a human-translated benchmark, Flores-200' (p.1)
[Q8] 'We evaluate performance using automatic metrics and compare the results of encoder-decoder models with decoder-only LLMs' (p.1)	[Q336] 'report detokenized BLEU on Flores-200 devtest' (p.41)
[Q44] 'Given that our created dataset can be used for sentence-level translation and as a baseline for document-level translation, we evaluate all models on the test splits for each domain.' (p.5)	[Q455] 'We use the chrF++ metric to compare the model performance' (p.53)

[Q51] 'we realigned sentence-level or pseudo-translation outputs into full documents, then computed BLEU... and chrF... to create document BLEU (d-BLEU) and document chrF (d-chrF).' (p.5)	[Q645] 'spBLEU, a BLEU metric based on a standardized SentencePiece model (SPM) covering 101 languages' (p.73)
[Q77] 'for Yorùbá, it often uses inconsistent or partial diacritics, resulting in inaccuracies.' (p.8)	[Q648] 'we primarily evaluate using chrF++ using the settings from sacrebleu' (p.73)
[Q119] 'Please do not delete double empty rows, as they serve to separate paragraphs.' (p.18)	[WEB-5] Ethiopia adult literacy ~51.8% (UNESCO via World Bank, 2017) — establishes oral-relay relevance for output use
[Q175] 'we used BLEU and chrF scores but excluded AfriCOMET due to its backbone model... having a context length of 512 tokens.' (p.22)	