

Comparative Assessment: LAILA vs. ASAP++

Deployment context: Arab High-School Arabic Essay Feedback (Qatar Primary)

Deployment Context

Use case: Support high school students (G10-12) in editing and revising their Arabic essay drafts before submitting them for grading.

Target population: high school students (G10-12) in Arab countries (especially Qatar)

Overall Risk Ratings

Benchmark	Risk	Avg. Score
LAILA	HIGH	2.2 / 5
ASAP++	HIGH	1.0 / 5

Dimension Scores Comparison

Dimension	LAILA	Rating	ASAP++	Rating	Δ
Input Ontology (IO)	2/5	Significant gaps	1/5	Serious concern	+1
Input Content (IC)	2/5	Significant gaps	1/5	Serious concern	+1
Input Form (IF)	4/5	Minor gaps	1/5	Serious concern	+3
Output Ontology (OO)	1/5	Serious concern	1/5	Serious concern	0
Output Content (OC)	3/5	Moderate gaps	1/5	Serious concern	+2
Output Form (OF)	1/5	Serious concern	1/5	Serious concern	0
Average	2.2		1.0		+1.2

Δ = LAILA score – ASAP++ score. Positive = regional benchmark scores higher.

Overall Summaries

LAILA (risk: HIGH)

LAILA is a high-quality Arabic AES dataset that establishes substantial within-Qatar validity for numeric trait scoring on persuasive and explanatory essays, but it is structurally misaligned with the proposed deployment along three of the four dimensions marked HIGH priority. (1) Output Form (1/5) and Output Ontology (1/5) are direct mismatches: the benchmark validates QWK agreement on ordinal scores while the deployment requires open-ended natural-language revision suggestions — an output category absent from the benchmark, and a field-wide gap in Arabic AES. (2) Input Ontology (2/5) covers only 2 of 5–7 deployment-required genres; literary analysis, religious/Quranic commentary, and cultural commentary are entirely absent and the CAST rubric is explicitly persuasive/argumentative. (3) Input Content (2/5) is single-country (Qatar) with explicit author acknowledgement of limited generalizability; cross-national rubric and curriculum alignment with Egypt, Saudi Arabia, Jordan, and Lebanon is empirically unstudied. Output Content (3/5) is strong within Qatar (qualified annotators, substantial IAA) but unvalidated cross-nationally. Input Form (4/5) aligns well as both contexts are digital Arabic text. The authors themselves prohibit operational deployment without independent validation [Q104, Q106].

ASAP++ (risk: HIGH)

ASAP++ is severely misaligned with the target deployment across all six validity dimensions, all of which were flagged HIGH priority in the elicitation. The benchmark provides English-only essays from native English-speaking US Grade 7–10 students, scored on

US-pedagogy-derived genres (argumentative, source-dependent, narrative) by India-based English-MA annotators anchored to US holistic scores, with numeric ordinal output evaluated by QWK. The deployment requires Arabic (MSA) Grade 10–12 essays under Qatar MOEHE and other Arab national curricula, covering literary analysis, religious-text commentary, and cultural commentary genres absent from the benchmark, evaluated against Arab curriculum stakeholders, and producing natural-language revision suggestions with explanatory rationale. There is no overlap in language, script, age band, curriculum, annotator expertise, genre coverage, or output modality. ASAP++ is at best a methodological reference (multi-trait scoring structure, QWK reliability template) for a future Arabic AES effort; it cannot serve as a validity anchor for the deployment. More directly relevant Arabic resources — LAILA (7,859 high school essays, seven-trait rubric, QRDI-funded), QAES/QCAW (195 Qatari argumentative essays), and Arwi (Arabic feedback-generation system) — should replace ASAP++ as the primary references.

Per-Dimension Justification Comparison

Each benchmark's full assessment justification and cited evidence, shown side by side.

Input Ontology (IO)

Whether the benchmark's test case categories cover the query types expected in deployment.

LAILA — 2/5 (Significant gaps) [high confidence]	ASAP++ — 1/5 (Serious concern) [high confidence]
The benchmark's task taxonomy is restricted to 8 prompts spanning only two genres — explanatory and persuasive [Q38] — with the CAST rubric explicitly designed for persuasive/argumentative writing [Q118]. The deployment requires coverage of literary analysis, religious/Quranic commentary, and cultural commentary genres common in Arab high-school curricula, which the authors explicitly acknowledge as out of scope: 'limiting genre diversity and potentially affecting model robustness across other styles, such as narrative or descriptive writing' [Q89]. Web research confirms this is a field-wide gap: TAQEEM 2025 and ZaQQ 2025 maintain the same explanatory/persuasive restriction [WEB-10, WEB-11]. The benchmark also organizes its evaluation around prompt-specific and cross-prompt scoring setups [Q14, Q46], with the entire model output taxonomy restricted to numeric trait scoring [Q48, Q181] — no benchmarked task category corresponds to natural-language revision suggestion generation that the deployment requires. Given IO is HIGH priority in elicitation and multiple genre categories are missing, this represents significant taxonomic underrepresentation.	The benchmark's category inventory consists of three essay types — argumentative/persuasive, source-dependent, and narrative/descriptive [Q12, Q13, Q14, Q15] — drawn from a US middle-school context. The deployment requires coverage of literary analysis, religious-text commentary, and cultural commentary genres that are standard in Arab high school curricula but entirely absent from the benchmark. The Qatar MOEHE curriculum confirms Arabic Literature as a distinct subject and Islamic Education as compulsory [WEB-11, WEB-5], yet no AES benchmark covers these genres in Arabic. The expert elicitation flagged IO as HIGH priority precisely because the benchmark's taxonomy omits the test-case types the deployment requires. Even within the persuasive genre that nominally overlaps, scoring logic is type-specific and embedded in English pedagogical norms [Q43, Q44, Q45].
LAILA evidence	ASAP++ evidence
[Q38] 'LAILA comprises 7,859 essays collected across 8 distinct prompts: 3 explanatory and 5 persuasive (3,446 and 4,413 essays, respectively).' (p.6)	[Q12] 'There are 3 types of essays in the dataset.' (p.2)
[Q89] 'LAILA includes only explanatory and persuasive writing prompts, limiting genre diversity and potentially affecting model robustness across other styles, such as narrative or descriptive writing.' (p.10)	[Q13] 'Argumentative / Persuasive essays - These are essays where the prompt is one in which the writer has to convince the reader about their stance for or against a topic...' (p.2)
[Q118] 'CAST Persuasive/Argumentative Writing Rubric - English Translation' (p.13)	[Q14] 'Source-dependent responses - These essays are responses to a source text...' (p.2)
[Q48] 'all models followed a multitask setup predicting holistic and trait scores jointly.' (p.7)	[Q15] 'Narrative / Descriptive essays - These are essays where the prompt requires us to describe / narrate a story.' (p.2)

[WEB-10] TAQEEEM 2025 shared task continues the explanatory/persuasive genre restriction	[Q43] 'For source-dependent essays, we found out that the most important feature for content was length, while for argumentative / persuasive essays, it was coherence and cohesion features, followed by length.' (p.4)
[WEB-11] ZaQQ 2025 dataset also confined to standard essay genres; no Arabic AES resource for religious/literary genres	[WEB-11] Qatar MOEHE 2025–2026 curriculum covers Arabic Language and Arabic Literature as separate subjects
	[WEB-5] Islamic Education is compulsory in Qatar national curriculum schools for Qatari pupils and Arab passport holders

Input Content (IC)

Whether individual datapoint content is culturally and contextually appropriate for the target region.

LAILA — 2/5 (Significant gaps) [high confidence]	ASAP++ — 1/5 (Serious concern) [high confidence]
Content was collected from 24 Qatari high schools with 4,372 students under exam-style supervision [Q3, Q27], with prompts deliberately designed to be 'developmentally appropriate for the target age groups, culturally relevant to the Arabic-speaking context, and free of sensitive topics' [Q99] and gender balance pursued [Q23, Q100]. However, the deployment targets Egypt, Saudi Arabia, Jordan, Lebanon and other Arab countries [elicitation A1], and the authors themselves acknowledge that single-country sampling 'may limit its generalizability to other Arabic-speaking populations due to diverse educational systems' [Q88]. The Qatar expatriate mitigation is partial only [Q88]. Web research confirms that secondary-curriculum Arabic writing rubrics in Egypt, Saudi Arabia, Jordan, and Lebanon are unverified against CAST [WEB-5, WEB-6, WEB-7], that Lebanese students show documented MSA-register avoidance [WEB-9], and that within Qatar itself ~60% of secondary students attend private schools across 23 curricula [WEB-1, WEB-3, WEB-4] with no documented school-type breakdown of the LAILA sample. Two prompts (P3, P4) are also under-represented [Q90], and average length is only 171 words [Q92], restricting extended-writing assessment. Given IC is HIGH priority, multiple substantive content gaps weigh heavily.	The benchmark's input instances are ~13,000 essays written entirely by native English-speaking US students in Grades 7–10 [Q3, Q10] across 8 prompts. The deployment targets Arabic-language Grade 10–12 students under the Qatar MOEHE curriculum (and other Arab curricula), writing to prompts grounded in Arab cultural, literary, and Islamic frames of reference [WEB-10, WEB-11]. There is no overlap in language, age band, nationality, curriculum, or cultural frame of reference. The anonymization scheme (e.g., @ORGANIZATION1, @PERSON1) was itself a source of interpretive difficulty even for the English annotators [Q41], and would compound transfer issues. No Arabic-language instances exist in the dataset.
LAILA evidence	ASAP++ evidence
[Q3] 'we visited 24 high schools in Qatar and collected essays directly from 4,372 students in authentic classroom settings.' (p.1)	[Q3] 'The ASAP AEG dataset comprises of approximately 13,000 essays, written across 8 prompts. The essays were written by students of class 7 to 10...' (p.1)
[Q88] 'LAILA was collected from students in a single country, Qatar, and specific grade levels, grades 10 to 12, which may limit its generalizability to other Arabic-speaking populations due to diverse educational systems.' (p.9)	[Q10] 'All the essays were written by native English speaking children from classes 7 to 10.' (p.2)
[Q99] 'Essay prompts were carefully designed to be developmentally appropriate for the target age groups, culturally relevant to the Arabic-speaking context, and free of sensitive topics or potentially harmful content.' (p.10)	[Q41] 'One of the major problems that the annotators faced was the fact that all the essays were anonymized. Named entities, like The New York Times would be referred to as @ORGANIZATION1, Donald Trump would be referred to as @PERSON1...' (p.4)
[Q90] 'two of the prompts (Prompts 3 and 4) have fewer essays than the rest, which could introduce imbalance' (p.10)	[WEB-10] Qatar MOE Arabic Language curriculum framework anchors Grades 1–12 writing standards

[Q92] 'with an average essay length of only 171 words, the dataset lacks extended writing samples, restricting the assessment of models on long-form academic tasks' (p.10)	[WEB-11] IMPACT-se review of Qatar 2025–2026 curriculum confirms Arabic Literature and politically/religiously contextualized content
[WEB-1] Qatar 2024–2025: 278 government schools (~137,048 students), 351 private schools (~228,488 students); ~60% of secondary students in private schools	
[WEB-4] 23 distinct curricula operate in Qatar; British curriculum leads at 45.8% of K-12 schools	
[WEB-9] Lebanese students documented to resist formal MSA literary register	
[WEB-5] Egypt's curriculum reform ongoing since 2018; writing assessment norms in transition	
[WEB-6] Saudi Arabia uses centrally produced rating tools managed by supervisors	

Input Form (IF)

Whether the input signal encoding (text, audio, image parameters) matches deployment conditions.

LAILA — 4/5 (Minor gaps) [high confidence]	ASAP++ — 1/5 (Serious concern) [high confidence]
The input form is digital Arabic text submitted via Microsoft Forms under a 65-minute timed condition [Q25, Q26], aligning well with the deployment's text-based formative-feedback interface. The script is native Arabic (RTL), eliminating Latin-script mismatch concerns. Length distribution is unimodal around 171 words [Q134, Q135] with a smaller cluster of short essays (11–60 words, 1,061 essays) [Q136], and very few essays exceed 500 words [Q138] — a meaningful but minor mismatch with deployment, where formative drafts may vary more in length and grade-level. Model-specific input formats (AraBERT pooling + 816 features [Q156, Q157], AraT5 instruction+text [Q148], LLMs with 4096-token context [Q184]) are reasonable for Arabic NLP. Internet penetration in Qatar is 99% [WEB-14], supporting digital deployment. IF was marked LOWER priority, and the form alignment is genuinely strong; the main residual concern is that benchmark essays are exam-style timed productions while deployment drafts are pre-submission and may differ in length/polish.	The benchmark operates on English plain text with a feature pipeline (Zesch et al. attribute-independent features, Barzilay & Lapata entity grids) extracted via Stanford CoreNLP [Q31, Q32, Q33]. The deployment operates in MSA, written in right-to-left Arabic script with root-and-pattern morphology, optional diacritics (tashkeel), Arabic-specific punctuation (., ؟), and potential Gulf Arabic dialect intrusion [writing_systems, languages.dialect_influence, WEB-3, WEB-4]. Stanford CoreNLP's English pipeline cannot be applied to Arabic text directly; the canonical Arabic tools (CAMEL Tools, Camel Morph MSA, CAMELBERT-MSA) are entirely separate ecosystems [WEB-7, WEB-8, WEB-9]. This is a complete signal-distribution mismatch, not a parameter mismatch.
LAILA evidence	ASAP++ evidence
[Q25] 'LAILA was designed to collect essays written under timed conditions using a digital submission platform.' (p.3)	[Q31] 'We used the attribute independent feature set provided by Zesch et al. (2015).' (p.3)
[Q26] 'we configured a Microsoft Form for digital essay submission, set with a 65-minute time limit' (p.5)	[Q32] 'In addition to those features, we also made use of entity grid features described in Barzilay and Lapata (2005).' (p.3)
[Q134] 'unimodal pattern centered around an average essay length of 171 words, with the majority of essays concentrated between 90 and 210 words.' (p.15)	[Q33] 'All the features were extracted using Stanford Core NLP (Manning et al., 2014).' (p.3)
[Q138] 'Beyond 420 words, the frequency drops sharply, with only a few essays exceeding 500 words' (p.15)	[WEB-7] CAMEL Tools provides MSA morphological analysis, POS, dialect ID and diacritization — the Arabic counterpart ecosystem to Stanford CoreNLP

[WEB-14] 99% internet penetration in Qatar as of 2024	[WEB-8] Camel Morph MSA is the largest open-source MSA morphological analyzer (LREC-COLING 2024)
	[WEB-9] CAMELBERT-MSA is the standard pre-trained Arabic encoder used in ZAEBUC, LAILA, and TAQEEEM AES systems
	[WEB-3] MSA/Gulf-Arabic code-mixing in student writing is empirically documented but unresourced for AES
	[WEB-4] ZAEBUC-Spoken confirms MSA/Gulf/English code-switching among Gulf students

Output Ontology (OO)

Whether the benchmark's output categories and scoring criteria reflect regionally valid decision boundaries.

LAILA — 1/5 (Serious concern) [high confidence]	ASAP++ — 1/5 (Serious concern) [high confidence]
This is the most severe misalignment. LAILA's output ontology is exclusively numeric ordinal trait scoring on the CAST rubric: 6 traits at 0–5 plus REL at 0–2 [Q34, Q36], with holistic computed as the sum [Q35, Q182]. All benchmarked models 'followed a multitask setup predicting holistic and trait scores jointly' [Q48]; LLMs produce trait scores in JSON format [Q181]; AraT5 formulates AES as text generation of trait names plus scores [Q147]. The deployment requires actionable natural-language revision suggestions explaining why points were lost [elicitation A3, regional YAML output_function], a qualitatively different output category that is entirely absent from the benchmark's output schema. Web research confirms this gap is field-wide: 'crucial aspects of writing evaluation, such as providing feedback, have often been overlooked' even in English [WEB-12], and no Arabic rubric-to-natural-language feedback system exists [WEB-13, WEB-23]. The CAST rubric is also explicitly genre-scoped to persuasive/argumentative writing [Q118], compounding the misalignment with deployment genres. The narrow REL scale (0–2) further limits granularity [Q91]. OO is HIGH priority and the mismatch is direct, structural, and acknowledged by the authors who 'explicitly prohibit its use for high-stakes educational decisions' [Q104].	The benchmark output ontology is a fixed integer score range [Q11] over a small set of attributes (content, organization, word choice, sentence fluency, conventions for narrative; analogous traits for other types) [Q16, Q17, Q18, Q19], with annotations described as gold standard for predicting attribute scores [Q50]. The deployment requires natural-language revision suggestions with explanatory rationale (per elicitation Q3/A3), which is a categorically different output space — not a numeric score taxonomy at all. Additionally, traits like 'word choice' and 'sentence fluency' are calibrated to English vocabulary range and syntactic variety; Arabic writing quality dimensions critical for this deployment include morphological correctness, classical rhetorical structure (????? / ?????), and dialect intrusion handling [cultural_norms_notes, writing_systems.morphological_considerations], which are not represented.
LAILA evidence	ASAP++ evidence
[Q34] 'The rubric covers 7 writing traits [D4]: Relevance (REL), Organization (ORG), Vocabulary (VOC), Style (STY), Development (DEV), Mechanics (MEC), and Grammar (GRA).' (p.5)	[Q11] 'We use the same score range as the overall score range of the essays.' (p.2)
[Q48] 'all models followed a multitask setup predicting holistic and trait scores jointly.' (p.7)	[Q16] 'The ASAP dataset already contains attribute scores for the narrative essays, namely content, organization, word choice, sentence fluency, conventions, etc.' (p.2)
[Q181] 'the LLM is tasked with generating trait scores in JSON format, given the prompt text, essay, and trait rubrics.' (p.18)	[Q18] 'Based on the types of essays, there are 2 sets of attributes.' (p.2)
[Q118] 'CAST Persuasive/Argumentative Writing Rubric - English Translation' (p.13)	[Q19] 'There are 5 attributes for narrative essays, namely 1. Content... 2. Organization... 3. Word Choice...' (p.2)

[Q91] 'the relevance trait is scored on a narrow scale from 0 to 2, which limits the granularity of evaluation and may cause QWK to underrepresent subtle differences in model predictions.' (p.10)	[Q50] 'These annotations can be used as a gold standard for future experiments in predicting different attribute scores.' (p.5)
[Q104] 'We explicitly prohibit its use for high-stakes educational decisions affecting individual students.' (p.10)	
[WEB-12] 2024 Springer survey: feedback generation explicitly identified as overlooked in AES literature	
[WEB-13] AAEE Saudi system generates corrected essay versions but not trait-grounded revision suggestions	

Output Content (OC)

Whether ground-truth labels align with the judgments of regional stakeholders.

LAILA — 3/5 (Moderate gaps) [high confidence]	ASAP++ — 1/5 (Serious concern) [high confidence]
Annotation quality within Qatar is a relative strength: 6 annotators + 3 supervisors, all Arabic-language teachers/lecturers with 5 holding advanced degrees [Q30, Q31], with detailed guidelines, training, moderation, and adjudication procedures [Q41, Q42, Q43]. Inter-annotator agreement is substantial, averaging QWK 0.66 (P5) to 0.75 (P1) [Q44]. However, P5 required third-annotator adjudication on 23.3% of essays [Q45], indicating interpretive instability on explanatory prompts. The decisive concern is institutional homogeneity: rubric, annotators, schools, and funding are all Qatar-anchored (QUTC, QRDl, Qatar MoE) [Q33, Q84, Q86], with no documented non-Qatari annotators and no cross-national norm validation (gap_id 1, gap_id 6). For deployment in Egypt, Saudi Arabia, Jordan, and Lebanon, it is unknown whether STY and DEV labels reflect pan-Arab teacher consensus or Qatar-specific institutional convention [WEB-5, WEB-6, WEB-7]. OC was rated MODERATE priority, and the in-region (Qatar) signal is genuinely positive — but the cross-national transferability is unsubstantiated. A score of 3 reflects strong within-Qatar content validity tempered by unvalidated transferability.	Ground-truth labels were produced by three India-based annotators with English MA-level training, English newspaper editorial experience, and high English exam/TOEFL scores [Q27, Q28, Q29, Q30, Q53], scoring single-pass [Q22, Q40] with anchoring to the original US ASAP raters' overall scores [Q23, Q25]. None of these qualifications confer expertise in Arabic writing assessment or Qatari/Arab curriculum norms. The anchoring mechanism [Q25, Q26] further embeds US rater judgments as the normative standard. Per the elicitation, OC was flagged HIGH priority because Qatar/Arab teachers may apply different MSA convention, organization, and style standards. Convergent validity with the deployment's regional stakeholders is essentially zero.
LAILA evidence	ASAP++ evidence
[Q30] 'The hired annotation team consisted of 6 annotators and 3 supervisors, all of whom were Arabic language teachers or lecturers with educational backgrounds.' (p.5)	[Q22] 'Unlike the ASAP AEG dataset in which every essay was annotated by 2 annotators, we use only 1 annotator here for each essay.' (p.3)
[Q31] 'Five members of the team hold advanced degrees (MSc or PhD) in the Arabic language.' (p.5)	[Q23] 'For the ground truth, we make use of the overall score of the essays given by the original annotators of the ASAP AEG dataset.' (p.3)
[Q44] 'IAA was substantial across prompts, with an average QWK ranging from 0.66 (P5) to 0.75 (P1)' (p.6)	[Q25] 'The final score that is chosen is the one from the annotator that is closest to the overall scores.' (p.3)
[Q45] 'P5 showed the lowest average agreement, noting that it has the highest percentage of essays that require a third annotator (A3: 23.3%)' (p.6)	[Q27] 'We made use of a total of three annotators to annotate the essays.' (p.3)

[Q84] 'we heartily thank our dedicated annotators ... Qatar University Testing Center, the Ministry of Education and Higher Education in Qatar' (p.9)	[Q28] 'Each of the annotators had competence in English, either by scoring quite high marks in their high school exams (over 90% in English), or scoring over 110 in ToEFL.' (p.3)
[Q43] 'essays with large discrepancies between annotators (≥ 6 points difference in the HOL score) were flagged and escalated to a supervisor' (p.6)	[Q30] 'All the annotators have either studied or are studying English at a Master of Arts (MA) level.' (p.3)
[WEB-5] Egypt's curriculum reform suggests evolving writing-assessment norms not aligned with CAST	[Q53] 'We thank the annotators of our task - Advait Jayakumar, Janice Pereira and Elaine Mathias...' (p.5)
[WEB-6] Saudi Arabia uses centrally produced rating tools managed by supervisors	
[WEB-7] Lebanon's CRDP authoritative but Arabic writing rubrics not in searchable form	

Output Form (OF)

Whether the expected output modality matches regional deployment needs and accessibility.

LAILA — 1/5 (Serious concern) [high confidence]	ASAP++ — 1/5 (Serious concern) [high confidence]
The output-form mismatch is the most direct and explicitly documented gap. LAILA evaluates models using QWK between human and model numeric scores [Q58] with hyperparameter optimization on average QWK [Q59, Q60, Q61]; LLMs are constrained to JSON-formatted scores via the outlines library [Q190]. The deployment requires open-ended natural-language revision suggestions in MSA accessible to G10–12 students [elicitation A3, regional YAML output_function]. QWK on numeric scores cannot capture pedagogical usefulness, actionability, or comprehensibility of generated feedback. The authors themselves recognize this limit by prohibiting high-stakes use [Q104] and recommending 'fairness auditing and stakeholder consultation before real-world application' [Q106]. Web research confirms no Arabic rubric-to-natural-language feedback evaluation framework exists [WEB-12, WEB-13, WEB-23]. OF was marked HIGH priority. Even basic infrastructure for evaluating feedback quality (rubrics for actionability, register-appropriateness, MSA fluency) is absent from the benchmark.	The benchmark's output form is integer ordinal scores per trait, evaluated by Quadratic Weighted Kappa using an Ordinal Class Classifier with Random Forest internals under 5-fold CV [Q34, Q35, Q36, Q37, Q38, Q39]. The deployment explicitly requires free-text natural-language revision suggestions with explanatory rationale (elicitation Q3/A3). QWK has no mechanism to evaluate the quality, actionability, or pedagogical appropriateness of generated feedback text. This is a complete output-form mismatch: numeric ordinal classification vs. open-ended generation. The benchmark cannot validate the deployment's output modality at all.
LAILA evidence	ASAP++ evidence
[Q58] 'We evaluate model performance using QWK to measure agreement between human and model scores.' (p.8)	[Q34] 'We evaluate each of the annotators using Cohen's Kappa, with quadratic weights - i.e. Quadratic Weighted Kappa (QWK) (Cohen, 1968).' (p.4)
[Q190] 'For all the LLM experiments, we used vLLM for inference with the outlines library to enforce the JSON output format.' (p.18)	[Q35] 'We chose this as the evaluation metric (as compared to accuracy and weighted F-Score) because of the following reasons: Unlike accuracy and F-Score, Kappa takes into account random agreement.' (p.4)
[Q104] 'We explicitly prohibit its use for high-stakes educational decisions affecting individual students.' (p.10)	[Q37] 'To the best of our knowledge, all of the papers using the ASAP dataset make use of this as the evaluation metric.' (p.4)
[Q106] 'We recommend that any system developed using LAILA undergo fairness auditing and stakeholder consultation before real-world application.' (p.10)	[Q38] 'We made use of the Ordinal Class Classifier (Frank and Hall, 2001) in Weka... We used the Random Forest classifier (Breiman, 2001) as the internal classifier.' (p.4)

[WEB-12] AES literature explicitly identifies feedback generation as an overlooked area	[Q39] 'We used 5-fold cross validation to get the results for each attribute for each prompt.' (p.4)
[WEB-13] AAEE Saudi system provides corrected essays but not trait-grounded revision suggestions	
[WEB-23] No Arabic system bridges trait scores to actionable natural-language feedback for secondary students	