

Validity Analysis: LAILA

Target context: Arab-Country G10–12 Arabic Teachers: AI Essay Grading Assistance

Overall risk: **HIGH**

Deployment Context

Use case: Assist high school (G10-12) Arabic teachers in Arab countries (teaching students who are native in Arabic) in grading Arabic essays (dataset was collected in Qatar but the writing is in modern standard Arabic which is standard across all Arab countries, although the standards for grading may differ from country to country).

Target population: Arabic teachers in Arab countries teaching high school students (G10-12)

Dimension Scores

Dimension	Score	Rating	Confidence
Input Ontology 	2	Significant gaps	high
Input Content 	2	Significant gaps	high
Input Form 	4	Minor gaps	high
Output Ontology 	2	Significant gaps	high
Output Content 	3	Moderate gaps	medium
Output Form 	1	Serious concern	high
Average	2.3		

 = highest concern  = strongest dimension

Dimension Key

Abbr.	Dimension	Definition
IO	Input Ontology	Whether the benchmark’s test case categories cover the query types expected in deployment.
IC	Input Content	Whether individual datapoint content is culturally and contextually appropriate for the target region.
IF	Input Form	Whether the input signal encoding (text, audio, image parameters) matches deployment conditions.
OO	Output Ontology	Whether the benchmark’s output categories and scoring criteria reflect regionally valid decision boundaries.
OC	Output Content	Whether ground-truth labels align with the judgments of regional stakeholders.
OF	Output Form	Whether the expected output modality matches regional deployment needs and accessibility.

Overall Summary

LAILA is a methodologically rigorous, well-documented Arabic AES benchmark that establishes the largest publicly available trait-annotated Arabic essay dataset to date [Q1, Q2]. For the target deployment — pan-Arab G10–12 Arabic teachers using AI for essay grading assistance — its validity is mixed. Input Form is well aligned (text-only MSA, matching modality). Output Content has strong annotation procedures but unresolved questions about annotator nationality diversity. The remaining four dimensions show significant deployment mismatches that map directly onto the HIGH-priority gaps flagged in elicitation: Input Ontology (Egypt’s holistic grading [WEB-5] and missing genres [Q89]), Input Content (Qatari-only prompts [Q88] with no cross-national degradation testing), Output Ontology (no confidence or rationale categories), and Output Form (discrete labels

only, with the benchmark explicitly prohibiting high-stakes operational use without independent validation [Q104, Q105]). The benchmark is best understood as a strong research resource that requires substantial supplementation — both empirical (cross-national validation) and methodological (rationale and uncertainty layers) — before responsible teacher-facing deployment.

ID	Evidence
Q1	“LAILA, the largest publicly available Arabic AES dataset to date, comprising 7,859 essays annotated with holistic and trait-specific scores on seven dimensions: relevance, organization, vocabulary, style, development, mechanics, and grammar.” (p.1)
Q2	“The dataset comprises 7,859 essays across 8 distinct prompts making it the most extensive publicly available Arabic resource of its kind.” (p.1)
Q88	“First, LAILA was collected from students in a single country, Qatar, and specific grade levels, grades 10 to 12, which may limit its generalizability to other Arabic-speaking populations due to diverse educational systems. However, this concern is partially mitigated by Qatar’s large population of Arab expatriates, resulting in student participants representing a wide variety of Arabic dialects and backgrounds.” (p.9)
Q89	“Second, LAILA includes only explanatory and persuasive writing prompts, limiting genre diversity and potentially affecting model robustness across other styles, such as narrative or descriptive writing.” (p.10)
Q104	“We explicitly prohibit its use for high-stakes educational decisions affecting individual students.” (p.10)
Q105	“Users must commit to: (1) preventing re-identification attempts, (2) avoiding demographic profiling or inference, and (3) refraining from deployment in operational assessment contexts without independent validation.” (p.10)
WEB-5	Egypt’s holistic grading culture suggests Egyptian teachers may apply different evaluative standards than Qatar-based annotators (https://grokipedia.com/page/academic_grading_in_egypt)

Practical Guidance

What This Benchmark Measures

LAILA measures how well models can predict CAST-rubric trait scores and a derived holistic score on explanatory and persuasive MSA essays produced by G10–12 Qatari students under timed digital conditions [Q14, Q34, Q38]. Its strongest dimensions for the target context are Input Form (text-only MSA matches deployment) and Output Content (rigorous annotation with substantial IAA [Q44]), making it useful for benchmarking baseline scoring agreement between models and qualified Arabic-educator judgment.

Construct Depth

The benchmark probes the construct of Arabic essay quality with substantial depth on the analytic-rubric dimension — fine-grained seven-trait scoring, double annotation, rigorous adjudication, and extensive baseline experiments across feature-based, encoder-based, and LLM models in both prompt-specific and cross-prompt setups [Q48, Q83]. However, depth is concentrated within a Qatari-curriculum, two-genre, discrete-label paradigm. It does not probe construct depth on dimensions central to the deployment: cross-national rubric portability (Egypt’s holistic grading, Levantine and Maghrebi conventions), narrative/descriptive genre coverage, uncertainty calibration, or rationale quality.

What Else You Need

For valid teacher-facing deployment beyond Qatar, supplementation is required across four dimensions: (IO) cross-national rubric mapping plus a holistic-only evaluation track for Egyptian-style use; (IC) cross-national prompt validation testing model scoring stability on Egyptian, Levantine, and Maghrebi locally-grounded prompts; (OO+OF) integration of an Arabic-language adaptation of rationale-generation frameworks such as RaDME or RMTS [WEB-13, WEB-12] plus calibrated uncertainty estimates evaluated against teacher disagreement; (OC) a cross-country annotator validation study — re-annotating a stratified sample of LAILA essays with Egyptian, Saudi, Jordanian, Lebanese, and Moroccan Arabic-language educators to quantify cross-country labeling drift on Style, Development, and Organization. The benchmark’s own ethics requirements [Q105, Q106] effectively mandate this supplementation before operational deployment.

ID	Evidence
Q14	“We establish baseline AES results for Arabic under two evaluation setups: prompt-specific and cross-prompt.” (p.2)

Q34	"The rubric covers 7 writing traits [D4]: Relevance (REL), Organization (ORG), Vocabulary (VOC), Style (STY), Development (DEV), Mechanics (MEC), and Grammar (GRA)." (p.5)
Q38	"LAILA comprises 7,859 essays collected across 8 distinct prompts: 3 explanatory and 5 persuasive (3,446 and 4,413 essays, respectively)." (p.6)
Q44	"We compute IAA using Quadratic Weighted Kappa (QWK) (Williamson et al., 2012) to assess agreement between the two main annotators (A1 and A2), prior to adjudication. Table 4 reports QWK per trait and prompt, with agreement strength following Landis and Koch (1977). Overall, IAA was substantial across prompts, with an average QWK ranging from 0.66 (P5) to 0.75 (P1), showing strong consistency." (p.6)
Q48	"The models fall into three categories: feature-based (FB), encoder-based, and large language models (LLMs). LLMs were evaluated in zero-shot and few-shot settings, and all models followed a multitask setup predicting holistic and trait scores jointly." (p.7)
Q83	"We also benchmarked LAILA using SOTA Arabic and English AES models in both prompt-specific and cross-prompt settings, showing strong baselines for future research." (p.9)
Q105	"Users must commit to: (1) preventing re-identification attempts, (2) avoiding demographic profiling or inference, and (3) refraining from deployment in operational assessment contexts without independent validation." (p.10)
Q106	"We recommend that any system developed using LAILA undergo fairness auditing and stakeholder consultation before real-world application." (p.10)
WEB-12	BERT-based Arabic AES system documented but without rationale generation (https://www.researchgate.net/publication/382302599_Automated_essay_scoring_in_Arabic_a_dataset_and_analysis_of_a_BERT-based_system)
WEB-13	RaDME generates trait scores with co-produced rationales using knowledge distillation — English only, no Arabic adaptation documented (https://arxiv.org/html/2502.20748v1)