

Validity Analysis: MRBENCH

Target context: Metropolitan India — Grades 1–8 Mathematics Tutoring (Student Population)

Overall risk: **HIGH**

Deployment Context

Use case: Deployment scenario: A GPT-4 model is implemented in a one-on-one tutoring session, where its responses are augmented with a human tutor’s input. We need to assesd whather Are the student stisfied with this newly GPT4 responses?

Domain: Educational tutoring

Setting: Mobile application / enterprise software

Target population: A student from a metropolitan city in India, in grades 1–8, who is fluent in English.

Dimension Scores

Dimension	Score	Rating	Confidence
Input Ontology 	2	Significant gaps	high
Input Content 	1	Serious concern	high
Input Form 	4	Minor gaps	high
Output Ontology 	1	Serious concern	high
Output Content 	2	Significant gaps	high
Output Form 	3	Moderate gaps	medium
Average	2.2		

 = highest concern  = strongest dimension

Dimension Key

Abbr.	Dimension	Definition
IO	Input Ontology	Whether the benchmark's test case categories cover the query types expected in deployment.
IC	Input Content	Whether individual datapoint content is culturally and contextually appropriate for the target region.
IF	Input Form	Whether the input signal encoding (text, audio, image parameters) matches deployment conditions.
OO	Output Ontology	Whether the benchmark's output categories and scoring criteria reflect regionally valid decision boundaries.
OC	Output Content	Whether ground-truth labels align with the judgments of regional stakeholders.
OF	Output Form	Whether the expected output modality matches regional deployment needs and accessibility.

Overall Summary

MRBench provides a methodologically rigorous, learning-sciences-grounded evaluation framework for LLM tutoring on mathematics mistake remediation, but it is materially misaligned with the Metropolitan India Grade 1–8 deployment on the dimensions the user has flagged as HIGH priority. The taxonomy (input ontology) lacks any category for NCERT alignment, Indian numeral conventions, or cultural register; the instance pool (input content) is drawn from Western-framed Bridge and MathDial datasets with no documented Indian content, against a deployment requirement (A1, A2) for NCERT-aligned, India-grounded problems; the output ontology has no dimension for cultural appropriateness or student satisfaction, with the

Tutor Tone dimension trivially collapsing to non-offensiveness [Q90] precisely on the axis where Indian classroom warmth norms (A4) would matter; and ground-truth labels (output content) come from four Computer Science post-graduates with undisclosed nationality and no teaching experience [Q31, Q32], with no Indian rater validation in either MRBench or the BEA 2025 Shared Task [WEB-17, WEB-18]. Input form and output form are well-matched as text-only English-medium evaluation, partially mitigating overall risk. The benchmark is useful as a partial diagnostic — particularly for per-LLM behaviors like answer revelation rate — but it cannot substitute for a culturally and curricularly validated evaluation of the deployment's primary success criterion.

ID	Evidence
Q31	"The annotation team consisted of two male and two female annotators, with all four annotators holding at least a post-graduate degree in Computer Science and being proficient in English." (p.5)
Q32	"We note that for this study, we do not require annotators to have direct teaching experience, as understanding of the mathematical tasks at the middle school level and being able to judge the responses from the perspective of a potential user of such AI tutors (or a student), rather than specifically a teacher, is sufficient." (p.5)
Q90	"When we combine these two labels into "Non-offensive", the DAMR score reaches 100% as we observe no offensive responses from any LLMs or humans." (p.8)
WEB-17	BEA 2025 Shared Task extended MRBench without Indian rater validation (https://arxiv.org/html/2507.10579v1)
WEB-18	NAACL 2025 BEA proceedings: no participating system documented Indian adaptation (https://aclanthology.org/2025.naacl-long.57.pdf)

Practical Guidance

What This Benchmark Measures

MRBench measures whether LLM tutor responses to student mathematical mistakes exhibit eight specific pedagogical behaviors: mistake identification, mistake-location, guidance, actionability, coherence, tone (operationalized as non-offensiveness), human-likeness, and not-revealing-the-answer. For the Indian deployment, the strongest signals come from input form (text-based English dialogue, fully matched to deployment) and output form (per-dimension DAMR diagnostics that directly support QA review by human tutor co-authors), particularly the answer-revelation rate which is a meaningful production-relevant signal.

Construct Depth

The benchmark probes pedagogical behaviors at the level of individual tutor turns, not learning outcomes or student satisfaction — limitations explicitly acknowledged by the authors [Q119, Q120, Q121]. It provides no evidence on cultural appropriateness, NCERT alignment, or Indian-context word problem handling, and the Tutor Tone dimension collapses trivially [Q90] on the very axis where Indian register differences would surface. DAMR's per-dimension granularity is useful but does not operationalize the deployment's primary success criterion (student satisfaction).

What Else You Need

Substantial supplementation is required: (1) NCERT-aligned content evaluation set (input_ontology, input_content) since no existing benchmark fills this [WEB-6, WEB-7]; (2) Indian-context word problem evaluation extending GSM8K-style robustness testing [WEB-8] to MRBench dialogues; (3) cultural-appropriateness output dimension with Indian-rater-validated desired labels (output_ontology, output_content); (4) student-satisfaction validation linking DAMR to actual engagement signals from the deployed system; (5) Indian rater pool (educators and metropolitan Grade 1–8 students/parents) to re-annotate a sample for tone, human-likeness, and any added cultural dimension. DPDP Act 2025 [WEB-23, WEB-24] imposes verifiable parental consent and bans on profiling that constrain how this supplementation can be done.

ID	Evidence
Q90	"When we combine these two labels into "Non-offensive", the DAMR score reaches 100% as we observe no offensive responses from any LLMs or humans." (p.8)
Q119	"The current taxonomy and annotation scheme focus on the appropriateness of the tutor responses." (p.9)
Q120	"However, one of the limitations is that it does not consider the tutoring dialogues' impact on the overall student learning." (p.9)

Q121	“Specifically, the annotation pertains to the individual tutor turns within educational dialogues, which restricts our understanding of broader implications on student learning processes and learning gains, typically observed after a conversation concludes.” (p.9)
WEB-6	No NCERT-aligned Grade 1–8 English-medium tutoring dialogue benchmark exists (https://arxiv.org/html/2404.13099v1)
WEB-7	HAWP is Hindi arithmetic word problems only, not tutoring dialogue (https://aclanthology.org/2022.lrec-1.373.pdf)
WEB-8	GSM8K robustness research shows Indian-context substitution causes measurable LLM performance drops (https://www.emergentmind.com/topics/gsm8k)
WEB-23	https://ksandk.com/data-protection-and-data-privacy/child-data-protection-under-dpdp-act-parental-consent-rules/
WEB-24	https://www.dpdpa.com/dpdpa2023/chapter-2/section9.html

Input Ontology (IO)

Score: 2/5 — Significant gaps | Confidence: high

Benchmark: MRBENCH | Context: Metropolitan India — Grades 1–8 Mathematics Tutoring (Student Population)

Input Ontology definition: Whether the benchmark’s test case categories cover the query types expected in deployment.

Each section below is followed by an evidence table. Three types of evidence are cited: [QN] — verbatim quotes extracted from the benchmark paper; [WEB-N] — web sources consulted during assessment; [DN] / [DATASET-DN] — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

MRBench's eight-dimension pedagogical taxonomy is rigorously derived from learning-sciences principles [Q1, Q15, Q129] and covers core mistake-remediation behaviors (mistake identification, guidance, not revealing the answer, tone, human-likeness) that are relevant to any tutoring deployment. However, for the Indian metropolitan Grade 1–8 deployment — explicitly flagged HIGH priority and confirmed in elicitation A1 to require NCERT-aligned curriculum coverage — the taxonomy contains no category for curriculum alignment, Indian numeral conventions (lakhs/crores), rupee-denominated content, or culturally appropriate register. The authors themselves acknowledge the taxonomy 'will need to be verified on and likely adapted if applied to other tasks ... and to subjects other than mathematics' [Q118] but do not extend this to non-Western educational settings. Web search confirms no NCERT Grade 1–8 tutoring dialogue benchmark exists [WEB-6, WEB-7], so the gap is not addressable by alternative resources. The taxonomy provides partial but inadequate coverage of the categories required for this deployment.

ID	Evidence
Q1	“we propose a unified evaluation taxonomy with eight pedagogical dimensions based on key learning sciences principles, which is designed to assess the pedagogical value of LLM-powered AI tutor responses grounded in student mistakes or confusions in the mathematical domain.” (p.1)
Q15	“In this section, we first present our approach, narrowing the evaluation taxonomy down to eight measurable dimensions aligned with key pedagogical strategies (Jurenka et al., 2024; Hennessy et al., 2016). These dimensions are most suitable for the student mistake remediation task and are based on the learning sciences principles.” (p.3)
Q118	“We acknowledge that the proposed taxonomy will need to be verified on and likely adapted if applied to other tasks such as concept learning, and to subjects other than mathematics.” (p.9)
Q129	“Through an iterative analysis of the taxonomy, we identify eight dimensions that comprehensively assess tutor response quality in the context of mistake remediation.” (p.12)
WEB-6	No NCERT-aligned Grade 1–8 English-medium tutoring dialogue benchmark exists (https://arxiv.org/html/2404.13099v1)
WEB-7	HAWP is Hindi arithmetic word problems only, not tutoring dialogue (https://aclanthology.org/2022.lrec-1.373.pdf)

Strengths

- Pedagogically grounded eight-dimension taxonomy derived through iterative validation [Q15, Q129, Q132]
- Explicit focus on mistake remediation in mathematics matches the deployment's task scope [Q12, Q117]
- Covers two curriculum levels (elementary/Bridge and middle school/MathDial) corresponding to Grades 1–8 banding [Q143, Q144]

ID	Evidence
Q12	“In this work, we focus on educational dialogues between a student and a tutor in the mathematical domain. Specifically, the conversations are grounded in students' mistakes or confusions, and the AI tutor aims to respond in order to remediate such mistakes or confusions.” (p.3)
Q15	“In this section, we first present our approach, narrowing the evaluation taxonomy down to eight measurable dimensions aligned with key pedagogical strategies (Jurenka et al., 2024; Hennessy et al., 2016). These dimensions are most suitable for the student mistake remediation task and are based on the learning sciences principles.” (p.3)
Q117	“The proposed taxonomy primarily focuses on the task of the student mistake remediation in the domain of mathematics.” (p.9)

Q129	"Through an iterative analysis of the taxonomy, we identify eight dimensions that comprehensively assess tutor response quality in the context of mistake remediation." (p.12)
Q132	"However, our validation pilot study (see Section 4.2) confirmed that the selected eight dimensions are both necessary and sufficient for evaluating tutor response quality in dialogues aimed at mistake remediation." (p.12)
Q143	"These aspects highlight that including both datasets in MRBench ensures diversity and provides for a good mix of easy and difficult mathematical problems, making the benchmark both comprehensive and challenging." (p.14)
Q144	"The problems covered in the Bridge dataset are at the elementary school level, whereas those in MathDial are at the middle school level." (p.14)

Checklist

Checklist item definitions:

ID	Assessment Question
IO-1	Identify test case categories required for regional deployment
IO-2	Check if taxonomy omits regionally relevant categories
IO-3	Check if taxonomy includes irrelevant categories
IO-4	Document category gaps harming content validity

Checklist responses:

IO-1: Required regional categories include NCERT-aligned topic coverage (place value with lakhs, fractions, integers, linear equations, mensuration), Indian numeral conventions, rupee-denominated word problems, and culturally appropriate register. Confirmed via elicitation A1 and NCERT syllabus sources [WEB-2, WEB-4, WEB-5].

IO-2: Yes. The benchmark's taxonomy omits curriculum alignment, Indian numeral conventions, and cultural-register categories. The authors acknowledge adaptation would be required for other educational settings [Q118, Q130], and no NCERT-aligned tutoring dialogue benchmark currently fills this gap [WEB-6, WEB-7].

IO-3: The taxonomy's pedagogical dimensions (mistake identification, guidance, actionability, coherence, tone, human-likeness, etc.) are broadly relevant; no included category is clearly irrelevant to the Indian deployment, though the operationalization of Tutor Tone trivially collapses to non-offensiveness [Q90].

IO-4: Documented gaps: (a) no NCERT curriculum-alignment dimension; (b) no Indian numeral/currency convention dimension; (c) no cultural-appropriateness/register dimension. These omissions harm content validity for the target deployment because A4 confirms cultural appropriateness is a necessary condition for satisfaction.

ID	Evidence
Q1	"we propose a unified evaluation taxonomy with eight pedagogical dimensions based on key learning sciences principles, which is designed to assess the pedagogical value of LLM-powered AI tutor responses grounded in student mistakes or confusions in the mathematical domain." (p.1)
Q15	"In this section, we first present our approach, narrowing the evaluation taxonomy down to eight measurable dimensions aligned with key pedagogical strategies (Jurenka et al., 2024; Hennessy et al., 2016). These dimensions are most suitable for the student mistake remediation task and are based on the learning sciences principles." (p.3)
Q90	"When we combine these two labels into "Non-offensive", the DAMR score reaches 100% as we observe no offensive responses from any LLMs or humans." (p.8)
Q117	"The proposed taxonomy primarily focuses on the task of the student mistake remediation in the domain of mathematics." (p.9)
Q118	"We acknowledge that the proposed taxonomy will need to be verified on and likely adapted if applied to other tasks such as concept learning, and to subjects other than mathematics." (p.9)
Q130	"However, other educational settings, particularly those involving tutorial dialogues beyond mistake remediation, may require modifications, as discussed in the limitations section." (p.12)
Q143	"These aspects highlight that including both datasets in MRBench ensures diversity and provides for a good mix of easy and difficult mathematical problems, making the benchmark both comprehensive and challenging." (p.14)
Q144	"The problems covered in the Bridge dataset are at the elementary school level, whereas those in MathDial are at the middle school level." (p.14)

WEB-2	NCERT Grade 1–5 mathematics syllabus (place value with lakhs, spiral sequencing) (https://ncert.nic.in/pdf/syllabus/06Math%20(I-V).pdf)
WEB-4	NCERT upper-primary syllabus emphasizing abstraction, fractions, integers, algebra (https://cbseportal.com/NCERT/syllabus/maths-secondary)
WEB-5	NCERT Class 8 chapter list (Rational Numbers, Linear Equations, Mensuration) (https://www.tiwariacademy.com/ncert-solutions/class-8/maths/)
WEB-6	No NCERT-aligned Grade 1–8 English-medium tutoring dialogue benchmark exists (https://arxiv.org/html/2404.13099v1)
WEB-7	HAWP is Hindi arithmetic word problems only, not tutoring dialogue (https://aclanthology.org/2022.lrec-1.373.pdf)

Information Gaps

- No documentation of whether MRBench's source datasets contain any NCERT-aligned content
- No empirical study comparing pedagogical taxonomy coverage across Indian vs. Western tutoring contexts

Requires Expert Verification

- Whether additional pedagogical dimensions (e.g., respect/hierarchy register, scaffolding to NCERT learning outcomes) are required for the Indian metropolitan cohort

Remediation

Gap 1: Taxonomy omits NCERT curriculum alignment, Indian numeral convention compliance, and cultural-register categories required by the deployment.

Recommendation: Extend the eight-dimension taxonomy with at least three additional categories (curriculum-alignment, numeral/currency-convention compliance, cultural-register/warmth) and validate them with Indian educators before any deployment QA decisions are made on benchmark scores.

Input Content (IC)

Score: 1/5 — Serious concern | Confidence: high

Benchmark: MRBENCH | Context: Metropolitan India — Grades 1–8 Mathematics Tutoring (Student Population)

Input Content definition: Whether individual datapoint content is culturally and contextually appropriate for the target region.

Each section below is followed by an evidence table. Three types of evidence are cited: [QN] — verbatim quotes extracted from the benchmark paper; [WEB-N] — web sources consulted during assessment; [DN] / [DATASET-DN] — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

Input content is a HIGH-priority dimension per elicitation, and the deployment requires Indian everyday contexts (cricket, rupees, local markets, school fees) with explicit rejection of Western contexts (A2). MRBench's 192 instances are drawn entirely from Bridge and MathDial [Q41, Q52], neither of which is documented as containing Indian curriculum content, Indian numeral conventions, or India-specific scenarios. The paper notes that Bridge problems are 'fundamental mathematical concepts, including operations such as multiplication, addition' [Q45] in short Western-style framings, and MathDial is 'middle school-level mathematical reasoning' [Q49] without regional grounding. Web search confirms no Indian-context tutoring dialogue dataset exists [WEB-6, WEB-7], and GSM8K robustness research (Tomar et al. 2025) directly demonstrates measurable LLM performance drops when Western contexts are replaced with Indian ones [WEB-8] — establishing that this is not cosmetic but a measurable validity threat. There is no evidence of any India-relevant content in MRBench's instance pool.

ID	Evidence
Q41	"We have compiled mistake remediation benchmark, MRBench, from the Bridge (Wang et al., 2024a) and MathDial (Macina et al., 2023) datasets." (p.5)
Q45	"The dialogue context is typically short (few turns) and predominantly focused on fundamental mathematical concepts, including operations such as multiplication, addition, and so on." (p.5)
Q49	"Specifically, these conversations are grounded in middle school-level mathematical reasoning questions." (p.5)
Q52	"Next, for the 192 instances in MRBench (60 from Bridge and 132 from MathDial), we generated appropriate subsequent responses based on the conversation history and the last utterance, which contained confusions or mistakes, using seven state-of-the-art LLMs." (p.6)
WEB-6	No NCERT-aligned Grade 1–8 English-medium tutoring dialogue benchmark exists (https://arxiv.org/html/2404.13099v1)
WEB-7	HAWP is Hindi arithmetic word problems only, not tutoring dialogue (https://aclanthology.org/2022.lrec-1.373.pdf)
WEB-8	GSM8K robustness research shows Indian-context substitution causes measurable LLM performance drops (https://www.emergentmind.com/topics/gsm8k)

Strengths

- Instance pool is rigorously filtered for high quality, with a clear criterion that the student's last utterance contains a mistake [Q46, Q47, Q51]
- Combines elementary and middle-school-level dialogues providing some grade-banding diversity within math [Q97, Q143, Q144]
- Fundamental mathematical operations covered (addition, multiplication) overlap with NCERT primary content even if framing differs [Q45]

ID	Evidence
Q45	"The dialogue context is typically short (few turns) and predominantly focused on fundamental mathematical concepts, including operations such as multiplication, addition, and so on." (p.5)
Q46	"The original dataset consists of a total of 700 dialogues; we filtered 60 high-quality instances for MRBench." (p.5)
Q47	"Among the various criteria for selecting high-quality dialogues, the key one was that the student's last utterance (or last few utterances) should exhibit an error or confusion." (p.5)

Q51	"To further ensure the reliability of our benchmark, we manually inspected the data in order to retain only high-quality examples, which resulted in 132 instances for MRBench." (p.6)
Q97	"Combining both types of data in MRBench makes it both challenging and comprehensive." (p.8)
Q143	"These aspects highlight that including both datasets in MRBench ensures diversity and provides for a good mix of easy and difficult mathematical problems, making the benchmark both comprehensive and challenging." (p.14)
Q144	"The problems covered in the Bridge dataset are at the elementary school level, whereas those in MathDial are at the middle school level." (p.14)

Checklist

Checklist item definitions:

ID	Assessment Question
IC-1	Do inputs require region-specific cultural/geographic/dialectal knowledge?
IC-2	Does culturally sensitive content align with target culture?
IC-3	Flag inputs assuming Western-specific knowledge
IC-4	Regional annotator recruitment for culturally sensitive instances
IC-5	Document content issues harming content validity

Checklist responses:

IC-1: Yes — the deployment requires Indian cultural, geographic, and curricular grounding (NCERT topics, rupees, cricket, local markets, train timetables, school fees, festival contexts). Confirmed via elicitation A1, A2 and NCERT syllabus sources [WEB-2, WEB-4].

IC-2: No — there is no evidence that MRBench's source datasets contain any India-aligned cultural content. Bridge and MathDial are sourced from Western/general-context tutoring corpora [Q41, Q44, Q48] with no documentation of Indian inclusion.

IC-3: Likely yes. The paper notes Bridge dialogues focus on 'fundamental operations' [Q45] in short Western-framed contexts, and MathDial's middle-school reasoning problems [Q49] are not documented as having Indian grounding. GSM8K robustness research directly demonstrates that Western-to-Indian context substitution degrades LLM performance [WEB-8], so Western framings in MRBench's instances do not transfer cleanly.

IC-4: Not performed by the benchmark authors. INSUFFICIENT DOCUMENTATION on whether any regional annotators reviewed instance content for cultural fit; the four annotators' nationalities are undisclosed [Q31].

IC-5: Documented harms to content validity: (a) no rupee-denominated problems; (b) no Indian numeral conventions in problem statements; (c) no Indian everyday scenarios; (d) Western framings that empirically degrade LLM performance when transposed [WEB-8]. The deployment's HIGH-priority requirement for Indian-context word problems is entirely unmet.

ID	Evidence
Q31	"The annotation team consisted of two male and two female annotators, with all four annotators holding at least a post-graduate degree in Computer Science and being proficient in English." (p.5)
Q41	"We have compiled mistake remediation benchmark, MRBench, from the Bridge (Wang et al., 2024a) and MathDial (Macina et al., 2023) datasets." (p.5)
Q44	"The Bridge dataset (Wang et al., 2024a) comprises partial dialogue interactions between real human tutors and students at the elementary level, featuring two distinct human tutor responses (novice and expert)." (p.5)
Q45	"The dialogue context is typically short (few turns) and predominantly focused on fundamental mathematical concepts, including operations such as multiplication, addition, and so on." (p.5)
Q48	"The dialogues in the MathDial dataset (Macina et al., 2023) consist of complete multi-turn conversations between a real human tutor and an LLM acting as a student, where the tutor aims to remediate the student's mistakes." (p.5)
Q49	"Specifically, these conversations are grounded in middle school-level mathematical reasoning questions." (p.5)
Q93	"However, they struggle to provide appropriate guidance without revealing the answer because the mistakes are generally related to quite basic operations like addition or multiplication, often in a one-step type of mathematical problems." (p.8)

WEB-2	NCERT Grade 1–5 mathematics syllabus (place value with lakhs, spiral sequencing) (https://ncert.nic.in/pdf/syllabus/06Math%20(I-V).pdf)
WEB-4	NCERT upper-primary syllabus emphasizing abstraction, fractions, integers, algebra (https://cbseportal.com/NCERT/syllabus/maths-secondary)
WEB-6	No NCERT-aligned Grade 1–8 English-medium tutoring dialogue benchmark exists (https://arxiv.org/html/2404.13099v1)
WEB-8	GSM8K robustness research shows Indian-context substitution causes measurable LLM performance drops (https://www.emergentmind.com/topics/gsm8k)

Information Gaps

- No instance-level content audit of Bridge/MathDial for cultural framings is published
- No data on whether any of the 192 instances incidentally contain India-applicable framings

Requires Expert Verification

- Sample-level review by Indian educators to confirm degree of cultural distance in actual instance content
- Empirical measurement of LLM performance drop on MRBench instances when transposed to Indian context

Remediation

Gap 1: All 192 instances are drawn from Western-context Bridge and MathDial datasets with no documented Indian content; GSM8K research shows context substitution causes measurable LLM performance drops.

Recommendation: Construct a parallel evaluation set of NCERT-aligned dialogues with Indian everyday contexts (cricket, rupees, local markets, school fees, train timetables) authored or reviewed by Indian primary/middle-school teachers, and re-evaluate the seven LLMs on this set to quantify the regional performance gap.

Gap 2: DPDP Act 2023 / Rules 2025 impose verifiable parental consent and ban profiling/tracking of under-18 users [WEB-23, WEB-24], constraining how Indian student data can be collected for any supplementation effort.

Recommendation: Design any Indian content-collection or re-annotation effort under DPDP-compliant verifiable parental consent workflows; avoid behavioural-monitoring data collection for under-18 users and document data fiduciary obligations before deployment.

ID	Evidence
WEB-23	https://ksandk.com/data-protection-and-data-privacy/child-data-protection-under-dpdp-act-parental-consent-rules/
WEB-24	https://www.dpdpa.com/dpdpa2023/chapter-2/section9.html

Input Form (IF)

Score: 4/5 — Minor gaps | Confidence: high

Benchmark: MRBENCH | Context: Metropolitan India — Grades 1–8 Mathematics Tutoring (Student Population)

Input Form definition: Whether the input signal encoding (text, audio, image parameters) matches deployment conditions.

Each section below is followed by an evidence table. Three types of evidence are cited: [QN] — verbatim quotes extracted from the benchmark paper; [WEB-N] — web sources consulted during assessment; [DN] / [DATASET-DN] — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

Input form is a LOWER-priority dimension. MRBench is text-based English dialogue [Q13, Q133], encoded as conversation history with tutor/student turns and lengths measured in characters [Q147]. The deployment is also text-only English-medium via mobile or enterprise software and explicitly targets fluent English speakers; no modality, script, or language mismatch exists. Mobile internet penetration in metropolitan India is high (75%+ of children aged 5–14 use mobile phones, with metro rates higher) [WEB-10, WEB-12], and Android dominates with 92.4% market share [WEB-13], so text-based delivery on mobile is well-supported. The minor concerns are that conversation lengths and turn structures differ between Bridge and MathDial [Q141, Q142] — but this represents internal diversity rather than a target-context mismatch.

ID	Evidence
Q13	“Formally, let’s define the conversation history between a tutor and a student as $H = \{(T1, S1), (T2, S2), \dots, (Tt, St)\}$, where Ti represents the i -th response from the tutor, and Si represents the i -th response from the student. Let Sk denote the student’s most recent k utterances, where $k \in [1, \dots, t]$, containing a mistake or confusion. Then the objective of the tutor is to provide the most appropriate response T_{t+1} to address this mistake or confusion.” (p.3)
Q133	“The prompt template used to generate responses from the seven considered LLMs for both the Bridge and MathDial datasets is shown in Figure 2.” (p.12)
Q141	“It can be observed that the conversation history and response lengths from different LLMs and humans are generally shorter in the Bridge dataset compared to the MathDial dataset.” (p.14)
Q142	“Additionally, the number of turns differs between them.” (p.14)
Q147	“In all cases, length is estimated using the number of characters.” (p.14)
WEB-10	Mobile internet broadly available in Indian metros for text delivery (https://datareportal.com/reports/digital-2024-india)
WEB-12	~65% of Indian children aged 5–14 used mobile phones in 2023 (https://www.dataforindia.com/comm-tech/)
WEB-13	Android dominates Indian smartphone OS market at 92.4% in 2025 (https://www.imarcgroup.com/india-smartphone-market)

Strengths

- Text-only English format perfectly matches the deployment’s text-based English-medium delivery [Q13]
- Conversation-history input format directly matches a mobile-app tutoring dialogue use case [Q13, Q133]
- Internal diversity in conversation length across Bridge and MathDial provides some signal-distribution variation [Q141, Q142]

ID	Evidence
Q13	“Formally, let’s define the conversation history between a tutor and a student as $H = \{(T1, S1), (T2, S2), \dots, (Tt, St)\}$, where Ti represents the i -th response from the tutor, and Si represents the i -th response from the student. Let Sk denote the student’s most recent k utterances, where $k \in [1, \dots, t]$, containing a mistake or confusion. Then the objective of the tutor is to provide the most appropriate response T_{t+1} to address this mistake or confusion.” (p.3)
Q133	“The prompt template used to generate responses from the seven considered LLMs for both the Bridge and MathDial datasets is shown in Figure 2.” (p.12)
Q141	“It can be observed that the conversation history and response lengths from different LLMs and humans are generally shorter in the Bridge dataset compared to the MathDial dataset.” (p.14)
Q142	“Additionally, the number of turns differs between them.” (p.14)

Checklist

Checklist item definitions:

ID	Assessment Question
IF-1	Compare signal distributions between source and target
IF-2	Check if regional infrastructure supports data capture specs
IF-3	Identify domain-specific form differences
IF-4	Document form mismatches harming external validity

Checklist responses:

IF-1: Signal distributions are aligned: both benchmark and deployment are English-medium text dialogues with conversation history and turn structure [Q13, Q133]. No script or language mismatch.

IF-2: Regional infrastructure supports the format. India has approximately 750M+ mobile internet users with strong metropolitan 5G coverage [WEB-10, WEB-11], and Android dominates at 92.4% market share [WEB-13].

IF-3: Domain-specific form is well-matched: tutor/student dialogue turns over text are the deployment's actual use case. No multimodal (audio, image, handwriting) requirement is specified.

IF-4: Minor gap: MathDial does not provide conversation topic labels [Q145], which could complicate routing/curriculum-tagging in deployment. Otherwise no form mismatch identified.

ID	Evidence
Q13	"Formally, let's define the conversation history between a tutor and a student as $H = \{(T1, S1), (T2, S2), \dots, (Tt, St)\}$, where Ti represents the i -th response from the tutor, and Si represents the i -th response from the student. Let Sk denote the student's most recent k utterances, where $k \in [1, \dots, t]$, containing a mistake or confusion. Then the objective of the tutor is to provide the most appropriate response T_{t+1} to address this mistake or confusion." (p.3)
Q133	"The prompt template used to generate responses from the seven considered LLMs for both the Bridge and MathDial datasets is shown in Figure 2." (p.12)
Q145	"The conversation topic is not provided in MathDial." (p.14)
WEB-10	Mobile internet broadly available in Indian metros for text delivery (https://datareportal.com/reports/digital-2024-india)
WEB-11	Mobile broadband dominates Indian internet access; 5G rollout in metros since 2022 (https://muftinternet.com/blog/usage-statistics-internet-and-mobile-users-in-india-2025/)
WEB-13	Android dominates Indian smartphone OS market at 92.4% in 2025 (https://www.imarcgroup.com/india-smartphone-market)

Information Gaps

- No data on actual deployment device distribution for Grade 1–8 metropolitan students specifically (smartphone vs. tablet vs. laptop share)

Requires Expert Verification

- Whether deployment intends to support voice/audio modalities later, which would require benchmark extension

Output Ontology (OO)

Score: 1/5 — Serious concern | Confidence: high

Benchmark: MRBENCH | Context: Metropolitan India — Grades 1–8 Mathematics Tutoring (Student Population)

Output Ontology definition: Whether the benchmark's output categories and scoring criteria reflect regionally valid decision boundaries.

Each section below is followed by an evidence table. Three types of evidence are cited: [QN] — verbatim quotes extracted from the benchmark paper; [WEB-N] — web sources consulted during assessment; [DN] / [DATASET-DN] — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

Output ontology is a HIGH-priority dimension per elicitation. MRBench scores responses across eight pedagogical dimensions using a three-tier label system [Q58, Q59] aggregated into DAMR [Q64]. Critically, the deployment's primary success criterion is student satisfaction, and elicitation A4 explicitly confirms that a pedagogically correct but culturally distant response would NOT count as satisfying — yet MRBench has no output category for cultural appropriateness, NCERT curriculum correctness, or Indian classroom register. The Tutor Tone dimension trivially collapses to a non-offensiveness boundary (DAMR reaches 100% when neutral and encouraging are merged into 'Non-offensive') [Q90], failing to operationalize the warmth/encouragement register required by the deployment. The authors explicitly acknowledge the taxonomy 'does not consider the tutoring dialogues' impact on overall student learning' [Q120, Q121] and treats dimensions as orthogonal/independent [Q30, Q86] despite real interdependencies [Q116]. A response could therefore receive a high DAMR score while being culturally unsuitable for metropolitan Indian students, with no flag in the output schema. Web search confirms no validation study links MRBench scores to Indian student satisfaction [WEB-9, WEB-17, WEB-18].

ID	Evidence
Q30	"Although there are inherent interdependencies among the proposed dimensions of the taxonomy (e.g., a response that reveals the answer is less likely to be actionable, and vice versa), we explicitly instructed all annotators to treat each dimension as independent and orthogonal to minimize confounding factors and potential biases during the annotation process." (p.4)
Q58	"Each dimension was annotated using a three-tier labeling system (see Figure 1 and Table 4)." (p.6)
Q59	"For instance, the 'mistake identification' dimension employed the following labels: (i) yes, (ii) to some extent, and (iii) no." (p.6)
Q64	"We utilize two key metrics to quantitatively assess the pedagogical effectiveness of LLMs and for comparative analysis: (1) Desired Annotation Match Rate (DAMR): This metric quantifies the percentage of responses from each human or LLM-based tutor that received the desired annotation labels." (p.6)
Q86	"This further demonstrates the need to treat the dimensions as independent." (p.8)
Q90	"When we combine these two labels into 'Non-offensive', the DAMR score reaches 100% as we observe no offensive responses from any LLMs or humans." (p.8)
Q116	"However, in practice, these dimensions may be inherently interrelated and may influence one another." (p.9)
Q120	"However, one of the limitations is that it does not consider the tutoring dialogues' impact on the overall student learning." (p.9)
Q121	"Specifically, the annotation pertains to the individual tutor turns within educational dialogues, which restricts our understanding of broader implications on student learning processes and learning gains, typically observed after a conversation concludes." (p.9)
WEB-9	LLM-in-education meta-analysis documents motivational effects but not for Indian school-age math (https://arxiv.org/html/2507.02180v1)
WEB-17	BEA 2025 Shared Task extended MRBench without Indian rater validation (https://arxiv.org/html/2507.10579v1)
WEB-18	NAACL 2025 BEA proceedings: no participating system documented Indian adaptation (https://aclanthology.org/2025.naacl-long.57.pdf)

Strengths

- Eight orthogonal pedagogical dimensions provide fine-grained per-behavior signal that supports targeted diagnosis (e.g., GPT-4 reveals answer 47% of time) [Q69]

- Three-tier labeling system [Q58, Q59] captures partial credit for ambiguous cases
- Explicit 'desired label' for each dimension makes the decision rule transparent and auditable [Q65, Q128]
- Domain-agnostic NLG metrics (BLEU, BERTScore) are explicitly rejected for being inappropriate to pedagogical evaluation [Q4, Q7, Q8, Q10] — a defensible methodological choice

ID	Evidence
Q4	"General domain-agnostic natural language generation (NLG) metrics (Lin, 2004; Popovic', 2017; Post, 2018; Gao et al., 2020; Liu et al., 2023) are not well-suited for this context, as most of them fail to account for pedagogical values and require gold references, which are often not available, especially in online interactions." (p.1)
Q7	"General domain-agnostic natural language generation (NLG) metrics like BLEU (Papineni et al., 2002), BERTScore (Lin, 2004), DialogRPT (Gao et al., 2020), and so on have been used as proxies to measure the coherence and human-likeness of AI tutor responses." (p.2)
Q8	"However, these metrics do not account for pedagogical values (Jurenka et al., 2024; Liu et al., 2024) and often require a ground truth answer to evaluate matching responses." (p.2)
Q10	"Additionally, these metrics can be easily manipulated; for instance, simple responses like "Hello" or "teacher:" (Baladón et al., 2023; Jurenka et al., 2024) can inflate scores." (p.2)
Q58	"Each dimension was annotated using a three-tier labeling system (see Figure 1 and Table 4)." (p.6)
Q59	"For instance, the 'mistake identification' dimension employed the following labels: (i) yes, (ii) to some extent, and (iii) no." (p.6)
Q65	"The desired labels for each dimension are detailed in Table 2." (p.6)
Q69	"Both these LLMs perform well in identifying students' mistakes and their exact location, with Llama-3.1-405B having a slight edge as GPT-4 reveals the answer approximately 47% of the time, making its responses less actionable and impacting student's learning experience." (p.7)
Q128	"The definitions, associated labels, and the desired labels for each dimension of the proposed taxonomy are provided in Table 4." (p.12)

Checklist

Checklist item definitions:

ID	Assessment Question
OO-1	Review output label categories for regional relevance
OO-2	Identify missing regionally specific categories
OO-3	Flag categories encoding non-regional values
OO-4	Consider stakeholder-driven taxonomy redesign
OO-5	Document taxonomy issues harming structural/content validity

Checklist responses:

OO-1: The eight output categories (mistake identification, mistake location, guidance, actionability, coherence, tone, human-likeness, revealing the answer) are pedagogically relevant for any tutoring context but are insufficient for the Indian deployment, which additionally requires cultural appropriateness and curriculum alignment categories (A4, A1).

OO-2: Missing categories: (a) cultural-register/warmth specific to Indian classroom norms; (b) NCERT curriculum correctness; (c) Indian numeral convention compliance; (d) student satisfaction proxy. Tutor Tone is operationalized as non-offensiveness [Q90], which does not capture warmth required by elicitation A4.

OO-3: Tutor Tone's collapse to 'Non-offensive' achieving 100% DAMR [Q90] encodes a Western default of register adequacy that does not match Indian classroom warmth norms (per A4). The 'human-likeness' dimension is similarly under-specified for cultural register.

OO-4: Stakeholder-driven taxonomy redesign is needed: at minimum, adding a cultural-appropriateness dimension and a curriculum-alignment dimension, with desired labels validated against Indian educators and students.

OO-5: Documented harms: structural validity is harmed because the deployment's success construct (student satisfaction including cultural fit) has structure not captured by the eight dimensions; content validity is harmed by missing categories;

external validity is harmed because high DAMR scores cannot be assumed to predict deployment-context satisfaction. Authors themselves acknowledge taxonomy does not address overall learning impact [Q120, Q121].

ID	Evidence
Q14	"The evaluation taxonomy detailed in Section 4 assesses the appropriateness of the Tt+1 response across eight key pedagogical dimensions." (p.3)
Q64	"We utilize two key metrics to quantitatively assess the pedagogical effectiveness of LLMs and for comparative analysis: (1) Desired Annotation Match Rate (DAMR): This metric quantifies the percentage of responses from each human or LLM-based tutor that received the desired annotation labels." (p.6)
Q89	"Our findings on the Tutor Tone align with those of Wang et al. (2024a) – in task-oriented conversations, AI tutors tend to be more Neutral than Encouraging." (p.8)
Q90	"When we combine these two labels into "Non-offensive", the DAMR score reaches 100% as we observe no offensive responses from any LLMs or humans." (p.8)
Q116	"However, in practice, these dimensions may be inherently interrelated and may influence one another." (p.9)
Q118	"We acknowledge that the proposed taxonomy will need to be verified on and likely adapted if applied to other tasks such as concept learning, and to subjects other than mathematics." (p.9)
Q119	"The current taxonomy and annotation scheme focus on the appropriateness of the tutor responses." (p.9)
Q120	"However, one of the limitations is that it does not consider the tutoring dialogues' impact on the overall student learning." (p.9)
Q121	"Specifically, the annotation pertains to the individual tutor turns within educational dialogues, which restricts our understanding of broader implications on student learning processes and learning gains, typically observed after a conversation concludes." (p.9)
WEB-9	LLM-in-education meta-analysis documents motivational effects but not for Indian school-age math (https://arxiv.org/html/2507.02180v1)
WEB-17	BEA 2025 Shared Task extended MRBench without Indian rater validation (https://arxiv.org/html/2507.10579v1)

Information Gaps

- No published correlation between DAMR scores and student satisfaction for any population
- No empirical specification of what Indian-classroom 'warmth register' looks like at the response level

Requires Expert Verification

- Definition and operationalization of a cultural-appropriateness output dimension by Indian educators and students
- Calibration of desired labels for Tutor Tone against Indian metropolitan student preferences

Remediation

Gap 1: Tutor Tone dimension collapses to non-offensiveness [Q90]; no output category captures cultural appropriateness or student satisfaction, both confirmed essential by elicitation A4.

Recommendation: Replace or supplement the Tutor Tone dimension with a multi-level cultural-register dimension calibrated to Indian classroom warmth norms; add a proxy output dimension for predicted student satisfaction validated against in-deployment engagement signals.

ID	Evidence
Q90	"When we combine these two labels into "Non-offensive", the DAMR score reaches 100% as we observe no offensive responses from any LLMs or humans." (p.8)

Output Content (OC)

Score: 2/5 — Significant gaps | Confidence: high

Benchmark: MRBENCH | Context: Metropolitan India — Grades 1–8 Mathematics Tutoring (Student Population)

Output Content definition: Whether ground-truth labels align with the judgments of regional stakeholders.

Each section below is followed by an evidence table. Three types of evidence are cited: [QN] — verbatim quotes extracted from the benchmark paper; [WEB-N] — web sources consulted during assessment; [DN] / [DATASET-DN] — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

Output content is MODERATE-priority per elicitation, with A3 indicating partial alignment expected between Indian student satisfaction and benchmark annotators but anticipating divergence on tone, warmth, and cultural fit. MRBench's annotators are four post-graduate Computer Science students proficient in English [Q31] with no disclosed nationality, cultural background, or teaching experience [Q32]; teaching experience was explicitly NOT required [Q32]. Inter-annotator agreement is substantial (Fleiss' kappa 0.65 pilot, Cohen's kappa 0.71 full-scale) [Q39, Q61], indicating internal label reliability — but reliability does not address whether labels reflect Indian stakeholder perspectives. No Indian educators, parents, or students contributed to ground-truth labels, and no cross-cultural validation has been performed; the BEA 2025 Shared Task extended MRBench without addressing this [WEB-17, WEB-18]. The deployment's actual judges (human tutor co-authors and student engagement signals per A3) are structurally different from the benchmark's annotator pool. LLM-as-judge alternatives are explicitly unreliable [Q107, Q109, Q114], a finding corroborated by 2024–2025 literature [WEB-19, WEB-20], so they cannot substitute for human re-annotation.

ID	Evidence
Q31	"The annotation team consisted of two male and two female annotators, with all four annotators holding at least a post-graduate degree in Computer Science and being proficient in English." (p.5)
Q32	"We note that for this study, we do not require annotators to have direct teaching experience, as understanding of the mathematical tasks at the middle school level and being able to judge the responses from the perspective of a potential user of such AI tutors (or a student), rather than specifically a teacher, is sufficient." (p.5)
Q39	"To measure inter-annotator agreement, we computed Fleiss' kappa value, which for this annotation experiment equals 0.65, indicating substantial agreement." (p.5)
Q61	"The annotators reached an average Cohen's kappa score of 0.71, which indicates substantial inter-annotator agreement (McHugh, 2012)." (p.6)
Q107	"Across both LLMs, it can be observed that most of the correlation scores are negative (except for the human-likeness dimension), indicating that the annotations from the LLMs are unreliable for the challenging pedagogical dimensions." (p.8)
Q109	"We believe both LLMs have a limited understanding of rich pedagogical concepts, as they were not specifically trained on pedagogically rich datasets." (p.8)
Q114	"We also assess the feasibility of LLMs as evaluators in this context by correlating their judgements with human annotations, indicating that they are often unreliable." (p.9)
WEB-17	BEA 2025 Shared Task extended MRBench without Indian rater validation (https://arxiv.org/html/2507.10579v1)
WEB-18	NAACL 2025 BEA proceedings: no participating system documented Indian adaptation (https://aclanthology.org/2025.naacl-long.57.pdf)
WEB-19	2025 MDPI: LLM evaluator inconsistency confirmed for educational evaluation (https://www.mdpi.com/2076-3417/15/2/671)
WEB-20	2025 study: LLM and human evaluator scores share little common signal (https://arxiv.org/pdf/2508.02442)

Strengths

- Strong inter-annotator agreement (Fleiss' kappa 0.65, Cohen's kappa 0.71) demonstrates internal label reliability [Q39, Q61]
- In-house annotation with comprehensive training and validation pilot avoids crowdsourcing quality issues [Q33, Q34, Q35]

- Expert and Novice human responses included as comparison baselines, with Expert treated as gold standard [Q68, Q78]
- Authors explicitly assess and report the unreliability of LLM critics [Q107, Q109, Q114] — providing honest evidence that automated re-annotation is not a substitute

ID	Evidence
Q33	"To control the annotation workflow and ensure quality, we opted not to use public annotation outsourcing platforms such as Prolific or MTurk, which allowed us to implement rigorous training protocols and a robust validation mechanism for the annotations." (p.5)
Q34	"First, we provided all annotators with comprehensive training, including an interactive training document (see Section C for more details) and oral instructions." (p.5)
Q35	"Following this, we conducted validation pilot study to evaluate the annotation quality and the annotators' understanding of the instructions before rolling out the large-scale human evaluation detailed in Section 5.2." (p.5)
Q39	"To measure inter-annotator agreement, we computed Fleiss' kappa value, which for this annotation experiment equals 0.65, indicating substantial agreement." (p.5)
Q61	"The annotators reached an average Cohen's kappa score of 0.71, which indicates substantial inter-annotator agreement (McHugh, 2012)." (p.6)
Q68	"We consider human-based evaluations as gold standard." (p.7)
Q78	"We also investigated the pedagogical value of human responses for both Novice and Expert." (p.7)
Q107	"Across both LLMs, it can be observed that most of the correlation scores are negative (except for the human-likeness dimension), indicating that the annotations from the LLMs are unreliable for the challenging pedagogical dimensions." (p.8)
Q109	"We believe both LLMs have a limited understanding of rich pedagogical concepts, as they were not specifically trained on pedagogically rich datasets." (p.8)
Q114	"We also assess the feasibility of LLMs as evaluators in this context by correlating their judgements with human annotations, indicating that they are often unreliable." (p.9)

Checklist

Checklist item definitions:

ID	Assessment Question
OC-1	Do ground-truth labels reflect regional stakeholder perspectives?
OC-2	Assess potential annotator-regional population disagreement
OC-3	Review annotator demographics from documentation
OC-4	Consider label re-annotation by regional annotators
OC-5	Review aggregation methods for minority perspective erasure
OC-6	Document label issues harming convergent/external validity

Checklist responses:

OC-1: No — labels reflect the perspectives of four Computer Science post-graduates [Q31] with undisclosed nationality and no teaching experience [Q32]. There is no documented Indian stakeholder representation in the annotator pool.

OC-2: Disagreement is anticipated, particularly on Tutor Tone and human-likeness dimensions where cultural register expectations differ. A3 indicates partial alignment expected; A4 confirms divergence on cultural appropriateness. The paper itself notes Tutor Tone's neutral/encouraging distinction collapses trivially [Q90], suggesting the dimension lacks discriminative power on the very axis where Indian register differences would matter.

OC-3: Annotator demographics partially disclosed: gender (2 male, 2 female), education (post-graduate Computer Science), and English proficiency [Q31]. Nationality, cultural background, and teaching experience are NOT disclosed; teaching experience explicitly not required [Q32]. This is below the level of detail recommended in Datasheets/Data Statements for cross-cultural validity assessment.

OC-4: Re-annotation by representative regional annotators (Indian educators, Indian metropolitan students or parents) is needed for the deployment, particularly on tone, human-likeness, and any added cultural-appropriateness dimension. No such

re-annotation exists [WEB-17].

OC-5: Aggregation method (majority across four annotators with Cohen's kappa reported [Q57, Q61]) is reasonable for internal reliability but provides no mechanism to detect minority cultural perspectives; with all four annotators from a single demographic profile, minority perspectives are structurally absent.

OC-6: Documented harms to convergent validity (labels not validated against Indian raters) and external validity (judgments may not generalize to Indian metropolitan student satisfaction). The deployment's reliance on human tutor co-authorship partially mitigates this for the deployed system, but the benchmark itself cannot be relied on as a satisfaction proxy.

ID	Evidence
Q31	"The annotation team consisted of two male and two female annotators, with all four annotators holding at least a post-graduate degree in Computer Science and being proficient in English." (p.5)
Q32	"We note that for this study, we do not require annotators to have direct teaching experience, as understanding of the mathematical tasks at the middle school level and being able to judge the responses from the perspective of a potential user of such AI tutors (or a student), rather than specifically a teacher, is sufficient." (p.5)
Q39	"To measure inter-annotator agreement, we computed Fleiss' kappa value, which for this annotation experiment equals 0.65, indicating substantial agreement." (p.5)
Q57	"A total of 192 instances were annotated, with 40 of those annotated independently by two annotators (10 instances from Bridge and 30 from MathDial) allowing us to calculate pairwise inter-annotator agreement." (p.6)
Q61	"The annotators reached an average Cohen's kappa score of 0.71, which indicates substantial inter-annotator agreement (McHugh, 2012)." (p.6)
Q90	"When we combine these two labels into "Non-offensive", the DAMR score reaches 100% as we observe no offensive responses from any LLMs or humans." (p.8)
Q107	"Across both LLMs, it can be observed that most of the correlation scores are negative (except for the human-likeness dimension), indicating that the annotations from the LLMs are unreliable for the challenging pedagogical dimensions." (p.8)
WEB-17	BEA 2025 Shared Task extended MRBench without Indian rater validation (https://arxiv.org/html/2507.10579v1)
WEB-18	NAACL 2025 BEA proceedings: no participating system documented Indian adaptation (https://aclanthology.org/2025.naacl-long.57.pdf)
WEB-19	2025 MDPI: LLM evaluator inconsistency confirmed for educational evaluation (https://www.mdpi.com/2076-3417/15/2/671)
WEB-20	2025 study: LLM and human evaluator scores share little common signal (https://arxiv.org/pdf/2508.02442)

Information Gaps

- Annotator nationality and cultural background not disclosed in the paper
- No cross-cultural validation study with Indian raters has been published
- No published comparison of MRBench labels to Indian human-tutor co-author judgments

Requires Expert Verification

- Re-annotation of a sample of MRBench instances by Indian educators and metropolitan Grade 1–8 students to quantify divergence on tone and human-likeness
- Calibration of the deployment's human tutor co-author judgments against MRBench gold labels

Remediation

Gap 1: Ground-truth labels reflect four Computer Science post-graduates with undisclosed nationality and no teaching experience; no Indian rater validation in MRBench or BEA 2025 Shared Task.

Recommendation: Recruit a panel of Indian educators (NCERT-familiar primary and middle-school math teachers) and a smaller panel of metropolitan Grade 1–8 students or parents to re-annotate at least the 40 double-annotated instances [Q57] across all dimensions, then quantify divergence from MRBench gold labels and report dimension-specific recalibration.

ID	Evidence
----	----------

Q57

“A total of 192 instances were annotated, with 40 of those annotated independently by two annotators (10 instances from Bridge and 30 from MathDial) allowing us to calculate pairwise inter-annotator agreement.” (p.6)

Output Form (OF)

Score: 3/5 — Moderate gaps | Confidence: medium

Benchmark: MRBENCH | Context: Metropolitan India — Grades 1–8 Mathematics Tutoring (Student Population)

Output Form definition: Whether the expected output modality matches regional deployment needs and accessibility.

Each section below is followed by an evidence table. Three types of evidence are cited: **[QN]** — verbatim quotes extracted from the benchmark paper; **[WEB-N]** — web sources consulted during assessment; **[DN]** / **[DATASET-DN]** — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

Output form is a LOWER-priority dimension. MRBench's evaluation outputs are categorical labels per dimension and aggregate DAMR/AC scores [\[Q64, Q66\]](#), not free-form text — providing actionable per-dimension diagnostics (e.g., GPT-4 reveals answer 47% of the time [\[Q69\]](#)) that map cleanly to deployment QA tasks. The deployment's actual product output is free-form English tutor text, which the benchmark does not directly score for free-form quality but indirectly assesses via the categorical taxonomy. Domain-agnostic NLG metrics (BLEU, BERTScore) are explicitly rejected as inappropriate [\[Q4, Q7, Q8, Q10\]](#) — a defensible methodological choice. The chief weakness is that DAMR does not operationalize student satisfaction or engagement, and the LLM-critic alternative is unreliable [\[Q107, Q109, Q114\]](#) — confirmed by 2024–2025 literature [\[WEB-19, WEB-20\]](#) — meaning automated re-evaluation of the deployed system using MRBench-style critics is not a valid substitute for human judgment. Partial mitigation comes from the deployment's human tutor co-authorship layer.

ID	Evidence
Q4	"General domain-agnostic natural language generation (NLG) metrics (Lin, 2004; Popovic', 2017; Post, 2018; Gao et al., 2020; Liu et al., 2023) are not well-suited for this context, as most of them fail to account for pedagogical values and require gold references, which are often not available, especially in online interactions." (p.1)
Q7	"General domain-agnostic natural language generation (NLG) metrics like BLEU (Papineni et al., 2002), BERTScore (Lin, 2004), DialogRPT (Gao et al., 2020), and so on have been used as proxies to measure the coherence and human-likeness of AI tutor responses." (p.2)
Q8	"However, these metrics do not account for pedagogical values (Jurenka et al., 2024; Liu et al., 2024) and often require a ground truth answer to evaluate matching responses." (p.2)
Q10	"Additionally, these metrics can be easily manipulated; for instance, simple responses like "Hello" or "teacher:" (Baladón et al., 2023; Jurenka et al., 2024) can inflate scores." (p.2)
Q64	"We utilize two key metrics to quantitatively assess the pedagogical effectiveness of LLMs and for comparative analysis: (1) Desired Annotation Match Rate (DAMR): This metric quantifies the percentage of responses from each human or LLM-based tutor that received the desired annotation labels." (p.6)
Q66	"(2) Annotation Correlation (AC): This metric is based on Pearson's correlation (Sedgwick, 2012), and it estimates the correlation between LLM-generated and human annotations (Kim et al., 2024), allowing us to assess the reliability of LLMs as evaluators in the context of student mistake remediation." (p.6)
Q69	"Both these LLMs perform well in identifying students' mistakes and their exact location, with Llama-3.1-405B having a slight edge as GPT-4 reveals the answer approximately 47% of the time, making its responses less actionable and impacting student's learning experience." (p.7)
Q107	"Across both LLMs, it can be observed that most of the correlation scores are negative (except for the human-likeness dimension), indicating that the annotations from the LLMs are unreliable for the challenging pedagogical dimensions." (p.8)
Q109	"We believe both LLMs have a limited understanding of rich pedagogical concepts, as they were not specifically trained on pedagogically rich datasets." (p.8)
Q114	"We also assess the feasibility of LLMs as evaluators in this context by correlating their judgements with human annotations, indicating that they are often unreliable." (p.9)
WEB-19	2025 MDPI: LLM evaluator inconsistency confirmed for educational evaluation (https://www.mdpi.com/2076-3417/15/2/671)
WEB-20	2025 study: LLM and human evaluator scores share little common signal (https://arxiv.org/pdf/2508.02442)

Strengths

- DAMR provides interpretable per-dimension percentages directly usable for model comparison and deployment QA [Q64, Q67]
- Per-dataset (Bridge vs. MathDial) reporting [Q149, Q150] enables grade-band-specific diagnostics roughly aligned with primary vs. upper-primary
- Explicit, well-reasoned rejection of NLG metrics that would be inappropriate for tutoring evaluation [Q4, Q7, Q8, Q10]
- Honest reporting of LLM-critic unreliability [Q107, Q109, Q114] prevents users from over-relying on automated re-evaluation

ID	Evidence
Q4	“General domain-agnostic natural language generation (NLG) metrics (Lin, 2004; Popovic´, 2017; Post, 2018; Gao et al., 2020; Liu et al., 2023) are not well-suited for this context, as most of them fail to account for pedagogical values and require gold references, which are often not available, especially in online interactions.” (p.1)
Q7	“General domain-agnostic natural language generation (NLG) metrics like BLEU (Papineni et al., 2002), BERTScore (Lin, 2004), DialogRPT (Gao et al., 2020), and so on have been used as proxies to measure the coherence and human-likeness of AI tutor responses.” (p.2)
Q8	“However, these metrics do not account for pedagogical values (Jurenka et al., 2024; Liu et al., 2024) and often require a ground truth answer to evaluate matching responses.” (p.2)
Q10	“Additionally, these metrics can be easily manipulated; for instance, simple responses like "Hello" or "teacher:" (Baladón et al., 2023; Jurenka et al., 2024) can inflate scores.” (p.2)
Q64	“We utilize two key metrics to quantitatively assess the pedagogical effectiveness of LLMs and for comparative analysis: (1) Desired Annotation Match Rate (DAMR): This metric quantifies the percentage of responses from each human or LLM-based tutor that received the desired annotation labels.” (p.6)
Q67	“Table 3 shows DAMR scores for each LLM across all eight dimensions.” (p.7)
Q107	“Across both LLMs, it can be observed that most of the correlation scores are negative (except for the human-likeness dimension), indicating that the annotations from the LLMs are unreliable for the challenging pedagogical dimensions.” (p.8)
Q109	“We believe both LLMs have a limited understanding of rich pedagogical concepts, as they were not specifically trained on pedagogically rich datasets.” (p.8)
Q114	“We also assess the feasibility of LLMs as evaluators in this context by correlating their judgements with human annotations, indicating that they are often unreliable.” (p.9)
Q149	“Pedagogical ability assessment of different LLMs using the DAMR scores (in %) across eight evaluation dimensions with human evaluation on the Bridge data.” (p.16)
Q150	“Pedagogical ability assessment of different LLMs using the DAMR scores (in %) across eight evaluation dimensions with human evaluation on the MathDial data.” (p.16)

Checklist

Checklist item definitions:

ID	Assessment Question
OF-1	Does output modality match regional deployment needs?
OF-2	Assess text-to-speech availability for speech output
OF-3	Consider literacy rates and accessibility requirements
OF-4	Document form mismatches harming external validity

Checklist responses:

OF-1: Output modality (categorical label scoring + aggregate metrics) is not the deployment's product output (free-form tutor text), but it serves as an evaluation signal that maps onto the deployment's per-dimension QA needs. The match is indirect but adequate as evaluation signal.

OF-2: Not applicable — neither benchmark nor deployment requires text-to-speech. INSUFFICIENT DOCUMENTATION on whether deployment plans audio extension; assumed text-only per regional context.

OF-3: Target population is described as fluent English speakers in metropolitan India [regional context]. Mobile internet penetration is high [WEB-10, WEB-12], so accessibility constraints around text output are minimal. INSUFFICIENT DOCUMENTATION on accessibility for students with disabilities in the deployment's target population.

OF-4: Documented form mismatch: DAMR does not operationalize student satisfaction or engagement, which is the deployment's primary success criterion. Automated LLM critics are unreliable [Q107, Q114] and cannot substitute for human judgment in production. External validity is partially harmed: benchmark scores will not transparently predict deployment satisfaction outcomes.

ID	Evidence
Q64	"We utilize two key metrics to quantitatively assess the pedagogical effectiveness of LLMs and for comparative analysis: (1) Desired Annotation Match Rate (DAMR): This metric quantifies the percentage of responses from each human or LLM-based tutor that received the desired annotation labels." (p.6)
Q67	"Table 3 shows DAMR scores for each LLM across all eight dimensions." (p.7)
Q107	"Across both LLMs, it can be observed that most of the correlation scores are negative (except for the human-likeness dimension), indicating that the annotations from the LLMs are unreliable for the challenging pedagogical dimensions." (p.8)
Q114	"We also assess the feasibility of LLMs as evaluators in this context by correlating their judgements with human annotations, indicating that they are often unreliable." (p.9)
WEB-10	Mobile internet broadly available in Indian metros for text delivery (https://datareportal.com/reports/digital-2024-india)
WEB-12	~65% of Indian children aged 5–14 used mobile phones in 2023 (https://www.dataforindia.com/comm-tech/)
WEB-19	2025 MDPI: LLM evaluator inconsistency confirmed for educational evaluation (https://www.mdpi.com/2076-3417/15/2/671)
WEB-20	2025 study: LLM and human evaluator scores share little common signal (https://arxiv.org/pdf/2508.02442)

Information Gaps

- No documented mapping between DAMR scores and student satisfaction outcomes
- No published deployment study using MRBench scores as a production QA signal

Requires Expert Verification

- Whether the deployment's human tutor co-authors can effectively use DAMR-style per-dimension reports for QA review
- Whether deployment satisfaction signals can be calibrated to MRBench dimensions via instrumentation

Remediation

Gap 1: DAMR does not operationalize student satisfaction; LLM critics are unreliable [Q107, Q114] and cannot substitute for human judgment, confirmed by 2024–2025 literature [WEB-19, WEB-20].

Recommendation: Treat MRBench scores as one input among several to a deployment QA dashboard that also incorporates human tutor co-author ratings and student engagement signals, rather than as a sufficient indicator of deployment readiness; do not use Prometheus2/Llama-3.1-8B critics for production re-evaluation.

ID	Evidence
Q107	"Across both LLMs, it can be observed that most of the correlation scores are negative (except for the human-likeness dimension), indicating that the annotations from the LLMs are unreliable for the challenging pedagogical dimensions." (p.8)
Q114	"We also assess the feasibility of LLMs as evaluators in this context by correlating their judgements with human annotations, indicating that they are often unreliable." (p.9)
WEB-19	2025 MDPI: LLM evaluator inconsistency confirmed for educational evaluation (https://www.mdpi.com/2076-3417/15/2/671)
WEB-20	2025 study: LLM and human evaluator scores share little common signal (https://arxiv.org/pdf/2508.02442)