

Output Content (OC)

Score: 1/5 — Serious concern | Confidence: high

Benchmark: CRISISLTLSUM | Context: Mexico City and Central America Bilingual Crisis Journalism

Output Content definition: Whether ground-truth labels align with the judgments of regional stakeholders.

Each section below is followed by an evidence table. Three types of evidence are cited: **[QN]** — verbatim quotes extracted from the benchmark paper; **[WEB-N]** — web sources consulted during assessment; **[DN]** / **[DATASET-DN]** — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

Ground-truth labels were determined by majority vote among three MTurk workers **[Q70]** recruited under QC1, which restricted workers to 'countries where native English speakers are most likely to be found' **[Q61, Q105]**. Reference summaries were written by the two workers who agreed most on extraction labels **[Q87, Q88]**. While inter-annotator agreement was high (90.06% **[Q77]**, 91.77% against experts **[Q78, Q79]**), no Spanish-speaking annotators, no journalists, and no individuals with Latin American crisis knowledge contributed to ground truth. The elicitation explicitly confirms that journalistic relevance judgments, attribution norms, and summary scope differ from US conventions (Q3/A3). Therefore the labels and reference summaries cannot be assumed to reflect Mexican/Central American journalistic stakeholder perspectives. The aggregation method (majority vote on binary labels **[Q70]**; selecting the two most-agreeing workers' summaries **[Q87, Q88]**) further erases minority perspectives — a particular concern when no annotator represents the target population at all. Documented Mexican-context phenomena like #Verificado19S verification posts and community solidarity posts **[WEB-23]** would likely be judged differently by Mexican journalists than by Anglophone MTurk workers.

ID	Evidence
Q3	"A timeline is a chronologically sorted set of posts, where each brings in new information or updates about an ongoing event (such as a fire, storm, or traffic incident)." (p.1)
Q61	"First, we use location restriction (QC1) to limit the pool of workers to countries where native English speakers are most likely to be found." (p.5)
Q70	"To make the ground-truth on whether each tweet is part of the timeline or not, we take the majority label among the three workers." (p.6)
Q77	"The average timeline-level agreement between those two workers is 90.06%." (p.6)
Q78	"We also explore a more profound analysis by comparing the workers' annotations on 20% of the timelines against the annotations of experts." (p.6)
Q79	"The average timeline-level agreement rate of this analysis is 91.77%." (p.6)
Q87	"Since each timeline in our dataset is annotated by three workers (§3.2), it includes three summaries." (p.7)
Q88	"To adapt the timeline summarization task to this setting, we pick the two summaries written by the two workers who agree the most based on the timeline extraction labels they assign to the tweets in each timeline." (p.7)
Q105	"Following a similar process in 3.2, we recruit a group of workers from MTurk with the same Location restriction (QC1), who can pass our pilot study as a qualification test." (p.8)
WEB-23	#Verificado19S citizen-verification genre would be judged differently by Mexican journalists than US MTurk workers (https://arxiv.org/abs/2007.05848)

Strengths

- Documentation of annotator quality control (QC1–QC4) is detailed **[Q60–Q66]**, providing transparency about who did label and enabling targeted re-annotation planning.
- High inter-annotator agreement (90.06% **[Q77]**) and expert validation (91.77% **[Q79]**) demonstrate internal consistency of labels within the Anglophone population.
- Compensation transparency (\$3/HIT, >\$10/hour effective rate) **[Q67, Q68]** is documented.

ID	Evidence
----	----------

Q60	"Following prior quality control practices (Briakou et al., 2021), we use multiple quality control (QC) steps to ensure the recruitment of high-quality annotators." (p.5)
Q61	"First, we use location restriction (QC1) to limit the pool of workers to countries where native English speakers are most likely to be found." (p.5)
Q62	"Next, we recruit annotators who pass our qualification test (QC2), where we ask them to annotate 3 timelines." (p.5)
Q63	"We run several small pilot tasks, each with a replication factor of nine." (p.5)
Q64	"We check annotators' performance on timeline extraction task against experts' labels and have experts manually review (QC3) annotators' summary quality." (p.5)
Q65	"Only workers passing all the quality control steps contribute to the final task." (p.5)
Q66	"During the final task, we perform regular quality checks (QC4), and only use workers who consistently perform well." (p.5)
Q67	"We compensate the workers at a rate of \$3 per HIT for the task." (p.5)
Q68	"Each batch of tasks is followed by a one-time bonus that makes the final rate over \$10 per hour." (p.5)
Q77	"The average timeline-level agreement between those two workers is 90.06%." (p.6)
Q79	"The average timeline-level agreement rate of this analysis is 91.77%." (p.6)

Checklist

Checklist item definitions:

ID	Assessment Question
OC-1	Do ground-truth labels reflect regional stakeholder perspectives?
OC-2	Assess potential annotator-regional population disagreement
OC-3	Review annotator demographics from documentation
OC-4	Consider label re-annotation by regional annotators
OC-5	Review aggregation methods for minority perspective erasure
OC-6	Document label issues harming convergent/external validity

Checklist responses:

OC-1: No — ground-truth labels reflect Anglophone MTurk worker judgments [Q61, Q70], not Mexican/Central American journalist perspectives. The elicitation (A3) confirms divergent relevance and attribution norms.

OC-2: Substantial disagreement is likely. Mexican journalists operating under AP Español norms apply distinct attribution standards (CENAPRED/SSN/CNPC) [WEB-1, WEB-12] and may judge community solidarity posts and verification posts [WEB-23] as more informative than US MTurk workers would.

OC-3: Annotator demographics are documented only at the level of 'countries where native English speakers are most likely to be found' [Q61, Q105]. No information on age, gender, occupation, or journalistic background is provided. No Datasheet or Data Statement is referenced.

OC-4: Re-annotation by a representative pool — Spanish-speaking journalists familiar with Mexican/Central American crisis contexts — is required for validity in the target deployment [elicitation A3].

OC-5: Aggregation by majority vote on binary labels [Q70] and selection of the two most-agreeing workers' summaries [Q87, Q88] erases minority perspectives. When the entire annotator pool is non-representative of the target population, this aggregation compounds rather than mitigates the validity gap.

OC-6: Convergent validity with regional stakeholders is unsupported; external validity to Mexican journalist judgments cannot be assumed.

ID	Evidence
Q61	"First, we use location restriction (QC1) to limit the pool of workers to countries where native English speakers are most likely to be found." (p.5)

Q70	"To make the ground-truth on whether each tweet is part of the timeline or not, we take the majority label among the three workers." (p.6)
Q87	"Since each timeline in our dataset is annotated by three workers (§3.2), it includes three summaries." (p.7)
Q88	"To adapt the timeline summarization task to this setting, we pick the two summaries written by the two workers who agree the most based on the timeline extraction labels they assign to the tweets in each timeline." (p.7)
Q105	"Following a similar process in 3.2, we recruit a group of workers from MTurk with the same Location restriction (QC1), who can pass our pilot study as a qualification test." (p.8)
WEB-1	CENAPRED is the authoritative Mexican source for crisis attribution — distinct from US norms (https://en.wikipedia.org/wiki/CENAPRED)
WEB-12	Central American national civil-protection agencies (CONRED, COPECO, DGPC, SINAPRED, CNE, SINAPROC) define the local source-attribution ecosystem (https://whndn.org/portfolio-item/disaster-response-services-tools/)
WEB-23	#Verificado19S citizen-verification genre would be judged differently by Mexican journalists than US MTurk workers (https://arxiv.org/abs/2007.05848)

Information Gaps

- Specific MTurk worker demographics beyond country of residence are undocumented
- The magnitude of disagreement between Anglophone MTurk workers and Mexican journalists is empirically unmeasured

Requires Expert Verification

- Pilot re-annotation study with Mexican journalists to quantify disagreement rates on representative Spanish-language crisis tweets
- AP Español editor review of summary references for register and attribution conformance

Remediation

Gap 1: Reference labels and summaries authored entirely by Anglophone MTurk workers under QC1 location restriction [Q61]; no journalist or Mexican stakeholder representation

Recommendation: Re-annotate a representative subset using Mexican Spanish-speaking journalists familiar with Mexican/Central American crisis contexts and AP Español. Quantify disagreement against the original labels and report convergent-validity statistics. Avoid majority-vote aggregation that erases minority perspectives when annotator pools are small.

ID	Evidence
Q61	"First, we use location restriction (QC1) to limit the pool of workers to countries where native English speakers are most likely to be found." (p.5)