

Output Form (OF)

Score: 4/5 — Minor gaps | Confidence: high

Benchmark: LTZGLUE | Context: Luxembourgish Professional Newsroom

Output Form definition: Whether the expected output modality matches regional deployment needs and accessibility.

Each section below is followed by an evidence table. Three types of evidence are cited: **[QN]** — verbatim quotes extracted from the benchmark paper; **[WEB-N]** — web sources consulted during assessment; **[DN]** / **[DATASET-DN]** — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

Output form aligns well with deployment needs (LOWER priority dimension). All tasks produce discrete label outputs evaluated by macro-F1 [\[Q93, Q96\]](#), directly compatible with the deployment's per-task significance-score-and-threshold pipeline [\[A4\]](#). Encoder results are averaged over three runs with reported standard deviations [\[Q89, Q96\]](#), enabling stability assessment. The label-output (not MCQA) paradigm [\[Q84\]](#) mirrors the deployment's classification framing. Class-imbalanced tasks use class-balanced loss [\[Q92\]](#). Caveats: LLM outputs failing label format are discarded [\[Q90\]](#), inflating reported scores; LLMs are evaluated only once with macro-F1 only [\[Q96\]](#) — meaning LLM performance is less stable, and the benchmark itself flags this as 'indicative rather than directly comparable' [\[Q86\]](#). The deployment's threshold-flagging logic [\[A4\]](#) requires confidence scores per task, which encoder fine-tuning supports natively but prompted LLM evaluation may not.

ID	Evidence
Q84	"We did not use a Multiple Choice Question Answering (MCQA)-setup, but provided the labels that should be used as output." (p.6)
Q86	"However, prompt-based evaluation introduces additional sources of variability, including prompt sensitivity and differences in instruction-following behaviour across models. As a result, performance should be interpreted as indicative rather than directly comparable to supervised results." (p.6)
Q89	"For encoder-based models, results are reported as averages over multiple runs (see Appendix 7.2 for more details)." (p.7)
Q90	"Prompted LLMs do not always produce well-formed outputs and may return an incorrect number of predictions for a given task; such outputs are discarded prior to evaluation." (p.7)
Q92	"For the supervised models, since the linguistic acceptability and sentiment analysis datasets are highly imbalanced, when fine-tuning on these tasks we use class-balanced loss based on effective size (Cui et al., 2019) with a beta of 0.99." (p.7)
Q93	"Table 6 shows F1 scores for all models across all tasks (see Appendix 7.9 for full results)." (p.7)
Q96	"Table 6: Test F1 scores across all ItzGLUE tasks. Encoder results are averaged over three runs with standard deviations as subscripts. Prompted LLMs were evaluated once; we report macro-F1 only." (p.8)

Strengths

- Discrete label outputs evaluated by macro-F1 [\[Q93, Q96\]](#) directly support the deployment's per-task threshold flagging [\[A4\]](#)
- Encoder results averaged over three runs with standard deviations [\[Q89, Q96\]](#) permit stability assessment for production deployment
- Class-balanced loss applied to imbalanced tasks (LA, SA) [\[Q92\]](#) addresses real distributional skew
- Label-output paradigm rather than MCQA [\[Q84\]](#) matches the deployment's classification framing

ID	Evidence
Q84	"We did not use a Multiple Choice Question Answering (MCQA)-setup, but provided the labels that should be used as output." (p.6)
Q89	"For encoder-based models, results are reported as averages over multiple runs (see Appendix 7.2 for more details)." (p.7)
Q92	"For the supervised models, since the linguistic acceptability and sentiment analysis datasets are highly imbalanced, when fine-tuning on these tasks we use class-balanced loss based on effective size (Cui et al., 2019) with a beta of 0.99." (p.7)
Q93	"Table 6 shows F1 scores for all models across all tasks (see Appendix 7.9 for full results)." (p.7)

Q96	"Table 6: Test F1 scores across all ItzGLUE tasks. Encoder results are averaged over three runs with standard deviations as subscripts. Prompted LLMs were evaluated once; we report macro-F1 only." (p.8)
------------	--

Checklist

Checklist item definitions:

ID	Assessment Question
OF-1	Does output modality match regional deployment needs?
OF-2	Assess text-to-speech availability for speech output
OF-3	Consider literacy rates and accessibility requirements
OF-4	Document form mismatches harming external validity

Checklist responses:

OF-1: Discrete label outputs [Q93, Q148, Q149, Q150, Q151, Q152, Q153] match the deployment's per-task scoring needs [A4]; the deployment's threshold flagging is a downstream wrapper compatible with macro-F1 evaluation.

OF-2: Not applicable — deployment is text-only [deployment_modality]; speech output is not required.

OF-3: Target users are professional journalists/editors with high literacy and trilingual proficiency [WEB-14, WEB-15]; no accessibility constraints introduced by the output form.

OF-4: Minor mismatches: (a) prompted LLM scores are single-run macro-F1 [Q96] less directly comparable to encoder runs [Q86]; (b) discarded malformed LLM outputs [Q90] inflate apparent performance — relevant if the deployment uses prompted LLMs rather than fine-tuned encoders.

ID	Evidence
Q86	"However, prompt-based evaluation introduces additional sources of variability, including prompt sensitivity and differences in instruction-following behaviour across models. As a result, performance should be interpreted as indicative rather than directly comparable to supervised results." (p.6)
Q90	"Prompted LLMs do not always produce well-formed outputs and may return an incorrect number of predictions for a given task; such outputs are discarded prior to evaluation." (p.7)
Q93	"Table 6 shows F1 scores for all models across all tasks (see Appendix 7.9 for full results)." (p.7)
Q96	"Table 6: Test F1 scores across all ItzGLUE tasks. Encoder results are averaged over three runs with standard deviations as subscripts. Prompted LLMs were evaluated once; we report macro-F1 only." (p.8)
Q148	"headline_classification: Decide if the given title/headline fits the text. Output True or False." (p.15)
Q149	"sentiment_analysis: Classify sentiment of the text. Allowed labels: positive, neutral, negative." (p.15)
Q150	"linguistic_acceptability_binary: Decide whether the sentence is linguistically acceptable in Luxembourgish. Output: 0 or 1." (p.15)
Q151	"linguistic_acceptability_multilabel: Detect if the sentence is correct or if some element is wrong. If the sentence is correct, Output: correct. If it is not, Output the label referencing the wrong element: syntax, verb, ortho or adj." (p.15)
Q152	"ner: Perform Named Entity Recognition on the given sequence of sentence tokens. Output tags as lists of ner_tags. Allowed Tags: O, B-LOC, I-LOC, B-PER, I-PER, B-DATE, I-DATE, B-ORG, I-ORG, B-MISC, I-MISC." (p.15)
Q153	"topic_classification: Classify topic of the document by title and text. Allowed category_names: sports, animals, business, culture, technology." (p.15)
WEB-14	Luxembourg trilingual education system produces newsroom professionals with high literacy in Luxembourgish/French/German plus English (https://men.public.lu/en/systeme-educatif/langues-ecole-luxembourgaise.html)
WEB-15	EU Education and Training Monitor confirms strong language education context for target users (https://op.europa.eu/webpub/eac/education-and-training-monitor/en/country-reports/luxembourg.html)

Information Gaps

- Whether confidence scores from fine-tuned encoders are well-calibrated for the deployment's threshold-based flagging logic

Requires Expert Verification

- Calibration of model confidence outputs against editorial flagging thresholds in production