

Validity Analysis: HOPE

Target context: US and South Asian Mental Health Counseling Practitioners (OMHC/Teletherapy)

Overall risk: **HIGH**

Deployment Context

Use case: A conversational AI system for mental health counseling classifies dialogue acts in real-time counseling conversations using the HOPE benchmark and any classification model. The system identifies the communicative intent (e.g., question, clarification) of each utterance by counselors and clients, and subsequently allows downstream tasks such as session quality assessment, automated counselor feedback, and training support in clinical and teletherapy settings.

Target population: Mental health counselors, trainee therapists, and clinical supervisors who conduct or review online counseling sessions in English. This includes practitioners and researchers working with teletherapy platforms and peer support services who need automated tools for counseling quality evaluation and dialogue understanding.

Dimension Scores

Dimension	Score	Rating	Confidence
Input Ontology 	2	Significant gaps	high
Input Content 	2	Significant gaps	high
Input Form	3	Moderate gaps	medium
Output Ontology 	3	Moderate gaps	high
Output Content	3	Moderate gaps	medium
Output Form 	4	Minor gaps	high
Average	2.8		

 = highest concern  = strongest dimension

Dimension Key

Abbr.	Dimension	Definition
IO	Input Ontology	Whether the benchmark's test case categories cover the query types expected in deployment.
IC	Input Content	Whether individual datapoint content is culturally and contextually appropriate for the target region.
IF	Input Form	Whether the input signal encoding (text, audio, image parameters) matches deployment conditions.
OO	Output Ontology	Whether the benchmark's output categories and scoring criteria reflect regionally valid decision boundaries.
OC	Output Content	Whether ground-truth labels align with the judgments of regional stakeholders.
OF	Output Form	Whether the expected output modality matches regional deployment needs and accessibility.

Overall Summary

HOPE provides a well-engineered dialogue-act classification benchmark for structured US clinical counseling (CBT, child, family therapy) with a thoughtfully designed 12-label taxonomy [Q17, Q27], substantial inter-annotator agreement [Q49], and rigorous evaluation methodology [Q63, Q87, Q88]. For the in-scope structured clinical therapy portion of the US deployment, the benchmark is reasonably valid. However, the deployment explicitly extends to peer support, crisis intervention, and South

Asian English contexts, and HOPE's source data, taxonomy, and annotators are predominantly aligned with US-based clinical therapy [Q98, Q99]. The HIGH-priority IO dimension shows a clear ontological gap: peer support communicative patterns and crisis-specific acts (safety-planning, risk stratification) have no dedicated representation and would default to overloaded categories like ACK or the residual GC catch-all [Q44, Q77]. No drop-in supplementary benchmark exists for crisis dialogue-act classification [WEB-2, WEB-20], and ESConv [WEB-18] / EMH [WEB-19] provide only structurally-different peer-support frameworks. Output form is well aligned (score 4); input form has bounded concerns from live-ASR fidelity (score 3); output ontology shows real boundary ambiguity even within scope (score 3); annotator representation lacks documented South Asian or peer/crisis presence (OC score 3). Input ontology and input content carry the largest validity risks (both score 2).

ID	Evidence
Q17	"We propose twelve dialogue-act labels that are aligned with mental-health counseling session." (p.2)
Q27	"Since the counseling conversations have inherent differences with the standard conversations (such as SwitchBoard dataset conversation), it demands a carefully designed set of dialogue-act labels capable of catering the requirements of counselling conversations. Hence, we, in consultation with therapist and counselling experts, design a set of 12 dialog-act labels that are arranged in a hierarchy." (p.3)
Q44	"General Chit-Chat (GC): Other utterances that do not belong to any of the above labels are tagged as GC, possibly because of the vagueness and the lack of sense in the context of the conversation." (p.4)
Q49	"We obtain the inter-rater agreement score of 0.7234 – which falls under the 'substantial' [7] category." (p.4)
Q63	"To measure the performances of SPARTA and other baseline systems, we compute macro-F1, weighted-F1, and accuracy scores." (p.6)
Q77	"We can observe that SPARTA consistently yields good scores for the majority of the dialogue-acts, except for the Acknowledgement (ACK) where it records F1-score of merely 47.86%." (p.7)
Q87	"Moreover, we also report the mean of the 3-fold cross-validation results for both SPARTA-MHA and SPARTA-TAA, and the results are consistent with the train-val-test split case." (p.7)
Q88	"We also perform statistical significance T-test comparing SPARTA-TAA and the best performing baseline (CASA). We observe that our results are significant with > 95% confidence across macro-F1 (p-value= 0.009), weighted-F1 (p-value= 0.014), and accuracy (p-value= 0.048) values." (p.7)
Q98	"Another important aspect of the current work is that the majority of the sessions in HOPE belong to the mental health professionals and patients based in United States." (p.8)
Q99	"Hence, the effectiveness of SPARTA on data from other geographical or demographical regions may vary." (p.8)
WEB-2	988 Lifeline three-phase crisis model uses communicative acts not represented in HOPE labels (https://988lifeline.org/wp-content/uploads/2023/02/FINAL_988_Suicide_and_Crisis_Lifeline_Suicide_Safety_Policy_-3.pdf)
WEB-18	ESConv 8-strategy framework for peer support is structurally distinct from HOPE (https://aclanthology.org/2021.acl-long.269/)
WEB-19	https://github.com/behavioral-data/Empathy-Mental-Health
WEB-20	Crisis Text Line research uses discourse analysis rather than HOPE-compatible dialogue-act labels; no crisis-specific dialogue-act benchmark identified (https://www.crisistextline.org/data-philosophy/research-impact/)

Practical Guidance

What This Benchmark Measures

HOPE measures the ability to assign one of 12 counseling-aligned dialogue-act labels to utterances in dyadic clinical therapy sessions, evaluating conversational-flow classification rather than clinical judgment [Q14, Q102]. For the deployment, this provides reasonable evidence about model performance on US-based structured clinical therapy (CBT/child/family) — particularly for tracking conversational trajectory in supervision and feedback workflows. Output Form (score 4) and Output Ontology design (score 3) are the strongest dimensions for this in-scope use case.

Construct Depth

The benchmark probes dialogue-act classification with reasonable depth for structured clinical therapy: 12.9K utterances across 212 sessions [Q13], substantial inter-annotator agreement [Q49], multi-tier baseline comparison [Q79, Q80], and per-label diagnostic information [Q76]. However, depth degrades sharply outside this scope: there is no empirical evidence about how labels behave on peer support register, no evaluation on crisis-context utterances, no per-session-type or per-label

Kappa breakdown, and no out-of-distribution diagnostic framework. The discarded emotion-annotation pilot [Q126, Q129] indicates the taxonomy intentionally excludes affective dimensions that matter for quality assessment.

What Else You Need

For peer support deployment: ESConv [WEB-18] and EMH [WEB-19] provide complementary annotation frameworks (support strategies, empathy mechanisms) that should be evaluated alongside HOPE-trained models — neither uses HOPE-compatible labels, so cross-framework validation is required. For crisis intervention deployment: no equivalent benchmark exists [WEB-20], so crisis-specific dialogue-act validation is a net-new task — the 988 Lifeline structured model [WEB-2] should inform a custom label extension. For South Asian deployment: expert elicitation from South Asian clinical psychologists is needed to validate label-boundary behavior under different formality/directiveness norms (gap_id 3, deferred). For real-time teletherapy: empirical comparison of YouTube vs. live-ASR transcripts is needed to quantify input-form drift (gap_id 5, deferred). For supervisory feedback applications: per-label safety-critical performance reporting (especially for IRQ/YNQ behavior on crisis-relevant utterances) should supplement macro-F1.

ID	Evidence
Q13	"The HOPE dataset contains ~ 12.9K utterances across 212 mental-health counseling sessions." (p.2)
Q14	"Each utterance in the dataset is tagged with one of the 12 counseling-aligned dialogue-act labels (c.f. Section 3)." (p.2)
Q49	"We obtain the inter-rater agreement score of 0.7234 – which falls under the 'substantial' [7] category." (p.4)
Q76	"We also present the label-wise performance of SPARTA in Table 4." (p.7)
Q79	"Based on the type of modelling, we categorize the baselines into three groups – utterance-driven (Ut), utterance+global context driven ($Ut + GC$), and utterance + global context + speaker-aware driven ($Ut + GC + SA$)." (p.7)
Q80	"Comparatively, SPARTA incorporates local context in addition to utterance, global context, and speaker dynamics ($Ut + LC + GC + SA$)." (p.7)
Q102	"It is important to note that we do not make any diagnostic claims." (p.8)
Q126	"In earlier phases of experiments, we also considered the role of emotion in dialogue act classification." (p.11)
Q129	"Due to such imbalance, this data was not used in the final version of our proposed architecture." (p.11)
WEB-2	988 Lifeline three-phase crisis model uses communicative acts not represented in HOPE labels (https://988lifeline.org/wp-content/uploads/2023/02/FINAL_988_Suicide_and_Crisis_Lifeline_Suicide_Safety_Policy_-3.pdf)
WEB-18	ESConv 8-strategy framework for peer support is structurally distinct from HOPE (https://aclanthology.org/2021.acl-long.269/)
WEB-19	https://github.com/behavioral-data/Empathy-Mental-Health
WEB-20	Crisis Text Line research uses discourse analysis rather than HOPE-compatible dialogue-act labels; no crisis-specific dialogue-act benchmark identified (https://www.crisistextline.org/data-philosophy/research-impact/)