

Validity Analysis: HOPE DAC

Target context: Global NLP/AI Research Community — Computational Mental Health (HOPE Benchmark Assessment)

Overall risk: **HIGH**

Deployment Context

Use case: An NLP/AI researcher uses the HOPE benchmark dataset to evaluate and improve dialogue-understanding for AI-assistive chatbots. The benchmark provides a evaluation protocol across 12 dialogue act categories in counseling conversations.

Target population: NLP/AI researchers in English speaking countries or willing to work with English data working on computational approaches to mental health, dialogue systems, and psycholinguistics. This includes PhD students, postdoctoral researchers, and industry practitioners building LLM-based systems for clinical NLP and digital mental health applications serving underrepresented regional communities.

Dimension Scores

Dimension	Score	Rating	Confidence
Input Ontology ✓	3	Moderate gaps	medium
Input Content △	2	Significant gaps	high
Input Form ✓	4	Minor gaps	high
Output Ontology	3	Moderate gaps	high
Output Content △	2	Significant gaps	high
Output Form	3	Moderate gaps	high
Average	2.8		

△ = highest concern ✓ = strongest dimension

Dimension Key

Abbr.	Dimension	Definition
IO	Input Ontology	Whether the benchmark's test case categories cover the query types expected in deployment.
IC	Input Content	Whether individual datapoint content is culturally and contextually appropriate for the target region.
IF	Input Form	Whether the input signal encoding (text, audio, image parameters) matches deployment conditions.
OO	Output Ontology	Whether the benchmark's output categories and scoring criteria reflect regionally valid decision boundaries.
OC	Output Content	Whether ground-truth labels align with the judgments of regional stakeholders.
OF	Output Form	Whether the expected output modality matches regional deployment needs and accessibility.

Overall Summary

HOPE provides a methodologically careful, transparent benchmark for dialogue-act classification in dyadic individual counseling — its task framing [Q7-Q9, Q22], 12-label scheme developed with expert consultation [Q27], substantial inter-annotator agreement [Q49], and statistically significant SPARTA results [Q88] make it a reasonable starting point for the elicited use case of evaluating chatbots in individual counseling-style conversations. The strongest validity dimensions are

Input Form (text-only English matches deployment exactly) and Input Ontology (well-scoped task framing for individual dyadic counseling). The most significant validity concerns are concentrated in the two HIGH-priority dimensions: Input Content (sessions are explicitly US-clinical [Q98] while downstream chatbots target South Asian, African American, and Latino communities whose communicative norms diverge in well-documented ways [WEB-5, WEB-9, WEB-11]) and Output Content (annotator cultural backgrounds undocumented [Q50], no label-stratified IAA, RoBERTa backbone inherits documented AAVE bias [WEB-8, WEB-9]). The Output Ontology and Output Form dimensions show acknowledged but unresolved limitations — single-label assignment [Q30] mismatches multi-intent reality (per elicitation A3), and no subgroup-disaggregated metrics are reported, which conflicts with emerging FDA equity expectations [WEB-20]. Many of these gaps are field-wide rather than HOPE-specific [WEB-17], but they materially limit the benchmark's external validity for the target downstream populations.

ID	Evidence
Q7	“Unlike general goal-oriented dialogues, a conversation between a patient and a therapist is considerably implicit, though the objective of the conversation is quite apparent.” (p.1)
Q8	“The nature of conversations in a social counselling setting is particularly distinct as compared to a conventional chit-chat or goal-oriented conversations.” (p.1)
Q9	“It follows a pattern which is different from both goal-oriented and general chit-chat based conversations.” (p.1)
Q22	“In this section, we present our dialogue-act classification dataset, called HOPE. In total, we annotated ~ 12.9K utterances with 12 dialogue-act labels carefully designed to cater to the requirements of a counseling session.” (p.3)
Q27	“Since the counseling conversations have inherent differences with the standard conversations (such as SwitchBoard dataset conversation), it demands a carefully designed set of dialogue-act labels capable of catering the requirements of counselling conversations. Hence, we, in consultation with therapist and counselling experts, design a set of 12 dialog-act labels that are arranged in a hierarchy.” (p.3)
Q30	“Sometimes, an utterance can have multiple dialog-acts; however, they are rare in the annotated dataset. Hence, for simplicity, we consider only one (primary) label for each utterance.” (p.3)
Q49	“We obtain the inter-rater agreement score of 0.7234 – which falls under the 'substantial' [7] category.” (p.4)
Q50	“Annotators are in the age group of 25-35, with 2-10 years of professional experience.” (p.4)
Q88	“We also perform statistical significance T-test comparing SPARTA-TAA and the best performing baseline (CASA). We observe that our results are significant with > 95% confidence across macro-F1 (p-value= 0.009), weighted-F1 (p-value= 0.014), and accuracy (p-value= 0.048) values.” (p.7)
Q98	“Another important aspect of the current work is that the majority of the sessions in HOPE belong to the mental health professionals and patients based in United States.” (p.8)
WEB-5	PMC review: only 9% of Asian Americans utilized mental health services vs. 18% of general US population (https://pmc.ncbi.nlm.nih.gov/articles/PMC3208524/)
WEB-8	ResearchGate: RoBERTa shows complex bias toward AAVE in masked language modeling — bias inherited by SPARTA (https://www.researchgate.net/publication/371085866_Quantifying_the_Bias_of_Transformer-Based_Language_Models_for_African_American_English_in_Masked_Language_Modeling)
WEB-9	npj Digital Medicine (2025): LLMs more racially biased in mental health than other health domains; AAVE dialect cues trigger inferior treatment recommendations (https://www.nature.com/articles/s41746-025-01746-4)
WEB-11	HHS Office of Minority Health: Latino adults 28% less likely than US adults overall to receive mental health treatment in 2024 (https://minorityhealth.hhs.gov/mental-and-behavioral-health-hispaniclatinos)
WEB-17	Translational Psychiatry (2023): only 4 of 102 NLP-MHI studies used external validation — subgroup performance reporting is field-wide rare (https://www.nature.com/articles/s41398-023-02592-2)
WEB-20	Orrick (Nov 2025): FDA DHAC signals sponsors will need 'evidence of equitable performance across populations and languages' for AI mental health devices (https://www.orrick.com/en/Insights/2025/11/FDAs-Digital-Health-Advisory-Committee-Considers-Generative-AI-Therapy-Chatbots-for-Depression)

Practical Guidance

What This Benchmark Measures

HOPE measures a model's ability to assign one of 12 counseling-domain dialogue-act labels [Q22] to utterances in dyadic individual-counseling text transcripts [Q26], with strongest signal on speaker-initiative questions and clear delivery acts and weakest signal on Acknowledgment [Q77] and on the speaker-initiative/responsive complementary pairs [Q51, Q91]. It is best suited to comparing dialogue-aware architectures' ability to model local context, global context, and speaker dynamics on

US-clinical English counseling conversations — the Input Ontology and Input Form dimensions support this well.

Construct Depth

The benchmark probes the dialogue-act classification construct moderately deeply on US-clinical English material — with hierarchical taxonomy design [Q29, Q32], expert-consulted label definitions [Q27], substantial annotator agreement [Q49], proper statistical testing [Q88], and ablations that disentangle local/global context and speaker-awareness contributions [Q115]. However, depth is constrained on the Output Content and Input Content dimensions: it does not probe whether the construct generalizes across cultural communicative norms, does not test multi-intent labeling [Q30], and provides no subgroup-disaggregated evidence for the populations the downstream chatbots are intended to serve.

What Else You Need

For valid use in evaluating chatbots targeting South Asian, African American, and Latino communities, supplementation is required across IC, OC, OO, and OF: (1) re-annotation or fresh annotation of a held-out evaluation sample by representative regional annotators (addresses OC); (2) collection or curation of supplementary counseling conversation samples from target communities (addresses IC); (3) multi-label or ranked-output evaluation protocols built on top of HOPE single-label predictions (addresses OO and OF, supports A3); (4) subgroup-disaggregated performance reporting consistent with emerging FDA equity expectations [WEB-20]; (5) AAVE-specific bias audits given RoBERTa backbone bias evidence [WEB-8, WEB-9]. AnnoMI [WEB-15] provides a methodological model for richer annotator documentation. As of May 2025 no culturally adapted counseling DA benchmark exists [WEB-23] to substitute for HOPE — supplementation rather than replacement is the realistic path.

ID	Evidence
Q22	"In this section, we present our dialogue-act classification dataset, called HOPE. In total, we annotated ~ 12.9K utterances with 12 dialogue-act labels carefully designed to cater to the requirements of a counseling session." (p.3)
Q26	"The data collection process provides us 12.9K utterances from 212 counseling therapy sessions – all of them are dyadic conversations only." (p.3)
Q27	"Since the counseling conversations have inherent differences with the standard conversations (such as SwitchBoard dataset conversation), it demands a carefully designed set of dialogue-act labels capable of catering the requirements of counselling conversations. Hence, we, in consultation with therapist and counselling experts, design a set of 12 dialog-act labels that are arranged in a hierarchy." (p.3)
Q29	"Each utterance in the dialogue belongs to one of the three categories – speaker initiative, speaker responsive, and general." (p.3)
Q30	"Sometimes, an utterance can have multiple dialog-acts; however, they are rare in the annotated dataset. Hence, for simplicity, we consider only one (primary) label for each utterance." (p.3)
Q32	"Our annotation scheme assigns three distinct dialogue-act labels to the first two categories, while the remaining four labels belong to the general category." (p.4)
Q49	"We obtain the inter-rater agreement score of 0.7234 – which falls under the 'substantial' [7] category." (p.4)
Q51	"The broad definitions of speaker-initiative and speaker-responsive dialogue-act label pairs, IRQ & ID; ORQ & OD, and CRQ & CD, seem complementary in nature." (p.4)
Q77	"We can observe that SPARTA consistently yields good scores for the majority of the dialogue-acts, except for the Acknowledgement (ACK) where it records F1-score of merely 47.86%." (p.7)
Q88	"We also perform statistical significance T-test comparing SPARTA-TAA and the best performing baseline (CASA). We observe that our results are significant with > 95% confidence across macro-F1 (p-value= 0.009), weighted-F1 (p-value= 0.014), and accuracy (p-value= 0.048) values." (p.7)
Q91	"We observe three pairs with significant error-rates ($\geq 25\%$) – YNQ:IRQ (26%), OD:ID (43%), ID:ACK (28%)." (p.7)
Q115	"The first set of results (i.e., SPARTA-BS or the baseline version of SPARTA) in Table 7 represents the absence of local context in SPARTA." (p.11)
WEB-8	ResearchGate: RoBERTa shows complex bias toward AAVE in masked language modeling — bias inherited by SPARTA (https://www.researchgate.net/publication/371085866_Quantifying_the_Bias_of_Transformer-Based_Language_Models_for_African_American_English_in_Masked_Language_Modeling)
WEB-9	npj Digital Medicine (2025): LLMs more racially biased in mental health than other health domains; AAVE dialect cues trigger inferior treatment recommendations (https://www.nature.com/articles/s41746-025-01746-4)

WEB-15	AnnoMI uses 10 annotators with post-annotation survey, demonstrating richer annotator documentation is achievable (https://www.mdpi.com/1999-5903/15/3/110)
WEB-20	Orrick (Nov 2025): FDA DHAC signals sponsors will need 'evidence of equitable performance across populations and languages' for AI mental health devices (https://www.orrick.com/en/Insights/2025/11/FDAs-Digital-Health-Advisory-Committee-Considers-Generative-AI-Therapy-Chatbots-for-Depression)
WEB-23	JMIR Mental Health (2025) confirms 'limited attention to linguistic or cultural adaptation' in LLM mental health studies (https://mental.jmir.org/2025/1/e78410)