

Output Content (OC)

Score: 1/5 — Serious concern | Confidence: high

Benchmark: SAMSUM | Context: South Asia and MENA — Clinical NLP Researchers and AI Practitioners

Output Content definition: Whether ground-truth labels align with the judgments of regional stakeholders.

Each section below is followed by an evidence table. Three types of evidence are cited: **[QN]** — verbatim quotes extracted from the benchmark paper; **[WEB-N]** — web sources consulted during assessment; **[DN]** / **[DATASET-DN]** — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

Reference summaries were authored by the same linguists who created the dialogues [\[Q23\]](#), creating data-label circularity, with a 50-conversation double-annotation validation showing 94% agreement [\[Q26\]](#) but limited generalization. Annotator demographics are not reported; acknowledged contributors have Polish-consistent names [\[Q75\]](#). The user elicitation explicitly confirms annotators were 'general mental health experts from mixed backgrounds (including India and the USA)' but 'were not specifically briefed on regional context when validating summaries' and that 'their salience judgments likely reflect generic CBT protocols rather than consciously incorporating regional norms.' Empirical evidence reveals concrete content failures: a stigmatizing self-description reproduced as factual content [\[DATASET-D6\]](#), a speaker attribution error in a simple two-turn exchange [\[DATASET-D21\]](#), peer distress treated as social transaction [\[DATASET-D33\]](#). HIGH priority weight on OC and confirmed norm-divergence risk produce a severe convergent validity failure.

ID	Evidence
Q23	"After collecting all of the conversations, we asked language experts to annotate them with summaries, assuming that they should (1) be rather short, (2) extract important pieces of information, (3) include names of interlocutors, (4) be written in the third person." (p.2)
Q26	"The results revealed that 94% of examined dialogues were classified by both annotators as good i.e. they do look like conversations from a messenger app and could be condensed in a reasonable way." (p.2)
Q75	"We would like to express our sincere thanks to Tunia Blachno, Oliwia Ebebenge, Monika Jędras and Małgorzata Krawentek for their huge contribution to the corpus collection – without their ideas, management of the linguistic task and verification of examples we would not be able to create this paper." (p.9)
D6	Pearl's boyfriend mistreats her but she loves him anyway. Patrick is trying to convince her to leave him. Pearl is fat and ugly, so it's not so easy for her. — Summary reproduces a derogatory self-description as factual without clinical reframing — illustrates embedded Western annotator salience norms
D21	Miranda will be ready in no more than 10 minutes. — Summary misattributes action (Danny said he was almost finished, not Miranda) — illustrates annotation quality ceiling
D33	Francine and Jessie talked yesterday. The talk has been helpful for Jessie. Francine says she's always there for Jessie. Jessie is very thankful. — Summary of peer emotional support — captures who did what but contains no clinical salience categories (distress type, risk level, intervention)

Strengths

- Documented validation procedure: double-annotation of 50 conversations with 94% pass rate [\[Q25, Q26\]](#)
- Conflict resolution mechanism for opposing –1/+1 ratings via third annotator [\[Q49\]](#)
- Reasonable inter-annotator agreement on dialogues (Cohen's kappa 0.506) [\[Q51\]](#)

ID	Evidence
Q25	"We asked two linguists to doubly annotate 50 conversations in order to verify whether the dialogues could appear in a messenger app and could be summarized (i.e. a dialogue is not too general or unintelligible) or not (e.g. a dialogue between two people in a shop)." (p.2)
Q26	"The results revealed that 94% of examined dialogues were classified by both annotators as good i.e. they do look like conversations from a messenger app and could be condensed in a reasonable way." (p.2)
Q49	"We noticed a few annotations (7 for news and 4 for dialogues) with opposite marks (i.e. one annotator judgement was –1, whereas the second one was 1) and decided to have them annotated once again by another annotator who had to resolve conflicts." (p.5)

Q51	"For news examples, we obtained agreement on the level of 0.371 and for dialogues – 0.506." (p.5)
------------	---

Checklist

Checklist item definitions:

ID	Assessment Question
OC-1	Do ground-truth labels reflect regional stakeholder perspectives?
OC-2	Assess potential annotator-regional population disagreement
OC-3	Review annotator demographics from documentation
OC-4	Consider label re-annotation by regional annotators
OC-5	Review aggregation methods for minority perspective erasure
OC-6	Document label issues harming convergent/external validity

Checklist responses:

OC-1: Ground truth labels do not reflect regional stakeholder perspectives. Annotators were not briefed on SA-MENA clinical context per elicitation; salience judgments reflect generic CBT/Western messenger norms. Empirical evidence: [DATASET-D6](#) reproduces stigmatizing self-description without clinical reframing.

OC-2: Disagreement between original annotators and SA-MENA mental health practitioners is highly likely. MENA practitioners would likely flag missing religious/fatalistic framing in distress summaries; South Asian practitioners would flag missing family-mediated resolution and izzat-sensitive framing. [WEB-13](#) corroborates that even Arabic mental health evaluation requires human-in-the-loop validation absent from SAMSum.

OC-3: Annotator demographics are not reported in the paper [[Q75 contributors have Polish-consistent names; otherwise undisclosed](#)]. No Datasheet or Data Statement is referenced. The elicitation supplied partial information (mixed Indian/US backgrounds, no regional briefing) but no formal demographic breakdown.

OC-4: Re-annotation by an SA-MENA-stratified clinical practitioner pool is required. [WEB-15](#) demonstrates feasibility (5 clinical health professionals validated MentalCLOUDS) but no analogous SA-MENA effort exists per [WEB-13](#), [WEB-17](#).

OC-5: Aggregation uses a single reference summary per dialogue [[Q24](#)] with no multi-reference design, structurally erasing minority/regional perspectives. Conflicts are resolved by a third annotator [[Q49](#)] rather than preserved as legitimate variation.

OC-6: Convergent validity is severely harmed: labels do not correlate with regional practitioner salience norms. External validity is severely harmed: original judgments will not generalize. The [DATASET-D21](#) attribution error additionally undermines the 94% quality-check ceiling.

ID	Evidence
Q23	"After collecting all of the conversations, we asked language experts to annotate them with summaries, assuming that they should (1) be rather short, (2) extract important pieces of information, (3) include names of interlocutors, (4) be written in the third person." (p.2)
Q24	"Each dialogue contains only one reference summary." (p.2)
Q26	"The results revealed that 94% of examined dialogues were classified by both annotators as good i.e. they do look like conversations from a messenger app and could be condensed in a reasonable way." (p.2)
Q49	"We noticed a few annotations (7 for news and 4 for dialogues) with opposite marks (i.e. one annotator judgement was –1, whereas the second one was 1) and decided to have them annotated once again by another annotator who had to resolve conflicts." (p.5)
Q51	"For news examples, we obtained agreement on the level of 0.371 and for dialogues – 0.506." (p.5)
Q75	"We would like to express our sincere thanks to Tunia Błachno, Oliwia Ebebenge, Monika Jędras and Małgorzata Krawentek for their huge contribution to the corpus collection – without their ideas, management of the linguistic task and verification of examples we would not be able to create this paper." (p.9)
D6	Pearl's boyfriend mistreats her but she loves him anyway. Patrick is trying to convince her to leave him. Pearl is fat and ugly, so it's not so easy for her. — Summary reproduces a derogatory self-description as factual without clinical reframing — illustrates embedded Western annotator salience norms

D21	Miranda will be ready in no more than 10 minutes. — Summary misattributes action (Danny said he was almost finished, not Miranda) — illustrates annotation quality ceiling
WEB-13	AraHealthQA 2025 found BERTScore 'failed to fully measure appropriateness or trustworthiness' for Arabic mental health QA — direct corroboration that automated metrics are insufficient (https://arxiv.org/html/2508.20047v2)
WEB-15	MentalCLOUDS supplements ROUGE/BERTScore with expert evaluation across six clinical parameters — shows the form of evaluation SAMSum lacks (https://mental.jmir.org/2024/1/e57306)
WEB-17	ACM survey confirms no LLM-grade contextualized models for low-resourced South Asian languages in mental health (https://dl.acm.org/doi/10.1145/3638761)

Information Gaps

- Exact regional breakdown of the 'mixed backgrounds (including India and the USA)' annotator pool described in elicitation
- Whether any annotators had clinical mental health training

Requires Expert Verification

- Magnitude of label-level disagreement between SA-MENA clinical practitioners and existing reference summaries on sampled distress-relevant examples

Remediation

Gap 1: Annotator demographics not regionally stratified; salience norms reflect generic CBT/Western messenger conventions and have produced concrete anti-clinical labels (e.g., DATASET-D6 reproducing stigmatizing self-description as fact).

Recommendation: Re-annotate any clinical-adjacent examples (or build new annotation infrastructure on a clinical corpus) using SA-MENA-stratified clinical practitioner panels with explicit briefing on regional context. Disclose demographics formally via a Datasheet/Data Statement. Preserve practitioner disagreement as legitimate variation rather than resolving via third-annotator override.

ID	Evidence
D6	Pearl's boyfriend mistreats her but she loves him anyway. Patrick is trying to convince her to leave him. Pearl is fat and ugly, so it's not so easy for her. — Summary reproduces a derogatory self-description as factual without clinical reframing — illustrates embedded Western annotator salience norms