

Validity Analysis: SAMSUM

Target context: South Asia and MENA — Clinical NLP Researchers and AI Practitioners

Overall risk: **HIGH**

Deployment Context

Use case: Researchers and developers building or evaluating AI-driven counseling support tools assess the capability of LLMs to produce faithful, coherent, and clinically appropriate summaries from counseling dialogues. The benchmark provides a standardized evaluation framework to compare multiple LLMs across summarization quality metrics relevant to the mental health domain.

Target population: NLP researchers and AI practitioners in South Asia and the Middle East and North Africa who are working on applied clinical NLP, digital mental health platforms, and LLM evaluation for low-resource or underrepresented languages and communities.

Dimension Scores

Dimension	Score	Rating	Confidence
Input Ontology 	1	Serious concern	high
Input Content 	1	Serious concern	high
Input Form 	2	Significant gaps	high
Output Ontology 	1	Serious concern	high
Output Content 	1	Serious concern	high
Output Form 	1	Serious concern	high
Average	1.2		

 = highest concern  = strongest dimension

Dimension Key

Abbr.	Dimension	Definition
IO	Input Ontology	Whether the benchmark's test case categories cover the query types expected in deployment.
IC	Input Content	Whether individual datapoint content is culturally and contextually appropriate for the target region.
IF	Input Form	Whether the input signal encoding (text, audio, image parameters) matches deployment conditions.
OO	Output Ontology	Whether the benchmark's output categories and scoring criteria reflect regionally valid decision boundaries.
OC	Output Content	Whether ground-truth labels align with the judgments of regional stakeholders.
OF	Output Form	Whether the expected output modality matches regional deployment needs and accessibility.

Overall Summary

SAMSum is a well-constructed informal English messenger-dialogue summarization benchmark that, by design and verifiable empirical sampling, contains no clinical content, no SA-MENA cultural content, no non-English language data, no clinical or regional output ontology, and uses metrics the authors themselves document as unreliable for dialogues — let alone for cross-lingual clinical evaluation. Five of six dimensions (IO, IC, OO, OC, OF) score 1; IF scores 2 because plain-text

speaker-attributed structure provides limited transferable infrastructure. The benchmark is being applied as a proxy for clinical counseling dialogue summarization quality assessment in SA-MENA contexts, which is a substantial domain, cultural, linguistic, and methodological mismatch on every priority-HIGH dimension identified in the elicitation. The empirical dataset analysis directly confirms claims made in the regional YAML and adds concrete evidence of anti-clinical annotator norms ([DATASET-D6](#), [DATASET-D10](#), [DATASET-D33](#)) that go beyond mere absence into active misalignment.

ID	Evidence
D6	Pearl's boyfriend mistreats her but she loves him anyway. Patrick is trying to convince her to leave him. Pearl is fat and ugly, so it's not so easy for her. — Summary reproduces a derogatory self-description as factual without clinical reframing — illustrates embedded Western annotator salience norms
D10	Mickey tries to calm her with the fact that more people commit suicide on the Golden Gate Bridge than in that forest. — Suicide statistics treated as social reassurance in summary — no clinical risk annotation or flagging
D33	Francine and Jessie talked yesterday. The talk has been helpful for Jessie. Francine says she's always there for Jessie. Jessie is very thankful. — Summary of peer emotional support — captures who did what but contains no clinical salience categories (distress type, risk level, intervention)

Practical Guidance

What This Benchmark Measures

SAMSum measures abstractive summarization of informal English messenger chitchat — meeting arrangements, gossip, social planning, workplace coordination — by Western linguists, scored against a single English reference summary using primarily ROUGE. Its strongest dimensions are Input Form (plain-text speaker-attributed multi-party structure that partially transfers to clinical pipelines) and the documented metacognitive finding that ROUGE is unreliable for dialogue summarization [[Q3](#), [Q67](#)]. It does not measure clinical counseling capability, cross-lingual capability, or culturally sensitive communication.

Construct Depth

The benchmark probes general English dialogue summarization with moderate depth (16,369 dialogues, multi-party variation, register diversity). For the deployment construct of interest — SA-MENA clinical counseling summarization — it provides essentially zero direct evidence. The structural multi-speaker tracking and informal-register robustness evidence transfers partially, but no input contains clinical structure, no label encodes clinical salience, and no metric is valid cross-lingually. The AraHealthQA 2025 finding ([WEB-13](#)) and MentalCLOUDS American-demographic caveat ([WEB-18](#)) jointly suggest that even the best English clinical analogues require regional adaptation that SAMSum cannot supply.

What Else You Need

For valid SA-MENA clinical deployment evaluation, practitioners need: (a) a clinical counseling dialogue corpus — MentalCLOUDS ([WEB-15](#)) or MEMO/HOPE provide the closest English-language starting points but require regional adaptation; (b) target-language reference summaries in Arabic, Hindi, Urdu, Farsi, Dari, Bengali — none exist per regional YAML; (c) a clinical output ontology covering symptom disclosure, risk, alliance, and region-specific categories (religious reframing, family-mediated resolution, stigma-sensitive omission); (d) regional clinical-practitioner annotators stratified by sub-region; (e) human-in-the-loop evaluation rubrics validated for the target population (per [WEB-13](#)'s explicit call for Arabic mental health). SAMSum can at best contribute structural format-handling signal; all other dimensions require external supplementation.

ID	Evidence
Q3	"We show that model-generated summaries of dialogues achieve higher ROUGE scores than the model-generated summaries of news – in contrast with human evaluators' judgement." (p.1)
Q67	"We show that the most popular summarization metric ROUGE does not reflect the quality of a summary." (p.8)
WEB-13	AraHealthQA 2025 found BERTScore 'failed to fully measure appropriateness or trustworthiness' for Arabic mental health QA — direct corroboration that automated metrics are insufficient (https://arxiv.org/html/2508.20047v2)
WEB-15	MentalCLOUDS supplements ROUGE/BERTScore with expert evaluation across six clinical parameters — shows the form of evaluation SAMSum lacks (https://mental.jmir.org/2024/1/e57306)

WEB-18

MentalCLOUDS itself notes American demographic limitation — even closest English clinical analogue carries cultural caveats (<https://arxiv.org/pdf/2402.19052>)

Input Ontology (IO)

Score: 1/5 — Serious concern | Confidence: high

Benchmark: SAMSUM | Context: South Asia and MENA — Clinical NLP Researchers and AI Practitioners

Input Ontology definition: Whether the benchmark's test case categories cover the query types expected in deployment.

Each section below is followed by an evidence table. Three types of evidence are cited: **[QN]** — verbatim quotes extracted from the benchmark paper; **[WEB-N]** — web sources consulted during assessment; **[DN]** / **[DATASET-DN]** — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

SAMSum's input ontology is explicitly scoped to informal messenger chitchat — chit-chats, gossip, meeting arrangements, political discussion, and university assignment consultation [\[Q20\]](#) — positioned against news/scientific summarization [\[Q7\]](#), not clinical counseling. The benchmark contains no task category for CBT-structured counseling components, symptom elicitation, risk assessment, or culturally specific counseling acts (religious reframing, family-mediated resolution, izzat/namus-mediated indirect disclosure). Empirical sampling confirms zero clinical instances across 76 examples; the closest analogues are peer emotional support [\[DATASET-D4\]](#) and interpersonal relationship advice [\[DATASET-D5\]](#), neither of which involve therapeutic technique. Given the HIGH priority weight on IO for this deployment and the user's confirmation that both generic CBT structure and region-specific counseling components matter, the ontology fails on both axes.

ID	Evidence
Q7	"Major research efforts have focused so far on summarization of single-speaker documents like news (e.g., Nallapati et al. (2016)) or scientific publications (e.g., Nikolov et al. (2018))." (p.1)
Q20	"It includes chit-chats, gossiping about friends, arranging meetings, discussing politics, consulting university assignments with colleagues, etc." (p.2)
D4	Jessie: i'm still pretty overwhelmed / Jessie: but I was able to calm down a little after we spoke / Jessie: thank you for checking up on me — Closest analogue to mental health support in the dataset — peer emotional support, not clinical counseling
D5	Patrick: He's a jerk. He mistreats you. / Pearl: He does but he also shows me lots of affection / Patrick: You need to value yourself more. — Interpersonal conflict and emotional support — secular, Western framing; no family honor, religious, or stigma-mediated framing

Strengths

- Multi-party dialogue structure (~25% of conversations involve 3+ interlocutors) exercises speaker-tracking abilities that have structural analogues to group counseling and multi-party clinical encounters ([DATASET-D9](#), [DATASET-D15](#)).

ID	Evidence
D9	Mickey: Aokigahara Forest in Japan. / Kelly: Isn't it call the Suicide Forest? / Mickey: more ppl commit suicide on the Golden Gate than in that Forest — Suicide discussed casually in entertainment/trivia register — no clinical safety framing
D15	Terry: I have a question. Does anybody know what's the youngest country in the world / Kate: South Sudan — Trivia exchange — no clinical relevance

Checklist

Checklist item definitions:

ID	Assessment Question
IO-1	Identify test case categories required for regional deployment
IO-2	Check if taxonomy omits regionally relevant categories
IO-3	Check if taxonomy includes irrelevant categories
IO-4	Document category gaps harming content validity

Checklist responses:

- IO-1:** Required categories for SA-MENA clinical deployment include CBT-structured counseling (symptom elicitation, psychoeducation, cognitive restructuring), clinical safety content (risk assessment, suicidality), region-specific counseling acts (religious reframing per tawakkul/qadar, family-mediated resolution, stigma-sensitive omission), and somatic distress presentation common in South Asian and MENA contexts.
- IO-2:** All required clinical categories are omitted. The benchmark's stated topic taxonomy [Q20] is explicitly informal messenger genres; no clinical or counseling component is named. Empirical sampling confirms zero clinical instances (DATASET-D1, DATASET-D13, DATASET-D15) with the closest analogue being peer support [DATASET-D4].
- IO-3:** The benchmark includes substantial categories irrelevant to clinical deployment — gossip [DATASET-D7], trivia [DATASET-D15], workplace coordination (DATASET-D14, DATASET-D17), consumer transactions, and Western leisure planning — which introduce construct-irrelevant variance for clinical evaluation use.
- IO-4:** Content validity is severely harmed: the construct (clinical counseling dialogue summarization) is entirely underrepresented, while the represented constructs (informal social chitchat) do not transfer. The authors themselves note 'no obvious baseline for the task of dialogues summarization' [Q38], underscoring that the ontology was not constructed with clinical reference.

ID	Evidence
Q7	"Major research efforts have focused so far on summarization of single-speaker documents like news (e.g., Nallapati et al. (2016)) or scientific publications (e.g., Nikolov et al. (2018))." (p.1)
Q20	"It includes chit-chats, gossiping about friends, arranging meetings, discussing politics, consulting university assignments with colleagues, etc." (p.2)
Q38	"There is no obvious baseline for the task of dialogues summarization." (p.3)
D1	Vicky will call Sonia to entertain her as she's bored. — Mundane chitchat arrangement summary — no clinical content
D4	Jessie: i'm still pretty overwhelmed / Jessie: but I was able to calm down a little after we spoke / Jessie: thank you for checking up on me — Closest analogue to mental health support in the dataset — peer emotional support, not clinical counseling
D7	Renee: Gawd, she looks like a horse face! / Sue: She is just the utter worst. / Renee: Even her hair is nasty! — Gossip/social ridicule exchange
D13	Trudy: A leather jacket / Trudy: My cat peed on it... — Trivial domestic topic — no clinical relevance
D14	Sharon: Hope it'll be done by tomorrow, the clients are anxious to close this soon. — Workplace coordination dialogue
D15	Terry: I have a question. Does anybody know what's the youngest country in the world / Kate: South Sudan — Trivia exchange — no clinical relevance
D17	Eva: Meeting today with clients at 11 am sharp.. / Jim: Yes sir / Jim: Yes sir I have — Semi-formal workplace exchange; conventional English register
WEB-12	Arabic NLP challenges include morphological richness, dialectal variation, and diacritics — unaddressed by SAMSum form (https://www.mdpi.com/2227-9032/13/9/963)
WEB-15	MentalCLOUDS supplements ROUGE/BERTScore with expert evaluation across six clinical parameters — shows the form of evaluation SAMSum lacks (https://mental.jmir.org/2024/1/e57306)

Information Gaps

- Whether validation/test splits contain any clinical-adjacent content not seen in the 76-example train sample

Requires Expert Verification

- Magnitude of taxonomy gap for specific sub-regional counseling traditions (e.g., Sufi-influenced counseling in Pakistan, Ayurvedic-adjacent mental health framing in India)

Remediation

Gap 1: No clinical or regional counseling categories represented in the benchmark's task taxonomy.

Recommendation: Substitute or supplement SAMSum with MentalCLOUDS (WEB-15) for English clinical structure, then commission a SA-MENA practitioner-led category extension covering religious reframing, family-mediated resolution, stigma-sensitive omission, and somatic distress presentations. Use SAMSum only as a structural-format pretraining signal, not as a clinical evaluation.

ID	Evidence
WEB-15	MentalCLOUDS supplements ROUGE/BERTScore with expert evaluation across six clinical parameters — shows the form of evaluation SAMSum lacks (https://mental.jmir.org/2024/1/e57306)

Input Content (IC)

Score: 1/5 — Serious concern | Confidence: high

Benchmark: SAMSUM | Context: South Asia and MENA — Clinical NLP Researchers and AI Practitioners

Input Content definition: Whether individual datapoint content is culturally and contextually appropriate for the target region.

Each section below is followed by an evidence table. Three types of evidence are cited: [QN] — verbatim quotes extracted from the benchmark paper; [WEB-N] — web sources consulted during assessment; [DN] / [DATASET-DN] — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

Concrete dialogue instances were synthetically authored by linguists fluent in English at Samsung R&D Institute Poland [Q17, Q19], reflecting their daily messenger conversation patterns. Empirical sampling confirms uniform European/North American cultural content: UK geography [DATASET-D29], US political content [DATASET-D11], US country music [DATASET-D35], Western consumer brands [DATASET-D28], Western dating norms [DATASET-D25], and a single 'Moroccan' reference framed orientalizzingly by a European speaker [DATASET-D8]. No Arabic, Urdu, Hindi code-switching, izzat/namus framing, religious distress framing (tawakkul, qadar), or stigma-mediated indirect disclosure is present. Combined with the user's HIGH priority weight on IC and confirmation that linguistic mismatch is the primary concern for MENA tooling, this constitutes a severe content validity failure.

ID	Evidence
Q17	"Our dialogue summarization dataset contains natural messenger-like conversations created and written down by linguists fluent in English." (p.2)
Q19	"We asked linguists to create conversations similar to those they write on a daily basis, reflecting the proportion of topics of their real-life messenger conversations." (p.2)
D8	Helen: You are so unsophisticated! It's a Moroccan casserole, cooked in a special pot with a pointy lid. — Tagine described from an external, orientalizzing perspective by a European speaker
D11	Kate: Trump is so awesome and benevolent and nice and good and amazing... Also i like girls. Make america great again. Viva Trump — US political content; explicitly American cultural referent
D25	Ruby: i'm considering meeting someone online / Grace: yeah, totally, go for it! / Grace: My sis met her boyfriend online you know?! — Online dating normalized; no stigma framing — Western social norm assumption
D28	Lee: the new Blackwidow / Archie: not a big fan of Razer myself / Lee: most of my peripherals are from Razer — Consumer electronics discussion; entirely Western brand ecosystem
D29	Sally: Widemouth Bay, Sandymouth Bay, Cambourne, Bodmin Railway, Seaton Trams, PeccoRama. — UK-specific geographic references embedded in natural dialogue
D35	Nathan: highway dont care / David: I hate that song / Nathan: I hate Taylor Swift but Tim McGraw is ok — US country music cultural references

Strengths

- Informal register diversity — emoticons, abbreviations, typos, kaomoji (DATASET-D18, DATASET-D27) — partially overlaps with realistic messaging registers used in some digital mental health platform contexts.

ID	Evidence
D18	Ammalee: This lasted over a month. 🤔 (• o •) 🤔 🤔 (• o •) 🤔 / Maryann: Ah! Hello Ma'am! Thank you for the good review! — Emoticons and Unicode kaomoji used naturally in dialogue
D27	Ava: cause I don't want to talk to you anymore, so leave me the fuck alone / Ava: I'm blocking you on fb. Goodbye Greg — Profanity, direct assertive communication, social-media-mediated conflict resolution

Checklist

Checklist item definitions:

ID	Assessment Question
IC-1	Do inputs require region-specific cultural/geographic/dialectal knowledge?
IC-2	Does culturally sensitive content align with target culture?
IC-3	Flag inputs assuming Western-specific knowledge
IC-4	Regional annotator recruitment for culturally sensitive instances
IC-5	Document content issues harming content validity

Checklist responses:

IC-1: Inputs require extensive Western cultural knowledge (US politics, UK geography, US music, Western consumer brands, Western social norms around dating and alcohol) that does not transfer to SA-MENA deployment context ([DATASET-D11](#), [DATASET-D29](#), [DATASET-D35](#), [DATASET-D2](#)).

IC-2: Culturally sensitive content is misaligned: alcohol normalization [[DATASET-D23](#)] conflicts with majority-Muslim deployment contexts; Western individualist conflict resolution norms [[DATASET-D26](#)] conflict with family-mediated norms documented in regional context YAML; secular framing throughout conflicts with religious frameworks central to MENA distress expression.

IC-3: Multiple inputs require Western-specific knowledge: 'went Dutch' idiom [[DATASET-D2](#)], Trump parody register [[DATASET-D11](#)], Taylor Swift/Tim McGraw recognition [[DATASET-D35](#)], UK coastal place names [[DATASET-D29](#)]. This introduces construct-irrelevant variance into summarization scoring for SA-MENA practitioners.

IC-4: INSUFFICIENT DOCUMENTATION — no regional annotator review of input cultural sensitivity has been conducted. Would require recruiting South Asian and MENA reviewers to systematically tag culturally distractor content.

IC-5: Content validity is severely harmed: the synthetic Polish-linguist origin [[Q17](#)] combined with empirical confirmation of monocultural Western content ([DATASET-D11](#), [DATASET-D16](#), [DATASET-D29](#)) means there is no representative input signal for the target deployment populations' actual discourse.

ID	Evidence
Q17	"Our dialogue summarization dataset contains natural messenger-like conversations created and written down by linguists fluent in English." (p.2)
Q19	"We asked linguists to create conversations similar to those they write on a daily basis, reflecting the proportion of topics of their real-life messenger conversations." (p.2)
D2	Eric: Okay, we watched a movie, then headed to a nice restaurant. Eric: We went Dutch of course. — Post-date debrief; culturally Western social norms
D11	Kate: Trump is so awesome and benevolent and nice and good and amazing... Also i like girls. Make america great again. Viva Trump — US political content; explicitly American cultural referent
D16	Jeniffer: diversity! / Jeniffer: something we don't have in Europe to that extend — European speaker perspective explicitly stated; US-centric tourism content
D23	Chris: He was so wasted that he puked down between the drywall! / June: He could've choked to death! — Party/drinking culture; casual treatment of intoxication risk
D26	Sam: he was saying he doesn't like being my roommate / Naomi: honestly if it's bothering you that much you should talk to him — Roommate conflict; individual direct communication advised — Western individualist framing
D29	Sally: Widemouth Bay, Sandymouth Bay, Cambourne, Bodmin Railway, Seaton Trams, PeccoRama. — UK-specific geographic references embedded in natural dialogue
D35	Nathan: highway dont care / David: I hate that song / Nathan: I hate Taylor Swift but Tim McGraw is ok — US country music cultural references
WEB-12	Arabic NLP challenges include morphological richness, dialectal variation, and diacritics — unaddressed by SAMSum form (https://www.mdpi.com/2227-9032/13/9/963)

Information Gaps

- Exact distribution of cultural distractor density across the full 14,731 training instances
- Whether any examples in dev/test splits contain code-switching or non-Latin script content

Requires Expert Verification

- Magnitude of construct-irrelevant variance attributable to Western cultural distractors when SA-MENA-trained models are evaluated

Remediation

Gap 1: Dialogue instances are monolingual English with European/North American cultural content; no SA-MENA discourse, code-switching, or help-seeking norms are represented.

Recommendation: Build a complementary corpus of de-identified clinical or simulated counseling dialogues authored or transcribed by regionally representative speakers in target languages (Arabic dialects, Urdu, Hindi, Farsi/Dari, Bengali), with explicit annotation of culturally specific framings (izzat, namus, tawakkul, qadar). MentalQA (WEB-16) is a starting point for Arabic but is QA-format and small; new dialogue-format collection is needed.

ID	Evidence
WEB-16	https://arxiv.org/abs/2405.12619

Input Form (IF)

Score: 2/5 — Significant gaps | Confidence: high

Benchmark: SAMSUM | Context: South Asia and MENA — Clinical NLP Researchers and AI Practitioners

Input Form definition: Whether the input signal encoding (text, audio, image parameters) matches deployment conditions.

Each section below is followed by an evidence table. Three types of evidence are cited: **[QN]** — verbatim quotes extracted from the benchmark paper; **[WEB-N]** — web sources consulted during assessment; **[DN]** / **[DATASET-DN]** — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

The plain-text turn-by-turn speaker-attributed format [\[Q29, Q33\]](#) is structurally compatible with clinical NLP pipelines and is empirically clean and consistent ([DATASET-D17](#), [DATASET-D19](#)). Speaker attribution and utterance separators have been shown to benefit models [\[Q65, Q66\]](#). However, all formatting decisions, tokenization, and length calibrations [\[Q42\]](#) assume English plain text exclusively. There is no provision for Arabic RTL rendering, Urdu Nastaliq, Devanagari, morphologically rich tokenization, or cross-script code-switching common in South Asian clinical contexts. Empirical sampling confirms 100% English Latin-script content with no non-Latin or RTL examples ([DATASET-D17](#), [DATASET-D19](#)). The MODERATE priority weight reflects partial transferability of structural conventions with substantive multilingual mismatches.

ID	Evidence
Q29	"Beforehand, we specified a format for written dialogues with summaries: a colon should separate an author of utterance from its content, each utterance is expected to be in a separate line." (p.2)
Q33	"Each utterance contains the name of the speaker." (p.3)
Q42	"In all our experiments, news and dialogues are truncated to 400 tokens, and summaries – to 100 tokens." (p.4)
Q65	"On the other hand, dialogues are divided into utterances, and for each utterance its author is assigned." (p.7)
Q66	"We demonstrate in experiments that the models benefit from the introduction of separators, which mark utterances for each person." (p.7)
D17	Eva: Meeting today with clients at 11 am sharp.. / Jim: Yes sir / Jim: Yes sir I have — Semi-formal workplace exchange; conventional English register
D19	Tom: Did anyone pick up prints? / Glen: I didn't have time, but called John? — Minimal workplace coordination; turn-by-turn structure with speaker attribution

Strengths

- Consistent plain-text format compatible with standard NLP preprocessing pipelines ([DATASET-D17](#), [DATASET-D19](#))
- Multi-party dialogue structure with named speaker attribution exercises clinically relevant speaker-tracking [\[DATASET-D15\]](#)
- Format tolerates emoticons, kaomoji, and informal register variation without structural disruption [\[DATASET-D18\]](#)

ID	Evidence
D15	Terry: I have a question. Does anybody know what's the youngest country in the world / Kate: South Sudan — Trivia exchange — no clinical relevance
D17	Eva: Meeting today with clients at 11 am sharp.. / Jim: Yes sir / Jim: Yes sir I have — Semi-formal workplace exchange; conventional English register
D18	Ammalee: This lasted over a month.👉 (· o ·)👈👉 (· o ·)👈 / Maryann: Ah! Hello Ma'am! Thank you for the good review! — Emoticons and Unicode kaomoji used naturally in dialogue
D19	Tom: Did anyone pick up prints? / Glen: I didn't have time, but called John? — Minimal workplace coordination; turn-by-turn structure with speaker attribution

Checklist

Checklist item definitions:

ID	Assessment Question
IF-1	Compare signal distributions between source and target
IF-2	Check if regional infrastructure supports data capture specs
IF-3	Identify domain-specific form differences
IF-4	Document form mismatches harming external validity

Checklist responses:

IF-1: Signal distribution is text-only English Latin script. Target deployment requires Arabic RTL, Urdu Nastaliq, Devanagari, Bengali, and code-switched content per regional context YAML; none are present.

IF-2: Regional infrastructure supports text-based pipelines, but Arabic NLP for mental health faces unique challenges (morphological richness, dialectal variation, diacritics) per WEB-12, none of which the benchmark form addresses. Token truncation at 400 [Q42] was calibrated for English and may interact poorly with morphologically agglutinative target languages.

IF-3: Domain-specific form gaps include: no provision for code-switching markers (Urdu-English, Hindi-English clinical contexts noted in regional YAML), no script-direction handling, no diacritic preservation conventions for Arabic clinical text.

IF-4: External validity is harmed for multilingual deployment: the benchmark's input form provides no signal for non-English clinical inputs, while structural conventions (speaker attribution, turn separators) do partially transfer.

ID	Evidence
Q29	"Beforehand, we specified a format for written dialogues with summaries: a colon should separate an author of utterance from its content, each utterance is expected to be in a separate line." (p.2)
Q33	"Each utterance contains the name of the speaker." (p.3)
Q42	"In all our experiments, news and dialogues are truncated to 400 tokens, and summaries – to 100 tokens." (p.4)
Q65	"On the other hand, dialogues are divided into utterances, and for each utterance its author is assigned." (p.7)
Q66	"We demonstrate in experiments that the models benefit from the introduction of separators, which mark utterances for each person." (p.7)
WEB-12	Arabic NLP challenges include morphological richness, dialectal variation, and diacritics — unaddressed by SAMSum form (https://www.mdpi.com/2227-9032/13/9/963)

Information Gaps

- Token length distribution interactions with Arabic/Urdu morphology
- Whether placeholder tokens for media [DATASET-D18] align with clinical platform conventions

Requires Expert Verification

- Impact of English-calibrated truncation thresholds on cross-lingual fine-tuning

Remediation

Gap 1: Format conventions and tokenization are calibrated for English Latin script with no provision for RTL Arabic, Urdu Nastaliq, Devanagari, or code-switching.

Recommendation: Establish multilingual preprocessing pipelines (RTL-aware rendering, morphology-aware tokenization for Arabic and Urdu, code-switching markers) before any cross-lingual application of SAMSum-derived models. Recalibrate length thresholds for morphologically richer target languages.

Output Ontology (OO)

Score: 1/5 — Serious concern | Confidence: high

Benchmark: SAMSUM | Context: South Asia and MENA — Clinical NLP Researchers and AI Practitioners

Output Ontology definition: Whether the benchmark's output categories and scoring criteria reflect regionally valid decision boundaries.

Each section below is followed by an evidence table. Three types of evidence are cited: [QN] — verbatim quotes extracted from the benchmark paper; [WEB-N] — web sources consulted during assessment; [DN] / [DATASET-DN] — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

Each dialogue receives exactly one reference summary [Q24] under a four-point style guide (short, information-extractive, interlocutor-naming, third-person) [Q23]. There are no clinical sub-categories — no symptom identification, risk indicators, therapeutic intervention markers, or empathy/alliance categories — and no representation of region-specific counseling outputs (religious reframing, family-mediated resolution plans, stigma-sensitive omissions). Empirical sampling reveals concrete failure modes: suicide discussed as social reassurance trivia [DATASET-D10], abuse-victim self-derogation reproduced as factual summary content [DATASET-D6], peer distress reduced to social transaction [DATASET-D33]. The single-reference design forecloses representation of regional practitioner norm divergence. Given HIGH priority weight on OO and user confirmation that region-specific salience matters, this is a severe structural validity violation.

ID	Evidence
Q23	"After collecting all of the conversations, we asked language experts to annotate them with summaries, assuming that they should (1) be rather short, (2) extract important pieces of information, (3) include names of interlocutors, (4) be written in the third person." (p.2)
Q24	"Each dialogue contains only one reference summary." (p.2)
D6	Pearl's boyfriend mistreats her but she loves him anyway. Patrick is trying to convince her to leave him. Pearl is fat and ugly, so it's not so easy for her. — Summary reproduces a derogatory self-description as factual without clinical reframing — illustrates embedded Western annotator salience norms
D10	Mickey tries to calm her with the fact that more people commit suicide on the Golden Gate Bridge than in that forest. — Suicide statistics treated as social reassurance in summary — no clinical risk annotation or flagging
D33	Francine and Jessie talked yesterday. The talk has been helpful for Jessie. Francine says she's always there for Jessie. Jessie is very thankful. — Summary of peer emotional support — captures who did what but contains no clinical salience categories (distress type, risk level, intervention)

Strengths

- The benchmark's style guide [Q23] does enforce interlocutor-naming, which transfers to clinical summarization where speaker attribution of symptoms and statements is clinically important [DATASET-D20].

ID	Evidence
Q23	"After collecting all of the conversations, we asked language experts to annotate them with summaries, assuming that they should (1) be rather short, (2) extract important pieces of information, (3) include names of interlocutors, (4) be written in the third person." (p.2)
D20	Jimmy is going to take medication for a month to cure his acute gastritis. — Summary extracts action item and medical fact; third-person, attribution-focused

Checklist

Checklist item definitions:

ID	Assessment Question
OO-1	Review output label categories for regional relevance
OO-2	Identify missing regionally specific categories

OO-3	Flag categories encoding non-regional values
OO-4	Consider stakeholder-driven taxonomy redesign
OO-5	Document taxonomy issues harming structural/content validity

Checklist responses:

- OO-1:** Output label categories are limited to a single free-form summary per dialogue with implicit salience rules (action items, attribution, factual outcomes). No clinical or regionally relevant categories exist.
- OO-2:** Missing clinical categories include: symptom disclosure markers, risk/safety flags, therapeutic intervention type, alliance markers, family-mediated resolution annotations, religious reframing markers, and stigma-sensitive framing tags. Empirical evidence: suicide content received no risk annotation [DATASET-D10], distress disclosure received no clinical marker [DATASET-D33].
- OO-3:** Embedded values include: action/outcome salience over emotional content, individual rather than family-mediated framing, secular rather than religious framing, direct disclosure rather than face-saving omission. These conflict with documented MENA fatalism/tawakkul framing and South Asian izzat-mediated indirect communication norms.
- OO-4:** Stakeholder-driven taxonomy redesign would be required for valid SA-MENA clinical deployment. WEB-15 (MentalCLOUDS) provides one English-language template with three clinical components (symptom/history, patient discovery, reflection) but is itself American-demographic-limited (WEB-18).
- OO-5:** Structural validity is severely harmed: the construct's structure (clinical summarization) is misrepresented as social transaction logging. Content validity is severely harmed: clinical and regional categories are entirely absent. External validity is harmed: scoring against this ontology will not generalize to clinical practitioner judgments.

ID	Evidence
Q23	"After collecting all of the conversations, we asked language experts to annotate them with summaries, assuming that they should (1) be rather short, (2) extract important pieces of information, (3) include names of interlocutors, (4) be written in the third person." (p.2)
Q24	"Each dialogue contains only one reference summary." (p.2)
D10	Mickey tries to calm her with the fact that more people commit suicide on the Golden Gate Bridge than in that forest. — Suicide statistics treated as social reassurance in summary — no clinical risk annotation or flagging
D33	Francine and Jessie talked yesterday. The talk has been helpful for Jessie. Francine says she's always there for Jessie. Jessie is very thankful. — Summary of peer emotional support — captures who did what but contains no clinical salience categories (distress type, risk level, intervention)
WEB-13	AraHealthQA 2025 found BERTScore 'failed to fully measure appropriateness or trustworthiness' for Arabic mental health QA — direct corroboration that automated metrics are insufficient (https://arxiv.org/html/2508.20047v2)
WEB-15	MentalCLOUDS supplements ROUGE/BERTScore with expert evaluation across six clinical parameters — shows the form of evaluation SAMSum lacks (https://mental.jmir.org/2024/1/e57306)
WEB-18	MentalCLOUDS itself notes American demographic limitation — even closest English clinical analogue carries cultural caveats (https://arxiv.org/pdf/2402.19052)

Information Gaps

- Whether regional practitioners would converge on a single replacement taxonomy or require sub-regional variants

Requires Expert Verification

- Specific clinical and regional categories that SA-MENA practitioners would prioritize in a redesigned ontology

Remediation

Gap 1: Single-reference summary design with messenger-chitchat salience norms; no clinical or regional categories.

Recommendation: Adopt or adapt MentalCLOUDS' three-component schema (symptom/history, patient discovery, reflection) per WEB-15 as a baseline, then add region-specific categories (religious reframing markers, family-mediated resolution, stigma-sensitive omission flags). Move from single-reference to multi-reference design to capture legitimate practitioner variation.

ID	Evidence
WEB-15	MentalCLOUDS supplements ROUGE/BERTScore with expert evaluation across six clinical parameters — shows the form of evaluation SAMSum lacks (https://mental.jmir.org/2024/1/e57306)

Output Content (OC)

Score: 1/5 — Serious concern | Confidence: high

Benchmark: SAMSUM | Context: South Asia and MENA — Clinical NLP Researchers and AI Practitioners

Output Content definition: Whether ground-truth labels align with the judgments of regional stakeholders.

Each section below is followed by an evidence table. Three types of evidence are cited: **[QN]** — verbatim quotes extracted from the benchmark paper; **[WEB-N]** — web sources consulted during assessment; **[DN]** / **[DATASET-DN]** — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

Reference summaries were authored by the same linguists who created the dialogues [\[Q23\]](#), creating data-label circularity, with a 50-conversation double-annotation validation showing 94% agreement [\[Q26\]](#) but limited generalization. Annotator demographics are not reported; acknowledged contributors have Polish-consistent names [\[Q75\]](#). The user elicitation explicitly confirms annotators were 'general mental health experts from mixed backgrounds (including India and the USA)' but 'were not specifically briefed on regional context when validating summaries' and that 'their salience judgments likely reflect generic CBT protocols rather than consciously incorporating regional norms.' Empirical evidence reveals concrete content failures: a stigmatizing self-description reproduced as factual content [\[DATASET-D6\]](#), a speaker attribution error in a simple two-turn exchange [\[DATASET-D21\]](#), peer distress treated as social transaction [\[DATASET-D33\]](#). HIGH priority weight on OC and confirmed norm-divergence risk produce a severe convergent validity failure.

ID	Evidence
Q23	"After collecting all of the conversations, we asked language experts to annotate them with summaries, assuming that they should (1) be rather short, (2) extract important pieces of information, (3) include names of interlocutors, (4) be written in the third person." (p.2)
Q26	"The results revealed that 94% of examined dialogues were classified by both annotators as good i.e. they do look like conversations from a messenger app and could be condensed in a reasonable way." (p.2)
Q75	"We would like to express our sincere thanks to Tunia Blachno, Oliwia Ebebenge, Monika Jędras and Małgorzata Krawentek for their huge contribution to the corpus collection – without their ideas, management of the linguistic task and verification of examples we would not be able to create this paper." (p.9)
D6	Pearl's boyfriend mistreats her but she loves him anyway. Patrick is trying to convince her to leave him. Pearl is fat and ugly, so it's not so easy for her. — Summary reproduces a derogatory self-description as factual without clinical reframing — illustrates embedded Western annotator salience norms
D21	Miranda will be ready in no more than 10 minutes. — Summary misattributes action (Danny said he was almost finished, not Miranda) — illustrates annotation quality ceiling
D33	Francine and Jessie talked yesterday. The talk has been helpful for Jessie. Francine says she's always there for Jessie. Jessie is very thankful. — Summary of peer emotional support — captures who did what but contains no clinical salience categories (distress type, risk level, intervention)

Strengths

- Documented validation procedure: double-annotation of 50 conversations with 94% pass rate [\[Q25, Q26\]](#)
- Conflict resolution mechanism for opposing –1/+1 ratings via third annotator [\[Q49\]](#)
- Reasonable inter-annotator agreement on dialogues (Cohen's kappa 0.506) [\[Q51\]](#)

ID	Evidence
Q25	"We asked two linguists to doubly annotate 50 conversations in order to verify whether the dialogues could appear in a messenger app and could be summarized (i.e. a dialogue is not too general or unintelligible) or not (e.g. a dialogue between two people in a shop)." (p.2)
Q26	"The results revealed that 94% of examined dialogues were classified by both annotators as good i.e. they do look like conversations from a messenger app and could be condensed in a reasonable way." (p.2)
Q49	"We noticed a few annotations (7 for news and 4 for dialogues) with opposite marks (i.e. one annotator judgement was –1, whereas the second one was 1) and decided to have them annotated once again by another annotator who had to resolve conflicts." (p.5)

Q51	"For news examples, we obtained agreement on the level of 0.371 and for dialogues – 0.506." (p.5)
------------	---

Checklist

Checklist item definitions:

ID	Assessment Question
OC-1	Do ground-truth labels reflect regional stakeholder perspectives?
OC-2	Assess potential annotator-regional population disagreement
OC-3	Review annotator demographics from documentation
OC-4	Consider label re-annotation by regional annotators
OC-5	Review aggregation methods for minority perspective erasure
OC-6	Document label issues harming convergent/external validity

Checklist responses:

OC-1: Ground truth labels do not reflect regional stakeholder perspectives. Annotators were not briefed on SA-MENA clinical context per elicitation; salience judgments reflect generic CBT/Western messenger norms. Empirical evidence: [DATASET-D6](#) reproduces stigmatizing self-description without clinical reframing.

OC-2: Disagreement between original annotators and SA-MENA mental health practitioners is highly likely. MENA practitioners would likely flag missing religious/fatalistic framing in distress summaries; South Asian practitioners would flag missing family-mediated resolution and izzat-sensitive framing. [WEB-13](#) corroborates that even Arabic mental health evaluation requires human-in-the-loop validation absent from SAMSum.

OC-3: Annotator demographics are not reported in the paper [[Q75 contributors have Polish-consistent names; otherwise undisclosed](#)]. No Datasheet or Data Statement is referenced. The elicitation supplied partial information (mixed Indian/US backgrounds, no regional briefing) but no formal demographic breakdown.

OC-4: Re-annotation by an SA-MENA-stratified clinical practitioner pool is required. [WEB-15](#) demonstrates feasibility (5 clinical health professionals validated MentalCLOUDS) but no analogous SA-MENA effort exists per [WEB-13](#), [WEB-17](#).

OC-5: Aggregation uses a single reference summary per dialogue [[Q24](#)] with no multi-reference design, structurally erasing minority/regional perspectives. Conflicts are resolved by a third annotator [[Q49](#)] rather than preserved as legitimate variation.

OC-6: Convergent validity is severely harmed: labels do not correlate with regional practitioner salience norms. External validity is severely harmed: original judgments will not generalize. The [DATASET-D21](#) attribution error additionally undermines the 94% quality-check ceiling.

ID	Evidence
Q23	"After collecting all of the conversations, we asked language experts to annotate them with summaries, assuming that they should (1) be rather short, (2) extract important pieces of information, (3) include names of interlocutors, (4) be written in the third person." (p.2)
Q24	"Each dialogue contains only one reference summary." (p.2)
Q26	"The results revealed that 94% of examined dialogues were classified by both annotators as good i.e. they do look like conversations from a messenger app and could be condensed in a reasonable way." (p.2)
Q49	"We noticed a few annotations (7 for news and 4 for dialogues) with opposite marks (i.e. one annotator judgement was –1, whereas the second one was 1) and decided to have them annotated once again by another annotator who had to resolve conflicts." (p.5)
Q51	"For news examples, we obtained agreement on the level of 0.371 and for dialogues – 0.506." (p.5)
Q75	"We would like to express our sincere thanks to Tunia Błachno, Oliwia Ebebenge, Monika Jędras and Małgorzata Krawentek for their huge contribution to the corpus collection – without their ideas, management of the linguistic task and verification of examples we would not be able to create this paper." (p.9)
D6	Pearl's boyfriend mistreats her but she loves him anyway. Patrick is trying to convince her to leave him. Pearl is fat and ugly, so it's not so easy for her. — Summary reproduces a derogatory self-description as factual without clinical reframing — illustrates embedded Western annotator salience norms

D21	Miranda will be ready in no more than 10 minutes. — Summary misattributes action (Danny said he was almost finished, not Miranda) — illustrates annotation quality ceiling
WEB-13	AraHealthQA 2025 found BERTScore 'failed to fully measure appropriateness or trustworthiness' for Arabic mental health QA — direct corroboration that automated metrics are insufficient (https://arxiv.org/html/2508.20047v2)
WEB-15	MentalCLOUDS supplements ROUGE/BERTScore with expert evaluation across six clinical parameters — shows the form of evaluation SAMSum lacks (https://mental.jmir.org/2024/1/e57306)
WEB-17	ACM survey confirms no LLM-grade contextualized models for low-resourced South Asian languages in mental health (https://dl.acm.org/doi/10.1145/3638761)

Information Gaps

- Exact regional breakdown of the 'mixed backgrounds (including India and the USA)' annotator pool described in elicitation
- Whether any annotators had clinical mental health training

Requires Expert Verification

- Magnitude of label-level disagreement between SA-MENA clinical practitioners and existing reference summaries on sampled distress-relevant examples

Remediation

Gap 1: Annotator demographics not regionally stratified; salience norms reflect generic CBT/Western messenger conventions and have produced concrete anti-clinical labels (e.g., DATASET-D6 reproducing stigmatizing self-description as fact).

Recommendation: Re-annotate any clinical-adjacent examples (or build new annotation infrastructure on a clinical corpus) using SA-MENA-stratified clinical practitioner panels with explicit briefing on regional context. Disclose demographics formally via a Datasheet/Data Statement. Preserve practitioner disagreement as legitimate variation rather than resolving via third-annotator override.

ID	Evidence
D6	Pearl's boyfriend mistreats her but she loves him anyway. Patrick is trying to convince her to leave him. Pearl is fat and ugly, so it's not so easy for her. — Summary reproduces a derogatory self-description as factual without clinical reframing — illustrates embedded Western annotator salience norms

Output Form (OF)

Score: 1/5 — Serious concern | Confidence: high

Benchmark: SAMSUM | Context: South Asia and MENA — Clinical NLP Researchers and AI Practitioners

Output Form definition: Whether the expected output modality matches regional deployment needs and accessibility.

Each section below is followed by an evidence table. Three types of evidence are cited: [QN] — verbatim quotes extracted from the benchmark paper; [WEB-N] — web sources consulted during assessment; [DN] / [DATASET-DN] — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

Primary evaluation uses ROUGE F1 (R-1, R-2, R-L) [Q47] supplemented by a -1/0/1 human rating scale [Q48]. The benchmark itself documents that ROUGE is misleading for dialogue summarization — high ROUGE corresponds with low human ratings [Q3, Q54, Q67, Q68] — and explicitly calls for new metrics [Q69, Q74] without providing them. For SA-MENA deployment, this dialogue-summarization unreliability is compounded by complete inapplicability to non-English outputs: ROUGE/BERTScore against English references provide no valid signal for Arabic, Hindi, Urdu, Farsi, or Dari per the regional YAML and user elicitation. WEB-13 (AraHealthQA 2025) directly corroborates that BERTScore fails to capture clinical appropriateness for Arabic mental health tasks. Empirical sampling confirms 100% English content (DATASET-D19, DATASET-D17), with no multilingual signal whatsoever. HIGH priority weight on OF and complete cross-lingual misalignment yield a severe external validity failure.

ID	Evidence
Q3	"We show that model-generated summaries of dialogues achieve higher ROUGE scores than the model-generated summaries of news – in contrast with human evaluators' judgement." (p.1)
Q47	"We evaluate models with the standard ROUGE metric (Lin, 2004), reporting the F1 scores (with stemming) for ROUGE-1, ROUGE-2 and ROUGE-L following previous works (Chen and Bansal, 2018; See et al., 2017)." (p.4)
Q48	"We asked two linguists to mark the quality of every summary on the scale of -1, 0, 1, where -1 means that a summarization is poor, extracts irrelevant information or does not make sense at all, 1 – it is understandable and gives a brief overview of the text, and 0 stands for a summarization that extracts only a part of relevant information, or makes some mistakes in the produced summary." (p.5)
Q54	"Models generating dialogue summaries can obtain high ROUGE results, but their outputs are marked as poor by human annotators." (p.5)
Q67	"We show that the most popular summarization metric ROUGE does not reflect the quality of a summary." (p.8)
Q68	"We performed an independent, manual analysis of summaries and we demonstrated that high ROUGE results, obtained for automatically-generated dialogue summaries, correspond with lower evaluation marks given by human annotators." (p.8)
Q69	"We conclude that when measuring the quality of model-generated summaries, the ROUGE metrics are more indicative for news than for dialogues, and a new metric should be designed to measure the quality of abstractive dialogue summaries." (p.8)
Q74	"In order to perform well, it may require designing dedicated tools, but also new, non-standard measures to capture the quality of abstractive dialogue summaries in a relevant way." (p.9)
D17	Eva: Meeting today with clients at 11 am sharp.. / Jim: Yes sir / Jim: Yes sir I have — Semi-formal workplace exchange; conventional English register
D19	Tom: Did anyone pick up prints? / Glen: I didn't have time, but called John? — Minimal workplace coordination; turn-by-turn structure with speaker attribution
WEB-13	AraHealthQA 2025 found BERTScore 'failed to fully measure appropriateness or trustworthiness' for Arabic mental health QA — direct corroboration that automated metrics are insufficient (https://arxiv.org/html/2508.20047v2)

Strengths

- The benchmark explicitly documents ROUGE's unreliability for dialogue summarization [Q3, Q67, Q68], providing authored justification for users to supplement or replace the metric

- Manual human rating procedure with conflict resolution [Q48, Q49] establishes a precedent for human-in-the-loop evaluation that aligns with WEB-13's recommendations for Arabic mental health

ID	Evidence
Q3	"We show that model-generated summaries of dialogues achieve higher ROUGE scores than the model-generated summaries of news – in contrast with human evaluators' judgement." (p.1)
Q48	"We asked two linguists to mark the quality of every summary on the scale of –1, 0, 1, where –1 means that a summarization is poor, extracts irrelevant information or does not make sense at all, 1 – it is understandable and gives a brief overview of the text, and 0 stands for a summarization that extracts only a part of relevant information, or makes some mistakes in the produced summary." (p.5)
Q49	"We noticed a few annotations (7 for news and 4 for dialogues) with opposite marks (i.e. one annotator judgement was –1, whereas the second one was 1) and decided to have them annotated once again by another annotator who had to resolve conflicts." (p.5)
Q67	"We show that the most popular summarization metric ROUGE does not reflect the quality of a summary." (p.8)
Q68	"We performed an independent, manual analysis of summaries and we demonstrated that high ROUGE results, obtained for automatically-generated dialogue summaries, correspond with lower evaluation marks given by human annotators." (p.8)
WEB-13	AraHealthQA 2025 found BERTScore 'failed to fully measure appropriateness or trustworthiness' for Arabic mental health QA — direct corroboration that automated metrics are insufficient (https://arxiv.org/html/2508.20047v2)

Checklist

Checklist item definitions:

ID	Assessment Question
OF-1	Does output modality match regional deployment needs?
OF-2	Assess text-to-speech availability for speech output
OF-3	Consider literacy rates and accessibility requirements
OF-4	Document form mismatches harming external validity

Checklist responses:

OF-1: Expected output modality is English-language plain text scored by ROUGE against single English reference. Regional deployment requires Arabic, Hindi, Urdu, Farsi, Dari, Bengali, Pashto outputs — none supported.

OF-2: Not directly applicable (text-to-text task), but downstream clinical deployment may require speech inputs/outputs; the benchmark provides no signal here.

OF-3: Literacy and accessibility considerations: regional context notes ~1 in 4 women in Pakistan/Bangladesh cite literacy as a mobile-connectivity barrier (WEB-2); the benchmark's text-only English form does not engage with this constraint.

OF-4: External validity is severely harmed: ROUGE is unreliable even for the English dialogue task by the authors' own admission [Q67], and provides zero signal for non-English clinical outputs. Users would need entirely separate multilingual evaluation infrastructure (regional YAML, user elicitation), which the benchmark does not provide.

ID	Evidence
Q3	"We show that model-generated summaries of dialogues achieve higher ROUGE scores than the model-generated summaries of news – in contrast with human evaluators' judgement." (p.1)
Q47	"We evaluate models with the standard ROUGE metric (Lin, 2004), reporting the F1 scores (with stemming) for ROUGE-1, ROUGE-2 and ROUGE-L following previous works (Chen and Bansal, 2018; See et al., 2017)." (p.4)
Q67	"We show that the most popular summarization metric ROUGE does not reflect the quality of a summary." (p.8)
Q69	"We conclude that when measuring the quality of model-generated summaries, the ROUGE metrics are more indicative for news than for dialogues, and a new metric should be designed to measure the quality of abstractive dialogue summaries." (p.8)
WEB-2	Literacy barriers affecting mental health access in Pakistan/Bangladesh — accessibility considerations not addressed by benchmark (https://datareportal.com/reports/digital-2024-deep-dive-the-state-of-internet-adoption)
WEB-13	AraHealthQA 2025 found BERTScore 'failed to fully measure appropriateness or trustworthiness' for Arabic mental health QA — direct corroboration that automated metrics are insufficient (https://arxiv.org/html/2508.20047v2)

WEB-15	MentalCLOUDS supplements ROUGE/BERTScore with expert evaluation across six clinical parameters — shows the form of evaluation SAMSum lacks (https://mental.jmir.org/2024/1/e57306)
--------	--

Information Gaps

- Whether any cross-lingual BERTScore variant has been validated specifically for clinical dialogue summarization in any target language (regional YAML notes this is unestablished)

Requires Expert Verification

- Which specific human evaluation rubric dimensions SA-MENA practitioners would prioritize when supplementing ROUGE

Remediation

Gap 1: ROUGE/BERTScore against English references provide no valid signal for non-English outputs and are documented as unreliable even for English dialogues by the authors themselves.

Recommendation: Replace ROUGE-only evaluation with a multi-tier framework: (i) target-language reference corpora with cross-lingual metrics (e.g., AraBERT-based BERTScore, IndicBERT) where validated; (ii) human-in-the-loop expert clinical evaluation per WEB-13 and WEB-15 across rubric dimensions including clinical appropriateness, cultural fit, risk preservation, and stigma-sensitive framing; (iii) error-typology audits for safety-critical content (suicide, abuse, self-harm) given DATASET-D6, DATASET-D10 demonstrate concrete failure modes.

ID	Evidence
D6	Pearl's boyfriend mistreats her but she loves him anyway. Patrick is trying to convince her to leave him. Pearl is fat and ugly, so it's not so easy for her. — Summary reproduces a derogatory self-description as factual without clinical reframing — illustrates embedded Western annotator salience norms
D10	Mickey tries to calm her with the fact that more people commit suicide on the Golden Gate Bridge than in that forest. — Suicide statistics treated as social reassurance in summary — no clinical risk annotation or flagging
WEB-13	AraHealthQA 2025 found BERTScore 'failed to fully measure appropriateness or trustworthiness' for Arabic mental health QA — direct corroboration that automated metrics are insufficient (https://arxiv.org/html/2508.20047v2)
WEB-15	MentalCLOUDS supplements ROUGE/BERTScore with expert evaluation across six clinical parameters — shows the form of evaluation SAMSum lacks (https://mental.jmir.org/2024/1/e57306)