Validity Analysis: AFRIDOC_MT

Target context: Amhara Region Agricultural & Microfinance Document Translation — Cooperative and Rural MFI End-Users

Overall risk: **HIGH**

Deployment Context

Use case: Document-level translation of agricultural input subsidy notices, land-use policy updates, and credit-scheme terms from federal and regional bureaus into Amharic

Target population: Farmer cooperative leaders and rural microfinance clients (e.g., ACSI, ADCSI borrowers) in Amhara interpreting government and lender document

Dimension Scores

Dimension	Score	Rating	Confidence
Input Ontology ⚠	2	Significant gaps	high
Input Content ⚠	1	Serious concern	high
Input Form ✓	4	Minor gaps	high
Output Ontology ⚠	2	Significant gaps	high
Output Content ⚠	2	Significant gaps	high
Output Form ✓	3	Moderate gaps	medium
Average	**2.3**		

⚠ = highest concern ✓ = strongest dimension

Dimension Key

Abbr.	Dimension	Definition
IO	Input Ontology	Whether the benchmark's test case categories cover the query types expected in deployment.
IC	Input Content	Whether individual datapoint content is culturally and contextually appropriate for the target region.
IF	Input Form	Whether the input signal encoding (text, audio, image parameters) matches deployment conditions.
OO	Output Ontology	Whether the benchmark's output categories and scoring criteria reflect regionally valid decision boundaries.
OC	Output Content	Whether ground-truth labels align with the judgments of regional stakeholders.
OF	Output Form	Whether the expected output modality matches regional deployment needs and accessibility.

Overall Summary

AFRIDOC-MT is a high-quality document-level MT benchmark for five African languages including Amharic in Ethiopic script, but its fitness for evaluating Amharic agricultural and microfinance document translation in Amhara region is severely limited by three convergent gaps: (1) total domain absence — no agricultural, land-tenure, cooperative governance, or credit-scheme content exists in the benchmark, confirmed empirically across all sampled examples and corroborated by web research showing no other Amharic MT benchmark covers these domains either; (2) source-language direction mismatch — the benchmark translates from English to Amharic and exhibits acknowledged translationese, whereas the deployment operates on documents natively authored in Amharic proclamation register; and (3) annotator-authority mismatch — the four named

Amharic translators were prepared for health/IT domains with no documented agricultural or legal Amharic expertise, while the user-specified preferred ground-truth authority is federal bureau translators. Output scoring is calibrated for narrative discourse coherence with no metric for terminological consistency — the critical coherence requirement for legal-administrative documents. The benchmark's strongest alignment is on Input Form (text-only Amharic in Ethiopic script with documented tokenization accommodations) and Output Form (text output modality match). Of the three HIGH-priority dimensions (IO, IC, OC), all score 1-2.

Practical Guidance

What This Benchmark Measures

AFRIDOC-MT can credibly measure: (a) general document-level MT capability for English-to-Amharic in Ethiopic script using NLLB-200, GPT-4o, and other current systems [Q4, Q41]; (b) basic discourse coherence in narrative health and IT news [Q16, Q55]; (c) fluency-vs-accuracy tradeoffs for Amharic at sentence-level vs. pseudo-document granularities [Q61, Q72]; (d) which model architectures struggle with non-Latin script and morphological richness [Q53, Q127]. The strongest dimensions for the deployment are Input Form (4) and Output Form (3), reflecting that the benchmark technically handles Amharic text in Ethiopic script and produces text outputs at relevant granularities.

Construct Depth

Construct depth is shallow for the deployment's specific construct (legal-administrative Amharic translation quality). The benchmark probes general document-level MT depth in news domains via a multi-evaluation stack (d-chrF + d-BLEU + GPT-4o-judge + human DA), but the GPT-4o judge is explicitly unreliable for African languages [Q67, Q85, Q219], leaving human DA with Amharic Krippendorff's alpha of only 0.46 [Q71] as the residual signal. No measurement instrument exists for terminological consistency across clauses [WEB-25, WEB-26 confirm metric gap], structural fidelity for numbered clauses, register-appropriateness against Ethiopian proclamation conventions, or accessibility for variable-literacy end-users. Performance on AFRIDOC-MT does not predict performance on deployment-domain documents.

What Else You Need

Critical supplementation: (1) construct a deployment-domain test set drawing from publicly available Ethiopian proclamation Amharic (e.g., Proclamation 985/2016 [WEB-10], Proclamation 456/2005 [WEB-12]) and ACSI/EABC notices [WEB-3, WEB-5] — closing the IO and IC gaps; (2) recruit federal Ministry of Agriculture or Amhara Regional Bureau translators as annotators for that test set — closing the OC gap; (3) implement a terminology-consistency metric adapted from Semenov-Bojar or HHI [WEB-25, WEB-26] for proclamation-derived term categories — closing the OO gap; (4) add structural-fidelity and plain-language accessibility scoring if format-preserving or oral-mediation use cases are confirmed — closing the OF gap; (5) collect Amharic-source-language examples to address the source-direction mismatch.

ID	Evidence
Q4	"Our results indicate that NLLB-200 achieves the best average performance among the standard NMT models, while GPT-4o outperforms general-purpose LLMs." (p.1)
Q16	"Both document-level and context-aware MT allow for the possibility of improving translation quality for context-dependent phenomena such as coreference resolution (Müller et al., 2018; Bawden et al., 2018; Voita et al., 2018; Herold and Ney, 2023), lexical disambiguation (Rios Gonzales et al., 2017; Martínez Garcia et al., 2019), and lexical cohesion (Wong and Kit, 2012; Garcia et al., 2014, 2017; Bawden et al., 2018; Voita et al., 2019)." (p.2)
Q41	"We present the result of 12 models in total, including the 1.2B version of Toucan, 1.3B and 3.3B versions of NLLB-200, 3B and 13B versions of MADLAD-400 and Aya-101 respectively. We also have the 8B instruction tuned version of LLama3.1 (LLama3.1-IT), 9B version of Gemma-2 (Gemma2-IT), and LLaMAX3-Alpaca." (p.4)
Q53	"We report d-chrF scores for the best prompt per model and language direction in the main text, as chrF better captures the morphological richness of African languages (Adelani et al., 2022), with full results provided in Appendix C." (p.5)
Q55	"Following Sun et al. (2025), we assess each translated document's fluency, content errors (CE), and cohesion errors—specifically lexical (LE) and grammatical (GE) errors—using GPT-4o, with evaluation limited to a few model outputs due to cost constraints (Appendix B.6)." (p.5)

Q61	“Pseudo-document translation is worse than sentence-level translation when translating into African languages Our results from pseudo-document translation show a performance drop across different models compared to sentence-level translation, especially when translating into African languages.” (p.6)
Q67	“These inconsistencies raise concerns about GPT-4o's reliability.” (p.7)
Q71	“We obtained Krippendorff's alpha values of ≥ 0.40, which are relatively low due to the fine granularity of the evaluation scale.” (p.8)
Q72	“Human evaluation results align closely with d-chrF, which favors sentence-level translations over pseudo-document translations when translating into African languages.” (p.8)
Q85	“However, our DA and qualitative analysis found GPT-4o's judgments inconsistent for African languages, and sentence-by-sentence translation proved more effective for some languages.” (p.9)
Q127	“Our analysis revealed that Amharic and Yorùbá—languages with unique characteristics such as non-Latin scripts and diacritics, respectively—had the largest average token counts across the tokenizers.” (p.19)
Q219	“These results suggest that using GPT-4o as a translation judge is not yet well-suited for low-resource languages.” (p.26)
WEB-3	EABC and primary cooperative unions are the institutional terminology anchors for deployment, distinct from health/IT translator preparation (https://addisstandard.com/sowing-under-shadows-of-violence-amhara-farmers-face-bleak-harvest-as-conflict-deepens-fertilizer-shortages-drives-up-input-prices/)
WEB-5	ACSI serves ~2 million clients across all woredas; key MFI in deployment (https://www.findevgateway.org/sites/default/files/publications/files/mfg-en-case-study-a-review-of-ethiopian-microfinance-institutions-and-their-role-in-poverty-reduction-a-case-study-on-amhara-credit-and-saving-institution-acsi-jan-2011.pdf)
WEB-10	Ministry of Justice provides authoritative Amharic terminology — federal bureau translators are the preferred authority unrepresented in the benchmark (https://justice.gov.et/en/law/cooperative-societies-proclamation-2/)
WEB-12	Land Administration Proclamation 456/2005 defines land-tenure terminology absent from benchmark (https://land.igad.int/index.php/countries/37-countries/ethiopia/38-ethiopia-profile?start=2)
WEB-25	Semenov & Bojar (WMT 2022) terminology-consistency metric exists but not validated for African languages (https://aclanthology.org/2022.wmt-1.41/)
WEB-26	HHI applied to legal corpora terminology consistency — not validated for Amharic (https://www.researchgate.net/publication/358368108_Measuring_Terminology_Consistency_in_Translated_Corpora_Implementation_of_the_Herfindahl-Hirshman_Index)

Input Ontology (IO)

Score: 2/5 — Significant gaps | Confidence: high

Benchmark: *AFRIDOC_MT* | ***Context:*** *Amhara Region Agricultural & Microfinance Document Translation — Cooperative and Rural MFI End-Users*

Input Ontology definition: Whether the benchmark's test case categories cover the query types expected in deployment.

> Each section below is followed by an evidence table. Three types of evidence are cited: **[QN]** — verbatim quotes extracted from the benchmark paper; **[WEB-N]** — web sources consulted during assessment; **[DN]** / **[DATASET-DN]** — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

The benchmark's task taxonomy (document-level and sentence-level MT between English and five African languages, including Amharic) is structurally relevant — the deployment also needs document-level Amharic MT, and the user agreed that document-level structural demands are broadly similar across genres. However, the test-case category distribution covers only health (WHO) and IT news (Techpoint Africa) genres [Q1, Q2, Q80], with zero coverage of agricultural input subsidy notices, land-use policy updates, credit-scheme terms, or cooperative governance documents — the four genres specified for the deployment. The only legally-binding-text instance in the entire dataset is a translation of the WHO International Health Regulations [DATASET-D17], which does not approach the register of Ethiopian federal proclamations. Web research confirms this is a sector-wide gap: no Amharic MT benchmark covers Ethiopian government proclamation, agricultural subsidy, cooperative governance, or legal-administrative domains [WEB-14, WEB-22, WEB-15]. Given the user's HIGH priority weight on IO and the total absence of deployment-relevant test categories, the score is low — but not 1, because the abstract task type (document-level Amharic MT) is correctly represented and supports both sentence-level and pseudo-document subtasks at multiple chunk sizes [Q3, Q35, Q49] that broadly match deployment document length variation.

ID	Evidence
Q1	“This paper introduces AFRIDOC-MT, a document-level multi-parallel translation dataset covering English and five African languages: Amharic, Hausa, Swahili, Yorùbá, and Zulu.” (p.1)
Q2	“The dataset comprises 334 health and 271 information technology news documents, all human-translated from English to these languages.” (p.1)
Q3	“We conduct document-level translation benchmark experiments by evaluating the ability of neural machine translation (NMT) models and large language models (LLMs) to translate between English and these languages, at both the sentence and pseudo-document levels, the outputs being realigned to form complete documents for evaluation.” (p.1)
Q35	“Given the AFRIDOC-MT data, we conducted both sentence- and document-level translation, evaluating two types of models: encoder-decoder and decoder-only models. While the majority of these models are open-source, we also evaluated two proprietary models of the same type.” (p.4)
Q49	“To select an appropriate chunk size, we conducted initial tests with k = 1 (sentence-level), 5, 10, and 25, choosing k = 10 for our experiments.” (p.5)
Q80	“We introduce AFRIDOC-MT, a document-level translation dataset in the health and tech domains for five African languages.” (p.9)
D17	አይኤችአር 194 የዓለም ጤና ድርጅት አባላትን ጨምሮ በ196 አገሮች ላይ በሕግ አስገዳጅ የሆነ የዓለም አቀፍ ሕግ መሣሪያ ነው — Amharic: International Health Regulations — legally-binding international law framework; closest genre to legal-administrative register in the dataset
WEB-14	No public Amharic legal-administrative glossary resource exists (https://github.com/EthioNLP/Ethiopian-Language-Survey)
WEB-15	EthioLLM resources do not include legal-administrative MT data (https://aclanthology.org/2024.lrec-main.561.pdf)
WEB-22	EthioMT parallel corpus is general-domain only (https://arxiv.org/html/2403.19365v1)

Strengths

- Document-level MT task is correctly framed and explicitly motivated by context-dependent phenomena (coreference, lexical disambiguation, lexical cohesion) [Q16] that also matter for cross-clause legal documents
- Multiple chunk-size configurations (k=5, 10, 25, full document, sentence-level) [DATASET-D5, Q49] enable evaluation at granularities relevant to bureau document lengths

- Twelve diverse model configurations evaluated [Q41], including African-centric (Toucan), broad-coverage (NLLB-200, MADLAD-400), and decoder-only LLMs, providing a useful map of MT system options for Amharic

ID	Evidence
Q16	"Both document-level and context-aware MT allow for the possibility of improving translation quality for context-dependent phenomena such as coreference resolution (Müller et al., 2018; Bawden et al., 2018; Voita et al., 2018; Herold and Ney, 2023), lexical disambiguation (Rios Gonzales et al., 2017; Martínez Garcia et al., 2019), and lexical cohesion (Wong and Kit, 2012; Garcia et al., 2014, 2017; Bawden et al., 2018; Voita et al., 2019)." (p.2)
Q41	"We present the result of 12 models in total, including the 1.2B version of Toucan, 1.3B and 3.3B versions of NLLB-200, 3B and 13B versions of MADLAD-400 and Aya-101 respectively. We also have the 8B instruction tuned version of LLama3.1 (LLama3.1-IT), 9B version of Gemma-2 (Gemma2-IT), and LLaMAX3-Alpaca." (p.4)
Q49	"To select an appropriate chunk size, we conducted initial tests with k = 1 (sentence-level), 5, 10, and 25, choosing k = 10 for our experiments." (p.5)
D5	ይህ የሚያሳየው የደብሊውፒቪ1 ስርጭት ከአፍጋኒስታን (ምስራቅ ክልል) እና ፓኪስታን (ደቡብ ኬፒ) ዞኖች — Amharic: polio surveillance language, mentions sub-national zones and infection control measures

Checklist

Checklist item definitions:

ID	Assessment Question
IO-1	Identify test case categories required for regional deployment
IO-2	Check if taxonomy omits regionally relevant categories
IO-3	Check if taxonomy includes irrelevant categories
IO-4	Document category gaps harming content validity

Checklist responses:

IO-1: Required deployment categories are: (1) agricultural input subsidy notices, (2) land-use policy updates, (3) credit-scheme terms / loan agreements, (4) cooperative membership and governance notices — all proclamation-derived formal administrative register documents in Amharic. Hybrid genres (e.g., MFI documents combining credit-scheme and subsidy terminology, since ACSI distributes both [WEB-7]) are also relevant.

IO-2: Yes, the benchmark omits all four required categories. Source corpus is exclusively Techpoint Africa (tech news) and WHO (health news) [Q18, Q22], confirmed empirically across all 61 sampled examples [DATASET-D7, DATASET-D8, DATASET-D9, DATASET-D13]. Web research confirms no other Amharic MT benchmark covers these categories either [WEB-14, WEB-22].

IO-3: The IT/tech news subset (sports betting commentary [DATASET-D28], crypto/Web3 news [DATASET-D9], African tech entrepreneurship interviews [DATASET-D27]) is largely irrelevant to the deployment context and represents construct-irrelevant variance for an agricultural/microfinance use case. Health news has marginal relevance only via its policy-register subset (IHR, TRIPS) [DATASET-D17, DATASET-D18].

IO-4: Documented gaps: (a) total absence of agricultural proclamation register; (b) total absence of cooperative governance documents per Proclamation 985/2016 [WEB-8, WEB-10]; (c) total absence of land-tenure documents per Proclamation 456/2005 and Amhara Revised Proclamation [WEB-12, WEB-13]; (d) total absence of MFI loan-agreement instances. These constitute construct underrepresentation for the deployment's task category.

ID	Evidence
Q1	"This paper introduces AFRIDOC-MT, a document-level multi-parallel translation dataset covering English and five African languages: Amharic, Hausa, Swahili, Yorùbá, and Zulu." (p.1)
Q2	"The dataset comprises 334 health and 271 information technology news documents, all human-translated from English to these languages." (p.1)
Q18	"We scraped English articles from the websites of Techpoint Africa and the World Health Organization (WHO). The articles cover different topics of different lengths with an average length of 30 and 37 sentences for health and tech respectively." (p.3)

Q22	“Finally, we selected 334 and 271 English articles/documents from the health and tech domains respectively, which represents 10k sentences each per domain.” (p.3)
Q80	“We introduce AFRIDOC-MT, a document-level translation dataset in the health and tech domains for five African languages.” (p.9)
D7	ፔይ ደይ ለመሸጥ ይፈልጋል...ናይጄሪያ ፕሬዚዳንት አዲስ CBN ገዥን ሾሙ — Amharic: African technology news digest — IT/fintech domain, no agricultural or legal content
D8	ደቡብ አፍሪካ ጎጂ ይዘቶችን በመስመር ላይ ማቆም ትፈልጋለች...ናላ ከ ዩኬ እና የአውሮፓ ህብረት ለናይጄሪያ ክፍያዎችን ይጀምራል — Amharic: African fintech and content moderation news
D9	ኔስት ኮይን (Nestcoin) የሂሳብ መዛግብቱን ለማጠናከር እና የአዲሱን ምርት በኦንቦርድ እድገት — Amharic: crypto/Web3 fintech startup news
D13	Sahel Capital's SEFAA (Social Enterprise Fund for Agriculture in Africa) Fund — Tech sentence: agriculture-sector investment fund — only agricultural adjacent content in dataset
D17	አይኤችአር 194 የዓለም ጤና ድርጅት አባላትን ጨምሮ በ196 አገሮች ላይ በሕግ አስገዳጅ የሆነ የዓለም አቀፍ ሕግ መሣሪያ ነው — Amharic: International Health Regulations — legally-binding international law framework; closest genre to legal-administrative register in the dataset
D18	ከ ቲአርአይፒኤስ ስምምነት ከፀደቀበት ጊዜ አንሥቶ የዓለም የጤና ስብሰባዎች ብዙ ውሳኔዎች — Amharic: TRIPS agreement, intellectual property law in health sector — international health-law language
D27	ፕሮዳክት ዳይቭ እንዴት ነው የተጀመረው...ወደ ምርት አስተዳደር ጉዞ — Amharic: long-form interview with tech entrepreneur — informal, personal narrative register
D28	In dissecting the social impacts, it becomes vital to delve deep into real-life cases that embody both the positive and negative facets of sports betting in Kenya. — English: sports betting social impacts — not directly relevant to deployment context
WEB-7	https://www.findevgateway.org/sites/default/files/publications/files/mfg-en-paper-subsidizing-microcredit-interest-how-important-is-it-to-the-poor-0-2004.pdf
WEB-8	Cooperative Societies Proclamation 985/2016 defines the operative tier vocabulary (https://leap.unep.org/en/countries/et/national-legislation/cooperative-societies-proclamation-no-9852016)
WEB-10	Ministry of Justice provides authoritative Amharic terminology — federal bureau translators are the preferred authority unrepresented in the benchmark (https://justice.gov.et/en/law/cooperative-societies-proclamation-2/)
WEB-12	Land Administration Proclamation 456/2005 defines land-tenure terminology absent from benchmark (https://land.igad.int/index.php/countries/37-countries/ethiopia/38-ethiopia-profile?start=2)
WEB-13	Amhara Revised Proclamation provides regional land-tenure framework (https://ukgreencitiesandinfrastructure.org/wp-content/uploads/2019/07/ICED-Ethiopia-land-administration-background-paper.pdf)
WEB-14	No public Amharic legal-administrative glossary resource exists (https://github.com/EthioNLP/Ethiopian-Language-Survey)
WEB-22	EthioMT parallel corpus is general-domain only (https://arxiv.org/html/2403.19365v1)

Information Gaps

- Whether deployment intermediaries can compose a custom evaluation set drawing from public Ethiopian proclamation Amharic texts to supplement AFRIDOC-MT
- Whether the deployment also includes oral-mediation use cases (extension worker reading aloud) that would alter the input ontology

Requires Expert Verification

- Confirmation from Amhara Regional Bureau staff that the four listed deployment genres exhaust the operative document categories

---

Remediation

Gap 1: All four required deployment genres (agricultural input subsidy notices, land-use policy updates, credit-scheme terms, cooperative governance notices) are absent from the benchmark, and no Amharic MT benchmark covers them.

Recommendation: Construct a supplementary evaluation set of 50–200 Amharic deployment-domain documents drawn from publicly available federal/regional proclamations (985/2016, 456/2005, Amhara Revised 2006), EABC subsidy notices, and ACSI loan-agreement templates. Treat AFRIDOC-MT scores as a general-MT baseline only and require deployment-domain test scores for go/no-go decisions.

Gap 2: Hybrid genres (e.g., MFI loan agreements bundled with input-subsidy distribution terms via ACSI's dual government-distribution role [WEB-7]) are not anticipated by either the benchmark or by single-genre supplementary sets.

Recommendation: Include hybrid ACSI-issued documents (MFI + subsidy) in the supplementary deployment-domain test set, sourced via direct consultation with ACSI or Amhara Regional Bureau partners; do not assume single-genre coverage suffices.

ID	Evidence
WEB-7	https://www.findevgateway.org/sites/default/files/publications/files/mfg-en-paper-subsidizing-microcredit-interest-how-important-is-it-to-the-poor-0-2004.pdf

Input Content (IC)

Score: 1/5 — Serious concern | Confidence: high

Benchmark: *AFRIDOC_MT* | ***Context:*** *Amhara Region Agricultural & Microfinance Document Translation — Cooperative and Rural MFI End-Users*

Input Content definition: Whether individual datapoint content is culturally and contextually appropriate for the target region.

> Each section below is followed by an evidence table. Three types of evidence are cited: **[QN]** — verbatim quotes extracted from the benchmark paper; **[WEB-N]** — web sources consulted during assessment; **[DN]** / **[DATASET-DN]** — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

The input content gap is total. All English source documents are scraped from Techpoint Africa and WHO [Q18, Q22], and the Amharic translations were produced from this English source material with explicit acknowledgment of translationese risk [Q103, Q104, Q105]. The deployment requires Amharic content drawn from Ethiopian federal and regional proclamations, agricultural subsidy notices, land-tenure documents, MFI loan agreements, and cooperative governance materials — none of which appear anywhere in the corpus. Critical proclamation-derived term categories (cooperative tier names, land-tenure category labels, subsidy scheme names, MFI credit-scheme terms) are entirely absent, confirmed empirically across all sampled examples [DATASET-D9, DATASET-D4]. The source-language direction is also mismatched: the benchmark translates FROM English [Q103], whereas the deployment operates on documents natively authored in Amharic administrative register [WEB-10]. No public Amharic legal-administrative glossary exists [WEB-14, WEB-15] to serve as an external terminology anchor. Given the user's HIGH priority weight on IC and the empirical confirmation of total content absence, this dimension is at the floor.

ID	Evidence
Q18	“We scraped English articles from the websites of Techpoint Africa and the World Health Organization (WHO). The articles cover different topics of different lengths with an average length of 30 and 37 sentences for health and tech respectively.” (p.3)
Q22	“Finally, we selected 334 and 271 English articles/documents from the health and tech domains respectively, which represents 10k sentences each per domain.” (p.3)
Q103	“A potential limitation of our dataset is the influence of translationese (Koppel and Ordan, 2011). Since all source material translated originates in English, translated sentences in African languages may exhibit patterns such as unnatural syntax or overly literal phrasing.” (p.10)
Q104	“Although we have not conducted an analysis to quantify these effects, prior work suggests that they can affect MT model performance, generalization and evaluation including direct assessment (Freitag et al., 2019; Edunov et al., 2020).” (p.10)
Q105	“Furthermore, AFRIDOC-MT may reflect a bias toward English in terms of structure, semantics, and cultural framing. We leave a deeper investigation of these issues to future work.” (p.10)
D4	ፀረ-ተሕዋስያንን መቋቋም(AMR) የሚከሰተው ባክቴሪያ፣ ቫይረሶች፣ ፈንገሶች እና ጥገኛ ተህዋሲያን — Amharic: antimicrobial resistance — health domain, WHO
D9	ኔስት ኮይን (Nestcoin) የሂሳብ መዛግብቱን ለማጠናከር እና የአዲሱን ምርት በኦንቦርድ እድገት — Amharic: crypto/Web3 fintech startup news
WEB-10	Ministry of Justice provides authoritative Amharic terminology — federal bureau translators are the preferred authority unrepresented in the benchmark (https://justice.gov.et/en/law/cooperative-societies-proclamation-2/)
WEB-14	No public Amharic legal-administrative glossary resource exists (https://github.com/EthioNLP/Ethiopian-Language-Survey)
WEB-15	EthioLLM resources do not include legal-administrative MT data (https://aclanthology.org/2024.lrec-main.561.pdf)

Strengths

- WHO health-policy subset includes some legally-binding international-law references (IHR, TRIPS) [DATASET-D17, DATASET-D18] — the only register-adjacent content, useful as a weak analog for numbered-clause discourse structure
- Amharic content is rendered fluently in Ethiopic script by native-speaker translators [Q23, Q24, DATASET-D2]

ID	Evidence

Q23	“We translated the extracted 10k English sentences to the 5 African languages through 4 expert translators per language.” (p.3)
Q24	“The translators were recruited through a language coordinator who is also a native speaker of the language.” (p.3)
D2	አካል ጉዳተኝነት ሰው የመሆን አካል ነው — Amharic translation: "Disability is part of being human" — WHO health article
D17	አይኤችአር 194 የዓለም ጤና ድርጅት አባላትን ጨምሮ በ196 አገሮች ላይ በሕግ አስገዳጅ የሆነ የዓለም አቀፍ ሕግ መሣሪያ ነው — Amharic: International Health Regulations — legally-binding international law framework; closest genre to legal-administrative register in the dataset
D18	ከ ቲአርአይፒኤስ ስምምነት ከፀደቀበት ጊዜ አንሥቶ የዓለም የጤና ስብሰባዎች ብዙ ውሳኔዎች — Amharic: TRIPS agreement, intellectual property law in health sector — international health-law language

Checklist

Checklist item definitions:

ID	Assessment Question
IC-1	Do inputs require region-specific cultural/geographic/dialectal knowledge?
IC-2	Does culturally sensitive content align with target culture?
IC-3	Flag inputs assuming Western-specific knowledge
IC-4	Regional annotator recruitment for culturally sensitive instances
IC-5	Document content issues harming content validity

Checklist responses:

IC-1: Yes — deployment inputs require region-specific knowledge of Ethiopian federal/regional proclamations (985/2016 cooperative; 456/2005 land; Amhara Revised 2006 + 252/2017 amendment) [WEB-8, WEB-10, WEB-12, WEB-13], Amhara-specific institutional vocabulary (ACSI, EABC, primary cooperative unions) [WEB-3, WEB-5], and Amharic administrative register conventions distinct from Addis Ababa standard. The benchmark contains none of this [DATASET concerns 1, 3].

IC-2: The benchmark's WHO and Techpoint sources do not align with deployment culture. The user noted contemporary documents draw primarily from federal/regional proclamation-derived terminology, which is absent. Cultural framing of the benchmark is described as English-biased [Q105]; the deployment operates within Ethiopian Orthodox Christian Amhara cultural context [WEB-11, WEB-20] with cooperative-mediated document interaction [WEB-3].

IC-3: Tech subset includes Western/global startup vocabulary (venture capital, crypto, Web3) [DATASET-D9, DATASET-D19] and sports betting [DATASET-D28] that do not transfer. Health subset is WHO-globalized rather than Ethiopia-grounded.

IC-4: Not done in the benchmark; required for deployment evaluation. Federal bureau translators are the preferred annotation authority [elicitation Q3] but were not represented in benchmark creation [gap_3 web search].

IC-5: Critical content issues: (a) source-direction mismatch (English→Amharic vs. Amharic-source) [Q103, CRITICAL Concern 2]; (b) total absence of proclamation-derived terminology [CRITICAL Concern 3]; (c) translationese in reference Amharic [Q103, Q104, MAJOR Concern 3]; (d) no ACSI/EABC/cooperative scheme content [WEB-3, WEB-5]; (e) no Amhara region zone-specific administrative vocabulary.

ID	Evidence
Q3	“We conduct document-level translation benchmark experiments by evaluating the ability of neural machine translation (NMT) models and large language models (LLMs) to translate between English and these languages, at both the sentence and pseudo-document levels, the outputs being realigned to form complete documents for evaluation.” (p.1)
Q23	“We translated the extracted 10k English sentences to the 5 African languages through 4 expert translators per language.” (p.3)
Q24	“The translators were recruited through a language coordinator who is also a native speaker of the language.” (p.3)
Q103	“A potential limitation of our dataset is the influence of translationese (Koppel and Ordan, 2011). Since all source material translated originates in English, translated sentences in African languages may exhibit patterns such as unnatural syntax or overly literal phrasing.” (p.10)
Q104	“Although we have not conducted an analysis to quantify these effects, prior work suggests that they can affect MT model performance, generalization and evaluation including direct assessment (Freitag et al., 2019; Edunov et al., 2020).” (p.10)

Q105	“Furthermore, AFRIDOC-MT may reflect a bias toward English in terms of structure, semantics, and cultural framing. We leave a deeper investigation of these issues to future work.” (p.10)
D9	ኔስት ኮይን (Nestcoin) የሂሳብ መዛግብቱን ለማጠናከር እና የአዲሱን ምርት በኦንቦርድ እድገት — Amharic: crypto/Web3 fintech startup news
D19	ቬንቸር ካፒታሊስቶች(Venture capitalists) የ ኩባንያዎችን ማንነት የሚያገኙት — Amharic: startup/entrepreneurship discourse — tech sector, no agricultural or cooperative content
D28	In dissecting the social impacts, it becomes vital to delve deep into real-life cases that embody both the positive and negative facets of sports betting in Kenya. — English: sports betting social impacts — not directly relevant to deployment context
WEB-3	EABC and primary cooperative unions are the institutional terminology anchors for deployment, distinct from health/IT translator preparation (https://addisstandard.com/sowing-under-shadows-of-violence-amhara-farmers-face-bleak-harvest-as-conflict-deepens-fertilizer-shortages-drives-up-input-prices/)
WEB-5	ACSI serves ~2 million clients across all woredas; key MFI in deployment (https://www.findevgateway.org/sites/default/files/publications/files/mfg-en-case-study-a-review-of-ethiopian-microfinance-institutions-and-their-role-in-poverty-reduction-a-case-study-on-amhara-credit-and-saving-institution-acsi-jan-2011.pdf)
WEB-8	Cooperative Societies Proclamation 985/2016 defines the operative tier vocabulary (https://leap.unep.org/en/countries/et/national-legislation/cooperative-societies-proclamation-no-9852016)
WEB-10	Ministry of Justice provides authoritative Amharic terminology — federal bureau translators are the preferred authority unrepresented in the benchmark (https://justice.gov.et/en/law/cooperative-societies-proclamation-2/)
WEB-11	https://en.wikipedia.org/wiki/East_Gojjam_Zone
WEB-12	Land Administration Proclamation 456/2005 defines land-tenure terminology absent from benchmark (https://land.igad.int/index.php/countries/37-countries/ethiopia/38-ethiopia-profile?start=2)
WEB-13	Amhara Revised Proclamation provides regional land-tenure framework (https://ukgreencitiesandinfrastructure.org/wp-content/uploads/2019/07/ICED-Ethiopia-land-administration-background-paper.pdf)
WEB-20	https://en.wikipedia.org/wiki/North_Wollo_Zone
CRITICAL Concern 3	The benchmark contains zero instances of the critical term categories identified by the user: input subsidy scheme names, land-tenure category labels (rist-derived terms, usufruct rights, holding certificates), cooperative membership tier designations from Proclamation 985/2016 (ቀዳሚ የሕብረት ሥራ ማኅበር, የሕብረት ሥራ ማህበራት ዩኒዮን), or MFI credit-scheme terms (ACSI group lending terminology). The Amharic vocabulary in the dataset is drawn from health (disease names, medical procedures, nutrition, epidemiology) and IT (startup names, fintech terms, crypto vocabulary) — entirely non-overlapping with the deployment's operative terminology.

Information Gaps

- Whether ACSI loan-agreement Amharic templates and EABC subsidy notice templates are publicly available for use in a supplementary evaluation set
- Whether Amhara Regional Bureau translation unit maintains an internal glossary that could anchor terminology evaluation

Requires Expert Verification

- Native-speaker assessment of register-distance between AFRIDOC-MT Amharic and federal proclamation Amharic
- Identification of specific Amharic terms most likely to be mistranslated by AFRIDOC-MT-finetuned systems

---

Remediation

Gap 1: Critical proclamation-derived terminology (cooperative tier names per 985/2016, land-tenure category labels, subsidy scheme names, MFI credit-scheme terms) is entirely absent; source-language direction is mismatched (English→Amharic vs. native Amharic-source documents).

Recommendation: Compile a proclamation-terminology glossary using Ministry of Justice Amharic texts [WEB-10] and IGAD land-administration profile [WEB-12] as anchors. Build deployment test items in the deployment's actual source-language direction (Amharic source) rather than relying on AFRIDOC-MT's English source provenance. Document and probe specifically for translationese artifacts when evaluating models that have been exposed to AFRIDOC-MT-like training data.

ID	Evidence

WEB-10	Ministry of Justice provides authoritative Amharic terminology — federal bureau translators are the preferred authority unrepresented in the benchmark (https://justice.gov.et/en/law/cooperative-societies-proclamation-2/)
WEB-12	Land Administration Proclamation 456/2005 defines land-tenure terminology absent from benchmark (https://land.igad.int/index.php/countries/37-countries/ethiopia/38-ethiopia-profile?start=2)

Input Form (IF)

Score: 4/5 — Minor gaps | Confidence: high

Benchmark: *AFRIDOC_MT* | ***Context:*** *Amhara Region Agricultural & Microfinance Document Translation — Cooperative and Rural MFI End-Users*

Input Form definition: Whether the input signal encoding (text, audio, image parameters) matches deployment conditions.

> Each section below is followed by an evidence table. Three types of evidence are cited: **[QN]** — verbatim quotes extracted from the benchmark paper; **[WEB-N]** — web sources consulted during assessment; **[DN]** / **[DATASET-DN]** — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

Input form is the most aligned dimension and the user assigned it LOWER priority. Both benchmark and deployment are text-only, and Amharic is explicitly handled in Ethiopic (Ge'ez) script with documented accommodations: token limits were specifically increased for Amharic [Q162] because its non-Latin script produced larger token counts than English [Q127]. The benchmark adopts pseudo-document chunking with k=10 to fit Amharic within model context limits [Q47, Q48, Q128], a practical choice that the deployment will likely also need. Empirically, all sampled Amharic content is rendered consistently in Ethiopic script across all configurations [DATASET-D1, DATASET-D5, Strength 1]. The remaining concern is that the deployment may evolve toward format-preserving or multimodal translation (scanned PDFs, structured documents with tables and signature blocks) — a possibility flagged but unconfirmed in the regional context. Some annotation-workflow infrastructure (CSV/Google Sheets with paragraph-marking empty rows) [Q119, Q122] partially preserves document structure at the corpus level. Score of 4 reflects strong technical alignment on script and tokenization, with a minor gap on structural fidelity that becomes relevant only if deployment expands to format-preserving use cases.

ID	Evidence
Q47	"An initial analysis showed that some models were unable to process entire documents due to input length limits, which were exceeded by token counts in some languages (Amharic and Yorùbá)." (p.5)
Q48	"To address this, we adopted a similar approach to Lee et al. (2022), splitting documents into fixed-size chunks of k sentences to fit within token limits; the final chunk may contain fewer than k sentences." (p.5)
Q119	"Please do not delete double empty rows, as they serve to separate paragraphs. Also, avoid deleting any rows, columns, or provided text." (p.18)
Q122	"The files are in .csv format, and you can open them using Google Sheets or Microsoft Excel (for offline work)." (p.18)
Q127	"Our analysis revealed that Amharic and Yorùbá—languages with unique characteristics such as non-Latin scripts and diacritics, respectively—had the largest average token counts across the tokenizers." (p.19)
Q128	"To accommodate both languages in our experiments, we chose pseudo-documents with k=10." (p.19)
Q162	"Therefore, we increased the token limit specifically for Amharic." (p.21)
D1	ሞቃታማ አውሎ ነፋሶችን እንዲሁም ሀይለኛ አውሎ ነፋስ ወይም አውሎ ነፋሶች በመባል የሚታወቁት በጣም አጥፊ የአየር ሁኔታ ክስተቶች ናቸው — Amharic translation of WHO health content on tropical cyclones — health domain, Ethiopic script
D5	ይህ የሚያሳየው የደብሊውፒቪ1 ስርጭት ከአፍጋኒስታን (ምስራቅ ክልል) እና ፓኪስታን (ደቡብ ኬፒ) ዞኖች — Amharic: polio surveillance language, mentions sub-national zones and infection control measures

Strengths

- Amharic Ethiopic script is correctly handled across all configurations with explicit token-limit accommodation [Q162, DATASET-D1]
- Multiple chunk-size configurations (k=5, 10, 25, full document, sentence-level) [Q49, DATASET-D5] enable matching deployment document lengths
- Annotation workflow preserves paragraph-separating empty rows at corpus level [Q119]
- Token-length statistics across languages [Q33] demonstrate empirical engagement with Amharic morphological characteristics

ID	Evidence
Q33	“Table 4 shows the average number of whitespace-separated tokens per sentence across domains and languages, including English. The health domain has more tokens than tech. Hausa and Yorùbá are longer than English, likely due to their descriptive nature, while Swahili is similar in length. Amharic and Zulu are relatively shorter, reflecting interesting linguistic properties.” (p.4)
Q49	“To select an appropriate chunk size, we conducted initial tests with k = 1 (sentence-level), 5, 10, and 25, choosing k = 10 for our experiments.” (p.5)
Q119	“Please do not delete double empty rows, as they serve to separate paragraphs. Also, avoid deleting any rows, columns, or provided text.” (p.18)
Q162	“Therefore, we increased the token limit specifically for Amharic.” (p.21)
D1	ሞቃታማ አውሎ ነፋሶችን እንዲሁም ሀይለኛ አውሎ ነፋስ ወይም አውሎ ነፋሶች በመባል የሚታወቁት በጣም አጥፊ የአየር ሁኔታ ክስተቶች ናቸው — Amharic translation of WHO health content on tropical cyclones — health domain, Ethiopic script
D5	ይህ የሚያሳየው የደብሊውፒቪ1 ስርጭት ከአፍጋኒስታን (ምስራቅ ክልል) እና ፓኪስታን (ደቡብ ኬፒ) ዞኖች — Amharic: polio surveillance language, mentions sub-national zones and infection control measures

Checklist

Checklist item definitions:

ID	Assessment Question
IF-1	Compare signal distributions between source and target
IF-2	Check if regional infrastructure supports data capture specs
IF-3	Identify domain-specific form differences
IF-4	Document form mismatches harming external validity

Checklist responses:

IF-1: Signal distribution match is good. Both benchmark and deployment are text-only Amharic in Ethiopic script. Deployment delivery channel is print/SMS/WhatsApp [regional context: technology_access], all of which are text representations the benchmark covers.

IF-2: Amharic Ethiopic-script text infrastructure is supported in the benchmark; rural Amhara connectivity issues affect distribution but not the input form the MT system processes [WEB-18 on internet blockages affects delivery, not signal format].

IF-3: Domain-specific form differences: (a) deployment may include scanned PDFs requiring OCR [regional context: multimodal_future NEEDS VERIFICATION] — not assessed by AFRIDOC-MT; (b) deployment documents may include numbered clause structures, tables, signature blocks — partially preserved at corpus level [Q119] but no metric assesses preservation [Q57]; (c) Amharic register-specific token patterns of formal proclamation Amharic may differ from health/IT news Amharic — not directly addressed.

IF-4: Form mismatches harming external validity: minor — text-only modality matches; structural fidelity gap is the only documented mismatch and is contingent on multimodal evolution. INSUFFICIENT DOCUMENTATION on whether deployment receives plain text vs. structured documents; would need consultation with Amhara Regional Bureau.

ID	Evidence
Q33	“Table 4 shows the average number of whitespace-separated tokens per sentence across domains and languages, including English. The health domain has more tokens than tech. Hausa and Yorùbá are longer than English, likely due to their descriptive nature, while Swahili is similar in length. Amharic and Zulu are relatively shorter, reflecting interesting linguistic properties.” (p.4)
Q47	“An initial analysis showed that some models were unable to process entire documents due to input length limits, which were exceeded by token counts in some languages (Amharic and Yorùbá).” (p.5)
Q57	“In Tables 5 and 6 we present d-chrF scores based on the realigned documents, created by merging the translated sentences into their corresponding documents.” (p.5)
Q119	“Please do not delete double empty rows, as they serve to separate paragraphs. Also, avoid deleting any rows, columns, or provided text.” (p.18)
Q127	“Our analysis revealed that Amharic and Yorùbá—languages with unique characteristics such as non-Latin scripts and diacritics, respectively—had the largest average token counts across the tokenizers.” (p.19)

Q162	“Therefore, we increased the token limit specifically for Amharic.” (p.21)
WEB-18	Internet blockages in Amhara region affect document delivery channels but not text-input signal format (https://www.aljazeera.com/features/2025/3/10/we-just-want-peace-a-pacifist-community-amid-ethiopias-amhara-conflict)

Information Gaps

- Whether deployment receives documents as plain text, structured digital files, or scanned PDFs requiring OCR
- Per-token-count statistics for Amharic proclamation register vs. AFRIDOC-MT health Amharic

Requires Expert Verification

- Empirical token-length comparison between AFRIDOC-MT Amharic and bureau proclamation Amharic to assess whether the chosen chunk size generalizes

Output Ontology (OO)

Score: 2/5 — Significant gaps | Confidence: high

Benchmark: *AFRIDOC_MT* | ***Context:*** *Amhara Region Agricultural & Microfinance Document Translation — Cooperative and Rural MFI End-Users*

Output Ontology definition: Whether the benchmark's output categories and scoring criteria reflect regionally valid decision boundaries.

> Each section below is followed by an evidence table. Three types of evidence are cited: **[QN]** — verbatim quotes extracted from the benchmark paper; **[WEB-N]** — web sources consulted during assessment; **[DN]** / **[DATASET-DN]** — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

The output ontology is calibrated for news discourse: GPT-4o-judged fluency (1–5), content errors, lexical cohesion errors, and grammatical cohesion errors [Q54, Q55, Q180, Q191], plus d-chrF and d-BLEU [Q51, Q52] and human direct assessment on a 0–100 scale [Q195]. The benchmark explicitly defines lexical cohesion as 'vocabulary usage, incorrect or missing synonyms' [Q189] — narrative discourse flow, not legal terminological consistency. No rubric component evaluates whether named entities such as cooperative tier designations, land-tenure category labels, or subsidy scheme names are rendered identically across all clauses of a document — the critical coherence requirement for the deployment. Web research confirms terminology-consistency metrics exist (Semenov & Bojar WMT 2022; HHI) but have not been validated for Amharic [WEB-25, WEB-26]. Compounding this, the benchmark itself reports that GPT-4o-as-judge is unreliable for African language translation directions [Q67, Q85, Q99, Q219], leaving human DA as the residual signal — which had Krippendorff's alpha of only 0.46 for Amharic [Q71, Q211]. Given the user's MODERATE priority on OO and the partial match (d-chrF captures morphological richness [Q53], cohesion definitions touch on relevant phenomena), score is 2 rather than 1.

ID	Evidence
Q51	"Hence, we realigned sentence-level or pseudo-translation outputs into full documents, then computed BLEU (Papineni et al., 2002) and chrF (Popović, 2015) to create document BLEU (d-BLEU) and document chrF (d-chrF)." (p.5)
Q52	"Metrics were computed using SacreBLEU (Post, 2018) with bootstrap resampling (n = 1000) to report 95% confidence intervals." (p.5)
Q53	"We report d-chrF scores for the best prompt per model and language direction in the main text, as chrF better captures the morphological richness of African languages (Adelani et al., 2022), with full results provided in Appendix C." (p.5)
Q54	"We use GPT-4o as a judge to evaluate translation outputs, following recent work showing LLMs' effectiveness in assessing translation quality and analyzing errors (Wu et al., 2024; Sun et al., 2025)." (p.5)
Q55	"Following Sun et al. (2025), we assess each translated document's fluency, content errors (CE), and cohesion errors—specifically lexical (LE) and grammatical (GE) errors—using GPT-4o, with evaluation limited to a few model outputs due to cost constraints (Appendix B.6)." (p.5)
Q67	"These inconsistencies raise concerns about GPT-4o's reliability." (p.7)
Q71	"We obtained Krippendorff's alpha values of ≥ 0.40, which are relatively low due to the fine granularity of the evaluation scale." (p.8)
Q85	"However, our DA and qualitative analysis found GPT-4o's judgments inconsistent for African languages, and sentence-by-sentence translation proved more effective for some languages." (p.9)
Q99	"GPT-4o was employed to assess and rate the translation outputs of four models. While its ratings were consistent for translations into English, the same was not observed for translations into African languages, likely due to the model's limited understanding of these languages." (p.10)
Q180	"GPT-4o is prompted to rate the fluency of a document on a scale from 1 to 5, where 5 indicates high fluency and 1 represents low fluency." (p.22)
Q189	"Lexical Cohesion Mistakes: Issues with vocabulary usage, incorrect or missing synonyms, or overuse of certain words that disrupt the flow." (p.23)
Q191	"Fluency can only have values between 1 and 5. However, the other metrics, including CE, GE, and LE, do not have a specific range and can take on any value because they are counts." (p.24)

Q195	“Each annotator was assigned 80 documents to evaluate, tasked with marking as many error spans as possible and rating the overall quality on a scale from 0 to 100.” (p.24)
Q211	“We obtained α scores of 0.46, 0.57, 0.40, and 0.81, and 0.54 for Amharic, Hausa, Swahili, Yorùbá, and Zulu respectively.” (p.26)
Q219	“These results suggest that using GPT-4o as a translation judge is not yet well-suited for low-resource languages.” (p.26)
WEB-25	Semenov & Bojar (WMT 2022) terminology-consistency metric exists but not validated for African languages (https://aclanthology.org/2022.wmt-1.41/)
WEB-26	HHI applied to legal corpora terminology consistency — not validated for Amharic (https://www.researchgate.net/publication/358368108_Measuring_Terminology_Consistency_in_Translated_Corpora_Implementation_of_the_Herfindahl-Hirshman_Index)

Strengths

- d-chrF is explicitly chosen as primary because it captures morphological richness of African languages [Q53]
- Output ontology engages with discourse-level phenomena (cohesion, fluency) [Q187, Q188] that have some overlap with cross-clause requirements of legal documents
- Benchmark transparently flags GPT-4o judging unreliability for African languages [Q67, Q85, Q99, Q219], enabling informed substitution

ID	Evidence
Q53	“We report d-chrF scores for the best prompt per model and language direction in the main text, as chrF better captures the morphological richness of African languages (Adelani et al., 2022), with full results provided in Appendix C.” (p.5)
Q67	“These inconsistencies raise concerns about GPT-4o's reliability.” (p.7)
Q85	“However, our DA and qualitative analysis found GPT-4o's judgments inconsistent for African languages, and sentence-by-sentence translation proved more effective for some languages.” (p.9)
Q99	“GPT-4o was employed to assess and rate the translation outputs of four models. While its ratings were consistent for translations into English, the same was not observed for translations into African languages, likely due to the model's limited understanding of these languages.” (p.10)
Q187	“Cohesion refers to how different parts of a text are connected using language structures like grammar and vocabulary.” (p.23)
Q188	“It ensures that sentences flow smoothly and the text makes sense as a whole.” (p.23)
Q219	“These results suggest that using GPT-4o as a translation judge is not yet well-suited for low-resource languages.” (p.26)

Checklist

Checklist item definitions:

ID	Assessment Question
OO-1	Review output label categories for regional relevance
OO-2	Identify missing regionally specific categories
OO-3	Flag categories encoding non-regional values
OO-4	Consider stakeholder-driven taxonomy redesign
OO-5	Document taxonomy issues harming structural/content validity

Checklist responses:

OO-1: Output label categories (fluency 1-5, CE, LE, GE, d-chrF, d-BLEU, DA 0-100) are partially relevant — fluency and content-error scoring map to deployment quality concerns, but the categories are calibrated to news adequacy rather than legal-administrative precision [Q189, Q190].

OO-2: Missing categories specific to the deployment: (a) terminological consistency across clauses for proclamation-derived terms — no metric exists [WEB-25, WEB-26 confirm metric gap for Amharic]; (b) structural fidelity preservation for numbered clauses; (c) register-appropriateness scoring against Ethiopian proclamation register; (d) cross-reference accuracy for proclamation article citations.

OO-3: Categories assume news-discourse decision rules — 'lexical cohesion mistakes involve... overuse of certain words that disrupt the flow' [Q189] explicitly penalizes repetition, which is the opposite of what is needed for legal terminological consistency where identical repeat-rendering is required.

OO-4: Stakeholder-driven taxonomy redesign is needed: deployment scoring should add (1) terminology-consistency (e.g., HHI-based [WEB-26]) for proclamation terms; (2) register-appropriateness assessed by federal bureau translators [elicitation Q3]; (3) structural-fidelity scoring if format preservation is required.

OO-5: Issues harming structural validity: scoring rubric structurally misrepresents the construct of legal-administrative translation quality (treats consistent term repetition as a defect rather than a virtue [Q189]). Issues harming content validity: missing terminology-consistency, register, and structural categories. External validity at risk: GPT-4o judging shown unreliable for African languages [Q99, Q219].

ID	Evidence
Q3	“We conduct document-level translation benchmark experiments by evaluating the ability of neural machine translation (NMT) models and large language models (LLMs) to translate between English and these languages, at both the sentence and pseudo-document levels, the outputs being realigned to form complete documents for evaluation.” (p.1)
Q53	“We report d-chrF scores for the best prompt per model and language direction in the main text, as chrF better captures the morphological richness of African languages (Adelani et al., 2022), with full results provided in Appendix C.” (p.5)
Q55	“Following Sun et al. (2025), we assess each translated document's fluency, content errors (CE), and cohesion errors—specifically lexical (LE) and grammatical (GE) errors—using GPT-4o, with evaluation limited to a few model outputs due to cost constraints (Appendix B.6).” (p.5)
Q67	“These inconsistencies raise concerns about GPT-4o's reliability.” (p.7)
Q99	“GPT-4o was employed to assess and rate the translation outputs of four models. While its ratings were consistent for translations into English, the same was not observed for translations into African languages, likely due to the model's limited understanding of these languages.” (p.10)
Q189	“Lexical Cohesion Mistakes: Issues with vocabulary usage, incorrect or missing synonyms, or overuse of certain words that disrupt the flow.” (p.23)
Q190	“Grammatical Cohesion Mistakes: Problems with pronouns, conjunctions, or grammatical structures that link sentences and clauses.” (p.23)
Q219	“These results suggest that using GPT-4o as a translation judge is not yet well-suited for low-resource languages.” (p.26)
WEB-25	Semenov & Bojar (WMT 2022) terminology-consistency metric exists but not validated for African languages (https://aclanthology.org/2022.wmt-1.41/)
WEB-26	HHI applied to legal corpora terminology consistency — not validated for Amharic (https://www.researchgate.net/publication/358368108_Measuring_Terminology_Consistency_in_Translated_Corpora_Implementation_of_the_Herfindahl-Hirshman_Index)
WEB-27	JUST-NLP 2025 legal MT shared task uses BLEU/METEOR/COMET only, not for African languages (https://aclanthology.org/2025.justnlp-main.3.pdf)

Information Gaps

- Whether HHI or Semenov-Bojar terminology-consistency metrics can be implemented for Amharic with available resources
- Whether federal bureau translation units have internal scoring rubrics that could be adopted

Requires Expert Verification

- Federal bureau translator validation of which output categories matter most for binding legal documents

---

Remediation

Gap 1: No metric evaluates terminological consistency across clauses for proclamation-derived terms; lexical cohesion definition treats term repetition as a defect, opposite to legal-document requirements; GPT-4o judge unreliable for African languages.

Recommendation: Implement a term-consistency metric (Semenov-Bojar WMT 2022 [WEB-25] or HHI-based [WEB-26]) over a curated list of 30–80 deployment-critical Amharic terms. Replace GPT-4o-as-judge with native-speaker review by federal bureau translators for any high-stakes evaluation. Add structural-fidelity scoring (numbered-clause preservation) if deployment evolves toward format-preserving translation.

ID	Evidence
WEB-25	Semenov & Bojar (WMT 2022) terminology-consistency metric exists but not validated for African languages (https://aclanthology.org/2022.wmt-1.41/)
WEB-26	HHI applied to legal corpora terminology consistency — not validated for Amharic (https://www.researchgate.net/publication/358368108_Measuring_Terminology_Consistency_in_Translated_Corpora_Implementation_of_the_Herfindahl-Hirshman_Index)

Output Content (OC)

Score: 2/5 — Significant gaps | Confidence: high

Benchmark: *AFRIDOC_MT* | ***Context:*** *Amhara Region Agricultural & Microfinance Document Translation — Cooperative and Rural MFI End-Users*

Output Content definition: Whether ground-truth labels align with the judgments of regional stakeholders.

> Each section below is followed by an evidence table. Three types of evidence are cited: **[QN]** — verbatim quotes extracted from the benchmark paper; **[WEB-N]** — web sources consulted during assessment; **[DN]** / **[DATASET-DN]** — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

Reference Amharic translations were produced by four named native-speaker translators [Q23, Q109] under a controlled workflow with QE filtering at threshold 0.65 [Q28, Q29] and language-coordinator manual inspection [Q27]. However, the user's preferred ground-truth authority is domain-trained federal bureau translators [elicitation Q3], with regional bureau staff or extension workers as acceptable surrogates — and no documentation of institutional affiliation or agricultural/legal domain expertise exists for the four Amharic translators [Q109; gap_3 web search NOT FOUND]. The translators were prepared for health/IT domains via a terminology workshop [Q26], not for proclamation-derived agricultural or financial vocabulary. Inter-annotator agreement for Amharic was Krippendorff's alpha = 0.46 [Q71, Q211], described by the authors as 'relatively low' [Q71], even within the health domain in which translators were specifically prepared. Reference translations also exhibit translationese due to English source provenance [Q103, Q104]. Human evaluation used three native-speaker annotators per language [Q194], primarily drawn from the translation team [Q194], 80 documents each with 30 shared [Q197]. No demographics on annotator regional origin (Amhara region vs. other Amharic-speaking areas) are documented. Given the user's HIGH priority weight on OC and the multiple convergent gaps (annotator domain mismatch, low IAA, translationese, no federal-bureau-translator anchor), score is 2.

ID	Evidence
Q3	“We conduct document-level translation benchmark experiments by evaluating the ability of neural machine translation (NMT) models and large language models (LLMs) to translate between English and these languages, at both the sentence and pseudo-document levels, the outputs being realigned to form complete documents for evaluation.” (p.1)
Q23	“We translated the extracted 10k English sentences to the 5 African languages through 4 expert translators per language.” (p.3)
Q26	“Due to the domain-specific nature of the task, before starting the translation process, we conducted a translation workshop, during which three translation experts shared their experiences in creating terminologies and they also shared existing resources with the translators including short translation guidelines (Appendix A.1).” (p.3)
Q27	“Quality control was conducted using automated quality estimation, followed by manual inspections by our language coordinators.” (p.3)
Q28	“We also used Quality Estimation (QE), specifically AfriCOMET (Wang et al., 2024a), Given a translated sentence in any African language and its corresponding source English sentence, AfriCOMET generates a score between 0 and 1, where 0 indicates poor quality and higher values signify better quality.” (p.3)
Q29	“The translators, in collaboration with the language coordinators, were tasked with reviewing instances that had quality estimation scores below 0.65.” (p.3)
Q71	“We obtained Krippendorff's alpha values of ≥ 0.40, which are relatively low due to the fine granularity of the evaluation scale.” (p.8)
Q103	“A potential limitation of our dataset is the influence of translationese (Koppel and Ordan, 2011). Since all source material translated originates in English, translated sentences in African languages may exhibit patterns such as unnatural syntax or overly literal phrasing.” (p.10)
Q104	“Although we have not conducted an analysis to quantify these effects, prior work suggests that they can affect MT model performance, generalization and evaluation including direct assessment (Freitag et al., 2019; Edunov et al., 2020).” (p.10)
Q109	“Also, we also appreciate the translators whose names we have listed below: 1. Amharic: Bereket Tilahun, Hana M. Tamiru, Biniam Asmlash, Lidetewold Kebede 2. Hausa: Junaid Garba, Umma Abubakar, Jibrin Adamu, Ruqayya Nasir Iro 3. Swahili: Mohamed Mwinyimkuu, Laanyu koone, Baraka Karuli Mgasa, Said Athumani Said 4. Yorùbá: Ifeoluwa Akinsola, Simeon Onaolapo, Ganiyat Dasola Afolabi, Oluwatosin Koya 5. Zulu: Busisiwe Pakade, Rooweither Mabuya, Tholakele Zungu, Zanele Thembani” (p.11)

Q194	"For the human evaluation, three native speakers of the African languages—primarily translators involved in the dataset creation—were recruited." (p.24)
Q197	"To assess consistency and inter-annotator agreement, 30 of the 80 documents were shared among all three annotators." (p.24)
Q211	"We obtained α scores of 0.46, 0.57, 0.40, and 0.81, and 0.54 for Amharic, Hausa, Swahili, Yorùbá, and Zulu respectively." (p.26)

Strengths

- Reference translations were produced by named native speakers under controlled workflow with QE [Q23, Q28]
- Pre-translation workshop shared terminology resources [Q26], demonstrating awareness of domain-terminology importance
- Translators worked with full document context [Q25], appropriate for document-level coherence requirements
- Translators were instructed not to use MT engines [Q120], reducing risk of MT-contaminated reference data
- Compensation was disclosed ($1,250 / 2,500 sentences for translators [Q31]; $55.15 for human evaluators [Q198]), enabling labor-conditions transparency

ID	Evidence
Q23	"We translated the extracted 10k English sentences to the 5 African languages through 4 expert translators per language." (p.3)
Q25	"The 10k sentences were distributed equally among the translators and the translations were done in-context (i.e. the translators translated on the sentence level but had access to the whole document)." (p.3)
Q26	"Due to the domain-specific nature of the task, before starting the translation process, we conducted a translation workshop, during which three translation experts shared their experiences in creating terminologies and they also shared existing resources with the translators including short translation guidelines (Appendix A.1)." (p.3)
Q28	"We also used Quality Estimation (QE), specifically AfriCOMET (Wang et al., 2024a), Given a translated sentence in any African language and its corresponding source English sentence, AfriCOMET generates a score between 0 and 1, where 0 indicates poor quality and higher values signify better quality." (p.3)
Q31	"Each translator was paid $1, 250 for 2, 500 sentences." (p.3)
Q120	"Use the language field to input the translations. It is essential not to rely on translation engines, as our quality assurance process can detect this. Depending on such tools may result in potential issues that you would need to address, leading to additional work on your part." (p.18)
Q198	"Each annotator was remunerated with $55.15" (p.24)

Checklist

Checklist item definitions:

ID	Assessment Question
OC-1	Do ground-truth labels reflect regional stakeholder perspectives?
OC-2	Assess potential annotator-regional population disagreement
OC-3	Review annotator demographics from documentation
OC-4	Consider label re-annotation by regional annotators
OC-5	Review aggregation methods for minority perspective erasure
OC-6	Document label issues harming convergent/external validity

Checklist responses:

OC-1: No — ground-truth labels reflect health/IT translation judgments. The user-specified preferred authority (federal bureau translators with agricultural/legal Amharic training) is unrepresented [elicitation Q3; Q109 lists translator names without institutional affiliation; gap_3 web search confirmed no public information].

OC-2: Likely substantial disagreement on domain-specific term choices. Amhara regional bureau staff and federal Ministry of Agriculture translators would apply proclamation-derived standardized terminology [WEB-10, WEB-8]; AFRIDOC-MT translators applied health/IT register translations. Disagreement is structurally expected; the magnitude is not directly measured.

OC-3: Datasheet-level demographics absent. The paper names translators [Q109] and notes they were native speakers recruited via a coordinator [Q23, Q24] but provides no institutional affiliation, regional origin (Amhara-region vs. other), education, or domain training.

OC-4: Required for deployment validity. Re-annotation by Amhara Regional Bureau translation staff or federal Ministry of Agriculture translators on a deployment-domain test set is needed.

OC-5: Aggregation: human DA averages three annotators per language [Q69]. Krippendorff's alpha 0.46 for Amharic [Q71, Q211] indicates substantial inter-annotator divergence; averaging may erase legitimate domain-driven disagreements. Zulu evaluation could not be conducted at all due to recruitment difficulties [Q100], a precedent that flags fragility of the annotator-recruitment pipeline.

OC-6: Issues harming convergent validity: translators' health/IT focus likely fails to correlate with federal bureau translator judgment on agricultural/legal terminology [WEB-14 confirms no Amharic legal-administrative resources exist for cross-validation]. External validity: judgments do not generalize to deployment ground-truth authority. Empirical: low IAA [Q71], translationese [Q103], and no documented domain expertise [DATASET-MAJOR Concern 2].

ID	Evidence
Q3	"We conduct document-level translation benchmark experiments by evaluating the ability of neural machine translation (NMT) models and large language models (LLMs) to translate between English and these languages, at both the sentence and pseudo-document levels, the outputs being realigned to form complete documents for evaluation." (p.1)
Q23	"We translated the extracted 10k English sentences to the 5 African languages through 4 expert translators per language." (p.3)
Q24	"The translators were recruited through a language coordinator who is also a native speaker of the language." (p.3)
Q69	"In Table 10 we report average direct assessment (DA) scores (on a scale from 0 to 100) from three annotators per language for the health domain, when translating into four African languages." (p.7)
Q71	"We obtained Krippendorff's alpha values of ≥ 0.40, which are relatively low due to the fine granularity of the evaluation scale." (p.8)
Q100	"Therefore, we conducted a human evaluation for translations from English to African languages, comparing only three models due to cost constraints. However, we were unable to recruit annotators for Zulu." (p.10)
Q103	"A potential limitation of our dataset is the influence of translationese (Koppel and Ordan, 2011). Since all source material translated originates in English, translated sentences in African languages may exhibit patterns such as unnatural syntax or overly literal phrasing." (p.10)
Q109	"Also, we also appreciate the translators whose names we have listed below: 1. Amharic: Bereket Tilahun, Hana M. Tamiru, Biniam Asmlash, Lidetewold Kebede 2. Hausa: Junaid Garba, Umma Abubakar, Jibrin Adamu, Ruqayya Nasir Iro 3. Swahili: Mohamed Mwinyimkuu, Laanyu koone, Baraka Karuli Mgasa, Said Athumani Said 4. Yorùbá: Ifeoluwa Akinsola, Simeon Onaolapo, Ganiyat Dasola Afolabi, Oluwatosin Koya 5. Zulu: Busisiwe Pakade, Rooweither Mabuya, Tholakele Zungu, Zanele Thembani" (p.11)
Q211	"We obtained α scores of 0.46, 0.57, 0.40, and 0.81, and 0.54 for Amharic, Hausa, Swahili, Yorùbá, and Zulu respectively." (p.26)
WEB-3	EABC and primary cooperative unions are the institutional terminology anchors for deployment, distinct from health/IT translator preparation (https://addisstandard.com/sowing-under-shadows-of-violence-amhara-farmers-face-bleak-harvest-as-conflict-deepens-fertilizer-shortages-drives-up-input-prices/)
WEB-8	Cooperative Societies Proclamation 985/2016 defines the operative tier vocabulary (https://leap.unep.org/en/countries/et/national-legislation/cooperative-societies-proclamation-no-9852016)
WEB-10	Ministry of Justice provides authoritative Amharic terminology — federal bureau translators are the preferred authority unrepresented in the benchmark (https://justice.gov.et/en/law/cooperative-societies-proclamation-2/)
WEB-14	No public Amharic legal-administrative glossary resource exists (https://github.com/EthioNLP/Ethiopian-Language-Survey)

Information Gaps

- Institutional affiliation and domain expertise of the four named Amharic translators
- Regional origin of Amharic translators (Amhara region vs. Addis Ababa vs. diaspora)
- Whether any of the human evaluators had agricultural or legal Amharic exposure

Requires Expert Verification

- Federal bureau translator judgment on AFRIDOC-MT reference Amharic register-appropriateness for agricultural/legal documents
- Magnitude of expected disagreement between AFRIDOC-MT translators and federal bureau translators on a deployment-domain test set

---

Remediation

Gap 1: Reference Amharic translations were produced by translators with no documented agricultural or legal Amharic expertise; user-specified preferred authority (federal bureau translators) is absent; Amharic IAA is 0.46 even within the prepared health domain.

Recommendation: For deployment validation, recruit at least two federal Ministry of Agriculture or Amhara Regional Bureau translators to re-annotate a deployment-domain test set. Document annotator institutional affiliation, agricultural/legal Amharic experience, and regional origin. Report Krippendorff's alpha disaggregated by term-category to surface terminology-specific disagreement rather than averaging it away.

Output Form (OF)

Score: 3/5 — Moderate gaps | Confidence: medium

Benchmark: *AFRIDOC_MT* | ***Context:*** *Amhara Region Agricultural & Microfinance Document Translation — Cooperative and Rural MFI End-Users*

Output Form definition: Whether the expected output modality matches regional deployment needs and accessibility.

> Each section below is followed by an evidence table. Three types of evidence are cited: **[QN]** — verbatim quotes extracted from the benchmark paper; **[WEB-N]** — web sources consulted during assessment; **[DN]** / **[DATASET-DN]** — sampled datapoints from the benchmark dataset, with an interpretation note.

Justification

Output form is text-only translation [Q8, Q44, Q201], matching the deployment's text-only requirement. Document-level outputs are produced by realigning sentence-level or pseudo-document translations into full documents [Q51, Q57], partially preserving paragraph structure via maintained empty rows [Q119]. However, no metric assesses structural fidelity (preservation of numbered clauses, tables, signature blocks) [regional context: structure_preservation_gap]. The benchmark also reports specific output-form pathologies for African-language translations: under-generation, over-generation, repetition of words and phrases, off-target translations [Q6], and inconsistent diacritics [Q77] — these failure modes would be especially harmful in a legal-administrative deployment where exact term repetition matters. Embedding-based document-level evaluation for African languages remains unavailable [Q50, Q96, Q175, Q176]. Greedy decoding is used [Q155]. Given the user's MODERATE priority on OF and the partial match (text-only ✓; structural fidelity gap; no plain-language accessibility scoring for variable-literacy end-users), score is 3.

ID	Evidence
Q6	"Furthermore, our analysis reveals that some LLMs exhibit issues such as under-generation, over-generation, repetition of words and phrases, and off-target translations, specifically for translation into African languages." (p.1)
Q8	"We evaluate performance using automatic metrics and compare the results of encoder-decoder models with decoder-only LLMs" (p.1)
Q44	"Given that our created dataset can be used for sentence-level translation and as a baseline for document-level translation, we evaluate all models on the test splits for each domain." (p.5)
Q50	"Evaluating document-level translation remains challenging, as existing automatic metrics struggle to capture improvements and account for discourse phenomena (Jiang et al., 2022; Dahan et al., 2024), and embedding-based metrics have not been explored in this context for African languages due to the lack of data." (p.5)
Q51	"Hence, we realigned sentence-level or pseudo-translation outputs into full documents, then computed BLEU (Papineni et al., 2002) and chrF (Popović, 2015) to create document BLEU (d-BLEU) and document chrF (d-chrF)." (p.5)
Q57	"In Tables 5 and 6 we present d-chrF scores based on the realigned documents, created by merging the translated sentences into their corresponding documents." (p.5)
Q77	"However, for Yorùbá, it often uses inconsistent or partial diacritics, resulting in inaccuracies." (p.8)
Q96	"Quality evaluation in MT is an open and ongoing area of research, especially for document-level translation. Recent works have proposed embedding-based metrics for evaluation at both the sentence and document levels. While this has been well explored for high-resource language pairs, it remains underexplored for African languages, although there is a tool, AfriCOMET, that works for sentence-level evaluation in African languages." (p.10)
Q119	"Please do not delete double empty rows, as they serve to separate paragraphs. Also, avoid deleting any rows, columns, or provided text." (p.18)
Q155	"In all cases, greedy decoding was used." (p.21)
Q175	"For evaluation, we used BLEU and chrF scores but excluded AfriCOMET due to its backbone model, AfroXLM-R-L (Alabi et al., 2022; Adelani et al., 2024a), having a context length of 512 tokens." (p.22)
Q176	"This made it impractical to compute COMET scores for document-level outputs." (p.22)
Q201	"Given that AFRIDOC-MT is a document-level translation dataset, and due to the limited context length of most translation models and LLMs, which makes it impossible to translate a full document at once, we opted to translate the sentences within the documents and then merge them back to form the complete document. For document-level evaluation, we split the documents into chunks of 10 sentences and translate these chunks using the different models." (p.24)

Strengths

- Output modality (text in Ethiopic script) directly matches deployment text-only requirement [Q8]
- Paragraph structure preserved at corpus level via empty-row convention [Q119]
- Sentence-to-document realignment supports document-level evaluation [Q51, Q57]
- Output failure modes (under/over-generation, repetition, off-target, diacritic inconsistency) [Q6, Q77] are documented, enabling targeted detection in deployment

ID	Evidence
Q6	"Furthermore, our analysis reveals that some LLMs exhibit issues such as under-generation, over-generation, repetition of words and phrases, and off-target translations, specifically for translation into African languages." (p.1)
Q8	"We evaluate performance using automatic metrics and compare the results of encoder-decoder models with decoder-only LLMs" (p.1)
Q51	"Hence, we realigned sentence-level or pseudo-translation outputs into full documents, then computed BLEU (Papineni et al., 2002) and chrF (Popović, 2015) to create document BLEU (d-BLEU) and document chrF (d-chrF)." (p.5)
Q57	"In Tables 5 and 6 we present d-chrF scores based on the realigned documents, created by merging the translated sentences into their corresponding documents." (p.5)
Q77	"However, for Yorùbá, it often uses inconsistent or partial diacritics, resulting in inaccuracies." (p.8)
Q119	"Please do not delete double empty rows, as they serve to separate paragraphs. Also, avoid deleting any rows, columns, or provided text." (p.18)

Checklist

Checklist item definitions:

ID	Assessment Question
OF-1	Does output modality match regional deployment needs?
OF-2	Assess text-to-speech availability for speech output
OF-3	Consider literacy rates and accessibility requirements
OF-4	Document form mismatches harming external validity

Checklist responses:

OF-1: Yes — text output modality matches deployment [Q8, Q44]. Both benchmark and deployment produce Amharic text translations.

OF-2: TTS not assessed by the benchmark. The deployment context indicates extension workers read documents aloud to less-literate cooperative members [regional context: oral_tradition_and_document_mediation], creating an implicit TTS-relevant accessibility requirement that AFRIDOC-MT does not address. INSUFFICIENT DOCUMENTATION on whether deployment system itself includes TTS.

OF-3: Literacy considerations: Ethiopia adult literacy is 51.8% (2017) [WEB-16, WEB-17]; Amhara-region-specific and rural-Amhara-specific rates are not publicly available [regional context: NOT FOUND]. The benchmark does not assess accessibility or comprehension of translated output for variable-literacy readers; reading-level / plain-language scoring is absent. End-users are mediated by cooperative leaders and extension workers, partially mitigating the literacy gap.

OF-4: Form mismatches harming external validity: (a) no structural-fidelity scoring for numbered clauses/tables/signature blocks [Q119 partial preservation only]; (b) no plain-language or readability scoring for variable-literacy readers; (c) documented LLM output pathologies (repetition, off-target) [Q6] would be especially harmful for legal terminology stability in deployment.

ID	Evidence

Q6	“Furthermore, our analysis reveals that some LLMs exhibit issues such as under-generation, over-generation, repetition of words and phrases, and off-target translations, specifically for translation into African languages.” (p.1)
Q8	“We evaluate performance using automatic metrics and compare the results of encoder-decoder models with decoder-only LLMs” (p.1)
Q44	“Given that our created dataset can be used for sentence-level translation and as a baseline for document-level translation, we evaluate all models on the test splits for each domain.” (p.5)
Q77	“However, for Yorùbá, it often uses inconsistent or partial diacritics, resulting in inaccuracies.” (p.8)
Q119	“Please do not delete double empty rows, as they serve to separate paragraphs. Also, avoid deleting any rows, columns, or provided text.” (p.18)
WEB-16	Ethiopia adult literacy 51.8% (2017) — variable-literacy end-users not addressed by benchmark scoring (https://www.theglobaleconomy.com/Ethiopia/Literacy_rate/)
WEB-17	World Bank confirms Ethiopia literacy figure (https://data.worldbank.org/indicator/SE.ADT.LITR.ZS?locations=ET)

Information Gaps

- Amhara-region and rural-Amhara literacy rates
- Whether deployment system includes TTS or remains text-only
- Whether deployment delivery format requires preservation of numbered-clause / table / signature-block structure

Requires Expert Verification

- Cooperative leader and extension worker assessment of translation comprehensibility for end-user mediation

---

Remediation

Gap 1: No structural-fidelity scoring for numbered clauses, tables, or signature blocks; no readability/accessibility scoring for variable-literacy end-users mediated by cooperative leaders and extension workers.

Recommendation: If deployment delivers structured documents, add a structural-fidelity check (template-preservation rate). Pilot a comprehension test with cooperative leaders and extension workers from one or two woredas (where conflict access permits [WEB-4, WEB-21]) to assess plain-language accessibility of translated output, especially for read-aloud mediation.

ID	Evidence
WEB-4	https://www.gov.uk/government/publications/ethiopia-country-policy-and-information-notes/country-policy-and-information-note-amhara-and-amhara-opposition-groups-ethiopia-june-2025-accessible
WEB-21	https://www.unocha.org/publications/report/ethiopia/ethiopia-situation-report-13-december-2024